

End-to-End Speech Recognition by Following my Development History

Shinji Watanabe

Center for Language and Speech Processing

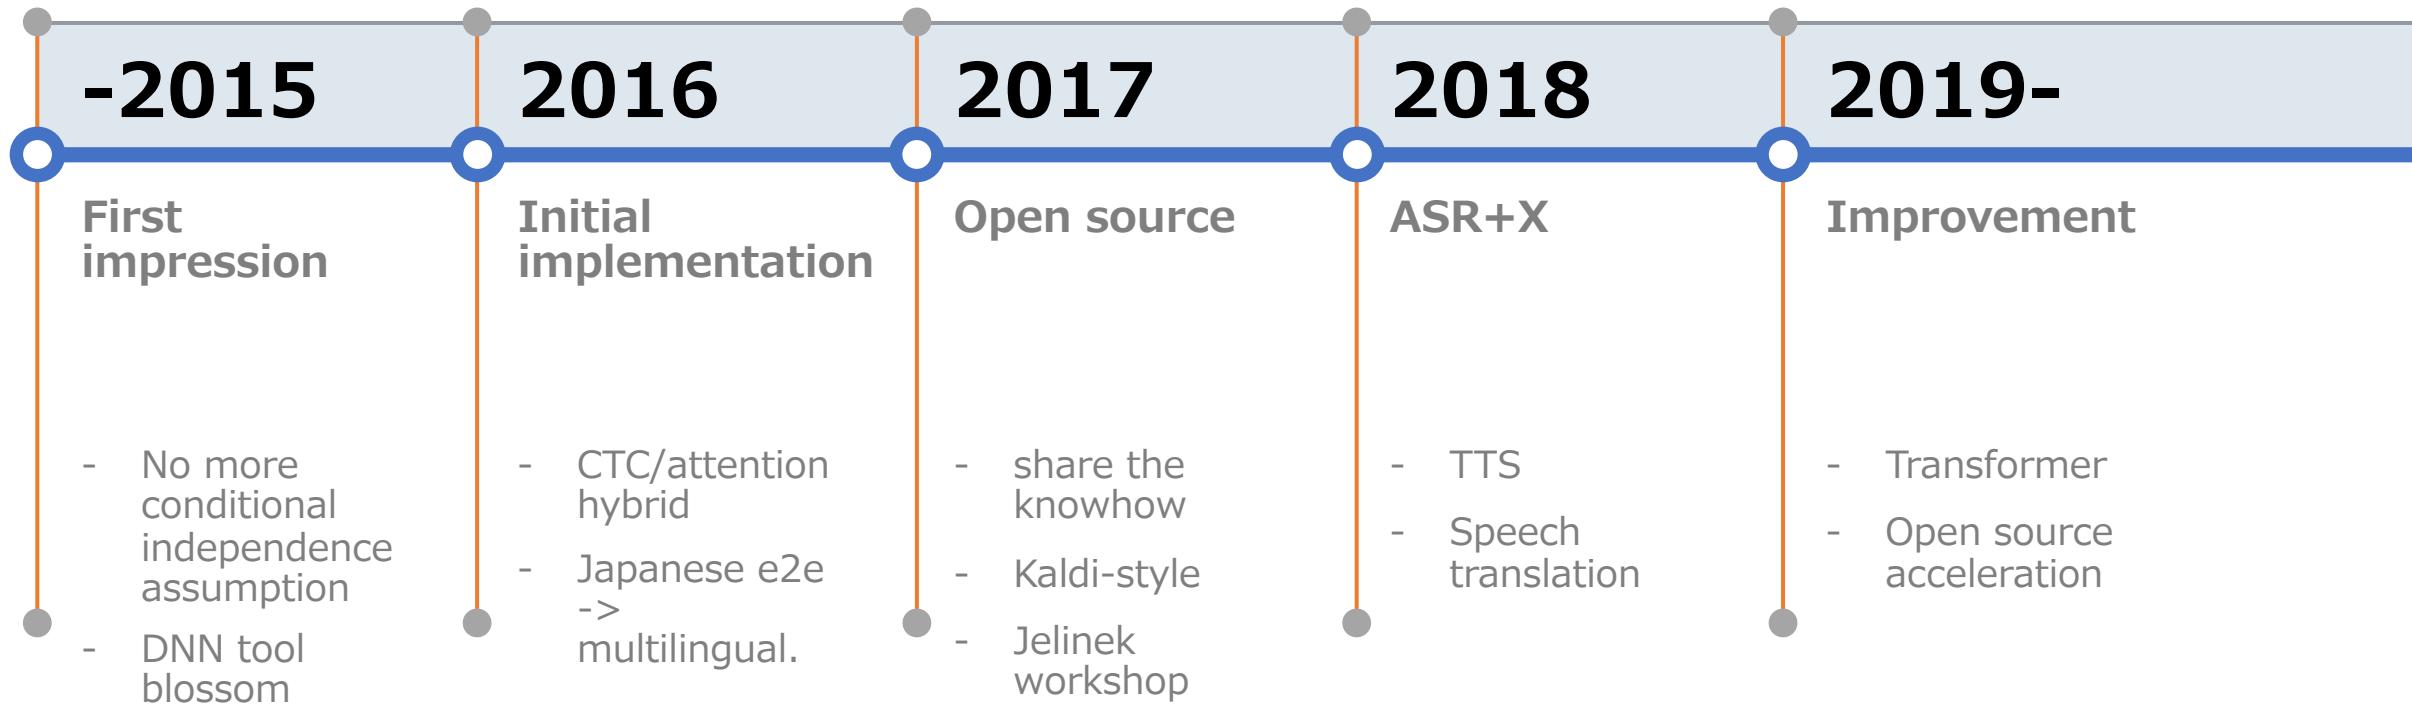
Johns Hopkins University

About this presentation

- This is based on my personal experience
- I re-order or re-structure several existing materials based on a chronological order
- I'm assuming people have some end-to-end ASR knowledge

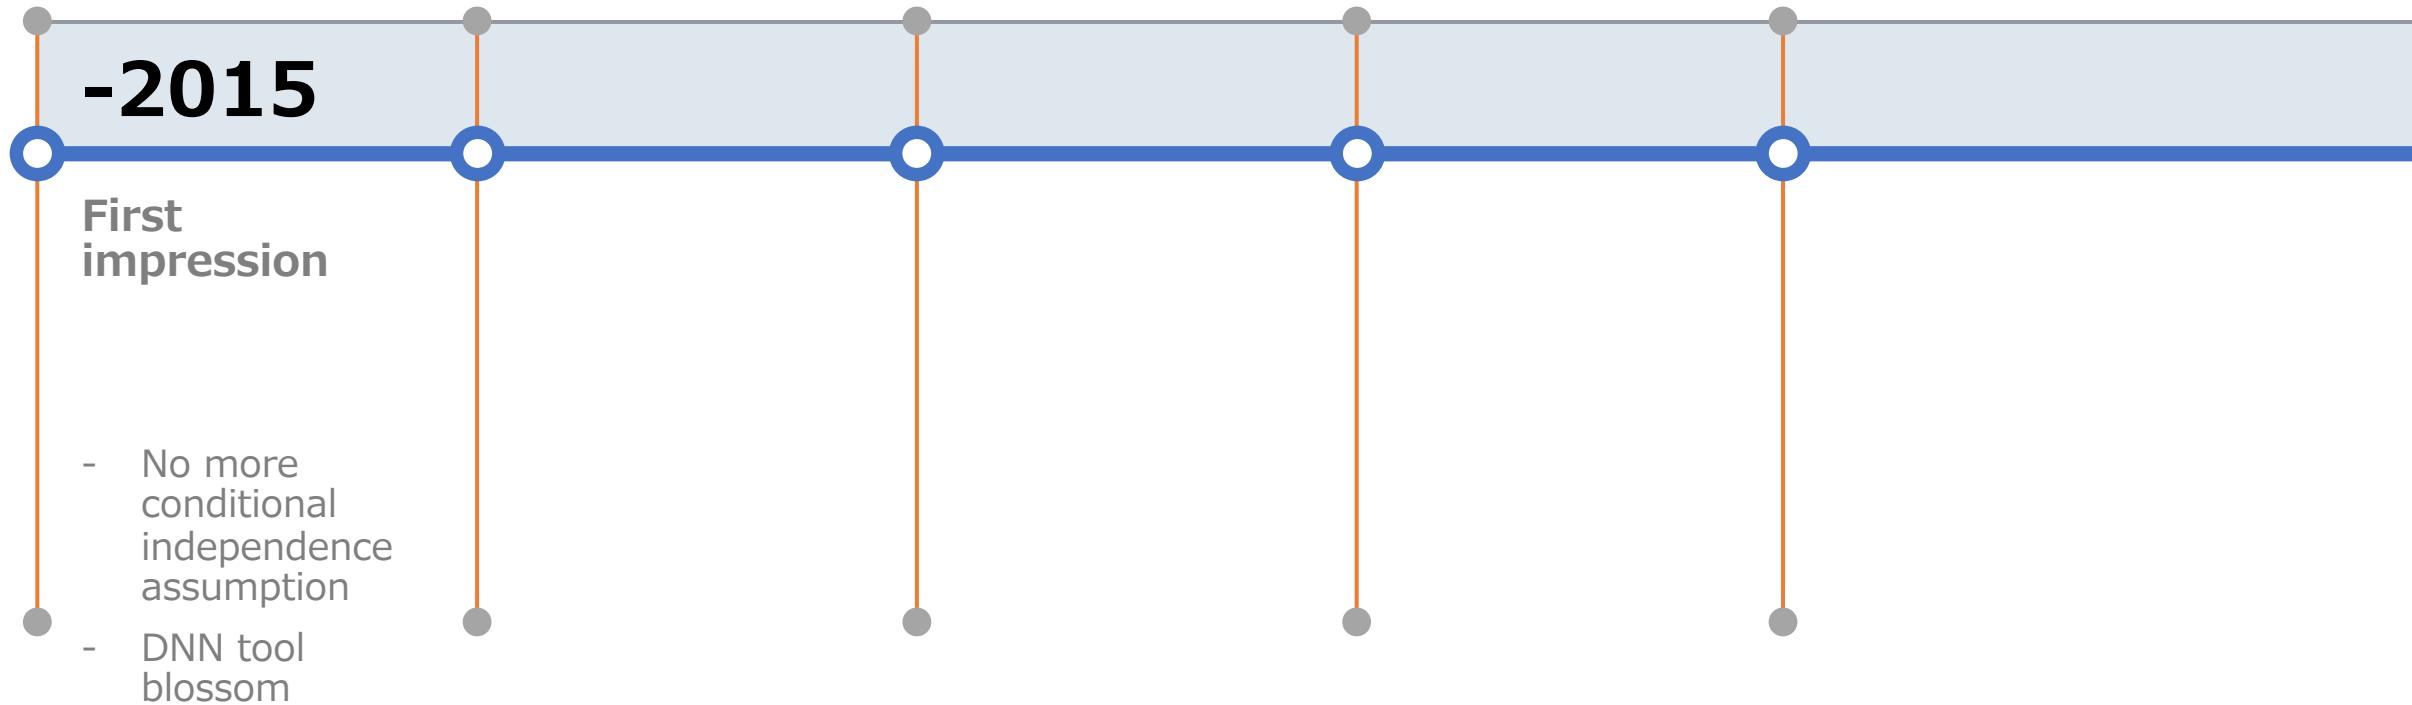
Timeline

Shinji's personal experience for end-to-end speech processing



Timeline

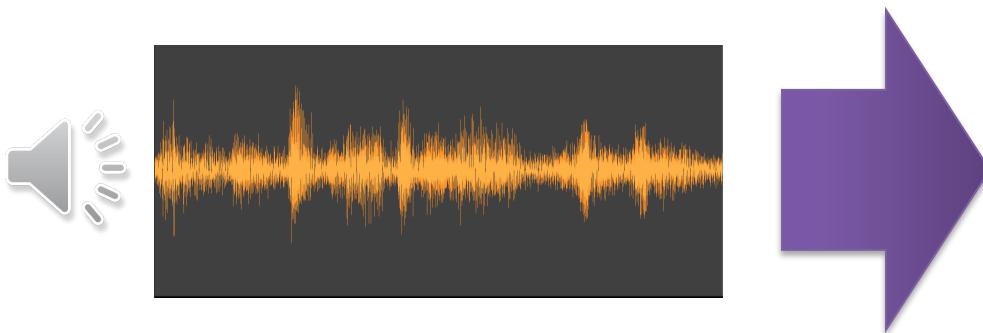
Shinji's personal experience for end-to-end speech processing



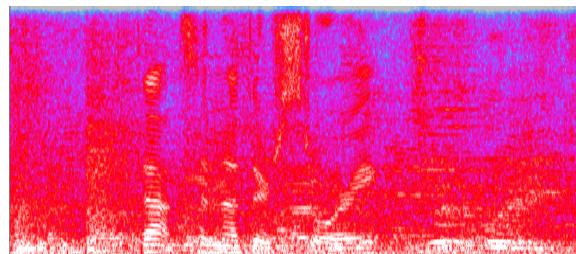
Noisy channel model (1970s-)

Noisy channel model (1970s-)

- Automatic Speech Recognition: Mapping *physical signal sequence* to *linguistic symbol sequence*



$$X = \{x_l \in \mathbb{Z} | l = 1, \dots, L\}$$
$$L = 43263$$



$$X = \{\mathbf{x}_t \in \mathbb{C}^D | t = 1, \dots, T\}$$
$$T = 268$$

"That's another story"

$$W = \{w_n \in \mathcal{V} | n = 1, \dots, N\}$$
$$N = 3$$

Noisy channel model (1970s-)

$$\arg \max_W p(W|X)$$

X : Speech sequence
 W : Text sequence

Noisy channel model (1970s-)

L : Phoneme sequence

$$\begin{aligned}\arg \max_W p(W|X) &= \arg \max_W p(X|W)p(W) \\ &\approx \arg \max_{W,L} p(X|L,W)p(L|W)p(W)\end{aligned}$$

- **Speech recognition**

- $p(X|L)$: Acoustic model (Hidden Markov model)
- $p(L|W)$: Lexicon
- $p(W)$: Language model (n-gram)

Noisy channel model (1970s-)

$$\begin{aligned}\arg \max_W p(W|X) &= \arg \max_W p(X|W)p(W) \\ &\approx \arg \max_{W,L} p(X|L,W)p(L|W)p(W)\end{aligned}$$

- **Speech recognition**

- $p(X|L)$: Acoustic model (Hidden Markov model)
- $p(L|W)$: Lexicon
- $p(W)$: Language model (n-gram)

- Factorization
- Conditional independence (Markov) assumptions

Noisy channel model (1970s-)

$$\arg \max_W p(W|X) = \arg \max_W p(X|W)p(W)$$

- **Machine translation**
 - $p(X|W)$: Translation model
 - $p(W)$: Language model

Noisy channel model (1970s-)

$$\begin{aligned}\arg \max_W p(W|X) &= \arg \max_W p(X|W)p(W) \\ &\approx \arg \max_{W,L} p(X|L,W)p(L|W)p(W)\end{aligned}$$

- **Speech recognition**
 - $p(X|L)$: Acoustic model (Hidden Markov model)
 - $p(L|W)$: Lexicon
 - $p(W)$: Language model (n-gram)
- Continued 40 years

Noisy channel model (1970s-)

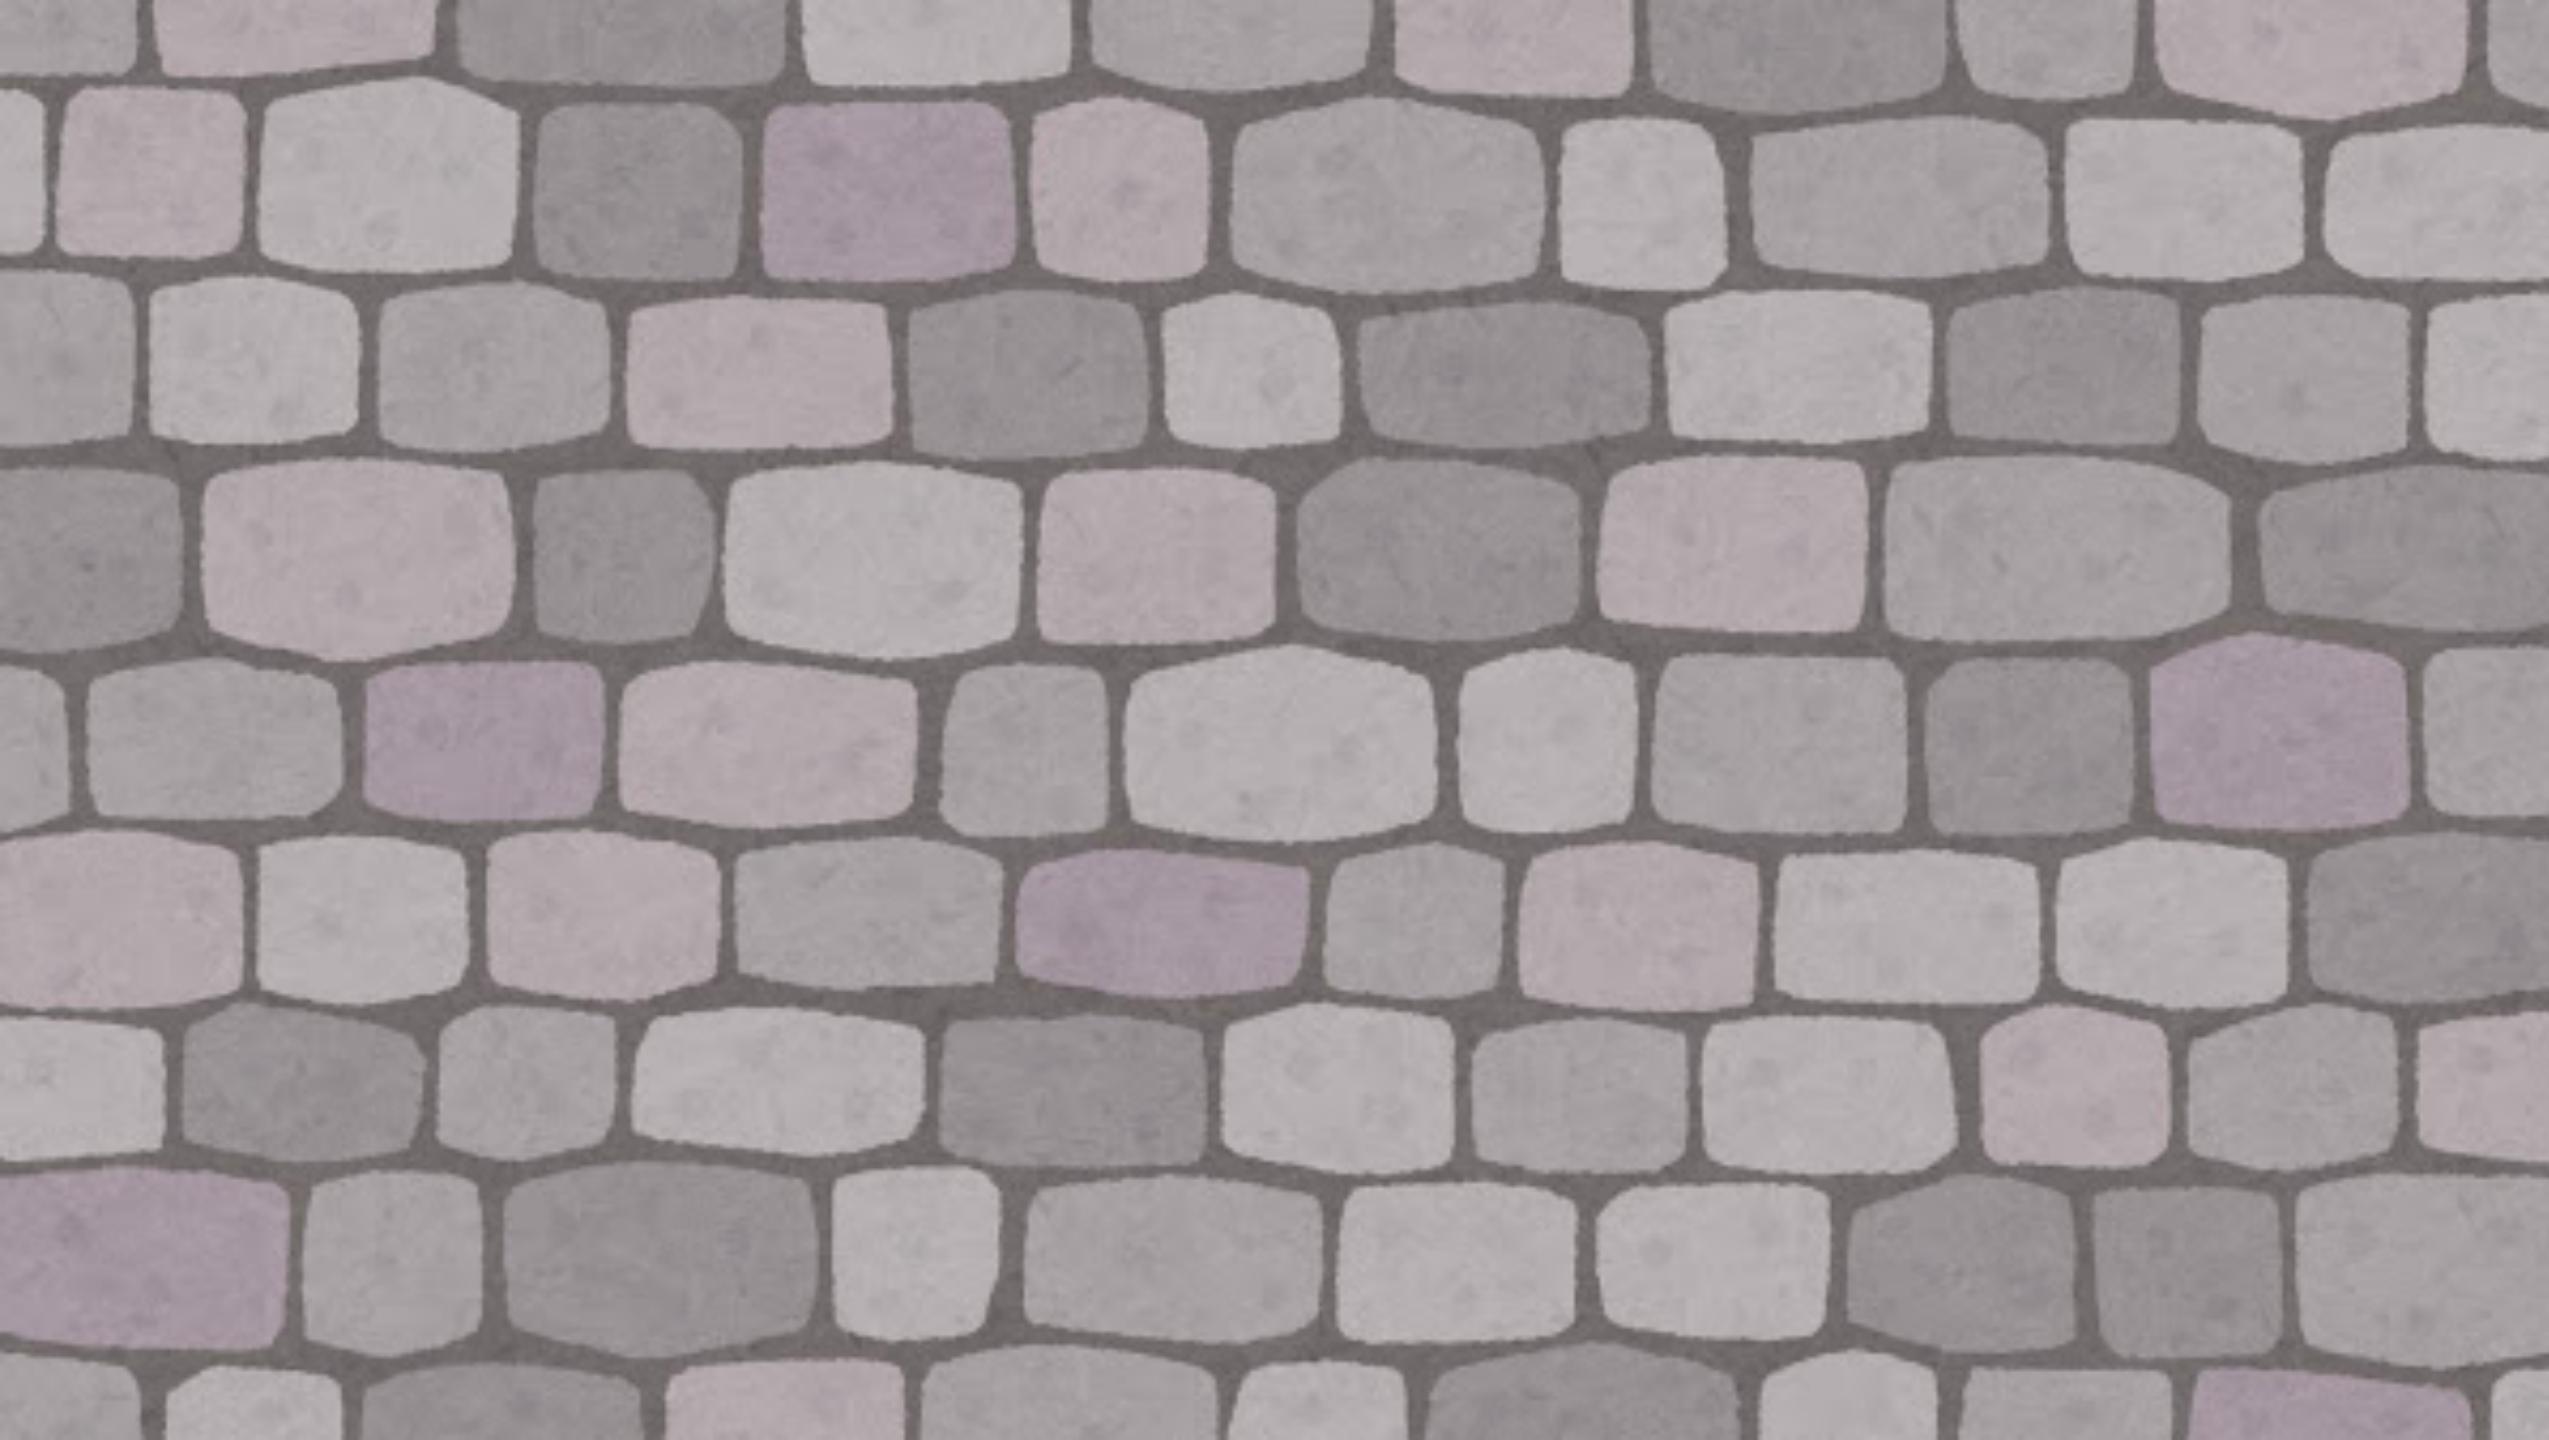
$$\begin{aligned}\arg \max_W p(W|X) &= \arg \max_W p(X|W)p(W) \\ &\approx \arg \max_{W,L} p(X|L,W)p(L|W)p(W)\end{aligned}$$

- **Speech recognition**
 - $p(X|L)$: Acoustic model
 - $p(L|W)$: Lexicon
 - $p(W)$: Language model
- Continued 40 years

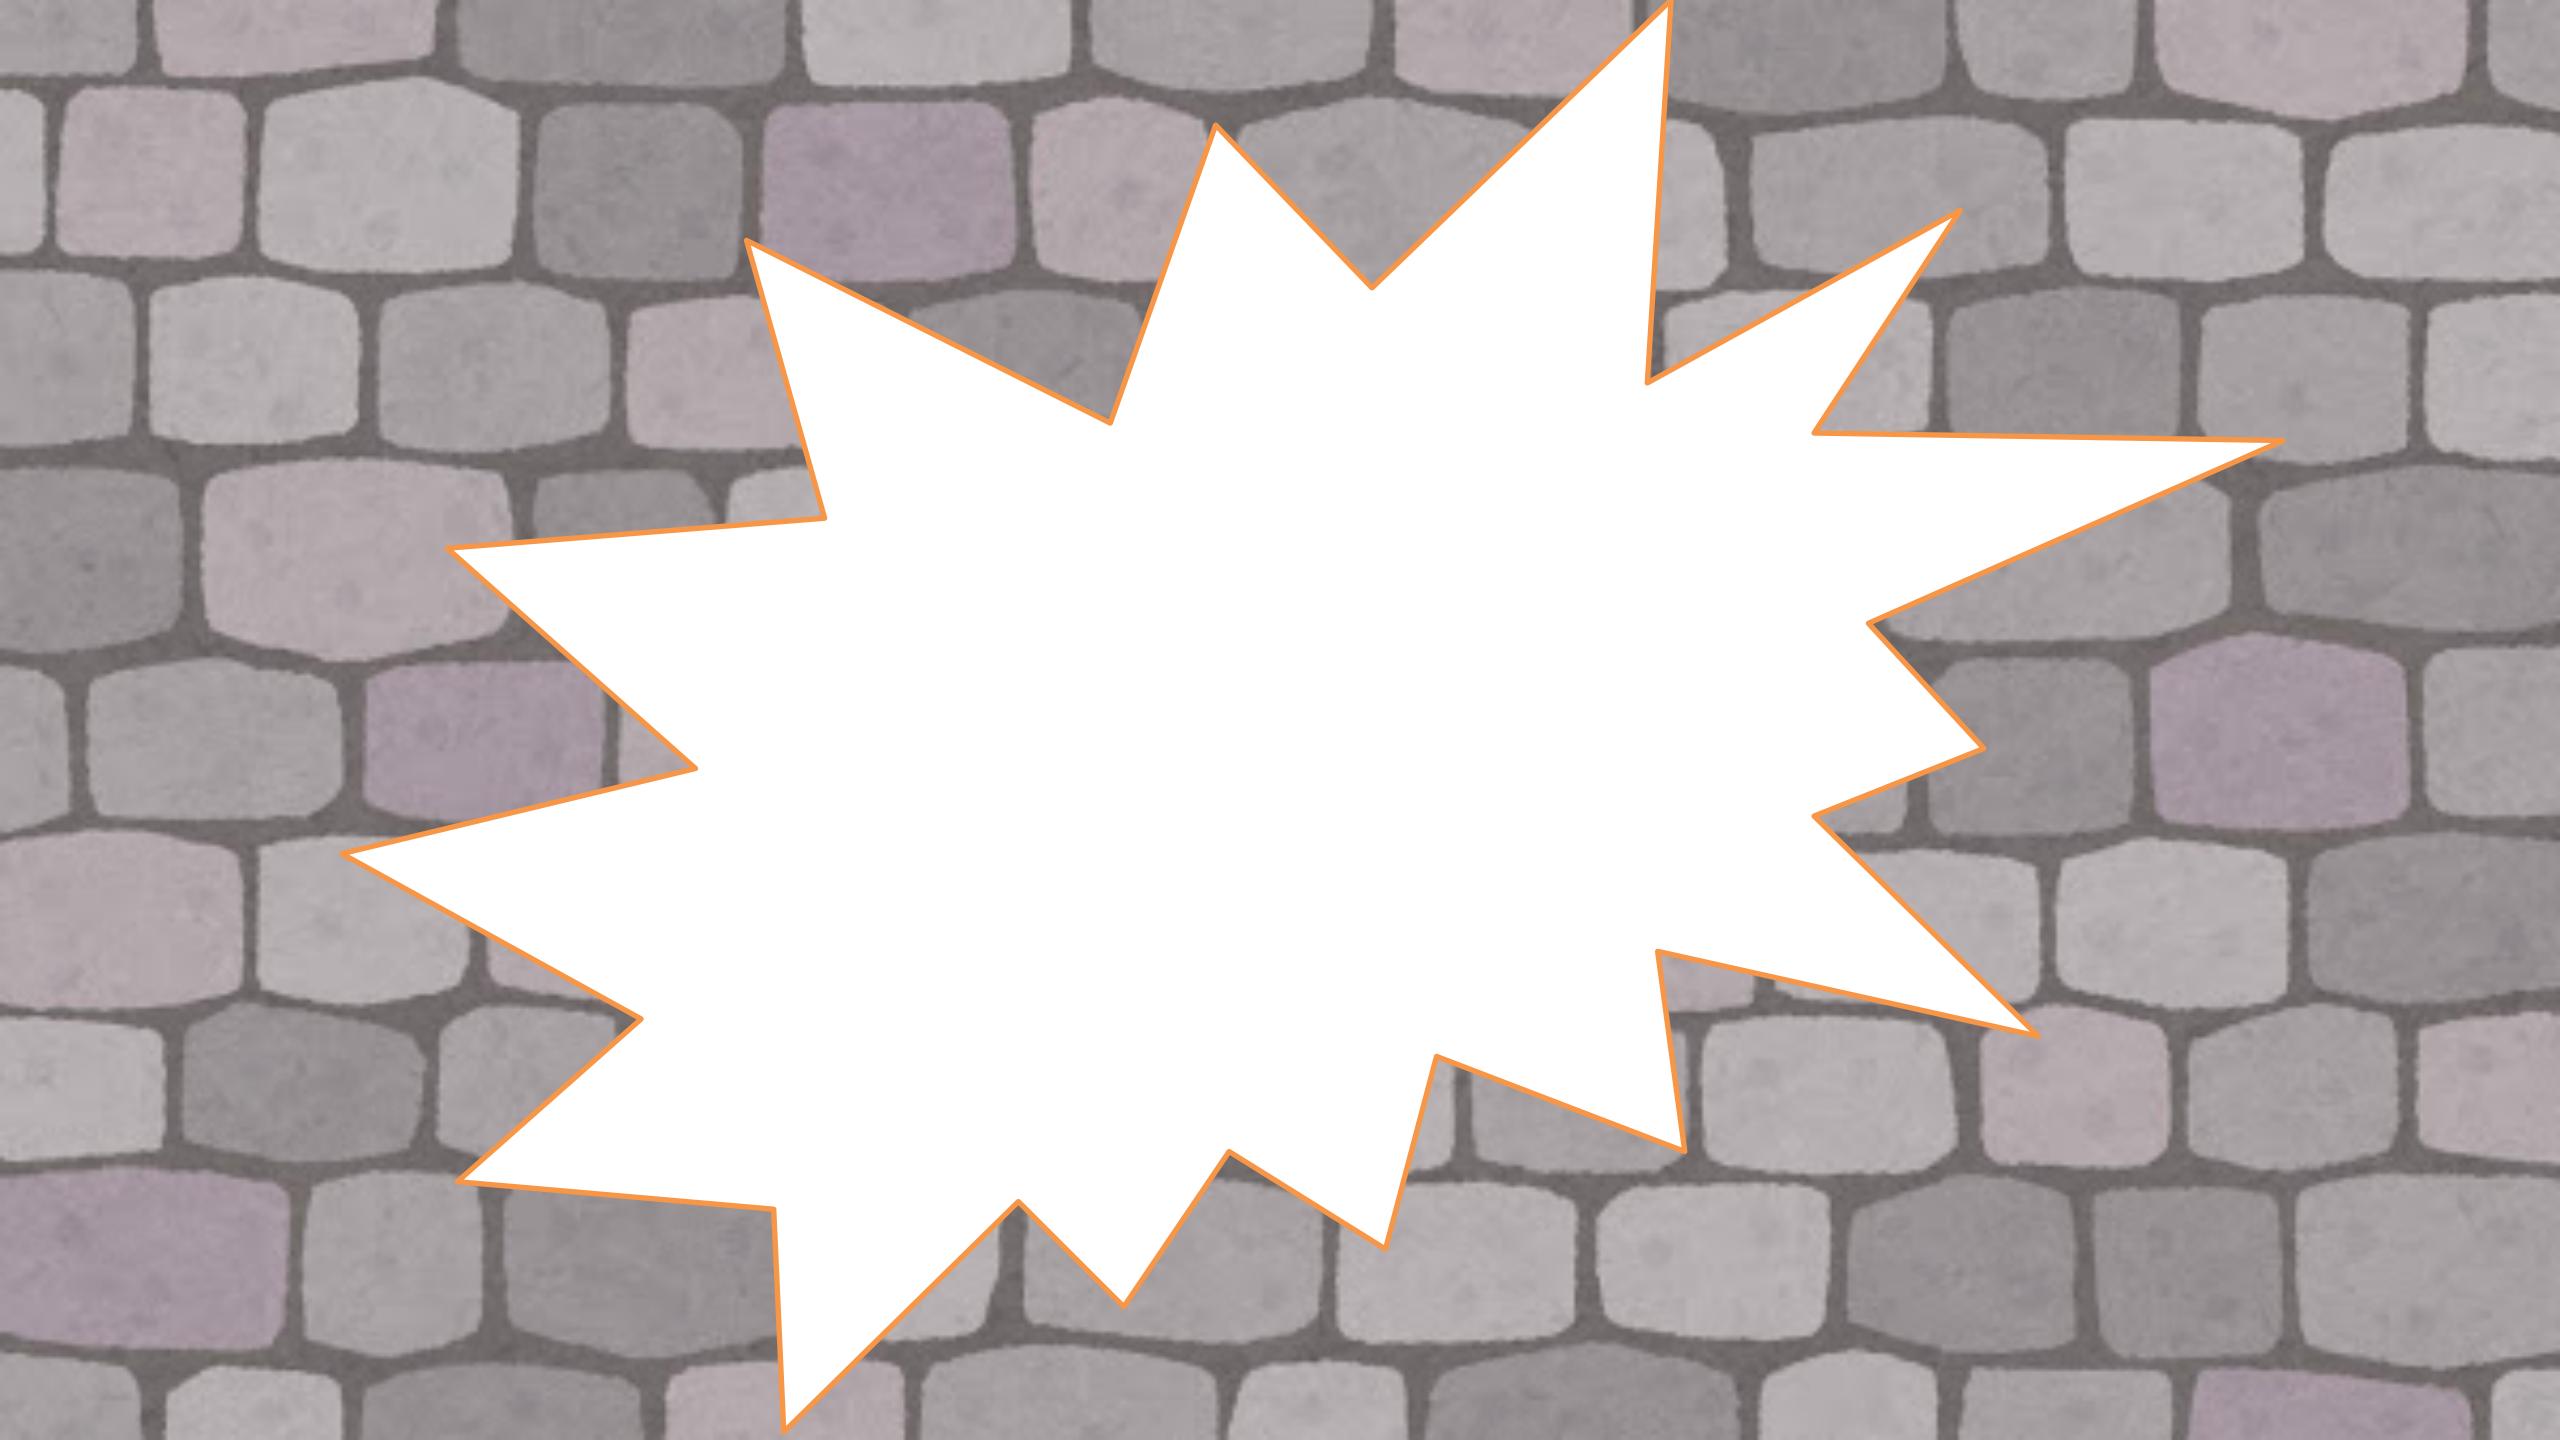


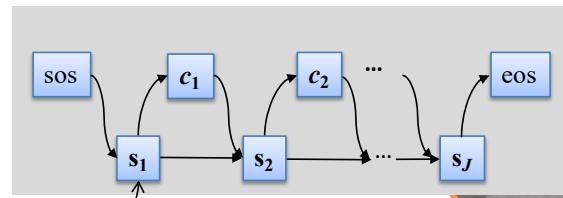
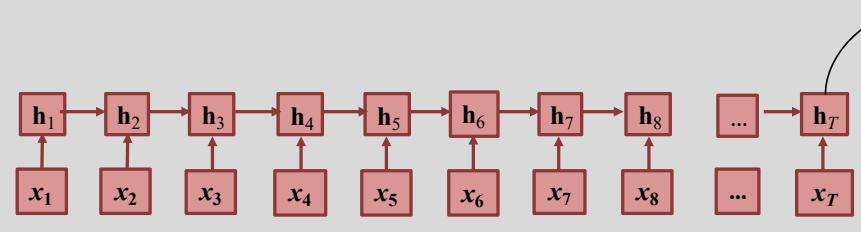
Big barrier:
noisy channel model
HMM
n-gram
etc.

However,

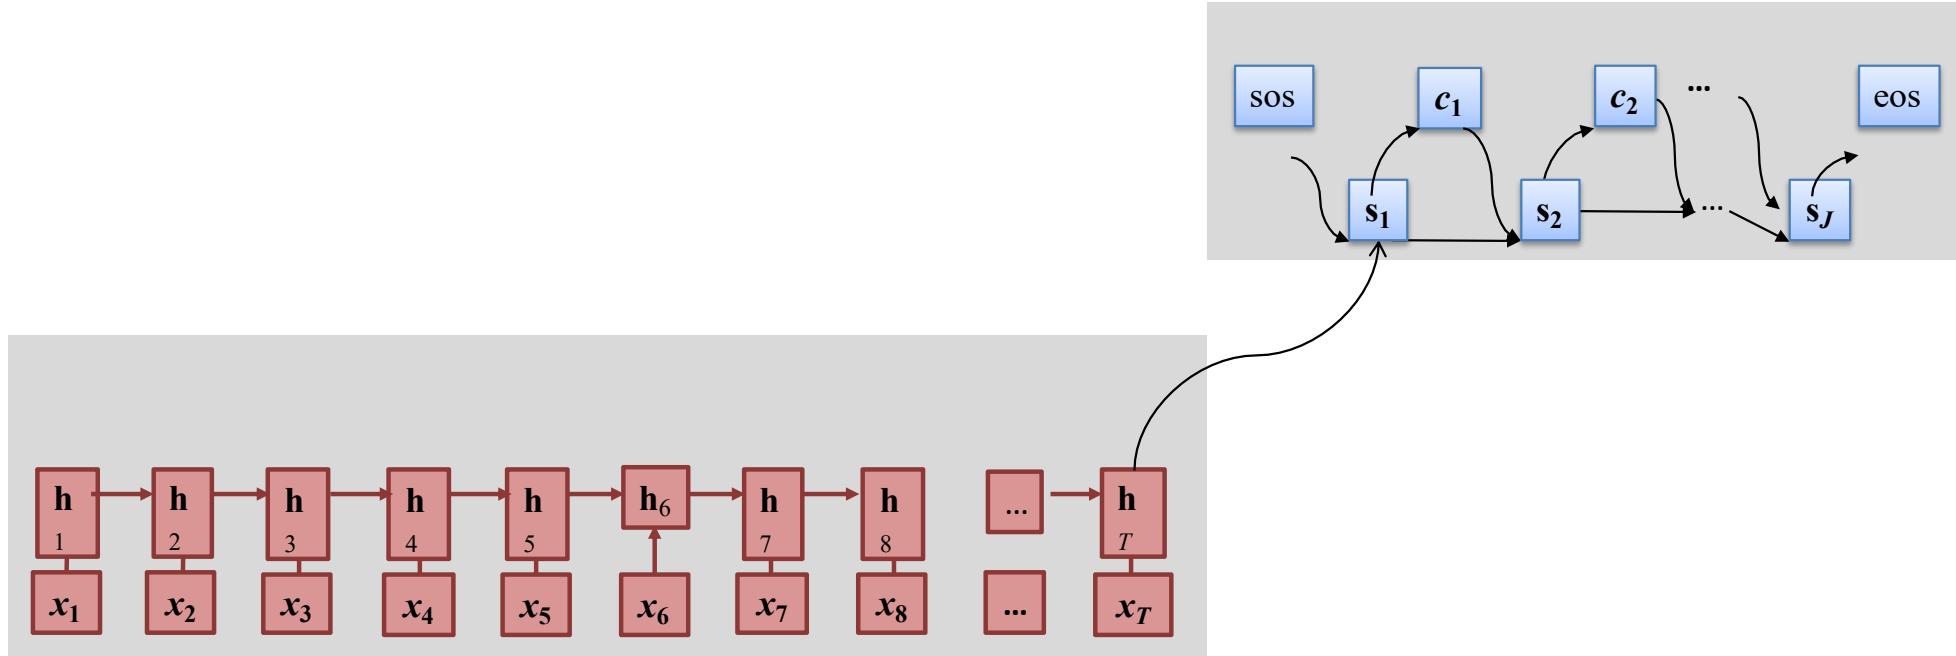








“End-to-End” Processing Using Sequence to Sequence



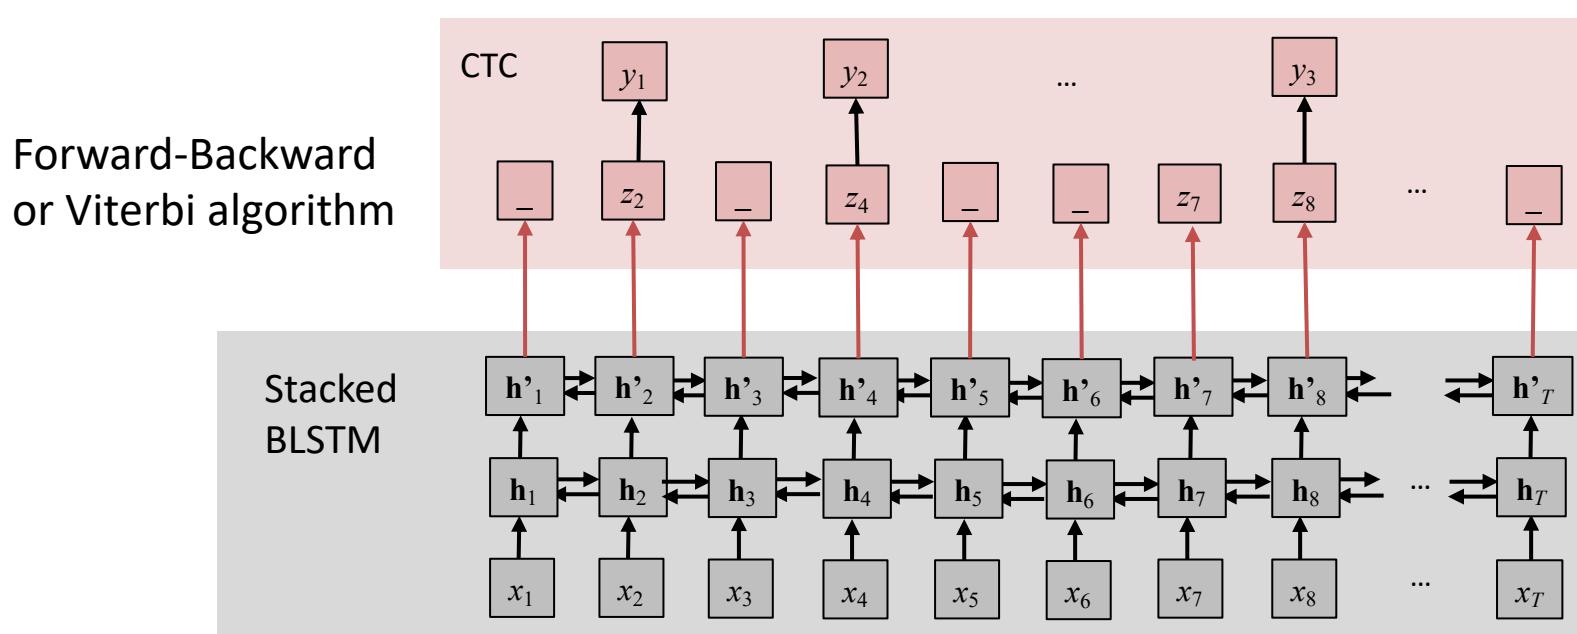
- Directly model $p(W|X)$ with a **single neural network**
 - **Integrate** acoustic $p(X|L)$, lexicon $p(L|W)$, and language $p(W)$ models
- Great success in neural machine translation

Connectionist temporal classification (CTC)

[Graves+ 2006,

Graves+ 2014, Miao+ 2015]

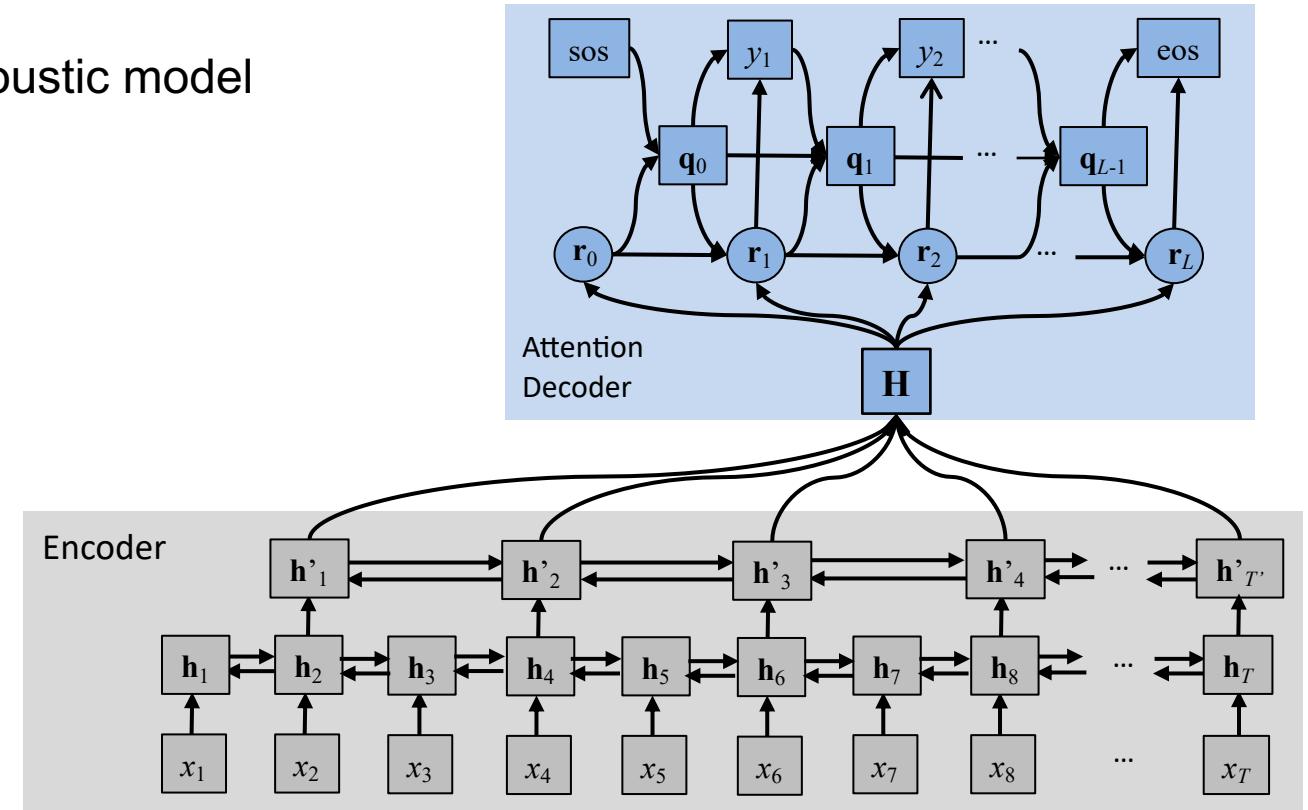
- Use bidirectional RNNs to predict frame-based labels including blanks
- Find alignments between X and Y using dynamic programming



Attention-based encoder decoder

[Chorowski+ 2014, Chan+ 2015]

- Combine acoustic and language models in a single architecture
 - Encoder: DNN part of acoustic model
 - Decoder: language model
 - Attention: HMM part of acoustic model



First impression in -2015

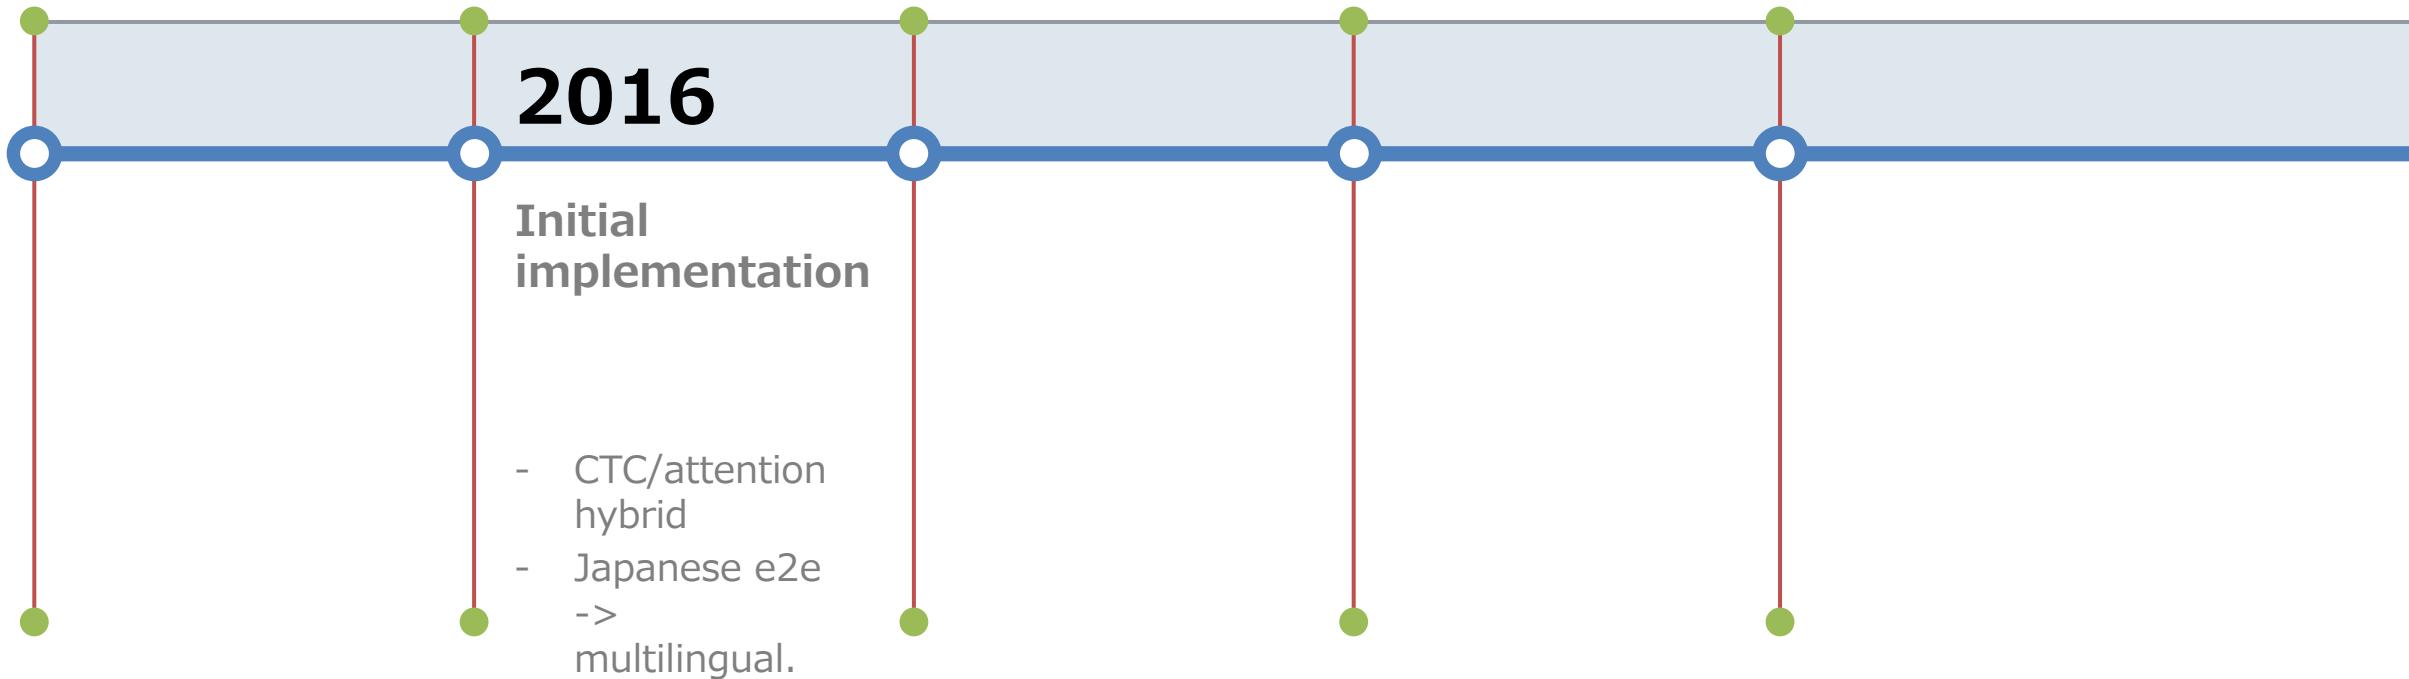
- Attention based encoder decoder

$$\arg \max_W p(W|X) = \arg \max_W \prod_j p(w_j|w_{<j}, X)$$

- No conditional independence assumption unlike HMM/CTC
 - More precise seq-to-seq model
 - This is what I have been struggling for 15 years!
- Attention mechanism allows too flexible alignments
 - Hard to train the model from scratch

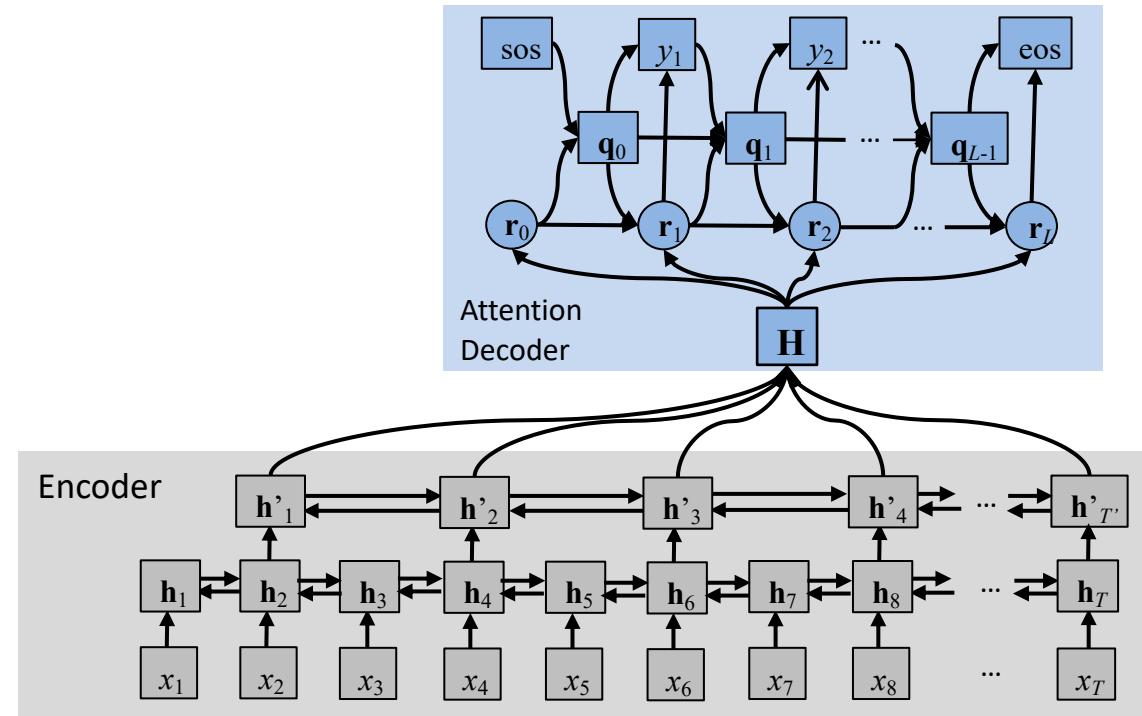
Timeline

Shinji's personal experience for end-to-end speech processing



Initial implementation in 2016

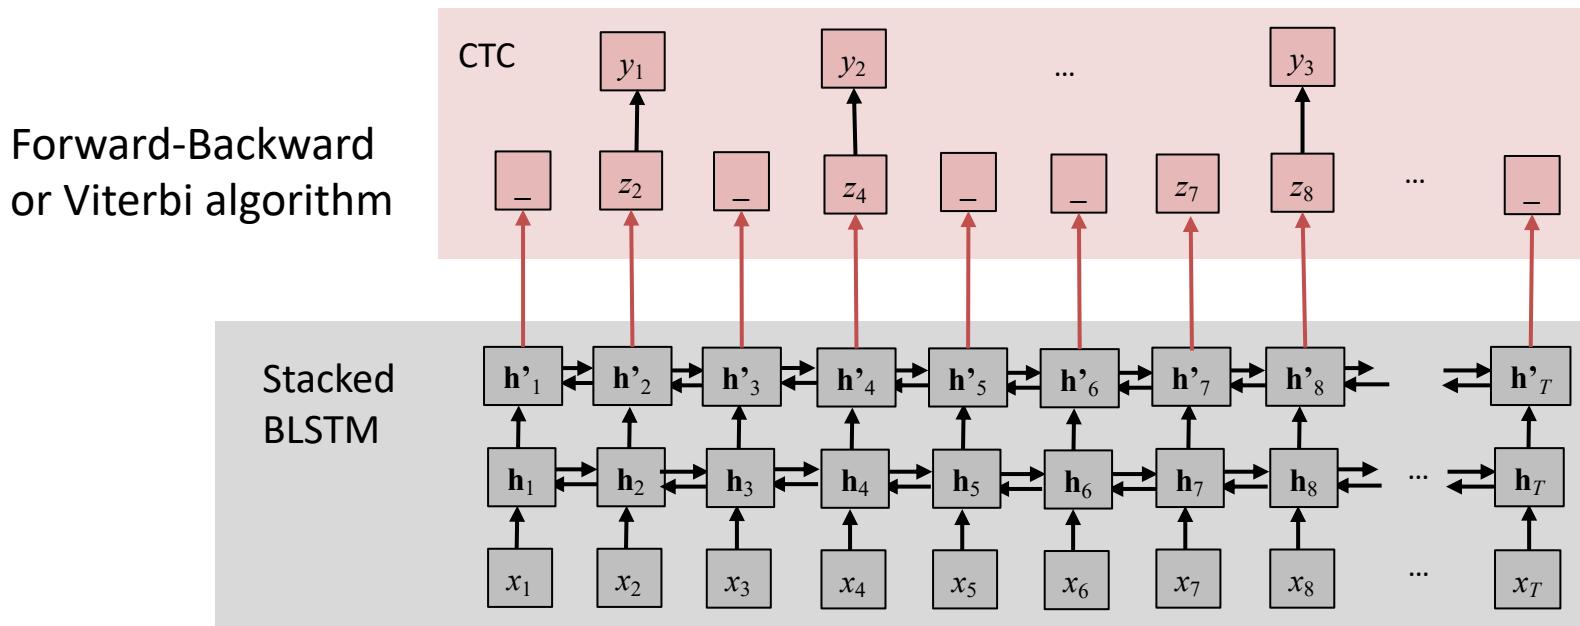
- Takaaki Hori, John Hershey, and I started an E2E project at MERL with some interns
- First, we implemented both
 - CTC
 - Attention-based encoder/decoder
- We found some pros. and cons.



Connectionist temporal classification (CTC)

[Graves+ 2006, Graves+ 2014, Miao+ 2015]

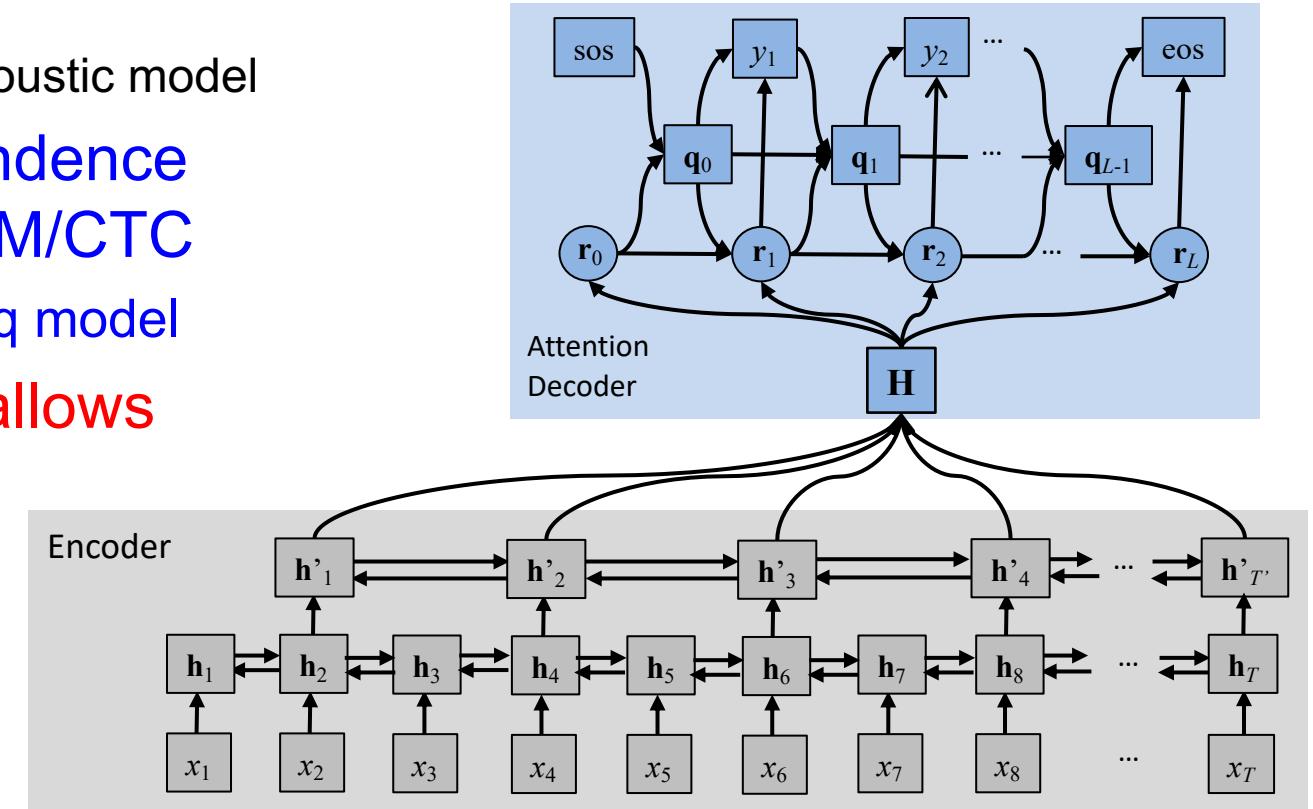
- Use bidirectional RNNs to predict frame-based labels including blanks
- Find alignments between X and Y using dynamic programming
- Relying on conditional independence assumptions (similar to HMM)
- Output sequence is not well modeled (no language model)



Attention-based encoder decoder

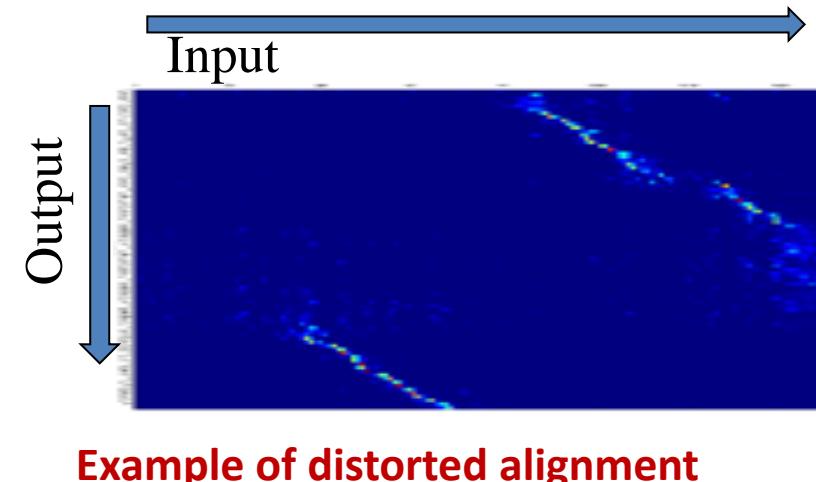
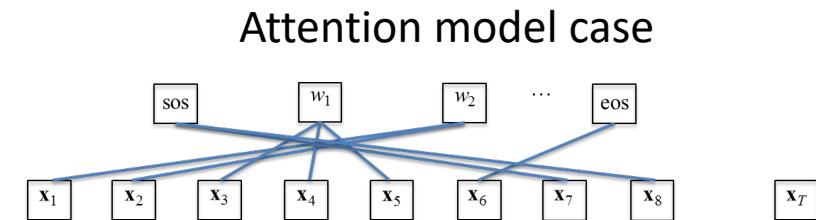
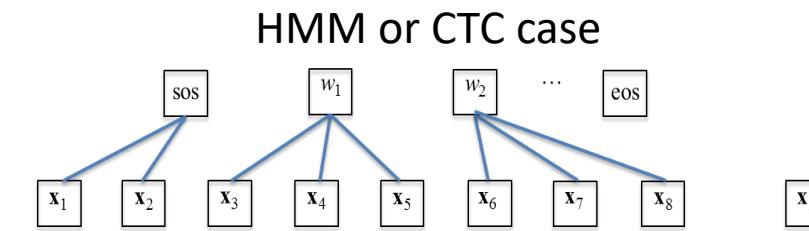
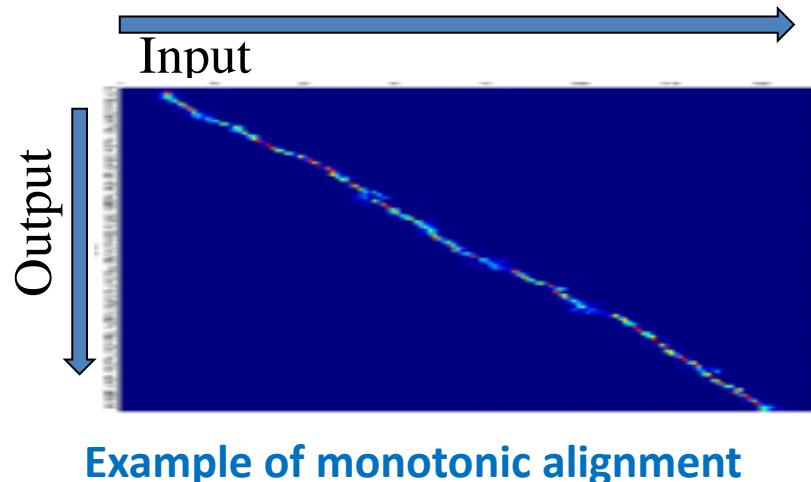
[Chorowski+ 2014, Chan+ 2015]

- Combine acoustic and language models in a single architecture
 - Encoder: DNN part of acoustic model
 - Decoder: language model
 - Attention: HMM part of acoustic model
- No conditional independence assumption unlike HMM/CTC
 - More precise seq-to-seq model
- Attention mechanism allows too flexible alignments
 - Hard to train the model from scratch



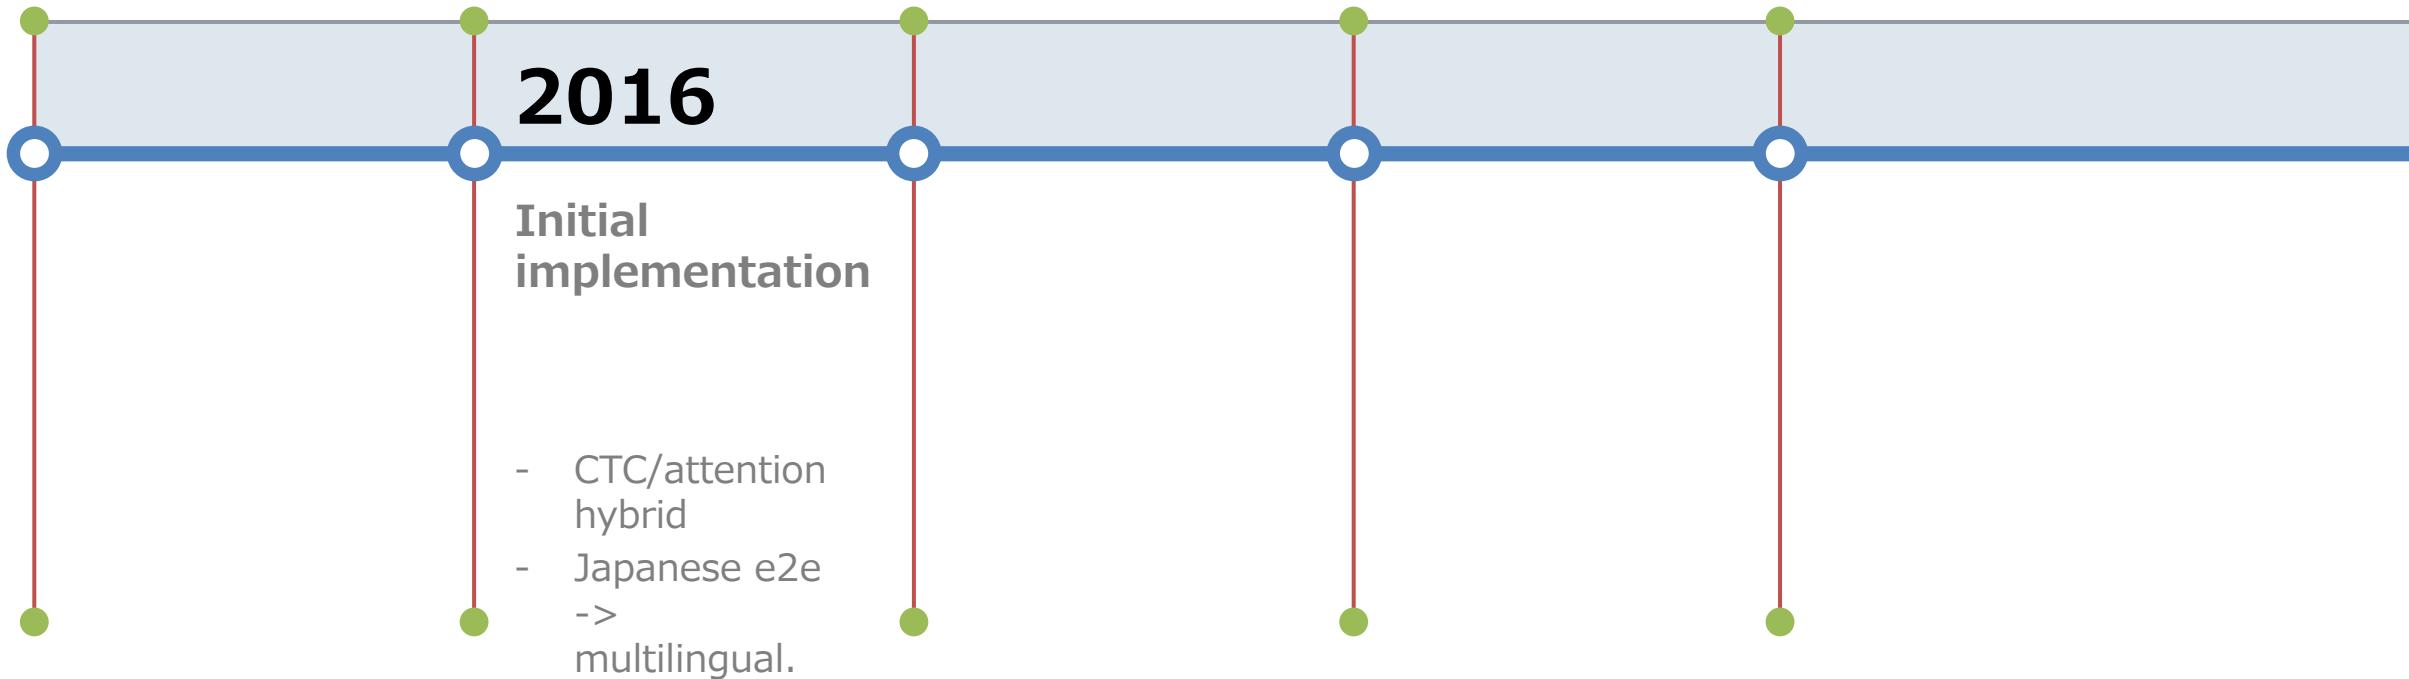
Input/output alignment by temporal attention

- Unlike CTC, attention model does not preserve order of inputs
- Our desired alignment in ASR task is **monotonic**
- Not regularized alignment makes the model **hard to learn** from scratch



Timeline

Shinji's personal experience for end-to-end speech processing



How to solve this unstable attention issues

It was **too unstable** to move to the next step...

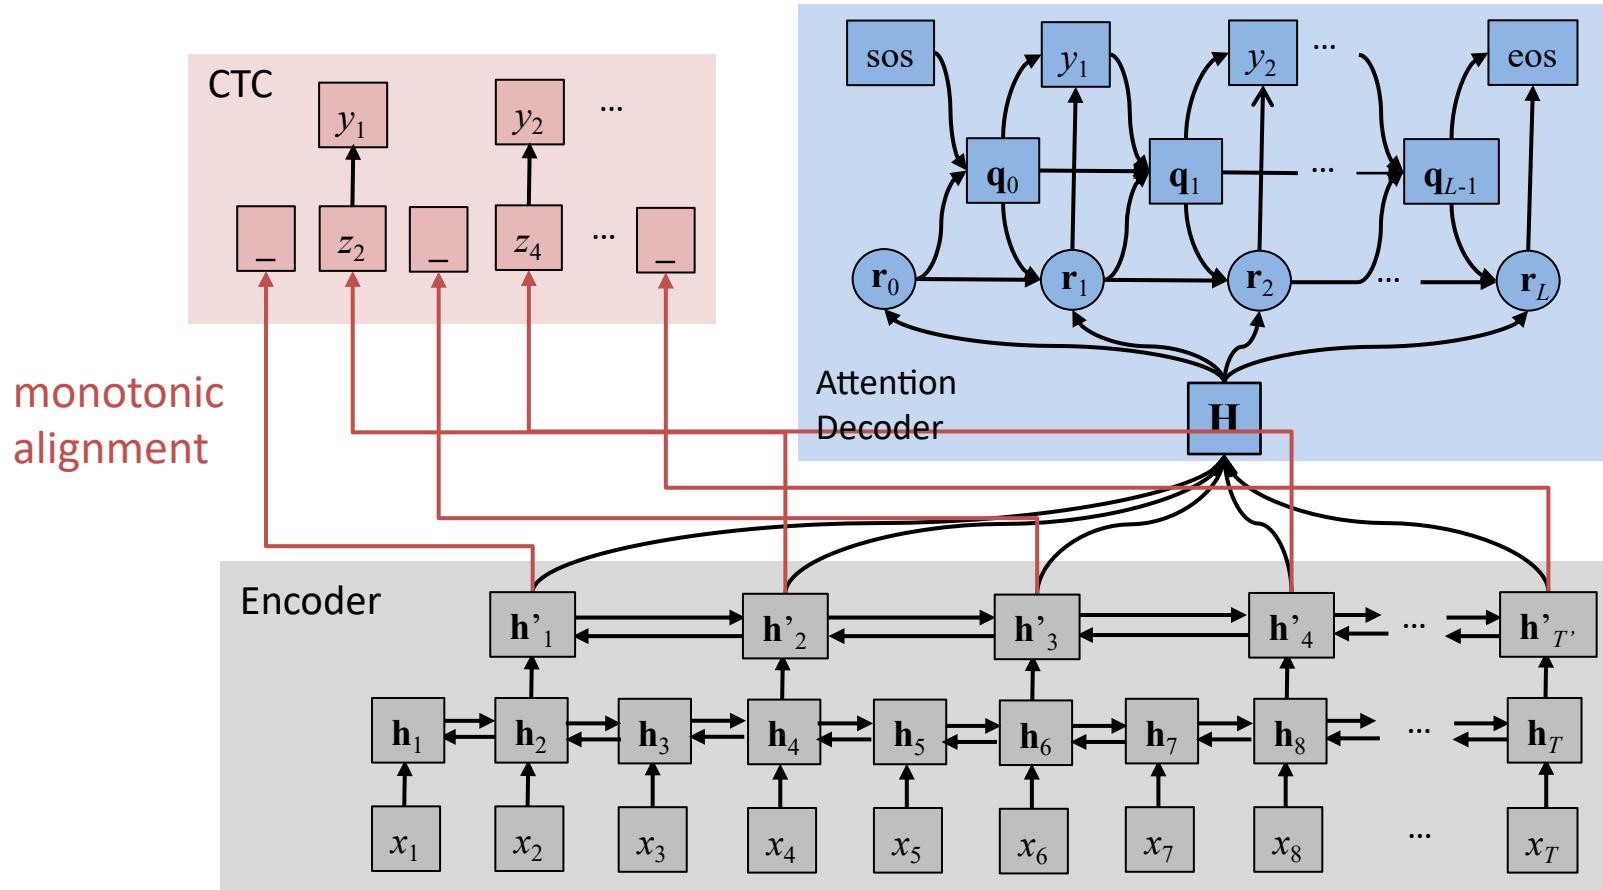
- We had a lot of ideas but those were pending due to that
- Probably we should try to use **both benefits of CTC and attention**

How to combine both?

- One possible solution: RNN transducer
 - Try to find another solution
 - Finally came up with a simple idea (or we decided to use this simple idea)
- **Hybrid CTC/attention**

Hybrid CTC/attention network [Kim+'17]

Multitask learning: $\mathcal{L}_{\text{MTL}} = \lambda \mathcal{L}_{\text{CTC}} + (1 - \lambda) \mathcal{L}_{\text{Attention}}$ λ : CTC weight

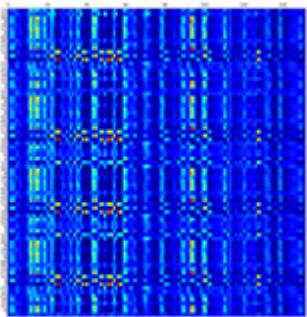


CTC guides attention alignment to be monotonic

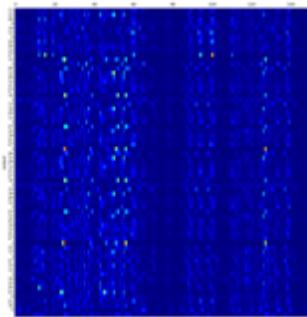
More robust input/output alignment of attention

- Alignment of one selected utterance from CHiME4 task

Attention Model

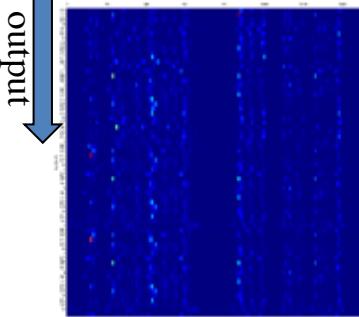


Epoch 1



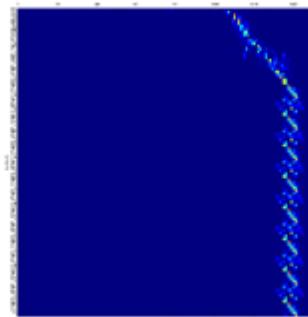
Epoch 3

Input →



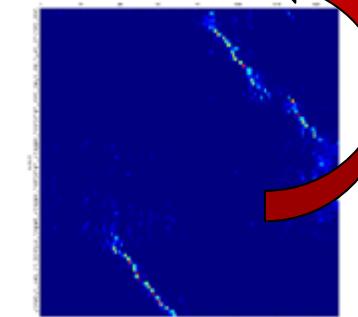
Epoch 5

↓ output



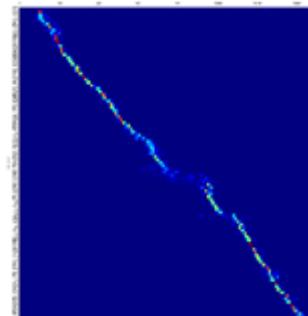
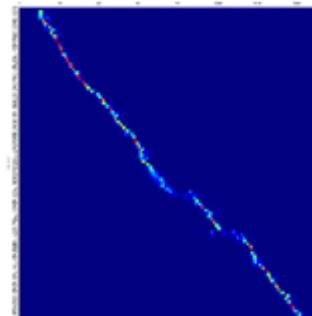
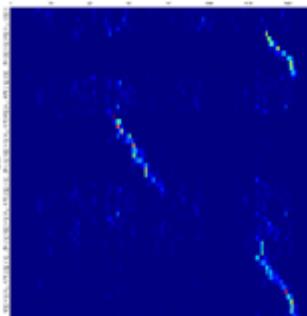
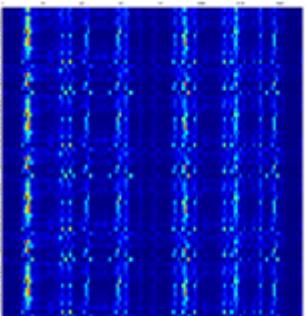
Epoch 7

Corrupted!

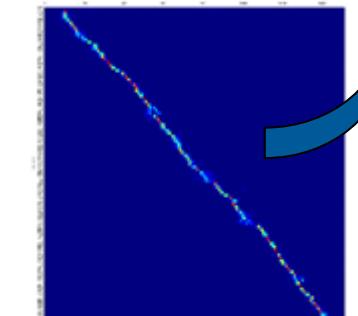


Epoch 9

Our joint CTC/attention model



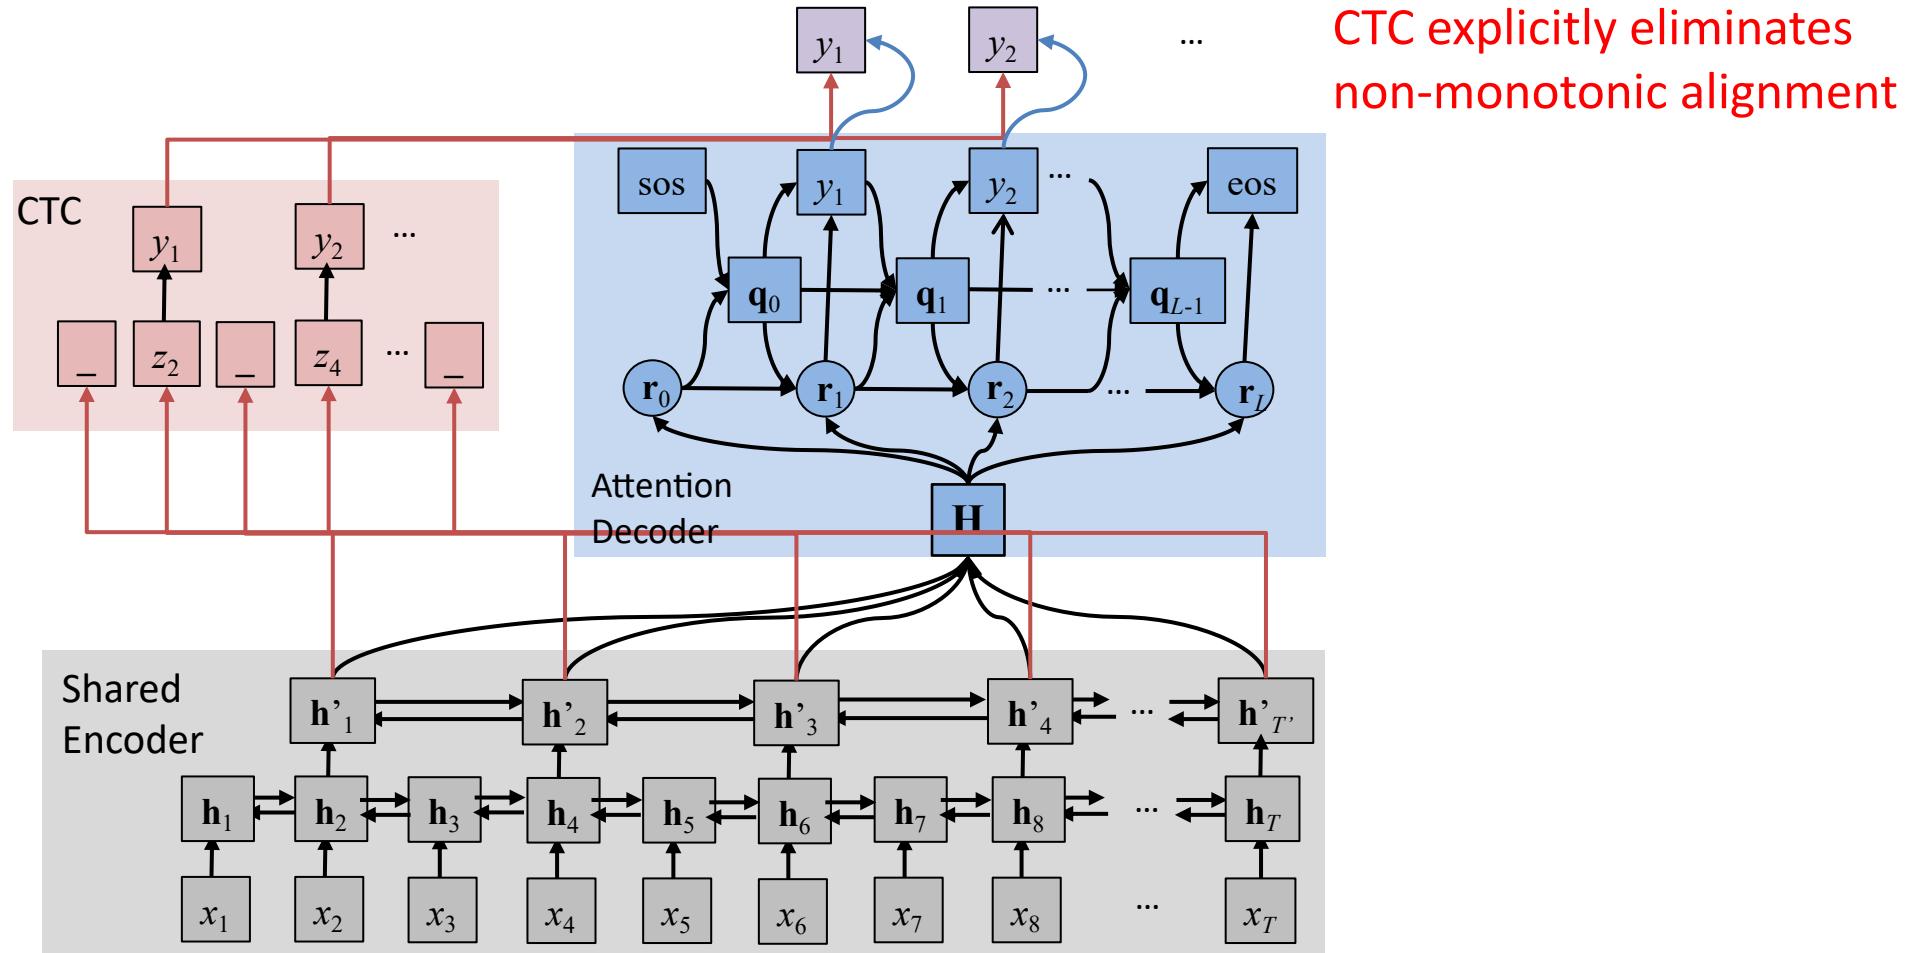
Monotonic!



Faster convergence

Joint CTC/attention decoding [Hori+'17]

Use CTC for decoding together with the attention decoder



Experimental Results

Character Error Rate (%) in **Mandarin Chinese Telephone Conversational** (HKUST, 167 hours)

Models	Dev.	Eval
Attention model (baseline)	40.3	37.8
CTC-attention learning (MTL)	38.7	36.6
+ Joint decoding	35.5	33.9

Character Error Rate (%) in **Corpus of Spontaneous Japanese** (CSJ, 581 hours)

Models	Task 1	Task 2	Task 3
Attention model (baseline)	11.4	7.9	9.0
CTC-attention learning (MTL)	10.5	7.6	8.3
+ Joint decoding	10.0	7.1	7.6

Example of recovering insertion errors (HKUST)

id: (20040717_152947_A010409_B010408-A-057045-057837)

Reference

但是如果你想如果回到了过去你如果带着这个现在的记忆是不是很痛苦啊

Hybrid CTC/attention (w/o joint decoding)

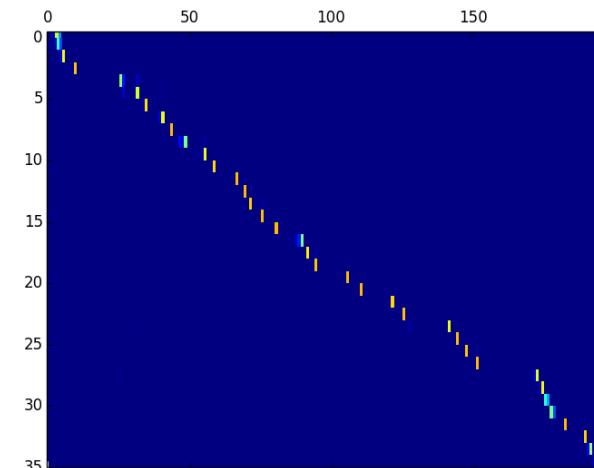
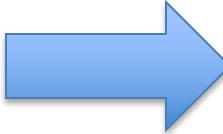
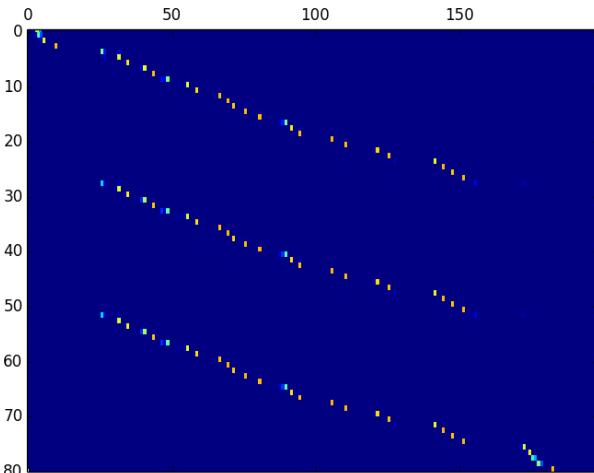
Scores: (#Correctness #Substitution #Deletion #Insertion) 28 2 3 45

但是如果你想如果回到了过去你如果带着这个现在的节如果你想如果回到了过去你如果带着这个现在的节如果你想如果回到了过去你如果带着这个现在的机是不是很 · · ·

w/ Joint decoding

Scores: (#Correctness #Substitution #Deletion #Insertion) 31 1 1 0

HYP: 但是如果你想如果回到了过去你如果带着这个现在的 · 机是不是很痛苦啊



Example of recovering deletion errors (CSJ)

id: (A01F0001_0844951_0854386)

Reference

またえ飛行時のエコーコーデーション機能をより詳細に解明する為に超小型マイクロホンおよび生体アンプをコウモリに搭載することを考えておりますそうすることによって

Hybrid CTC/attention (w/o joint decoding)

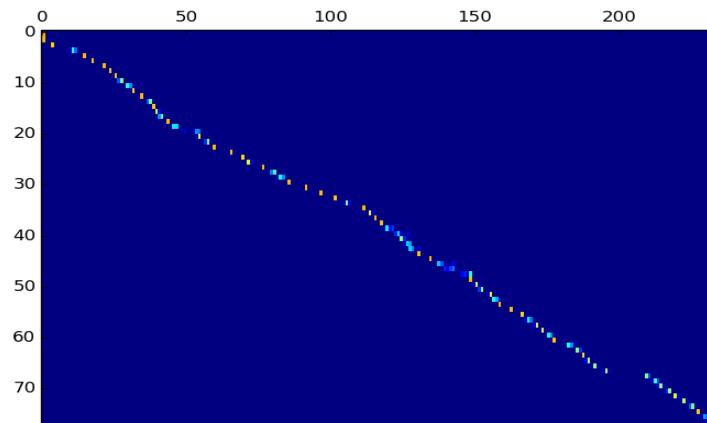
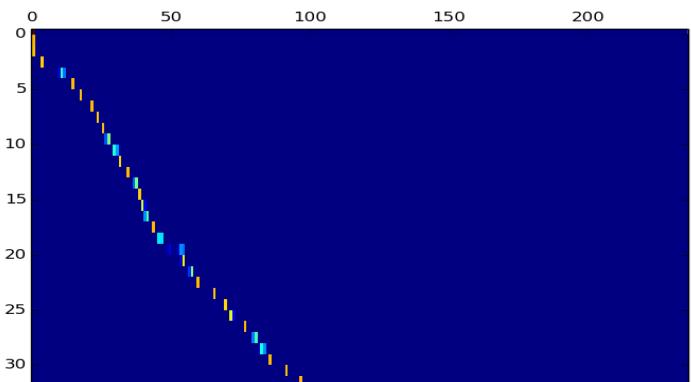
Scores: (#Correctness #Substitution #Deletion #Insertion) 30 0 47 0

またえ飛行時のエコーコーデーション機能をより詳細に解明する
為 …
… … … … … … … に … …

w/ Joint decoding

Scores: (#Correctness #Substitution #Deletion #Insertion) 67 9 10

またえ飛行時のエコーコーデーション機能をより詳細に解明する為に長国型マイ
クロホンお・いく声単位方をコウモリに登載することを考えておりますそうす
ることによって

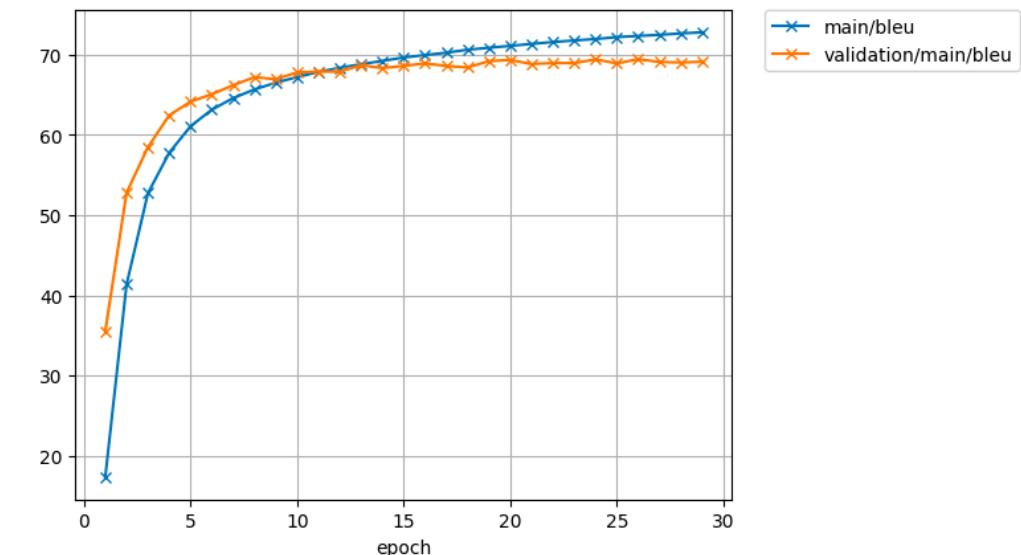
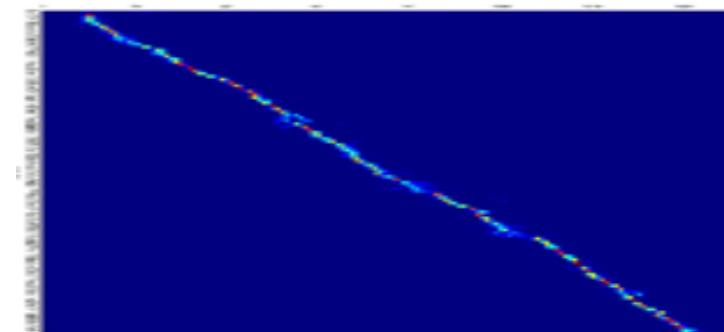


Discussions

- Hybrid CTC/attention-based end-to-end speech recognition
 - Multi-task learning during training
 - Joint decoding during recognition
- Now we have a good end-to-end ASR tool
 - **Make use of both benefits, completely solve alignment issues**
- **NOTE:** This can be solved by large amounts of training data and a lot of tuning. This is one solution (but quite academia friendly)

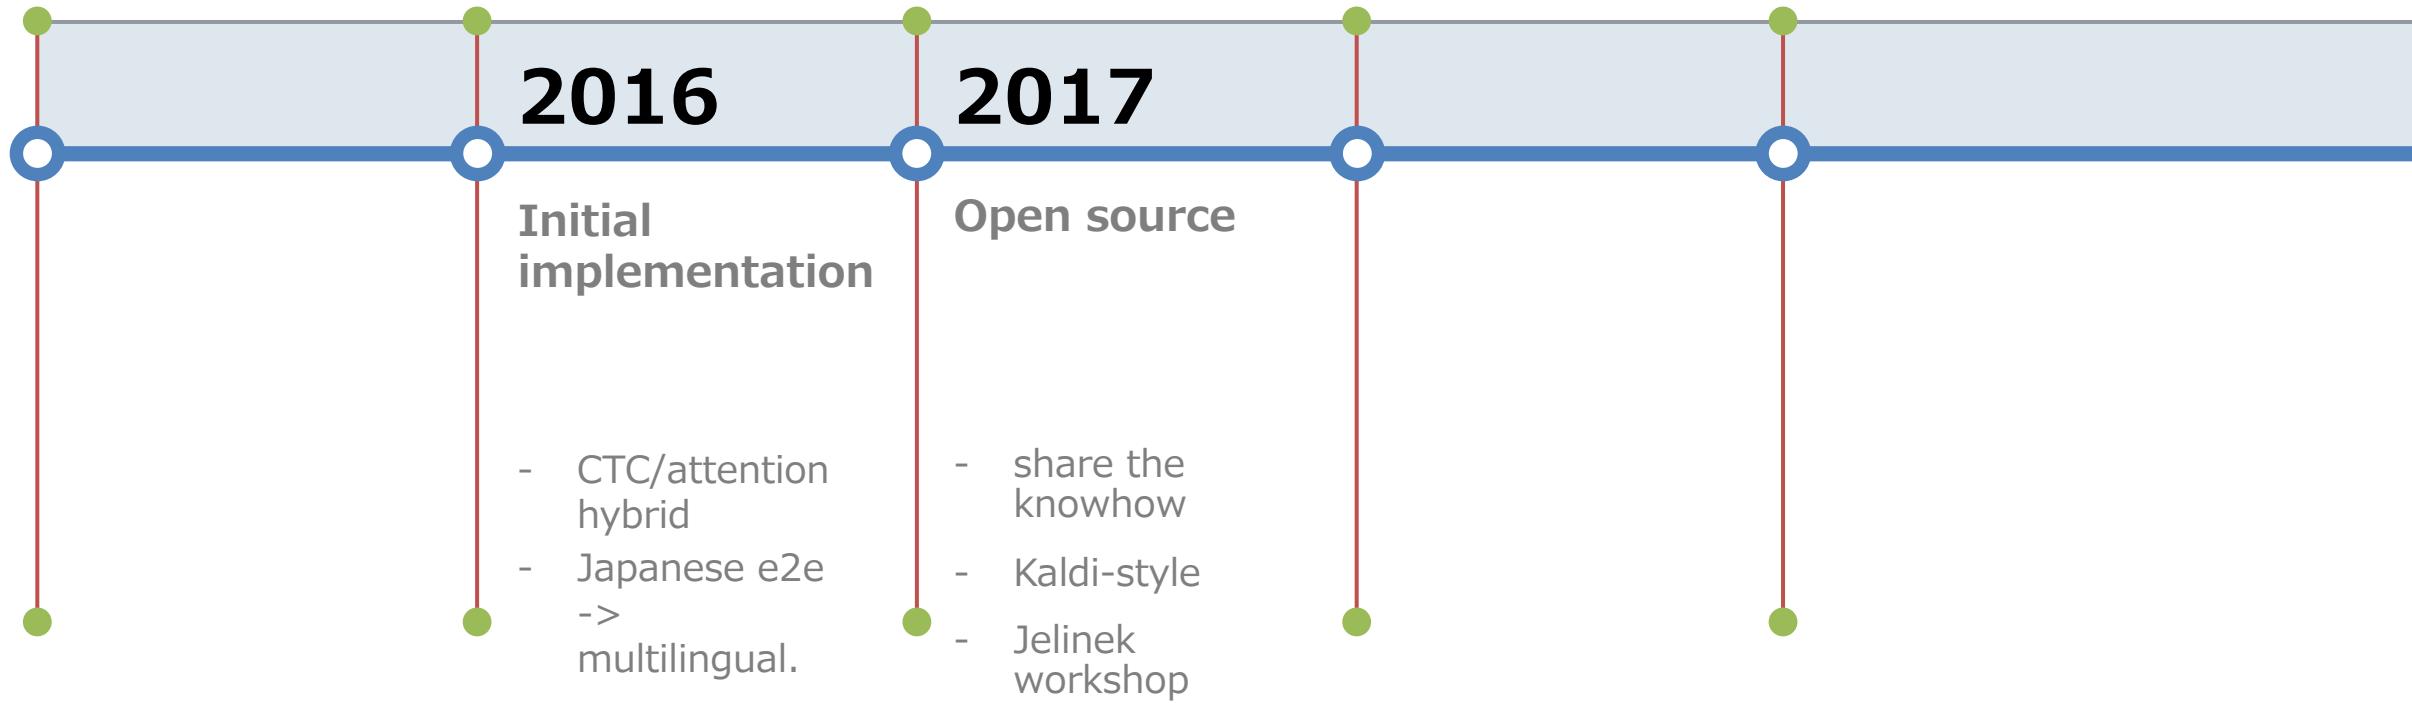
FAQ

- How to debug attention-based encoder/decoder?
- Please check
 - Attention pattern!**
 - Learning curves!**
- It gives you a lot of intuitive information!

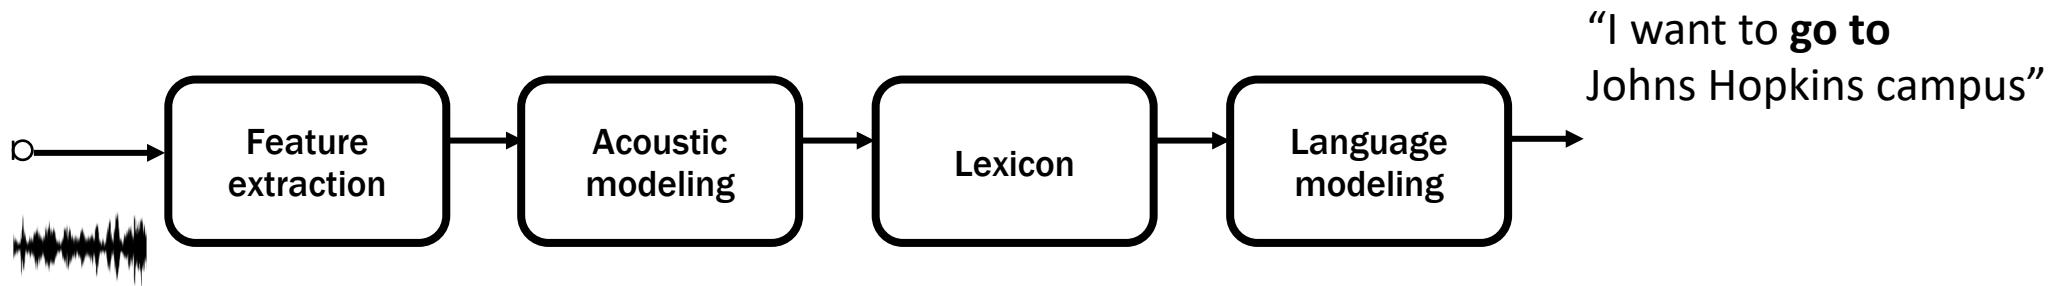


Timeline

Shinji's personal experience for end-to-end speech processing

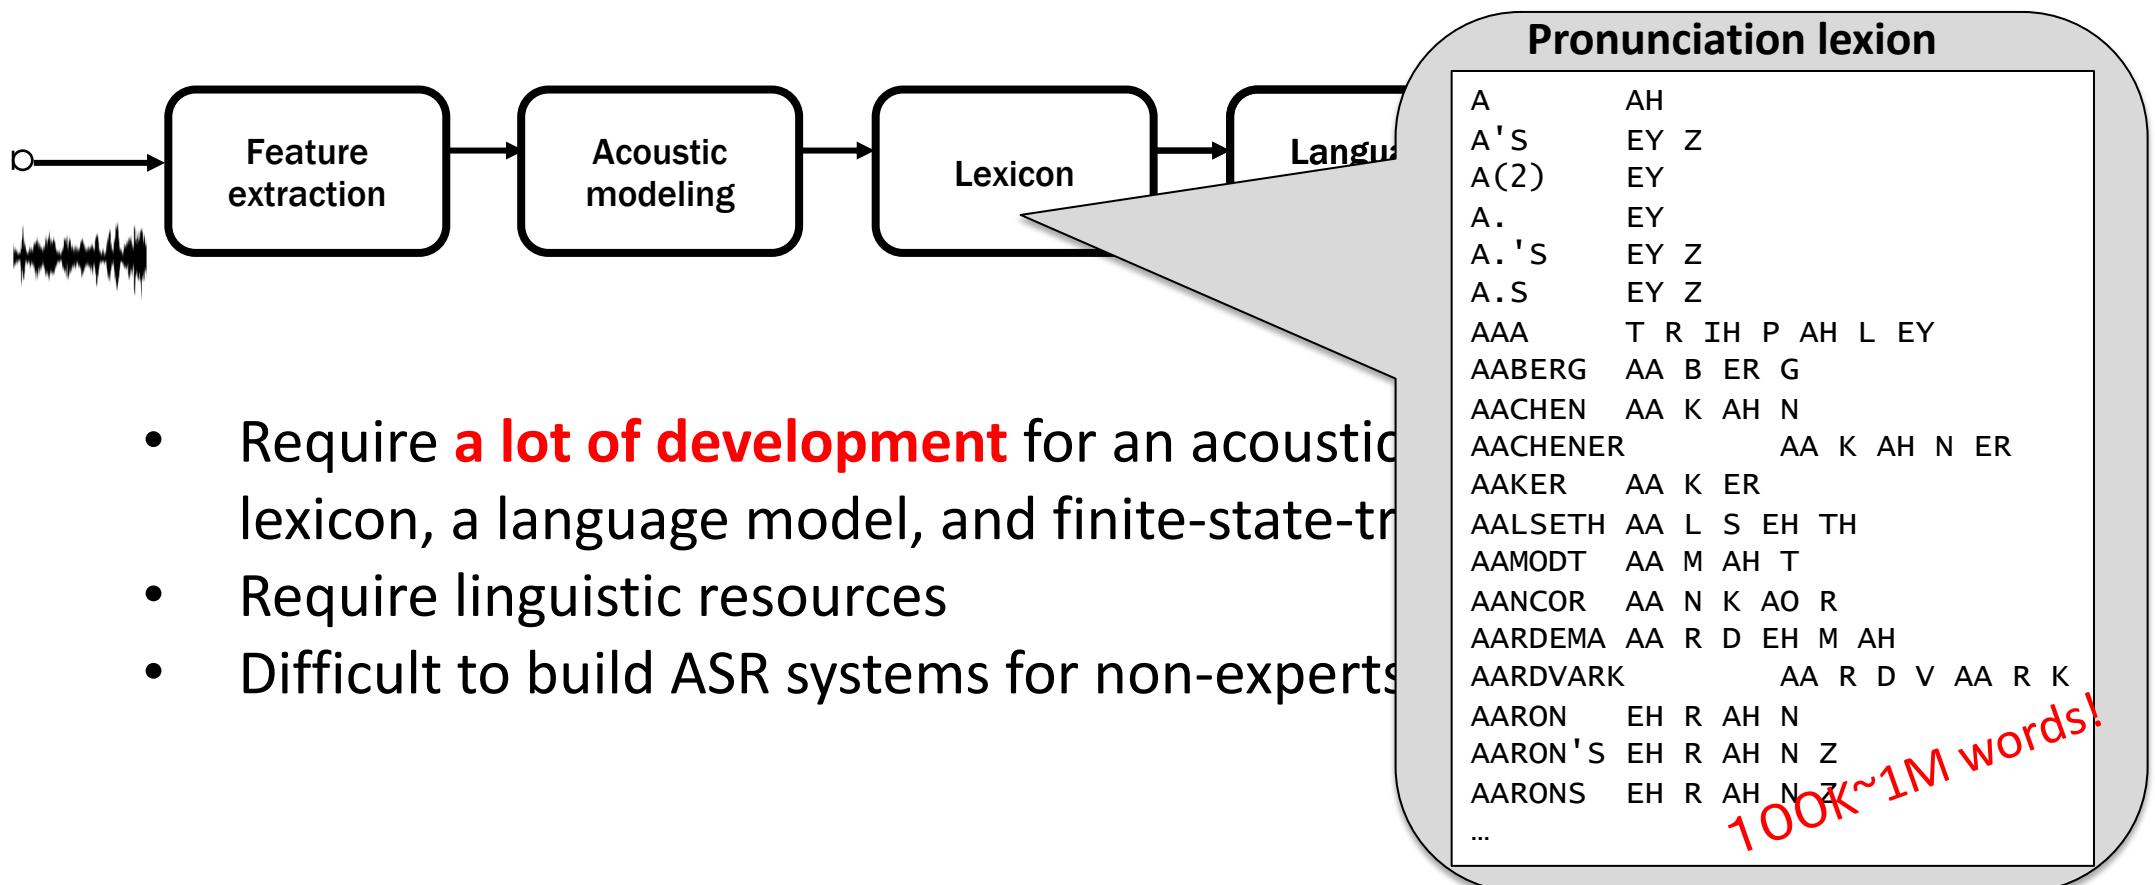


Speech recognition pipeline

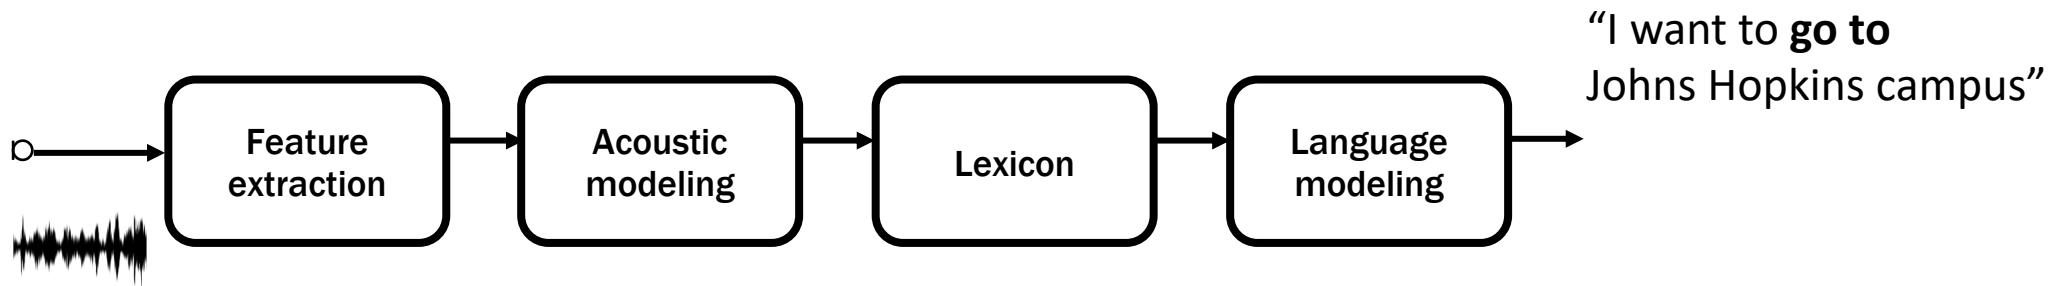


- Require **a lot of development** for an acoustic model, a pronunciation lexicon, a language model, and finite-state-transducer decoding
- Require linguistic resources
- Difficult to build ASR systems for non-experts

Speech recognition pipeline

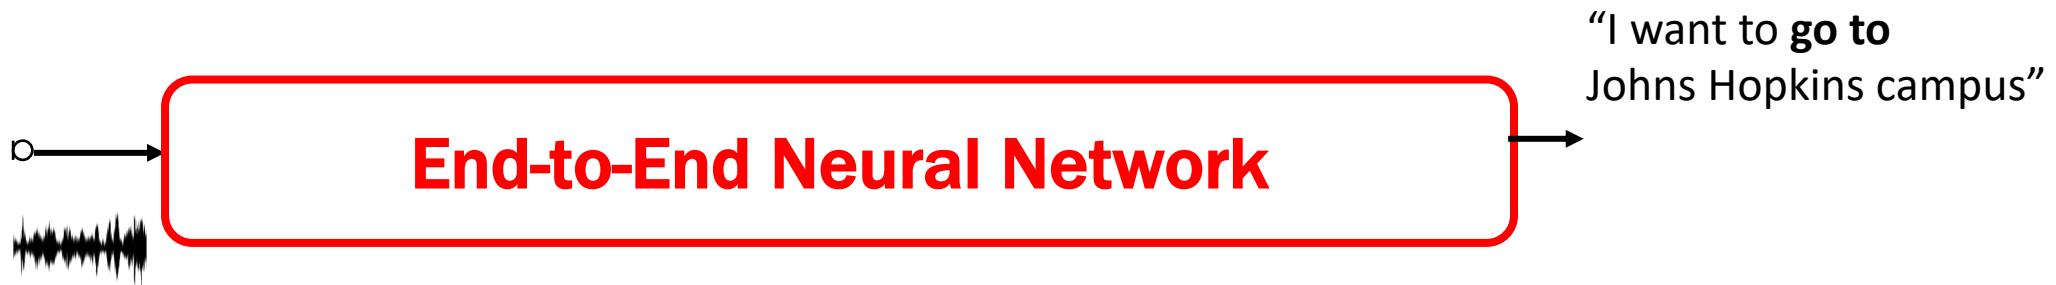


Speech recognition pipeline



- Require **a lot of development** for an acoustic model, a pronunciation lexicon, a language model, and finite-state-transducer decoding
- Require linguistic resources
- Difficult to build ASR systems for **non-experts**

From pipeline to integrated architecture



- Train a deep network that directly maps speech signal to the target letter/word sequence
- Greatly simplify the complicated model-building/decoding process
- Easy to build ASR systems for new tasks **without expert knowledge**
- Potential to outperform conventional ASR by **optimizing the entire network** with a single objective function

Japanese is a very ASR unfriendly language

“二つ目の要因は計算機資源・音声データの増加及びKaldiやTensorflowなどのオープンソースソフトウェアの普及である”

- **No word boundary**
- **Mix of 4 scripts** (Hiragana, Katakana, Kanji, Roman alphabet)
- Frequent **many to many pronunciations**
 - A lot of homonym (same pronunciations but different chars.)
 - A lot of multiple pronunciations for each char
- **Very different phoneme lengths per character**
 - “ン”: /n/, “侍”: /s/ /a/ /m/ /u/ /r/ /a/ /i/ (from 1 to 7 phonemes per character!)

We need very accurate **tokenizer** (chasen, mecab) to solve the above problems **jointly**

My attempt (2016)

- Japanese NLP/ASR: always go through NAIST Matsumoto lab's tokenizer

言葉の理解を
コンピューターで深める

情報科学研究科 自然言語処理学研究室
松本 裕治 教授 ケビン・デュー 助教

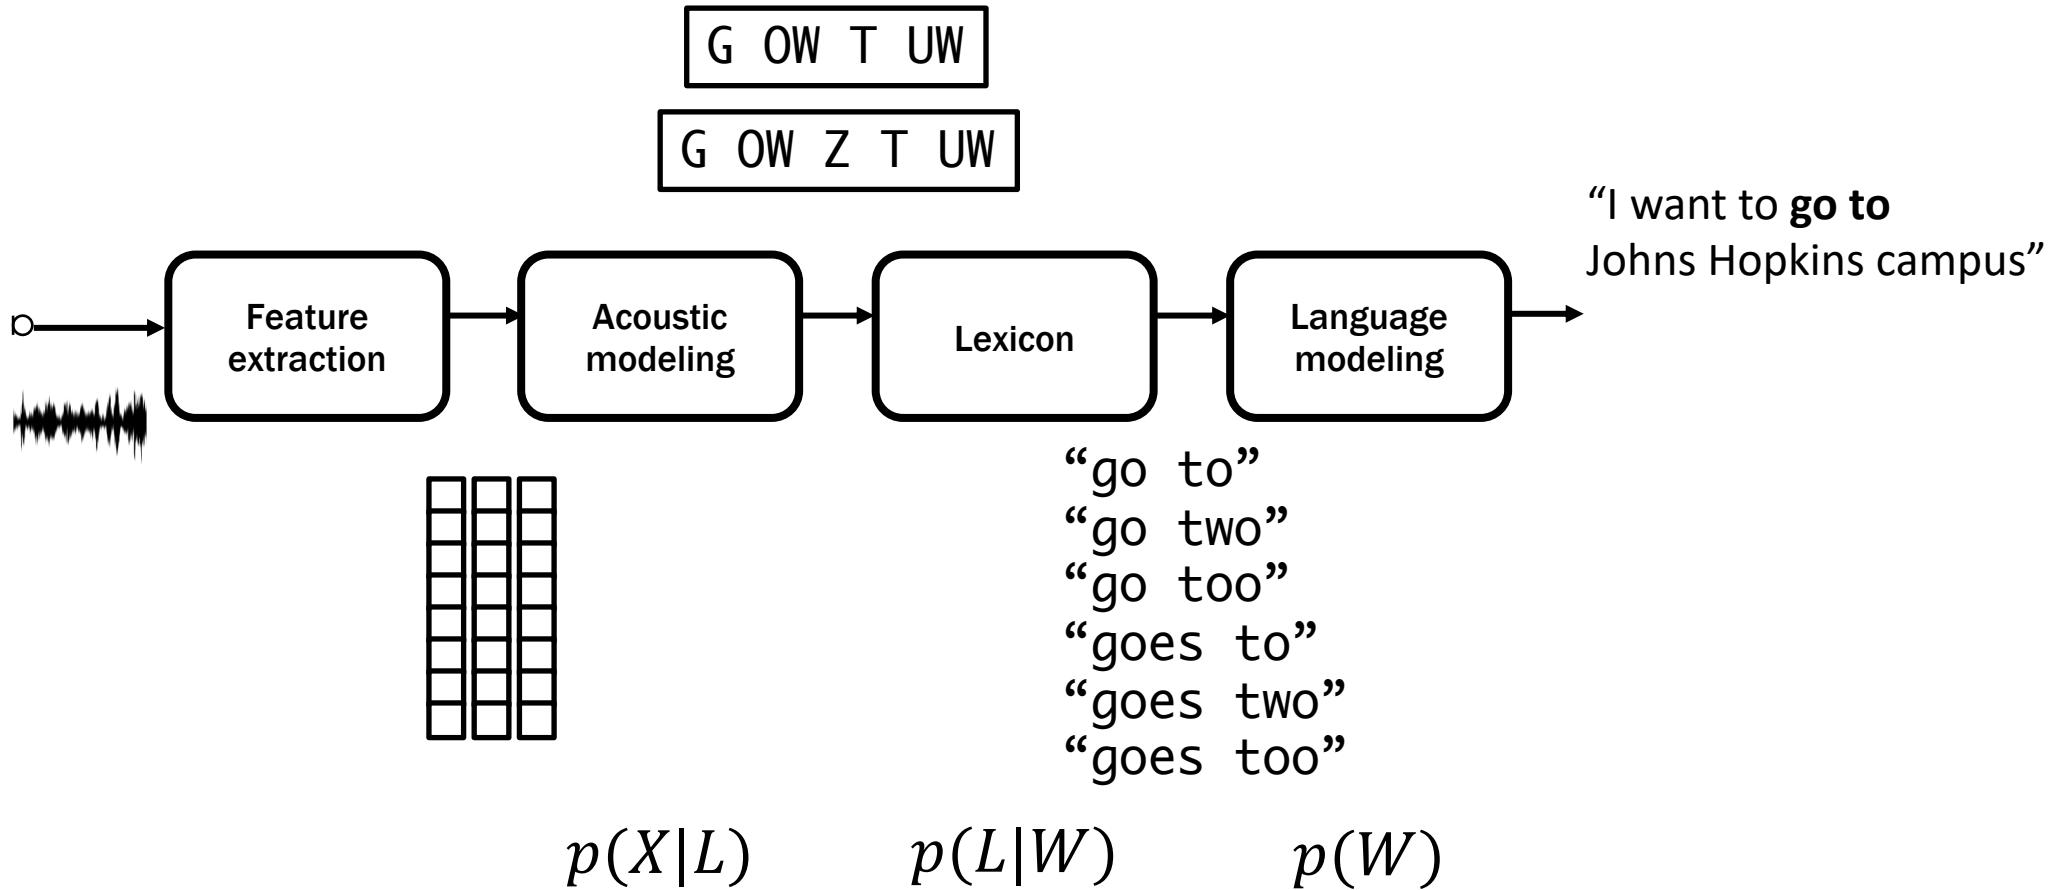


- My goal: remove the tokenizer (sorry Kevin)
- Directly predict Japanese text only from audio
- Surprisingly working very well. Our initial attempt reached Kaldi state-of-the-art with a tokenizer (CER~10% (2016) cf. ~5% (2020))
- This was the first Japanese ASR without using tokenizer (one of my dreams)

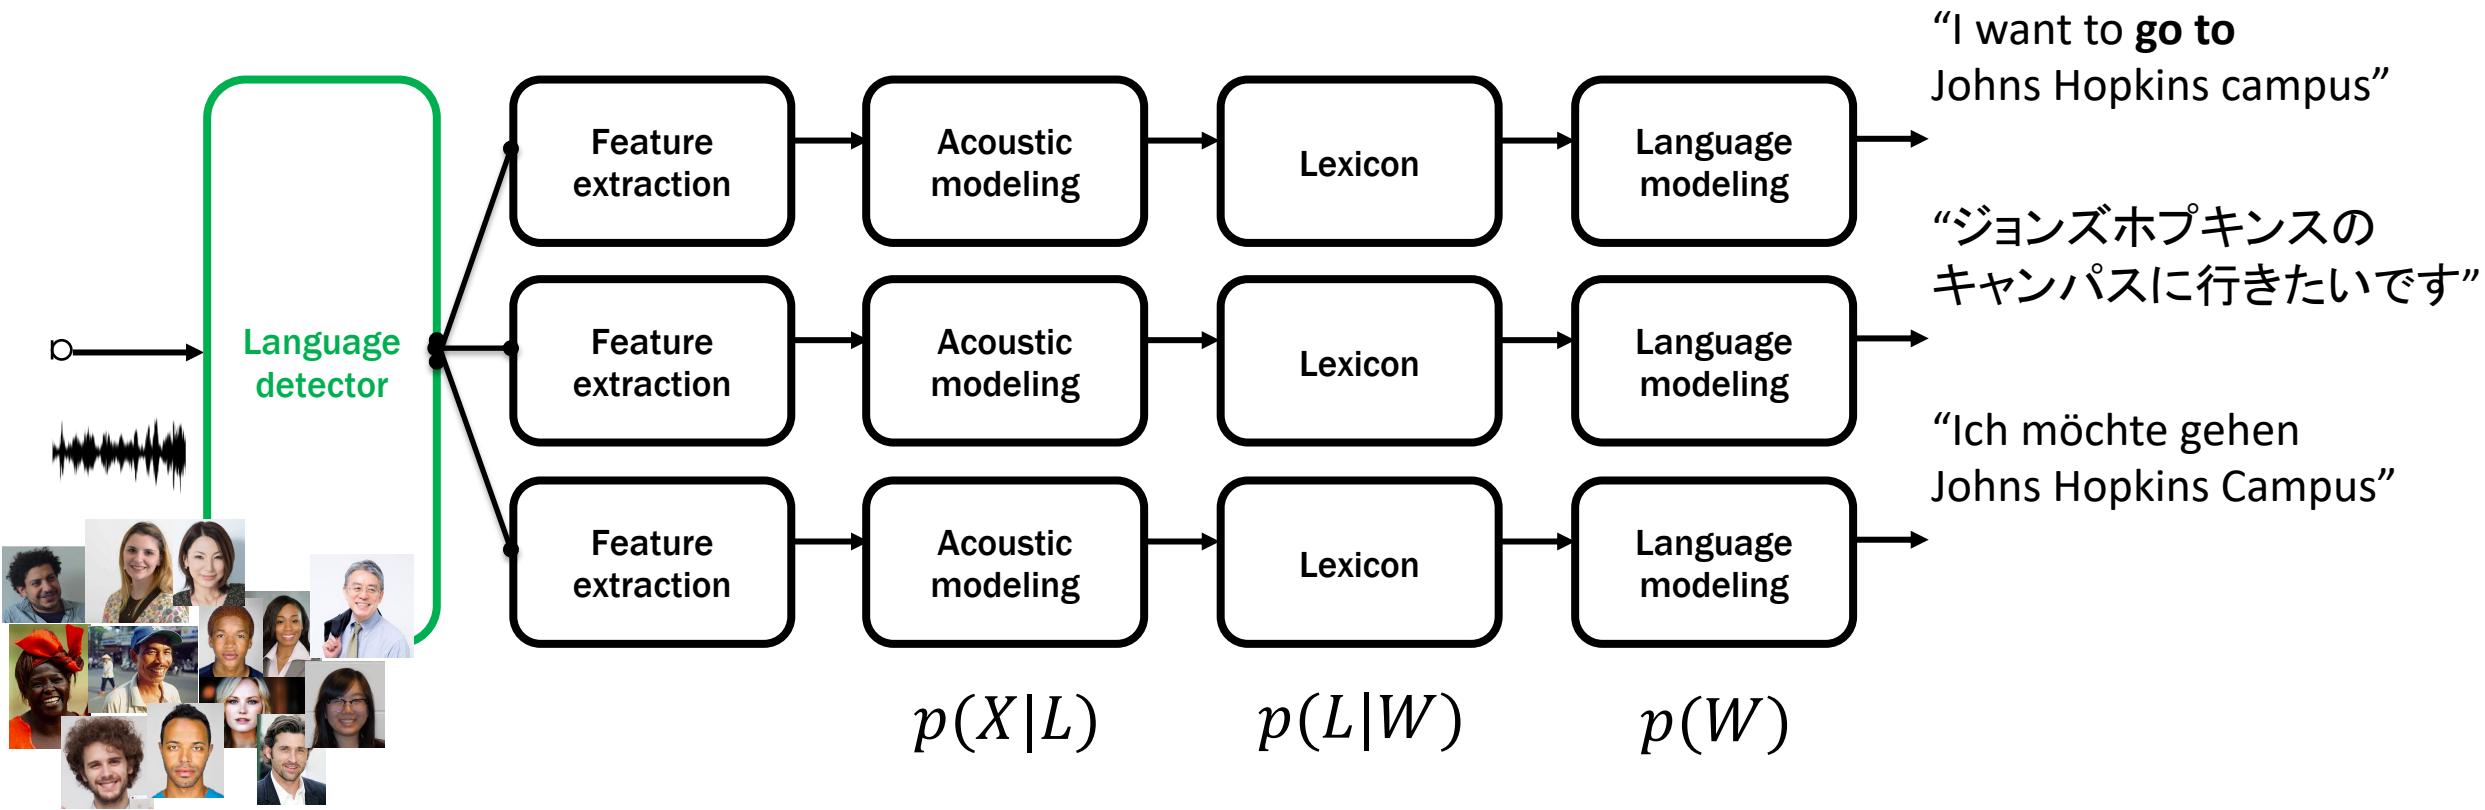
Multilingual e2e ASR

- Given the Japanese ASR experience, I thought that e2e ASR can handle mixed languages with a single architecture
 - ➔ Multilingual e2e ASR (2017)
 - ➔ Multilingual code-switching e2e ASR (2018)

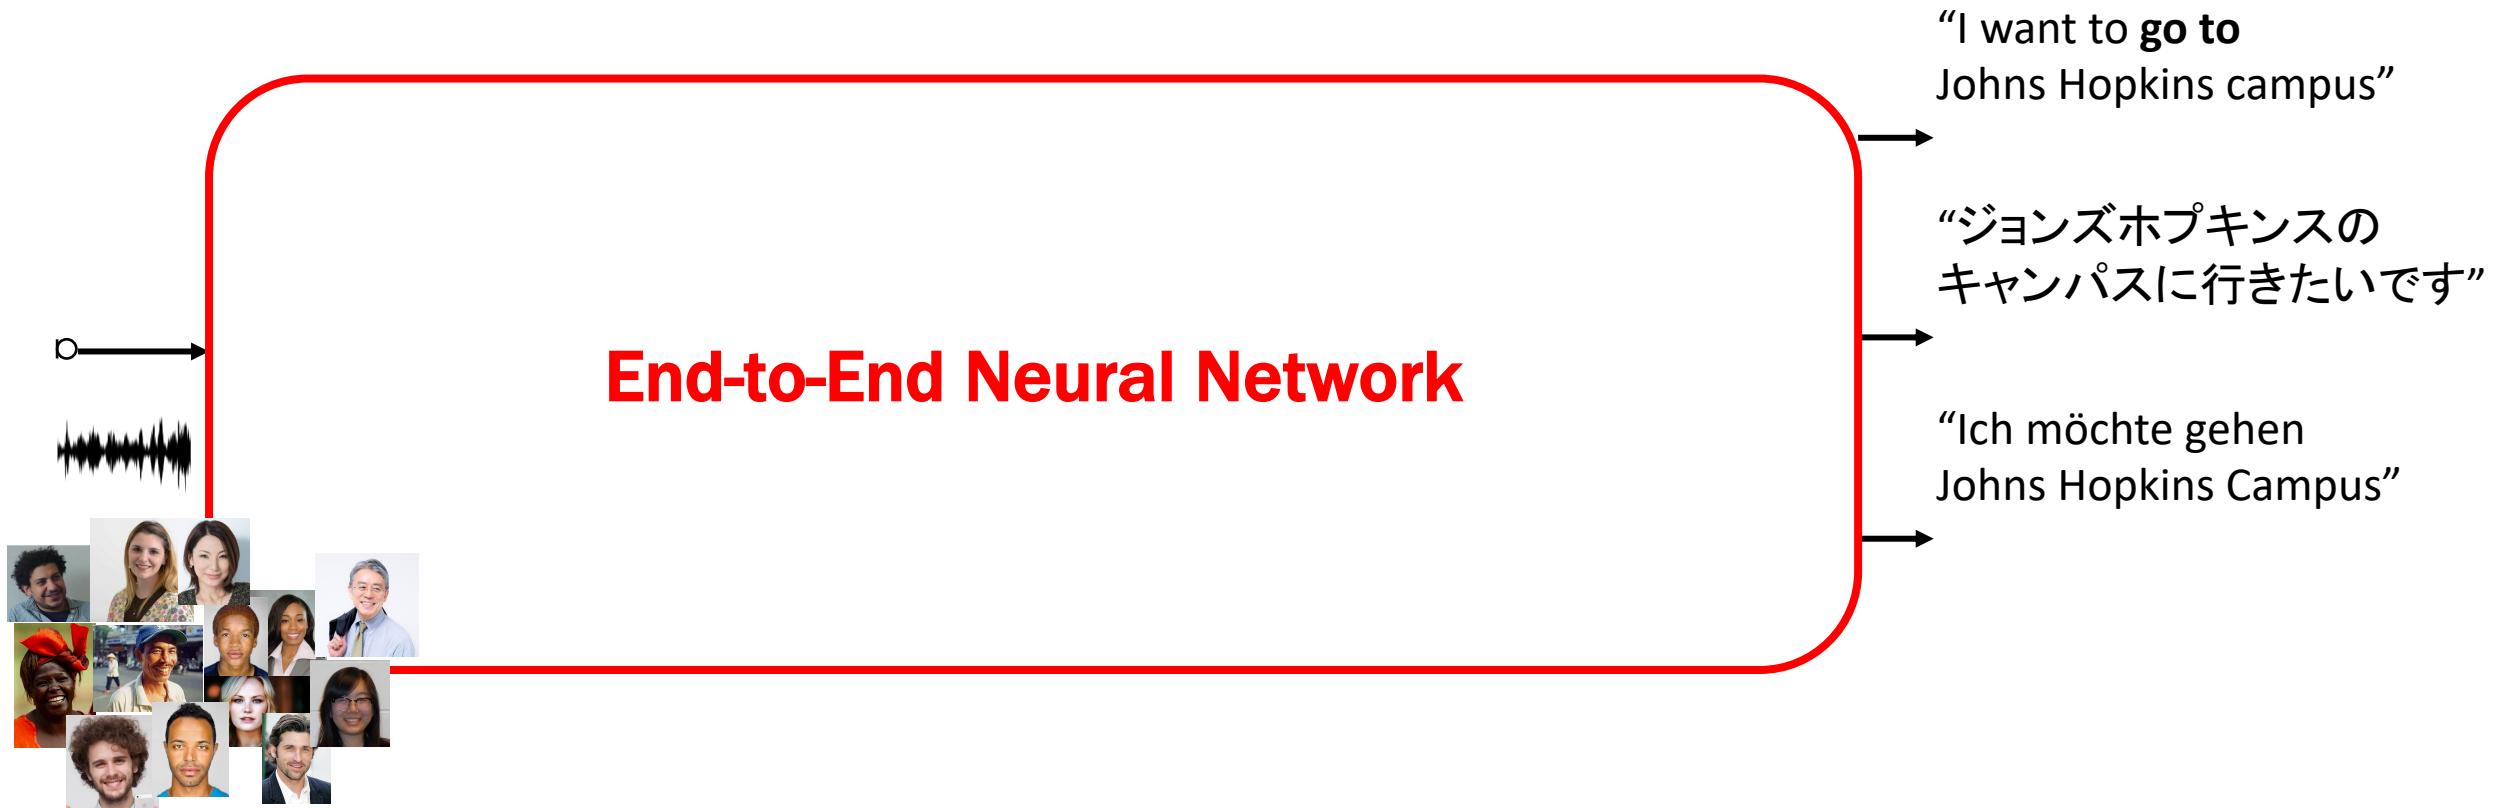
Speech recognition pipeline



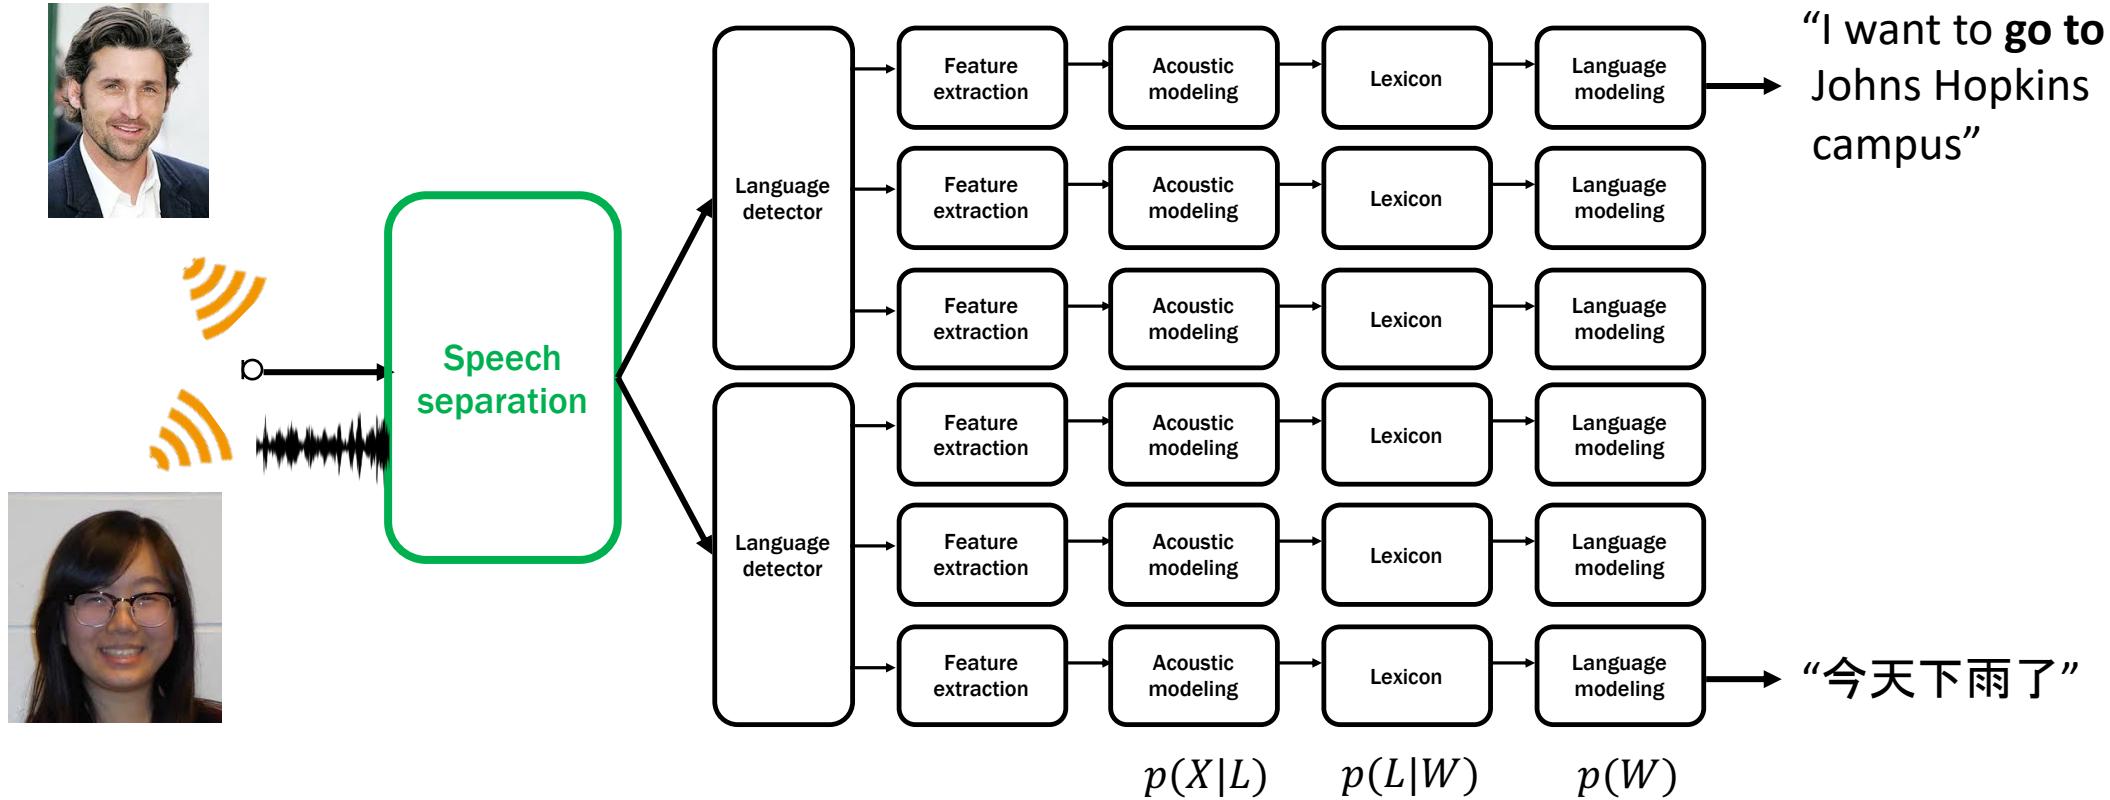
Multilingual speech recognition pipeline



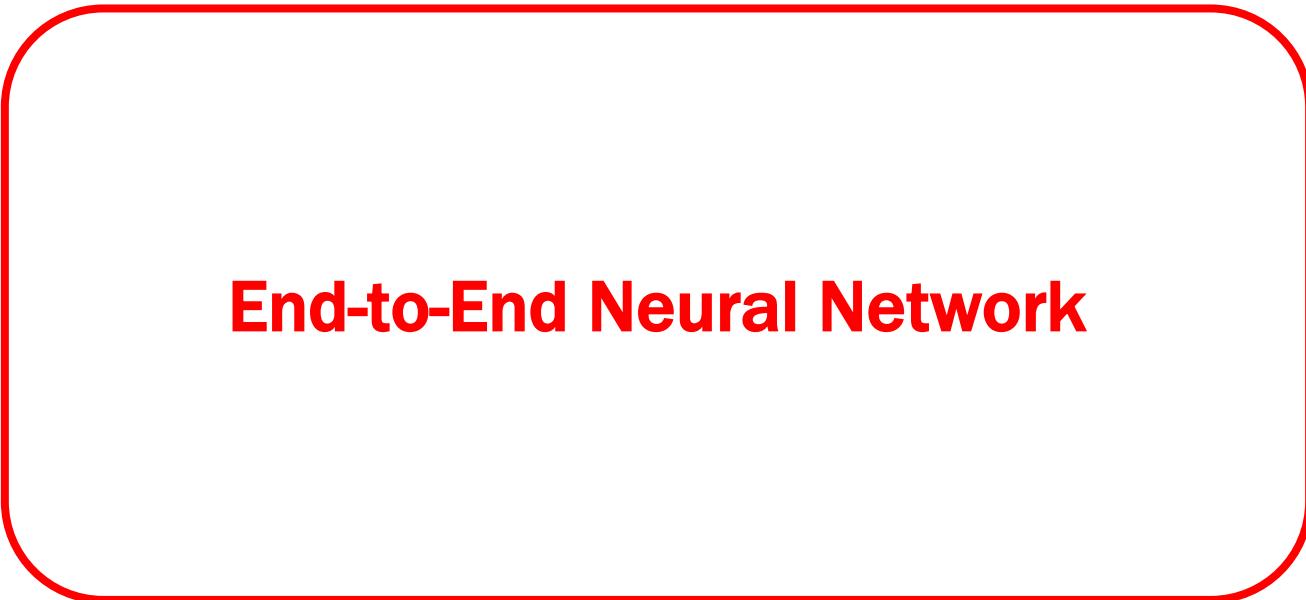
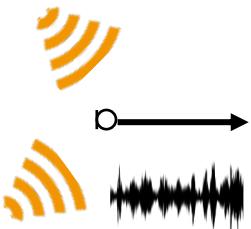
Multilingual speech recognition pipeline



Multi-speaker multilingual speech recognition pipeline



Multi-speaker multilingual speech recognition pipeline



→ “I want to go to
Johns Hopkins
campus”

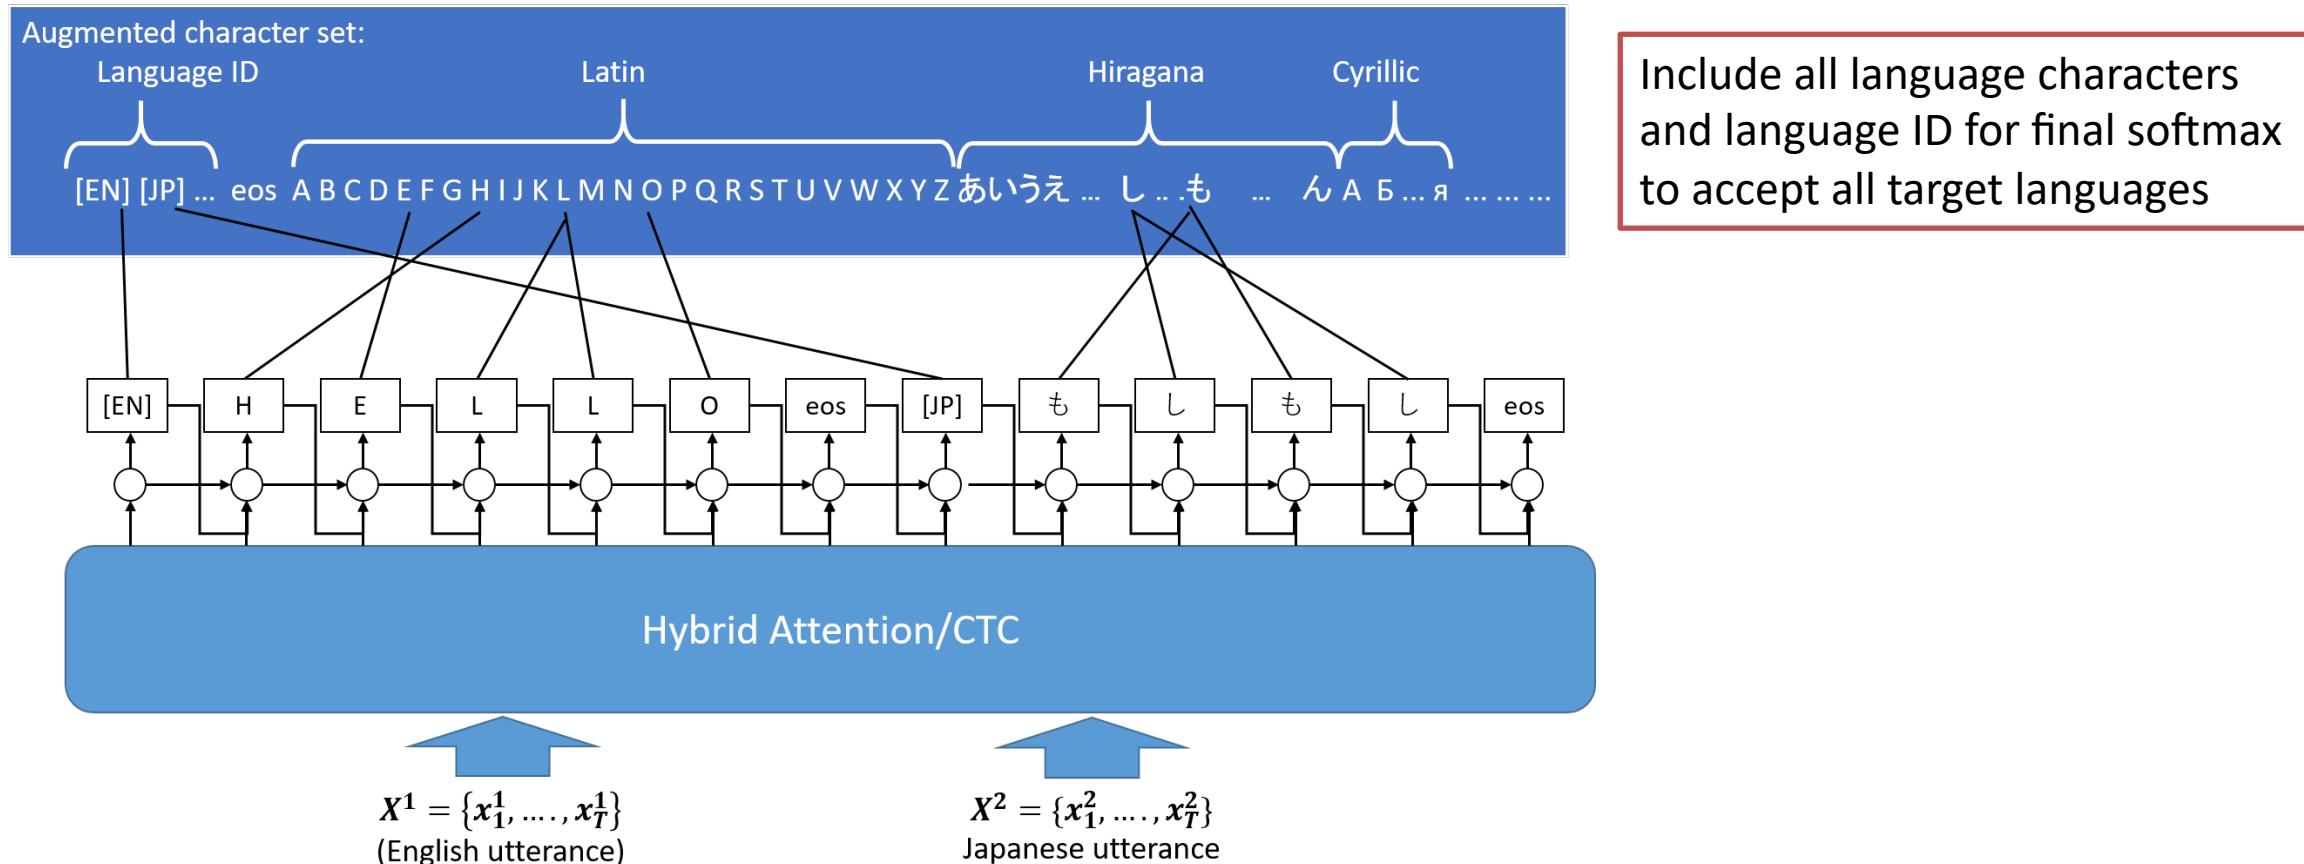
→ “今天下雨了”



Multi-lingual end-to-end speech recognition

[Watanabe+'17, Seki+'18]

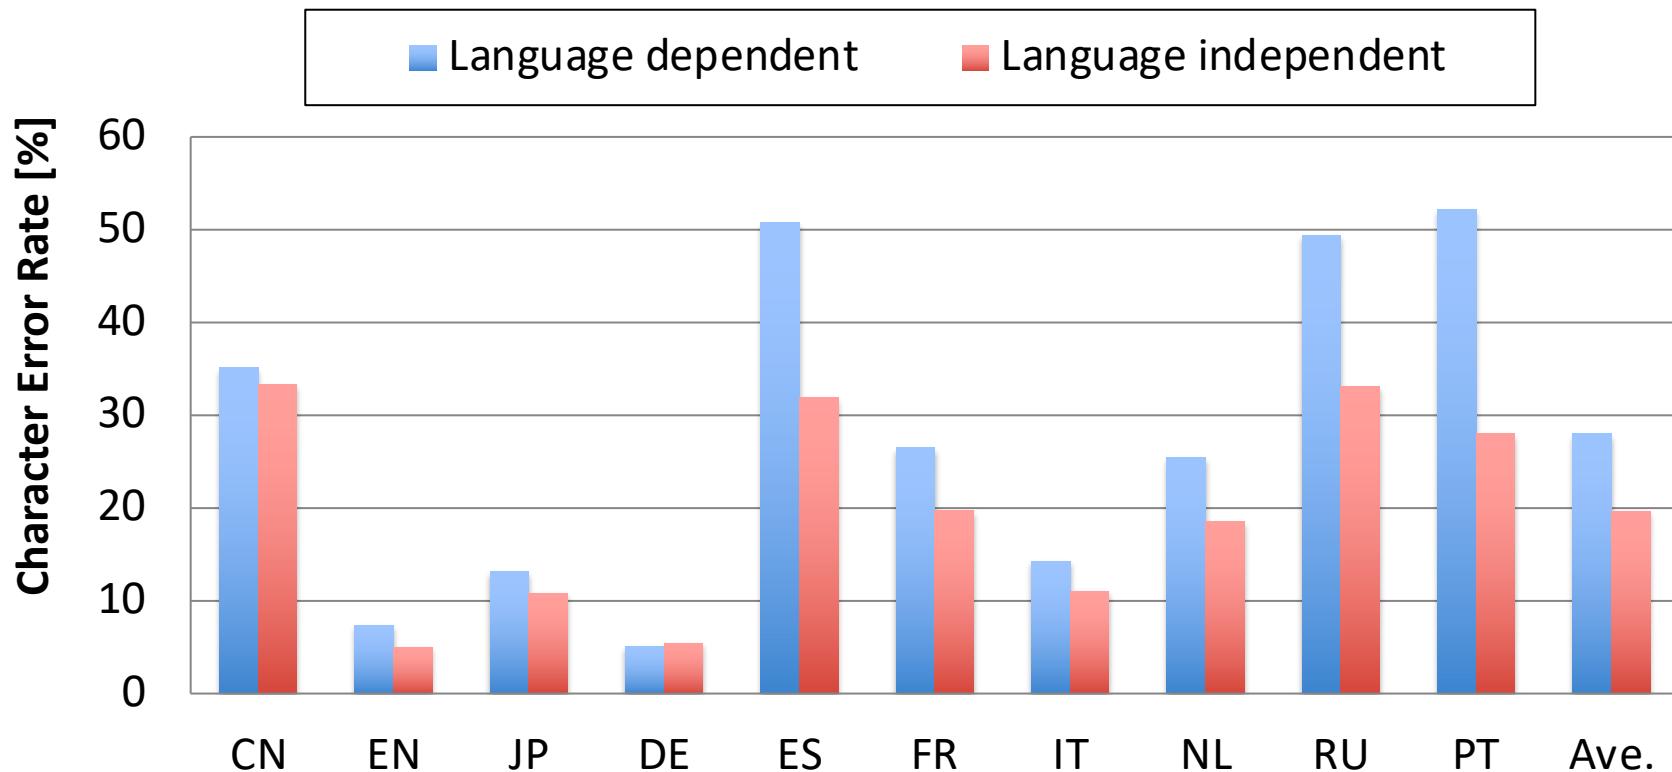
- Learn a single model with multi-language data (10 languages)
- **Integrates** language identification and 10-language speech recognition systems
- **No pronunciation lexicons**





ASR performance for 10 languages

- Comparison with language dependent systems
- Language-independent single end-to-end ASR works well!



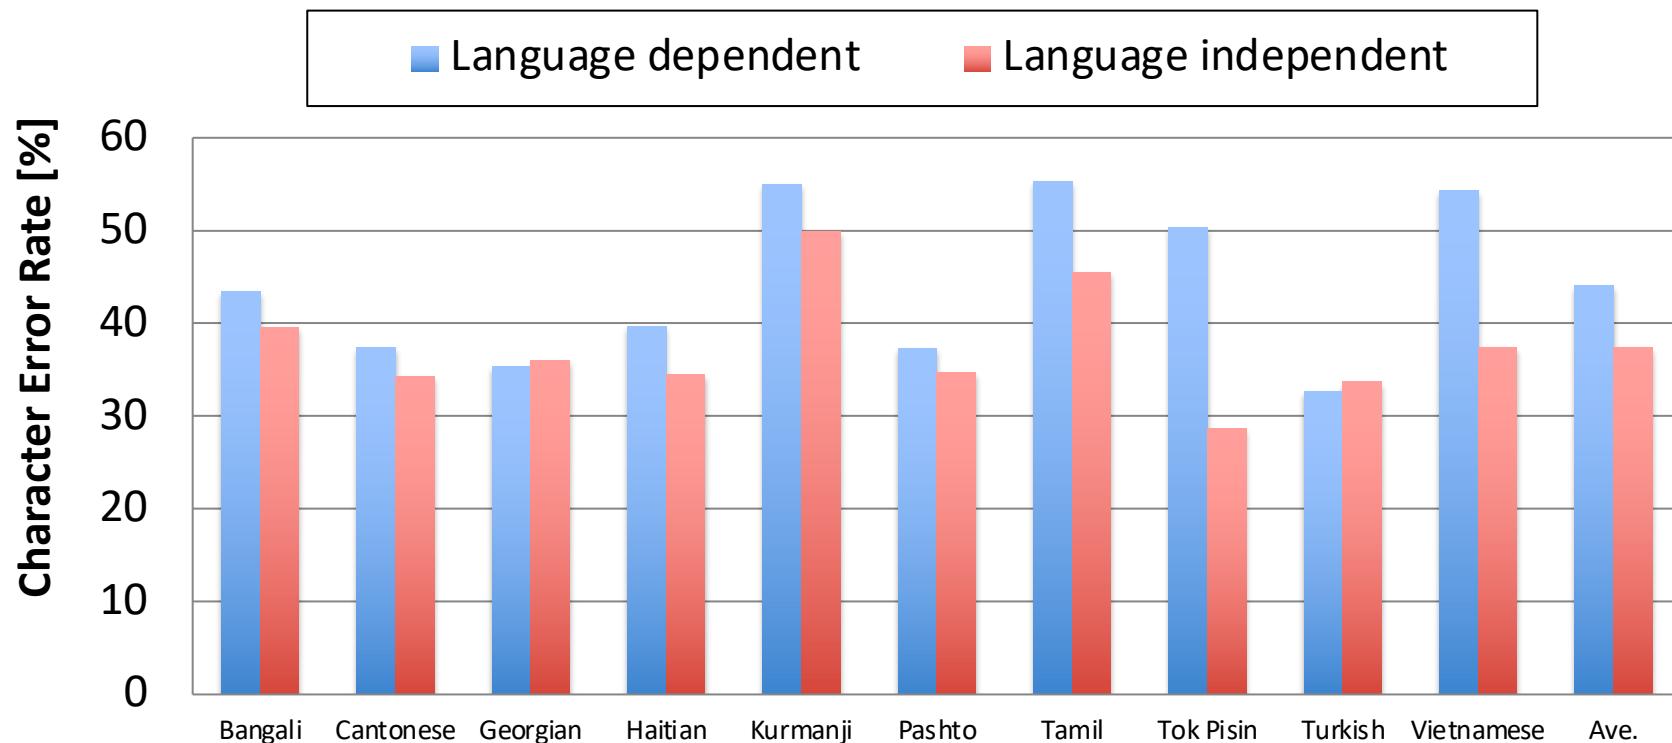
你好
Hello
こんにちは
Hallo
Hola
Bonjour
Ciao
Hallo
Привет
Olá

Language recognition performance

		CH	EN	JP	DE	ES	FR	IT	NL	RU	PT
CH	train_dev	100.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
	dev	100.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
EN	test_eval92	0.0	100.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
	test_dev93	0.0	100.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
JP	eval1_jpn	0.0	0.0	100.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
	eval2_jpn	0.0	0.0	100.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
DE	eval3_jpn	0.0	0.0	99.9	0.0	0.0	0.0	0.1	0.0	0.0	0.0
	et_de	0.0	0.0	0.0	99.7	0.0	0.0	0.0	0.3	0.0	0.0
ES	dt_de	0.0	0.0	0.0	99.7	0.0	0.0	0.0	0.3	0.0	0.0
	dt_es	0.0	0.0	0.0	0.0	67.9	0.0	31.9	0.0	0.0	0.2
FR	et_es	0.0	0.0	0.0	0.1	91.1	0.0	8.4	0.1	0.0	0.2
	dt_fr	0.0	0.0	0.0	0.1	0.0	99.4	0.0	0.2	0.0	0.3
IT	et_fr	0.0	0.0	0.0	0.1	0.0	99.5	0.0	0.1	0.0	0.3
	dt_it	0.0	0.0	0.0	0.0	0.3	0.4	99.1	0.0	0.0	0.3
NL	et_it	0.0	0.0	0.0	0.0	0.4	0.4	98.3	0.2	0.1	0.7
	dt_nl	0.0	0.0	0.0	1.3	0.0	0.1	0.1	97.2	0.0	1.3
RU	et_nl	0.0	0.0	0.0	1.0	0.0	0.2	0.2	97.6	0.0	0.9
	dt_ru	0.2	0.0	0.0	0.0	0.2	0.6	0.5	0.0	97.9	0.8
PT	et_ru	0.0	0.0	0.0	0.2	0.2	0.3	4.3	0.0	94.7	0.3
	dt_pt	0.0	0.0	0.0	0.3	0.3	2.6	1.7	3.4	0.6	91.2
PT	et_pt	0.0	0.3	0.0	0.3	0.0	0.0	3.9	3.6	0.3	91.5

ASR performance for law-resource 10 languages

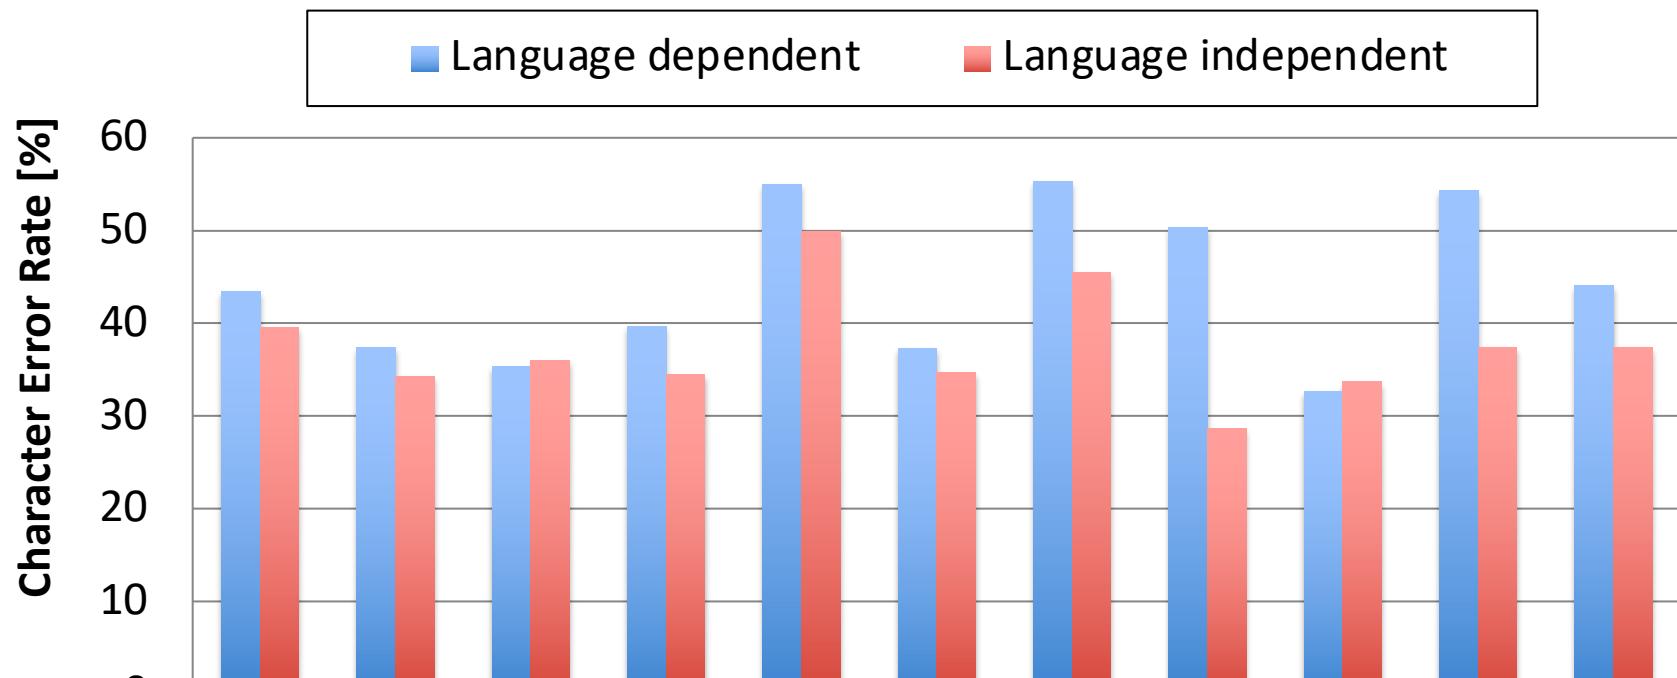
- Comparison with language dependent systems



ଶ୍ୟାଳୋ
你好
ଘାମାର୍ଜନବା
hello
???
سلام
வணக்கம்
???
Merhaba
xin chào

ASR performance for law-resource 10 languages

- Comparison with language dependent systems



~100 languages with CMU Wilderness Multilingual Speech Dataset
[Adams+(2019)]

ଶ୍ୟାଳୋ
你好
ଧାର୍ମାର୍ଜନବା
hello
???

سلام
வணக்கம்
???

Merhaba
xin chào

Actually it was one of the easiest studies in my work

Q. How many people were involved in the development?

A. 1 person

Q. How long did it take to build a system?

A. Totally ~1 or 2 day efforts with bash and python scripting (no change of main e2e ASR source code), **then I waited 10 days to finish training**

Q. What kind of linguistic knowledge did you require?

A. Unicode (because python2 Unicode treatment is tricky. If I used python3, I would not even have to consider it)

ASRU'17 best paper **candidate (not best paper ☹)**

Multi-lingual ASR

(Supporting 10 languages: CN, EN, JP, DE, ES, FR, IT, NL, RU, PT)

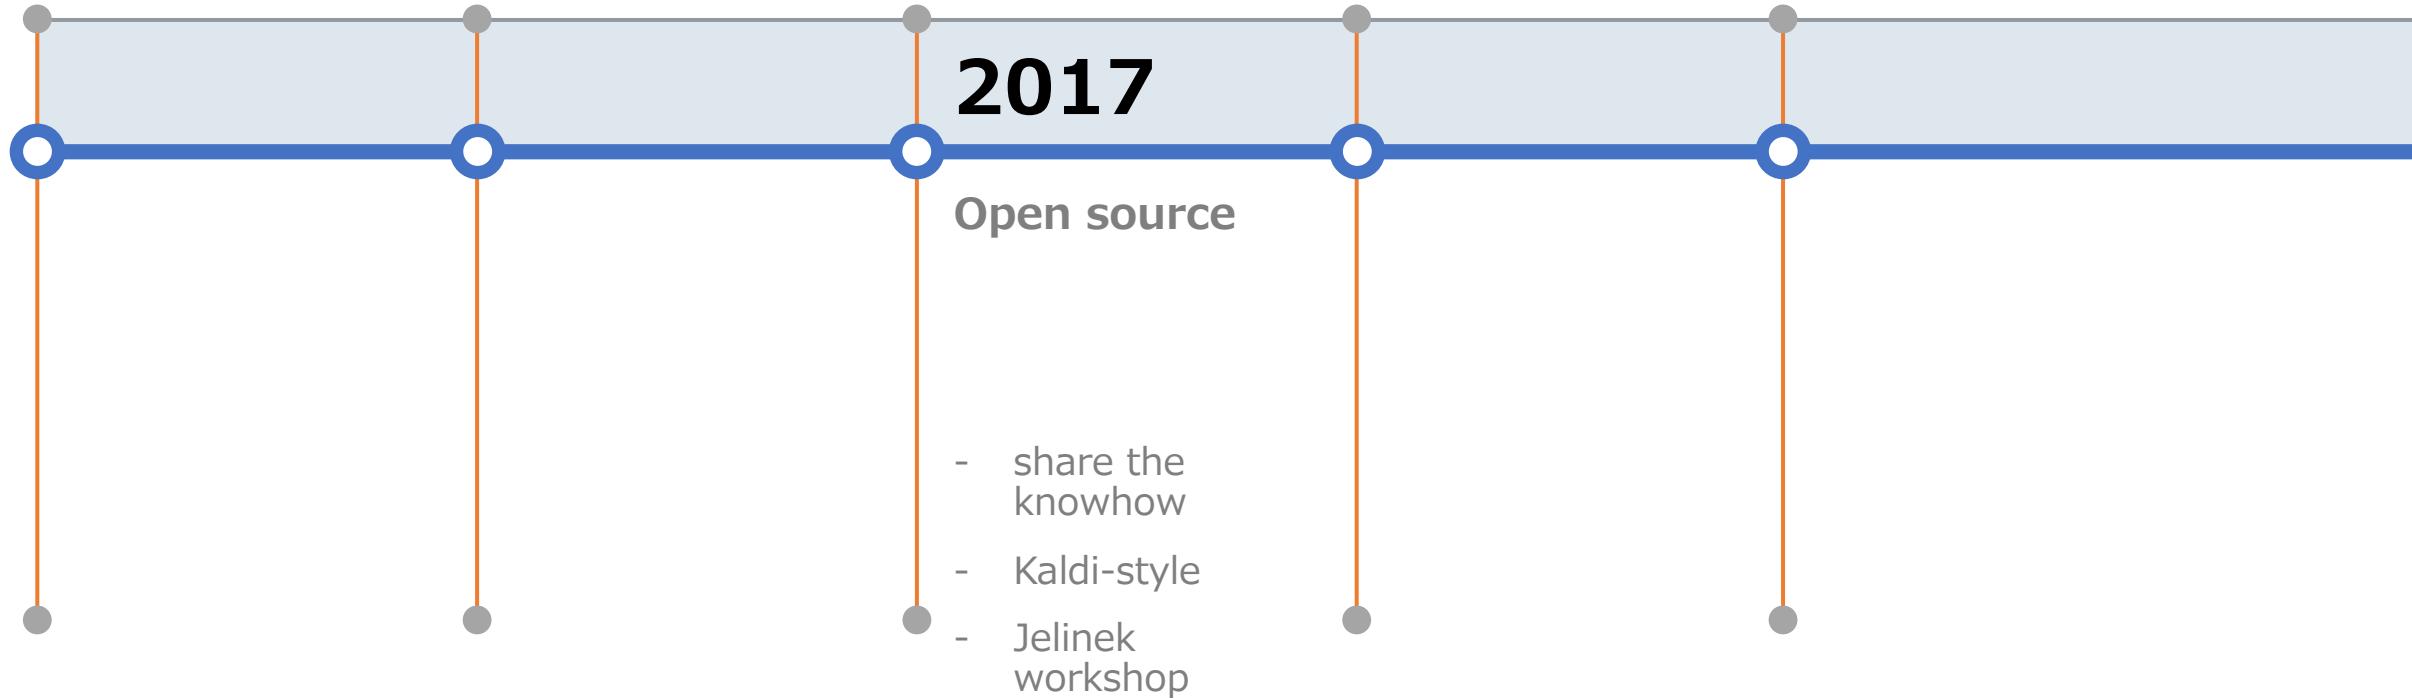
ID	a04m0051_0.352274410405	
	REF: [DE] bisher sind diese personen rundherum versorgt worden [EN] u. s. exports rose in the month but not nearly as much as imports ASR: [DE] bisher sind diese personen rundherum versorgt worden [EN] u. s. exports rose in the month but not nearly as much as imports	

ID	csj-eval:s00m0070-0242356-0244956:voxforge-et-fr:mirage59-20120206-njp-fr-sb-570	
	REF: [JP] 日本でもニュースになったと思いますが [FR] le conseil supérieur de la magistrature est présidé par le président de la république ASR: [JP] 日本でもニュースになったと思いますが [FR] le conseil supérieur de la magistrature est présidé par le président de la république	

ID	voxforge-et-pt:insinfo-20120622-orb-209:voxforge-et-de:guenter-20140127-usn-de5-069:csj-eval:a01m0110-0243648-0247512	
	REF: [PT] segunda feira [DE] das gilt natürlich auch für bestehende verträge [JP] え同一人物による異なるメッセージを示しております ASR: [PT] segunda feira [DE] das gilt natürlich auch für bestehende verträge [JP] え同一人物による異なるメッセージを示しております	

Timeline

Shinji's personal experience for end-to-end speech processing





ESPnet: End-to-end speech processing toolkit

Shinji Watanabe

Center for Language and Speech Processing

Johns Hopkins University

Joint work with Takaaki Hori , Shigeki Karita, Tomoki Hayashi, Jiro Nishitoba, Yuya Unno, Nelson Enrique Yalta Soplin, Jahn Heymann, Matthew Wiesner, Nanxin Chen, Adithya Renduchintala, Tsubasa Ochiai,

and more and more

ESPnet

- Open source (Apache2.0) end-to-end speech processing toolkit developed at Frederick Jelinek Memorial Summer Workshop 2018
- Major concept

Reproducible end-to-end speech processing studies for speech researchers

Keep simplicity

- Follows the **Kaldi style**
 - Data processing, feature extraction/format
 - Recipes to provide a complete setup for speech processing experiments

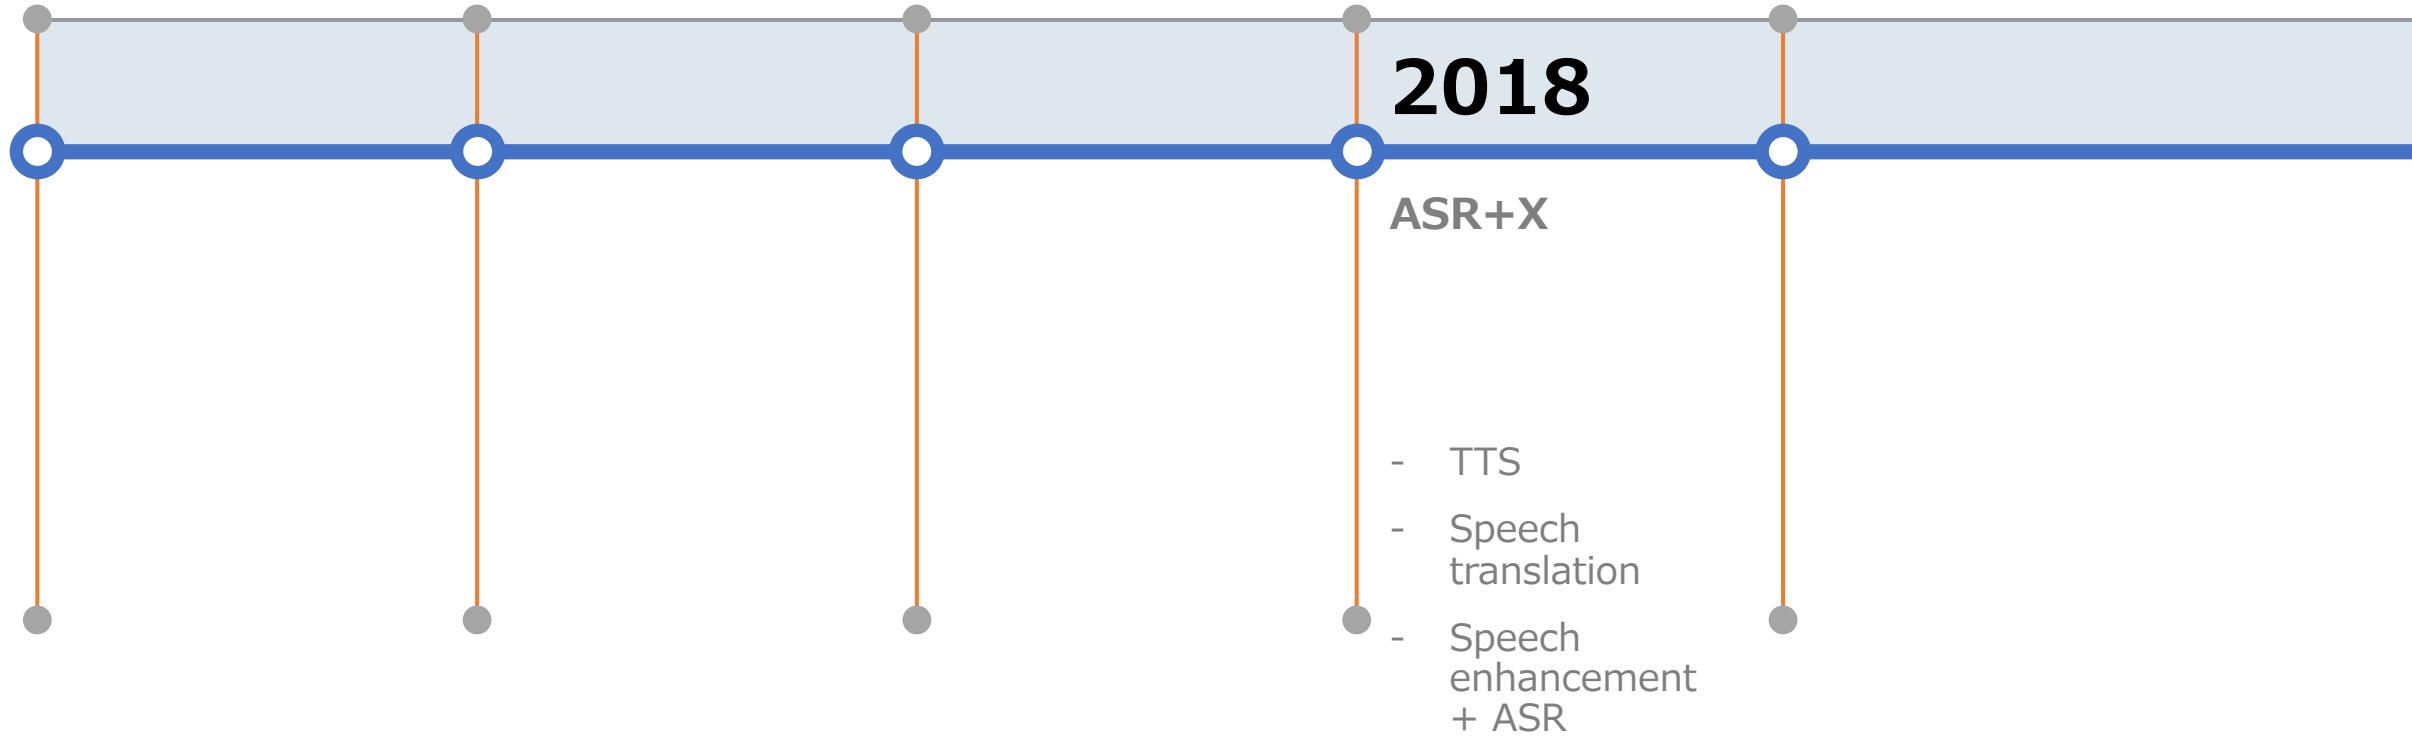
I personally don't like pre-training fine-tuning strategies (but I'm changing my mind)

Functionalities

- Kaldi style data preprocessing
 - 1) fairly comparable to the performance obtained by Kaldi hybrid DNN systems
 - 2) easily porting the Kaldi recipe to the ESPnet recipe
- Attention-based encoder-decoder
 - Subsampled BLSTM and/or VGG-like encoder and location-based attention (+10 attentions)
 - beam search decoding
- CTC
 - WarpCTC, beam search (label-synchronous) decoding
- **Hybrid CTC/attention**
 - Multitask learning
 - Joint decoding with label-synchronous hybrid CTC/attention decoding (solve monotonic alignment issues)
- RNN transducder
 - Warptransducer, beam search (label-synchronous) decoding
- Use of language models
 - Combination of RNNLM/n-gram trained with external text data (shallow fusion)

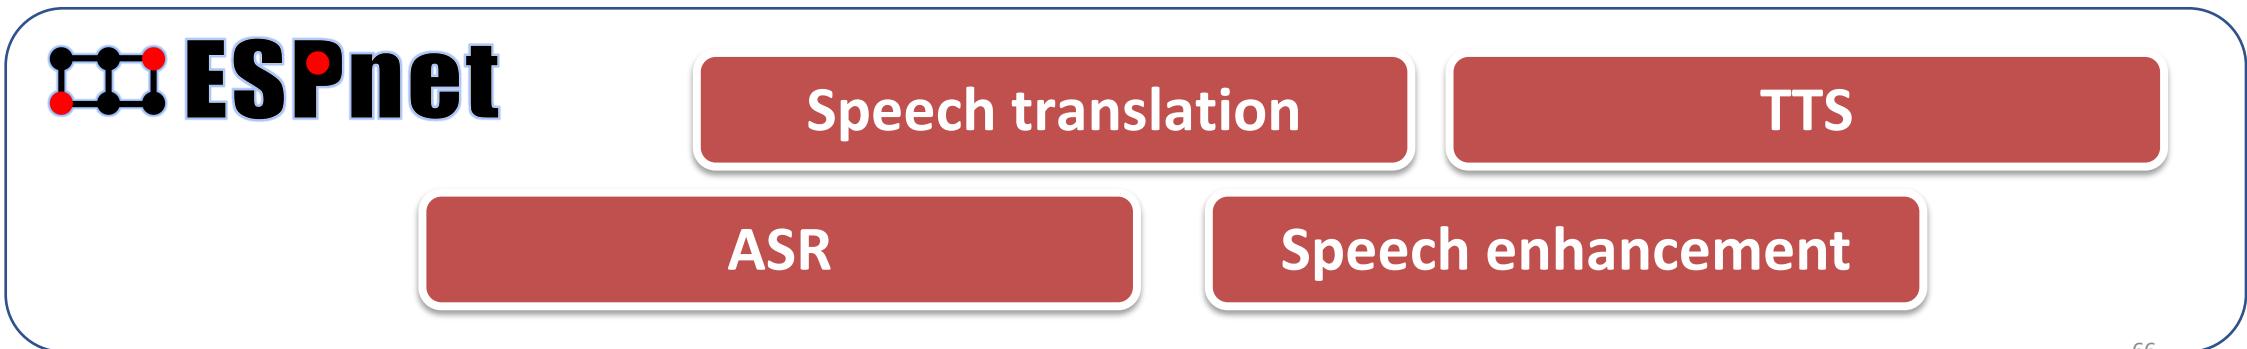
Timeline

Shinji's personal experience for end-to-end speech processing



ASR+X (Sanjeev called it)

- This toolkit (**ASR+X**) covers the following topics complementally



66

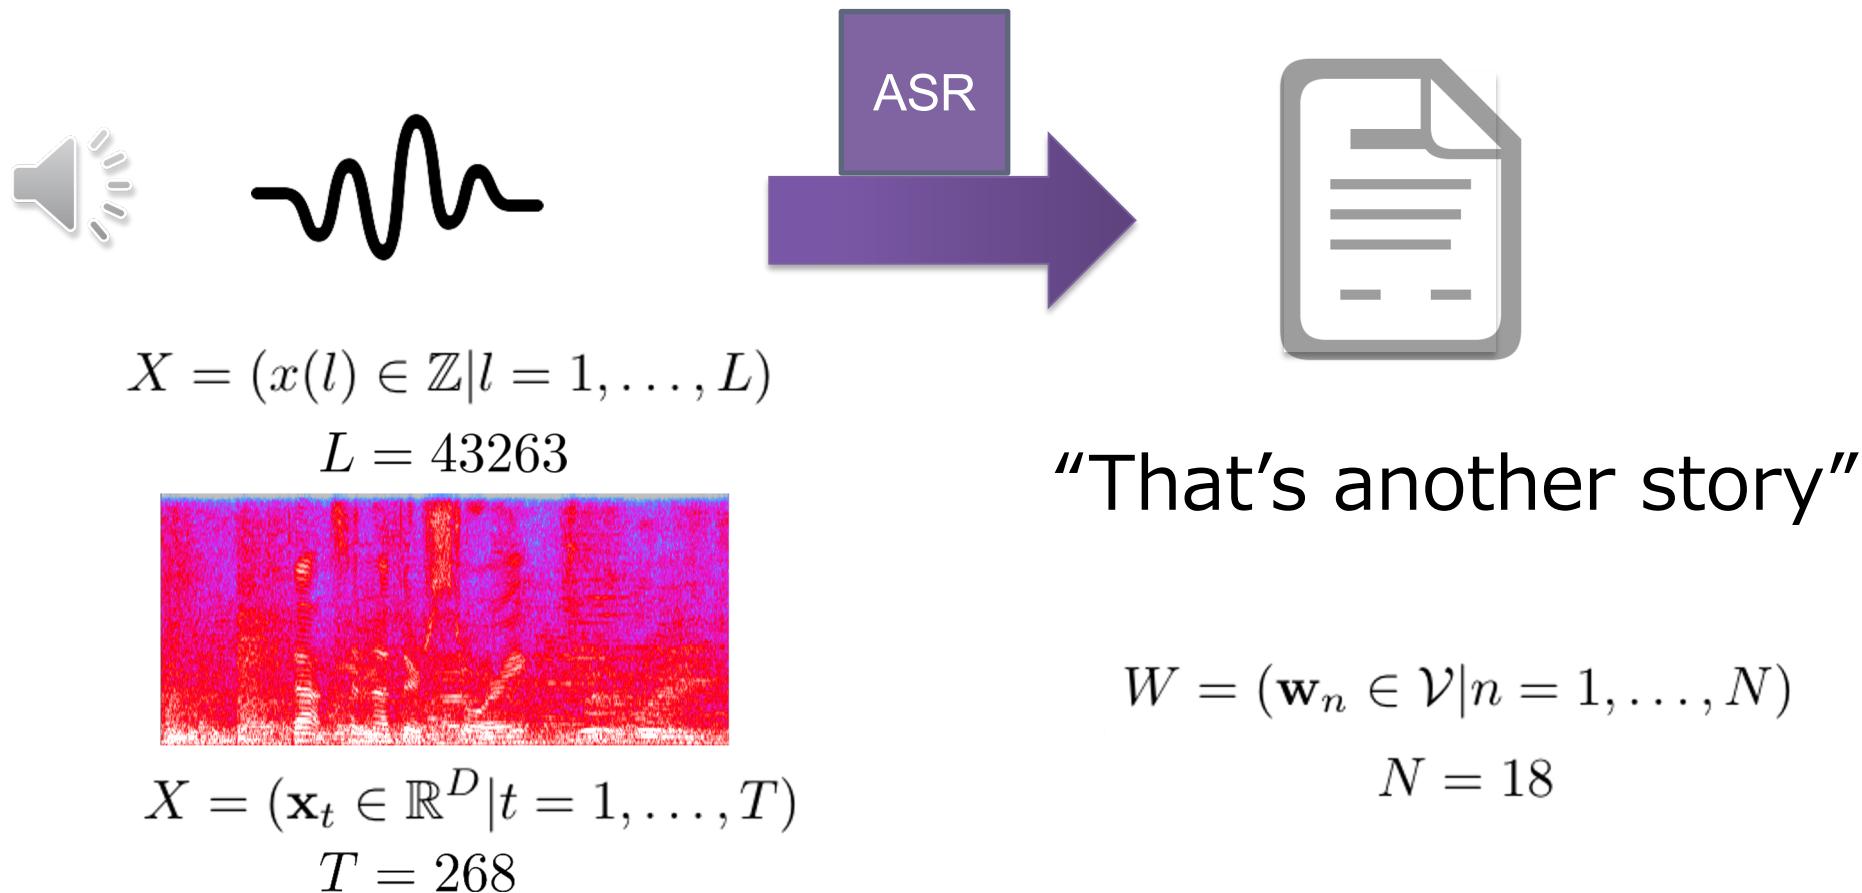
- Why we can support such wide-ranges of applications?

High-level benefit of e2e neural network

- **Unified** views of multiple speech processing applications based on end-to-end neural architecture
- **Integration** of these applications in a single network
- **Implementation** of such applications and their integrations based on an open source toolkit like ESPnet, nemo, espresso, ctc++, fairseq, opennmtpy, lingvo, etc. etc., in an unified manner

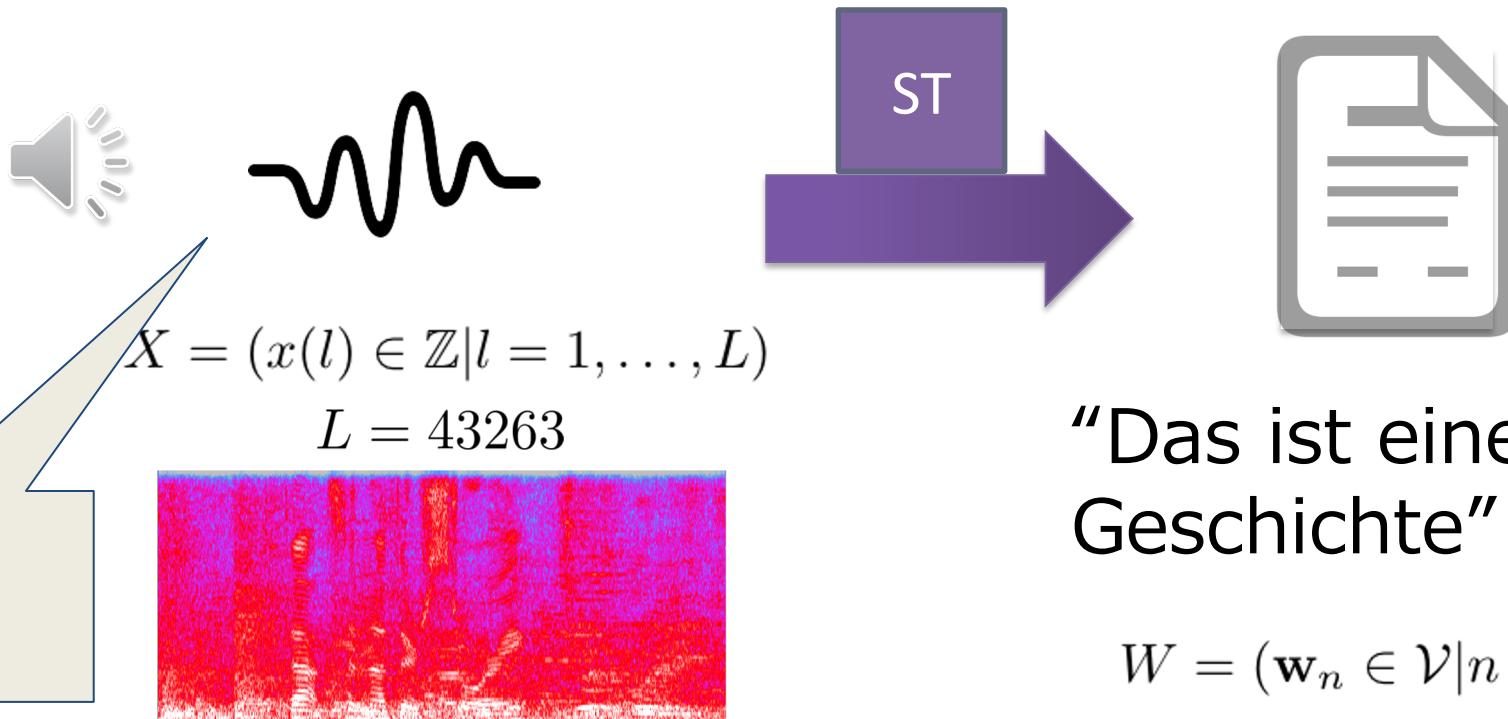
Automatic speech recognition (ASR)

- Mapping **speech** sequence to **character** sequence



Speech to text translation (ST)

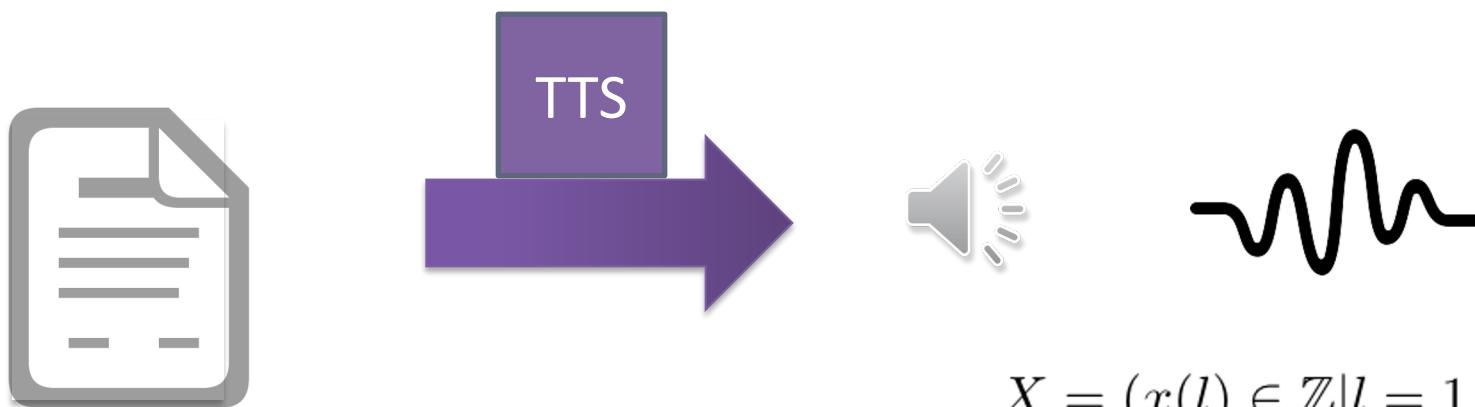
- Mapping **speech** sequence in a **source** language to **character** sequence in a **target** language



“Das ist eine andere
Geschichte”

Text to speech (TTS)

- Mapping **character** sequence to **speech** sequence

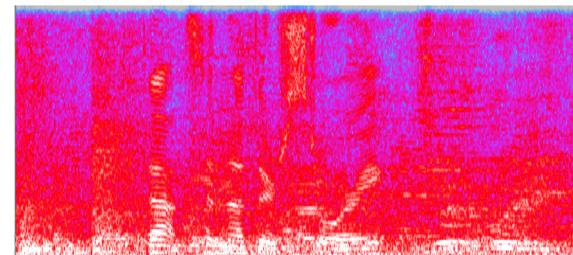


$$X = (x(l) \in \mathbb{Z} | l = 1, \dots, L)$$
$$L = 43263$$

“That’s another story”

$$W = (\mathbf{w}_n \in \mathcal{V} | n = 1, \dots, N)$$

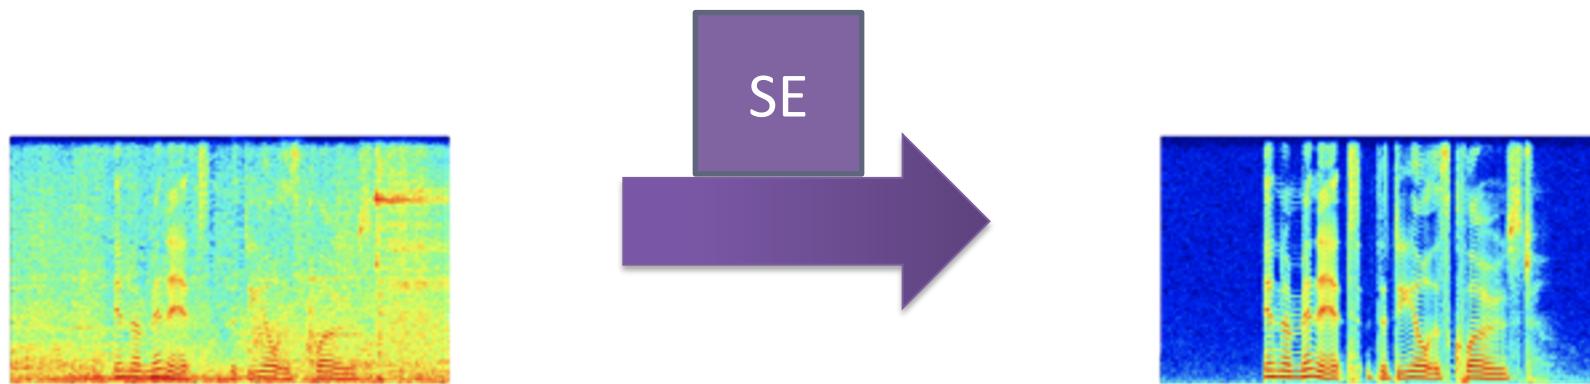
$$N = 18$$



$$X = (\mathbf{x}_t \in \mathbb{R}^D | t = 1, \dots, T)$$
$$T = 268$$

Speech enhancement (SE)

- Mapping **noisy** speech sequence to **clean** speech sequence



$$X = (\mathbf{x}_t \in \mathbb{R}^D | t = 1, \dots, T)$$
$$T = 268$$

$$X' = (\mathbf{x}'_t \in \mathbb{R}^D | t = 1, \dots, T)$$
$$T = 268$$

All of the problems

$$X = (x_1, x_2, \dots, x_T) \xrightarrow{f} Y = (y_1, y_2, \dots, y_N)$$

Unified view with sequence to sequence

- All the above problems: find a mapping function from *sequence to sequence* (**unification**)

$$X = (x_1, x_2, \dots, x_T) \xrightarrow{f} Y = (y_1, y_2, \dots, y_N)$$

- ASR: $X = \text{Speech}$, $Y = \text{Text}$
- TTS: $X = \text{Text}$, $Y = \text{Speech}$
- ST: $X = \text{Speech (EN)}$, $Y = \text{Text (JP)}$
- Speech Enhancement: $X = \text{Noisy speech}$, $Y = \text{Clean speech}$
- Mapping function $f(\cdot)$
 - Sequence to sequence (seq2seq) function
 - ASR as an example

Seq2seq end-to-end ASR

$$X = (x_1, x_2, \dots, x_T) \xrightarrow{f} Y = (y_1, y_2, \dots, y_N)$$

Mapping seq2seq function $f(\cdot)$

1. Connectionist temporal classification (CTC)
2. Attention-based encoder decoder
3. Joint CTC/attention (Joint C/A)
4. RNN transducer (RNN-T)
5. Transformer

Unified view

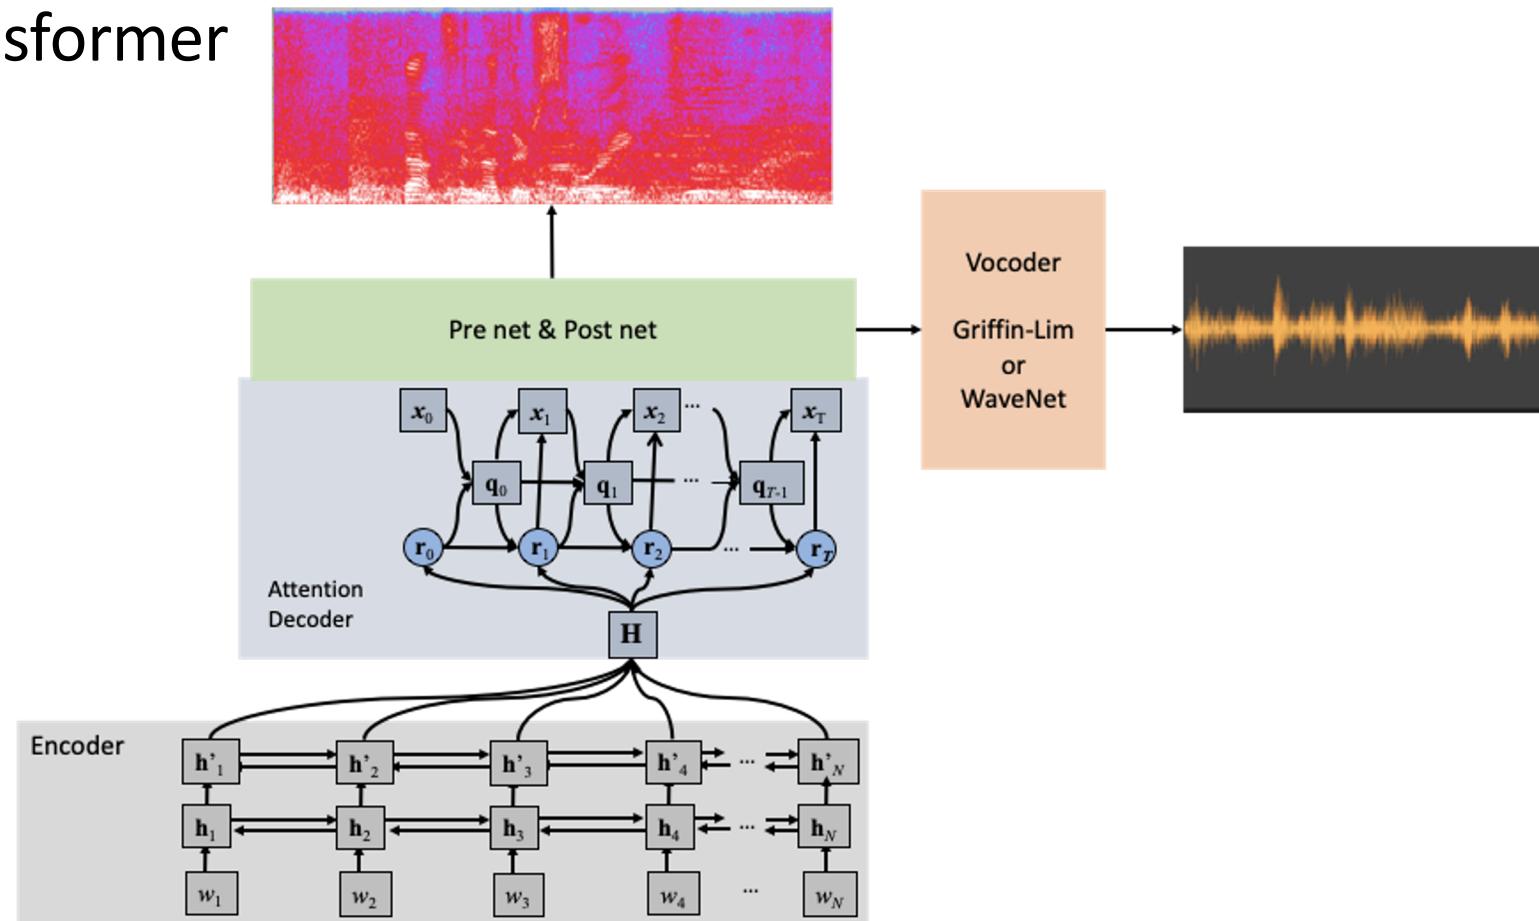
- Target speech processing problems: find a mapping function from *sequence to sequence* (**unification**)

$$X = (x_1, x_2, \dots, x_T) \xrightarrow{f} Y = (y_1, y_2, \dots, y_N)$$

- ASR: X = Speech, Y = Text
- TTS: X = Text, Y = Speech
- ...
- Mapping function (f)
 - Attention based encoder decoder
 - Transformer
 - ...

Seq2seq TTS (e.g., Tacotron2) [Shen+ 2018]

- Use seq2seq generate a spectrogram feature sequence
- We can use either attention-based encoder decoder or transformer



Unified view → Unified software design

We design a new speech processing toolkit based on

$$X = (x_1, x_2, \dots, x_T) \xrightarrow{f} Y = (y_1, y_2, \dots, y_N)$$

Unified view → Unified software design

We design a new speech processing toolkit based on

$$X = (x_1, x_2, \dots, x_T)$$

$$Y = (y_1, y_2, \dots, y_N)$$

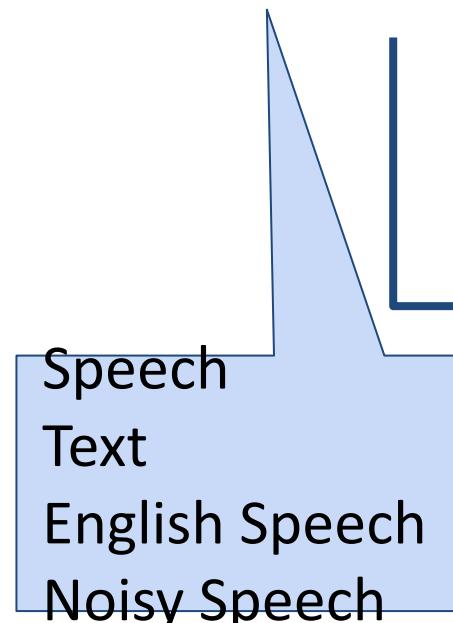


$$f(\cdot)$$

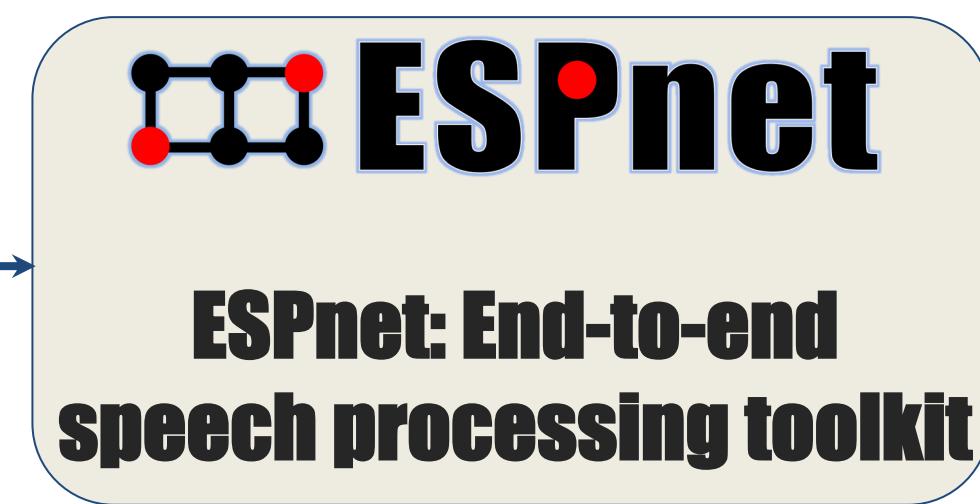
Unified view → Unified software design

We design a new speech processing toolkit based on

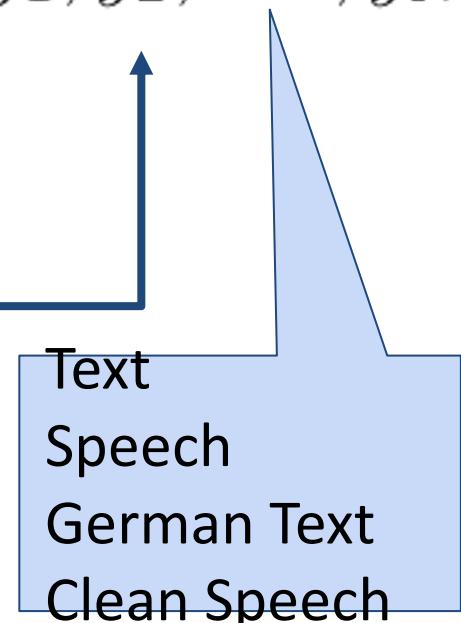
$$X = (x_1, x_2, \dots, x_T)$$



$$Y = (y_1, y_2, \dots, y_N)$$



$$f(\cdot)$$

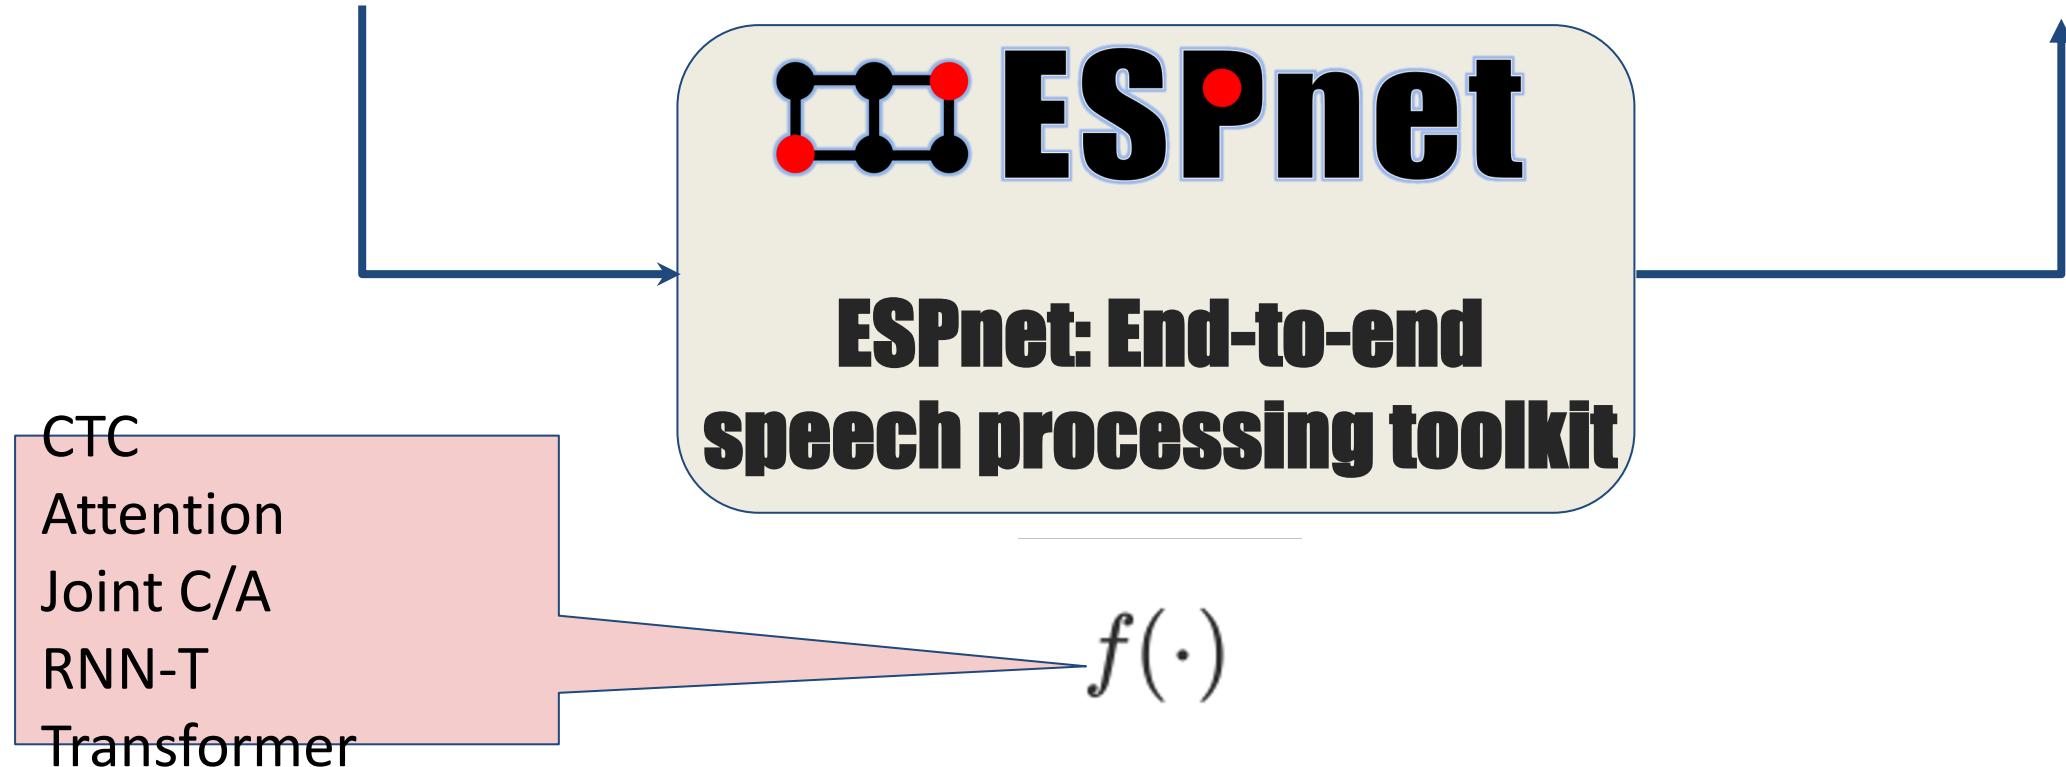


Unified view → Unified software design

We design a new speech processing toolkit based on

$$X = (x_1, x_2, \dots, x_T)$$

$$Y = (y_1, y_2, \dots, y_N)$$

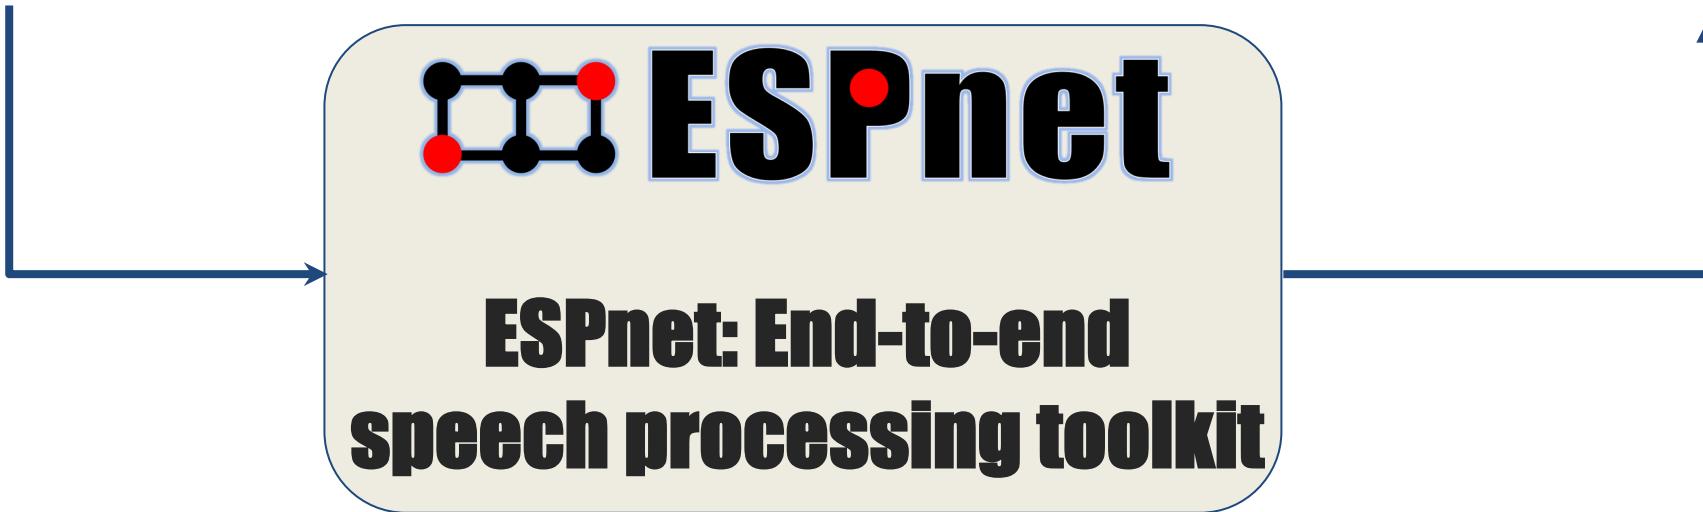


Unified view → Unified software design

We design a new speech processing toolkit based on

$$X = (x_1, x_2, \dots, x_T)$$

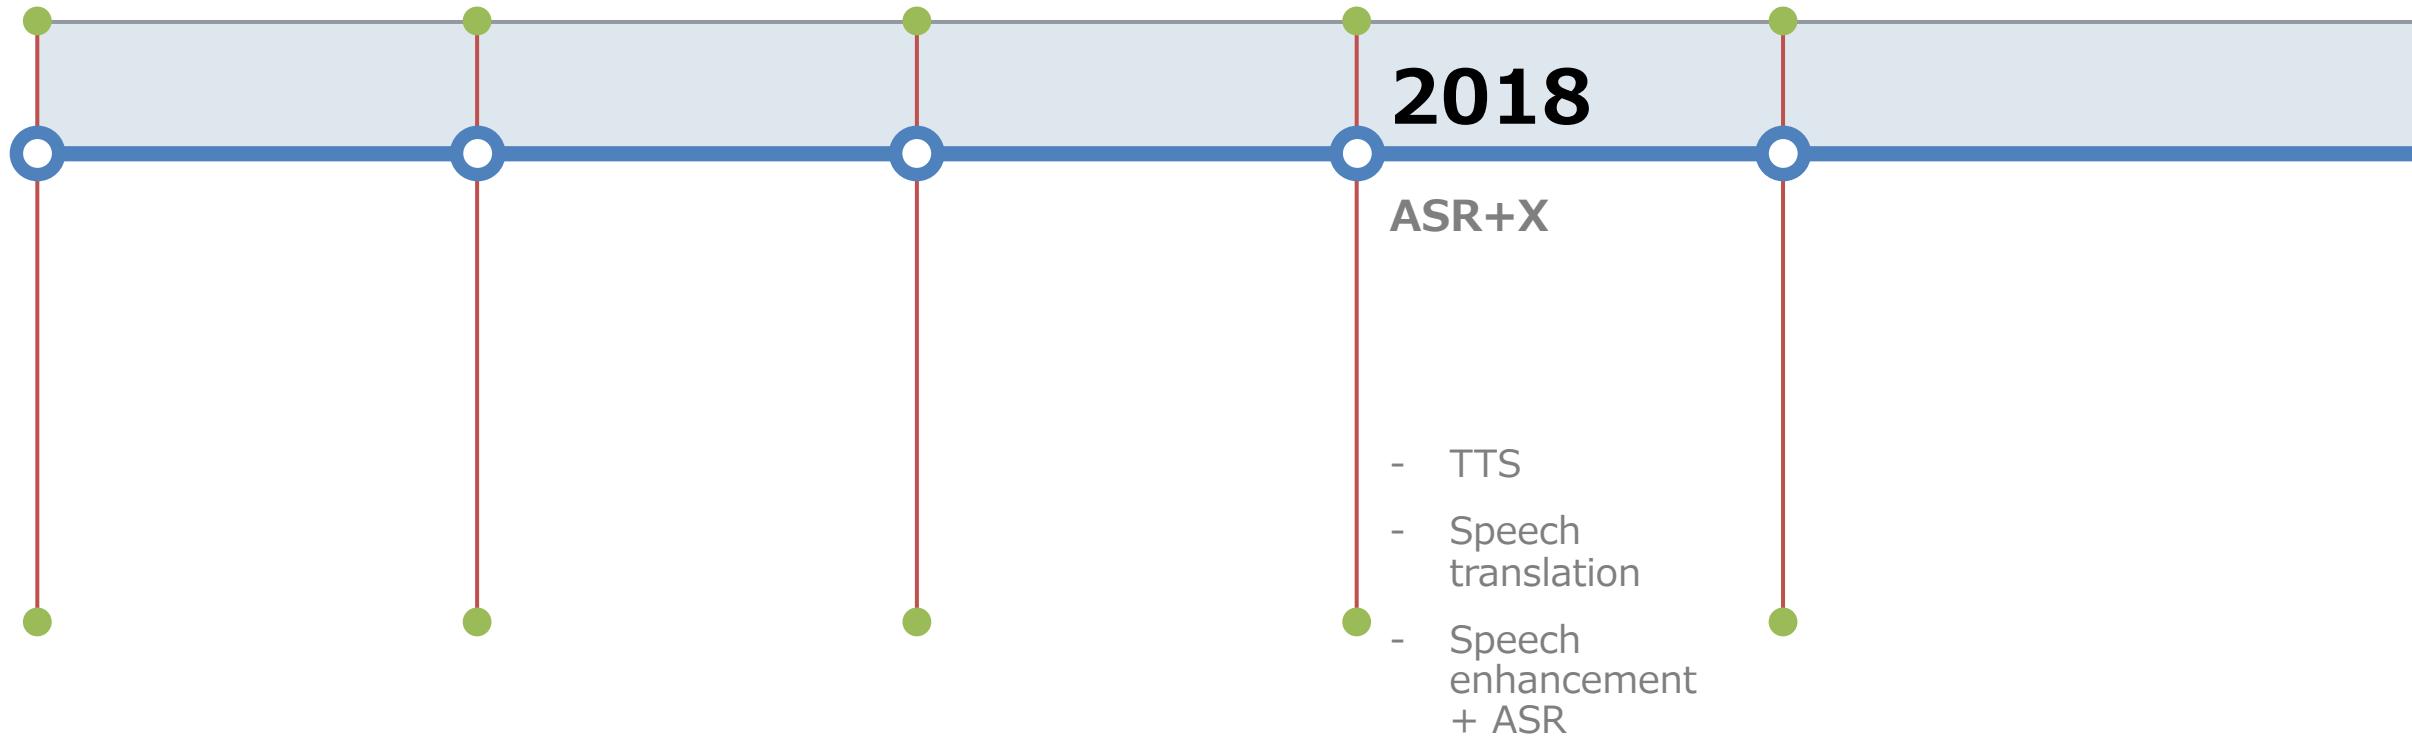
$$Y = (y_1, y_2, \dots, y_N)$$



- Many speech processing applications can be **unified** based on seq2seq
- Again, **Espresso**, **Nemo**, **Fairseq**, **Lingvo** and other toolkits also fully make use of these functions.

Timeline

Shinji's personal experience for end-to-end speech processing



Examples of integrations

Dereverberation + beamforming + ASR

□ Multichannel end-to-end ASR framework

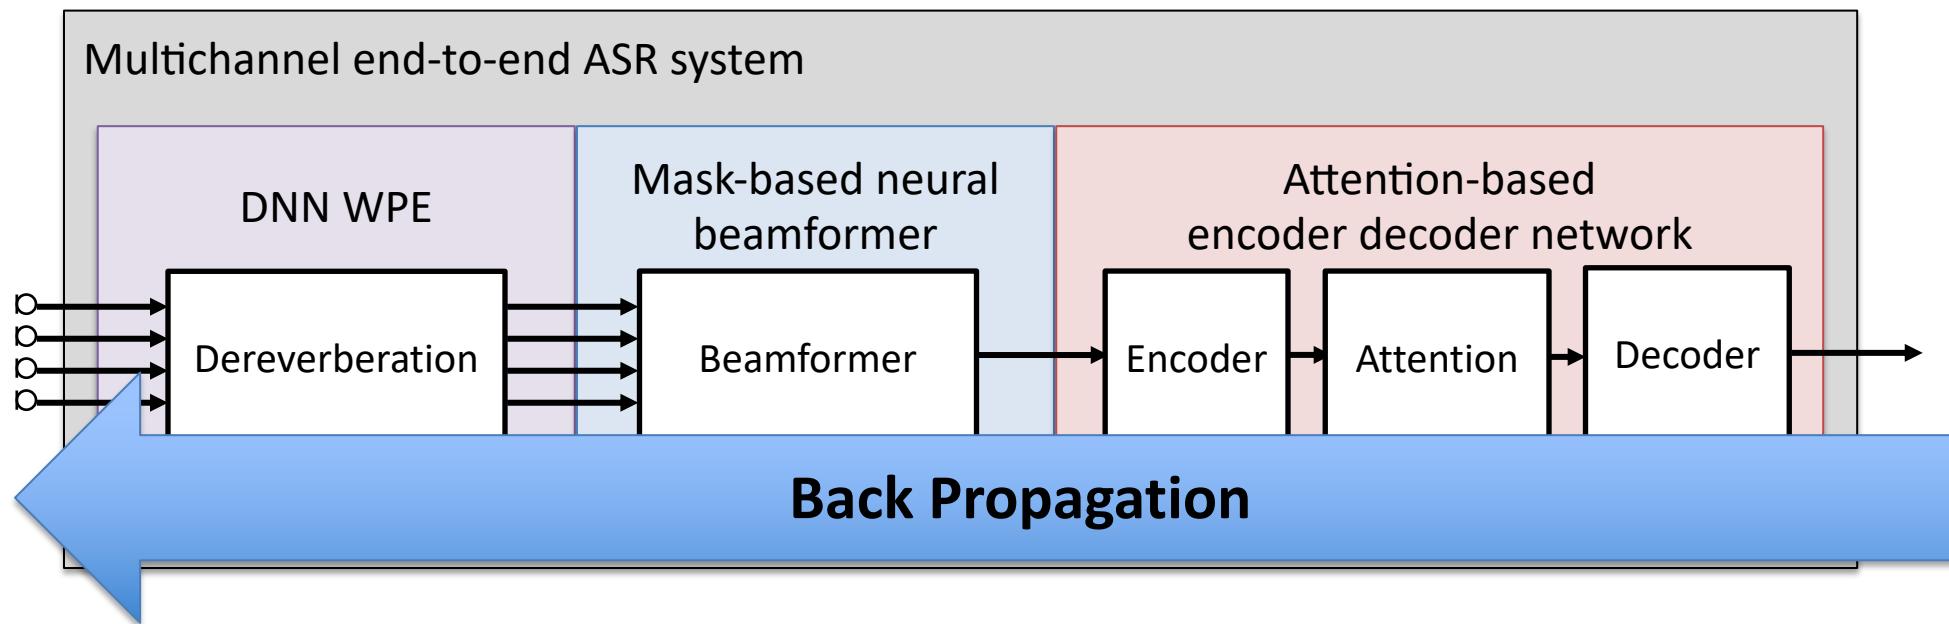
- integrates entire process of **speech dereverberation (SD)**, **beamforming (SB)** and **speech recognition (SR)**, by single neural-network-based architecture



SD : DNN-based weighted prediction error (DNN-WPE) [Kinoshita et al., 2016]

SB : Mask-based neural beamformer [Erdogan et al., 2016]

SR : Attention-based encoder-decoder network [Chorowski et al., 2014]

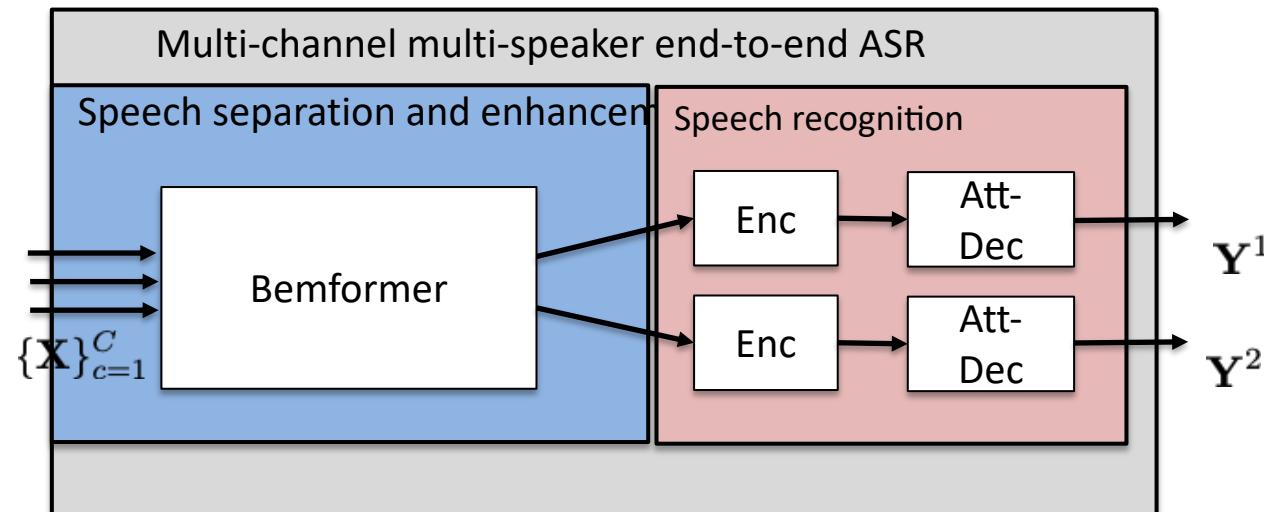


Beamforming + separation + ASR

[Xuankai Chang., 2019, ASRU]

- Multi-channel (MI) multi-speaker (MO) end-to-end architecture
 - Extend our previous model to **multispeaker end-to-end network**
 - Integrate the **beamforming-based speech enhancement and separation networks** inside the neural network

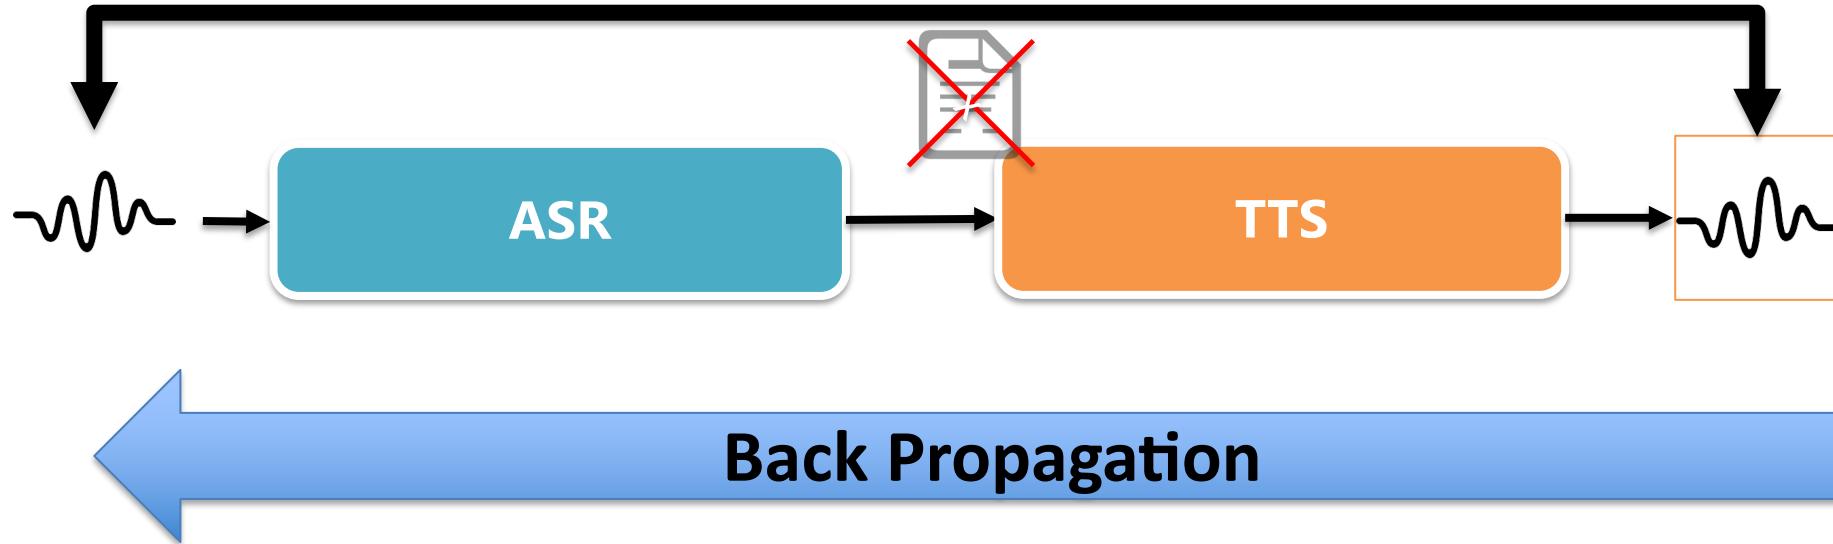
We call it **MIMO speech**



Back Propagation

ASR + TTS feedback loop

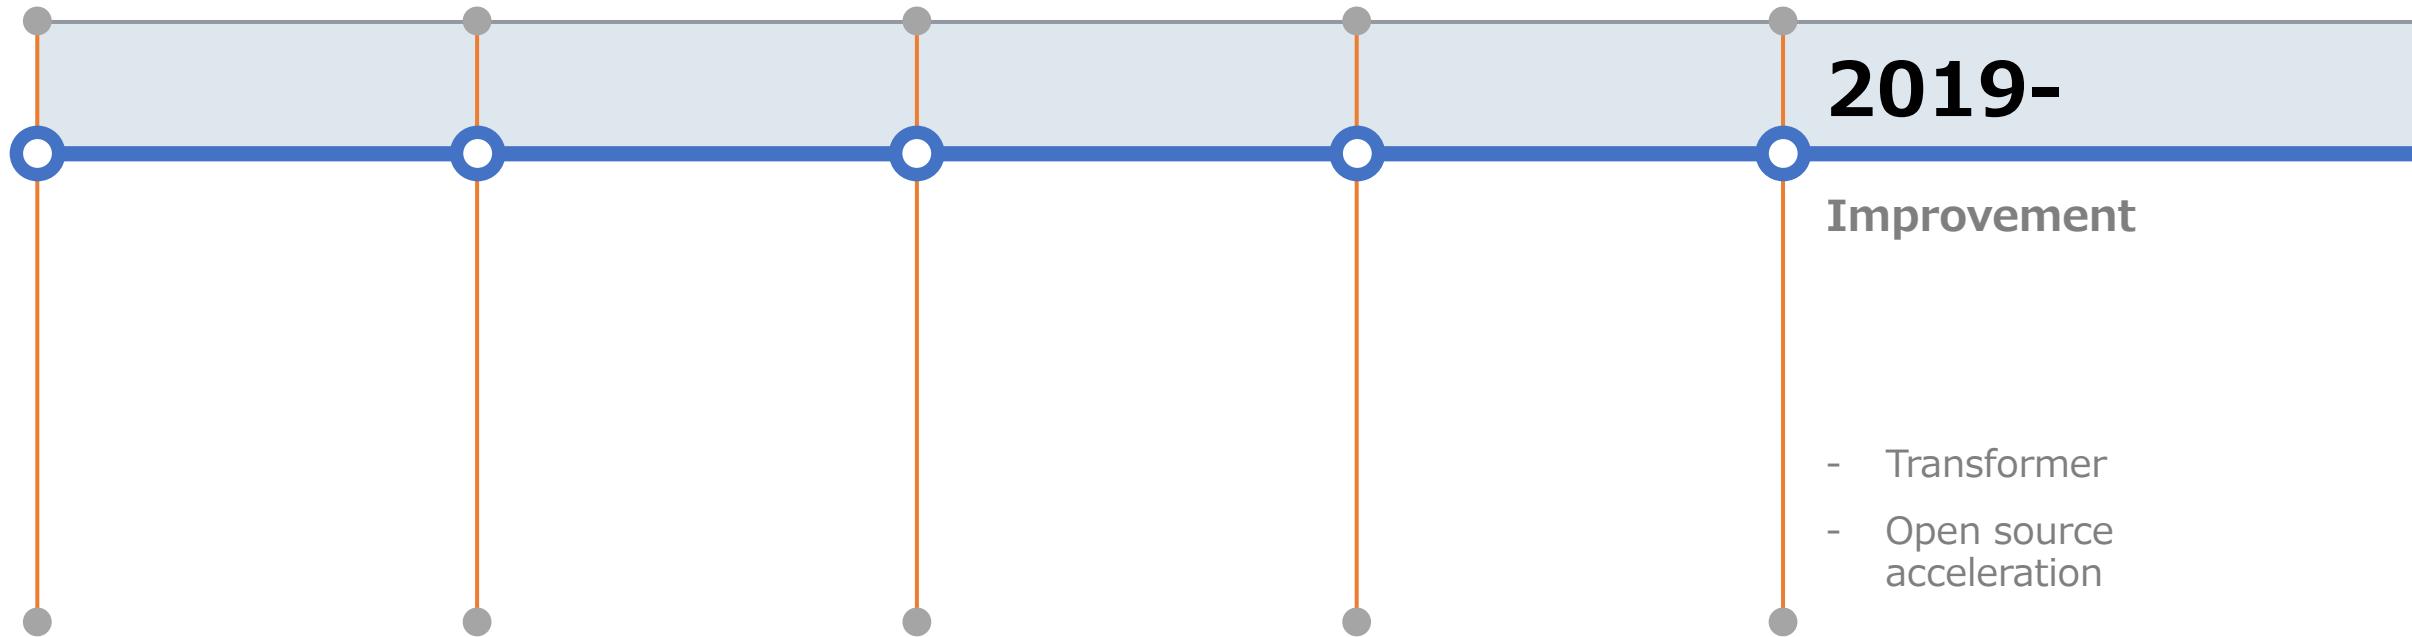
→ Unpaired data training



Only audio data to train both ASR and TTS
We do not need a pair data!!!

Timeline

Shinji's personal experience for end-to-end speech processing



Experiments (~ 1000 hours) Librispeech (Audio book)

Toolkit	dev_clean	dev_other	test_clean	test_other
Facebook wav2letter++	3.1	10.1	3.4	11.2
RWTH RASR	2.9	8.8	3.1	9.8
Nvidia Jasper	2.6	7.6	2.8	7.8
Google SpecAug.	N/A	N/A	2.5	5.8

- Very impressive results by Google

Experiments (~ 1000 hours)

Librispeech

Toolkit	dev_clean	dev_other	test_clean	test_other
Facebook wav2letter++	3.1	10.1	3.4	11.2
RWTH RASR	2.9	8.8	3.1	9.8
Nvidia Jasper	2.6	7.6	2.8	7.8
Google SpecAug.	N/A	N/A	2.5	5.8
ESPnet	2.2	5.6	2.6	5.7

- Reached Google's best performance by community-driven efforts (on September 2019)



GAFAM





Good example of “Collapetition”
= Collaboration + Competition

Experiments (~ 1000 hours)

Librispeech

Toolkit	dev_clean	dev_other	test_clean	test_other
Facebook wav2letter++	3.1	10.1	3.4	11.2
RWTH RASR	2.9	8.8	3.1	9.8
Nvidia Jasper	2.6	7.6	2.8	7.8
Google SpecAug.	N/A	N/A	2.5	5.8
ESPnet	2.2	5.6	2.6	5.7
MS Semantic Mask (ESPnet)	2.1	5.3	2.4	5.4
Facebook wav2letter Transformer	2.1	5.3	2.3	5.6

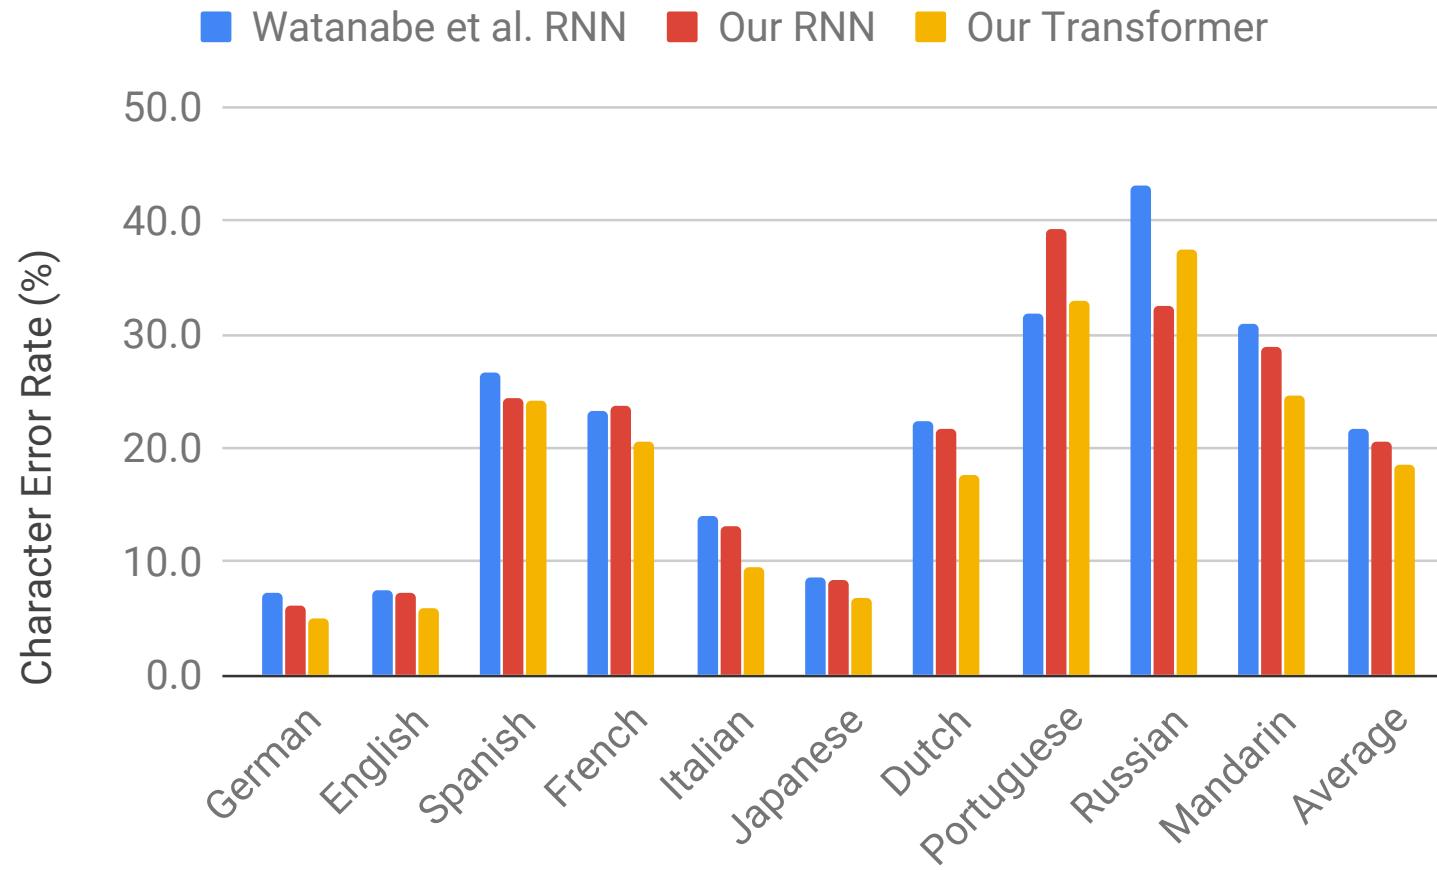
Experiments (~ 1000 hours)

Librispeech

End-to-End

Toolkit	dev_clean	dev_other	test_clean	test_other
Facebook wav2letter++	3.1	10.1	3.4	11.2
RWTH RASR	2.9	8.8	3.1	9.8
Nvidia Jasper	2.6	7.6	2.8	7.8
Google SpecAug.	N/A	N/A	2.5	5.8
ESPnet	2.2	5.6	2.6	5.7
MS Semantic Mask (ESPnet)	2.1	5.3	2.4	5.4
Facebook wav2letter Transformer	2.1	5.3	2.3	5.6
Kaldi (Pipeline)	1.8	5.8	2.2	5.8

Transformer is powerful for multilingual ASR



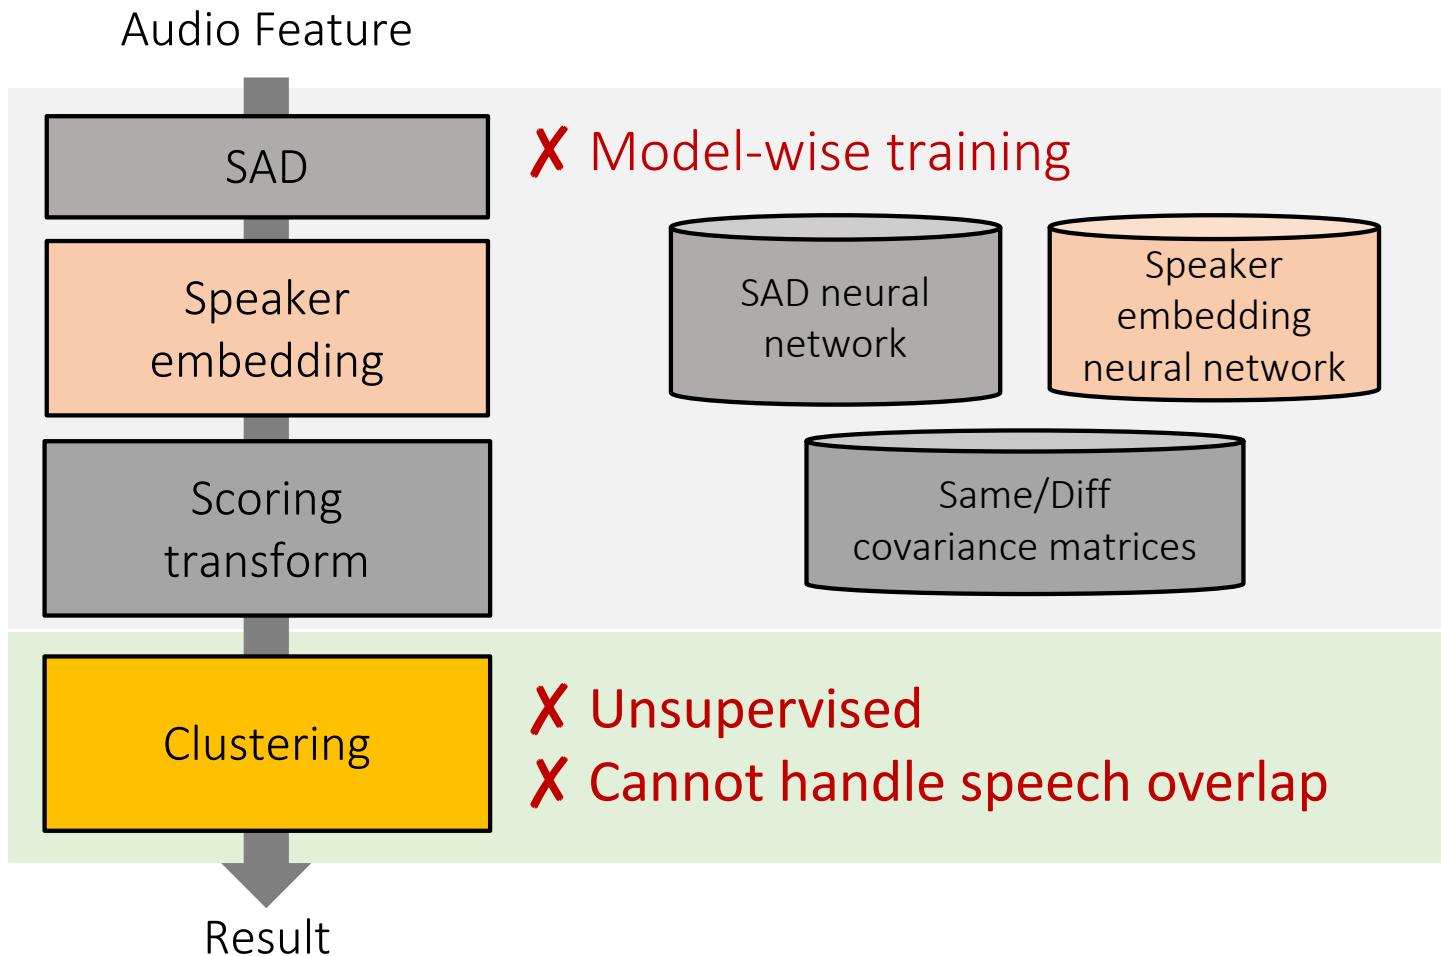
One of the most stable and biggest gains compared with other multilingual ASR techniques



By Philipp Koehn

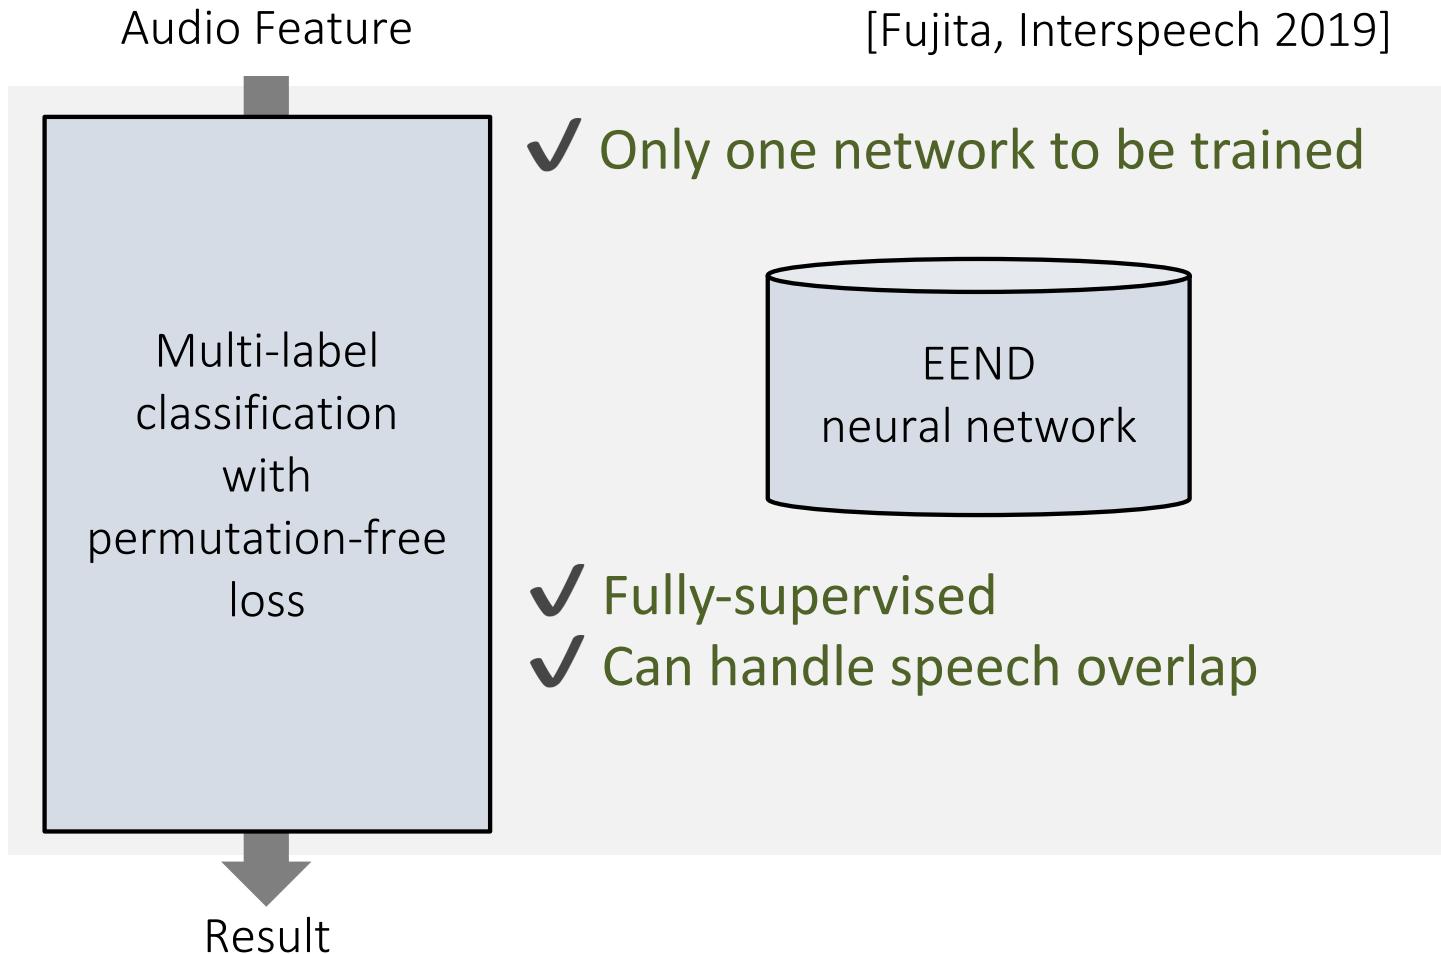
Self-Attentive End-to-End Diarization

[Fujita+(2019)]



Self-Attentive End-to-End Diarization

[Fujita+(2019)]



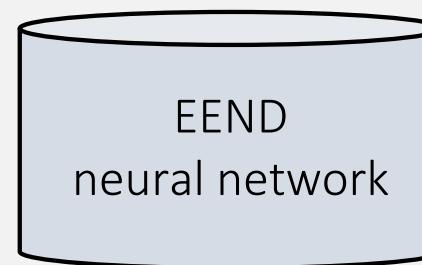
Self-Attentive End-to-End Diarization

[Fujita+(2019)]



[Fujita, Interspeech 2019]

✓ Only one network to be trained



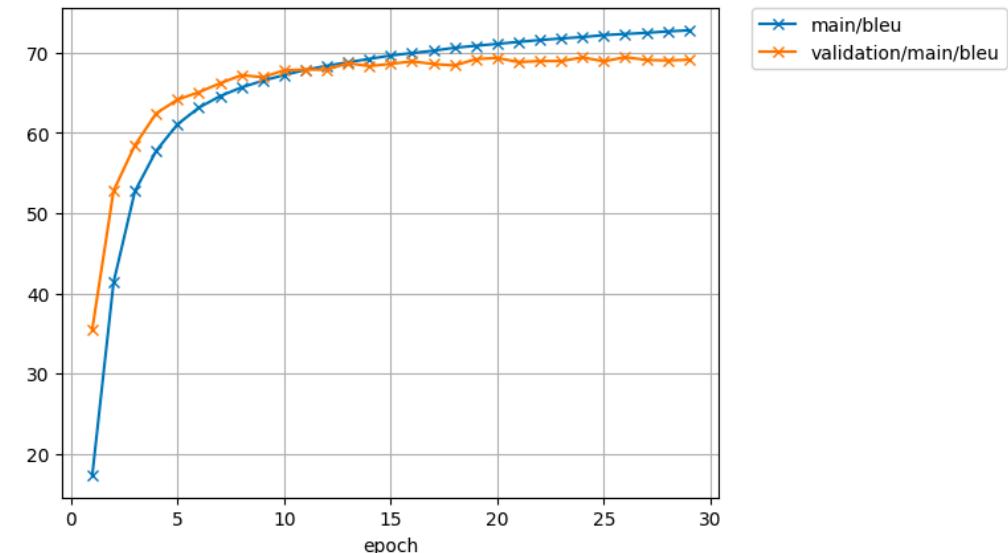
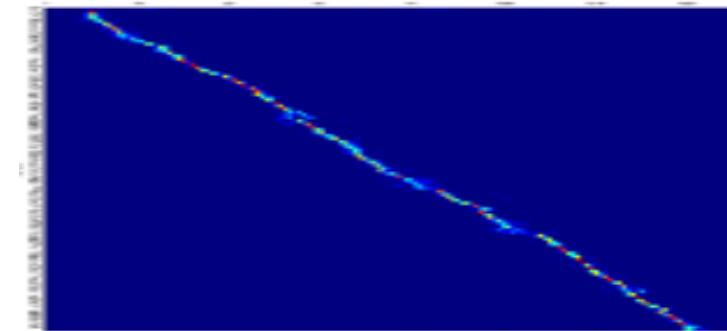
- ✓ Fully-supervised
- ✓ Can handle speech overlap

	CALL HOME DER (%)	CSJ EDR (%)
x-vector	11.53	22.96
EEND BLSTM	23.07	25.37
EEND Self-attention	9.54	20.48

- Outperform the state-of-the-art x-vector system!
- Check <https://github.com/hitachi-speech/EEND>

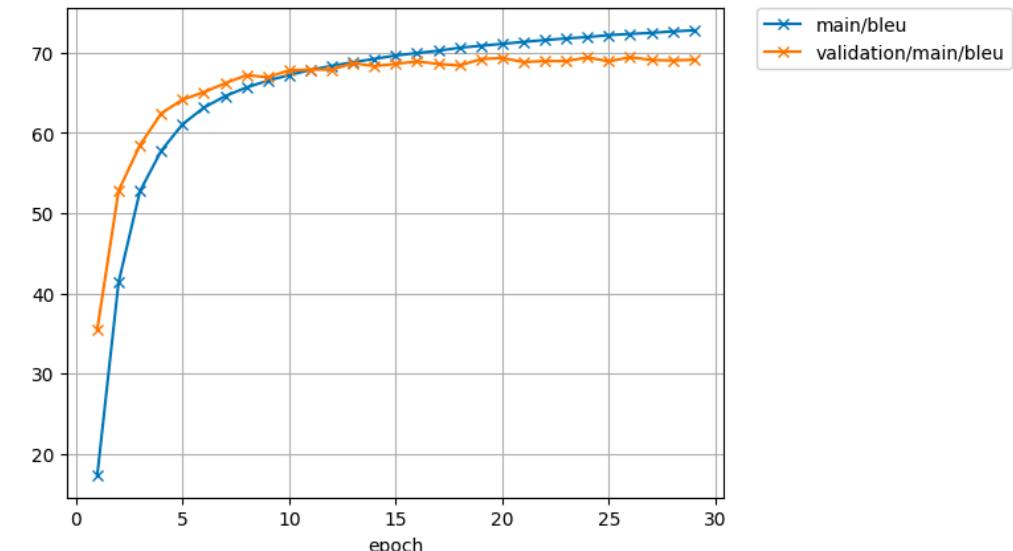
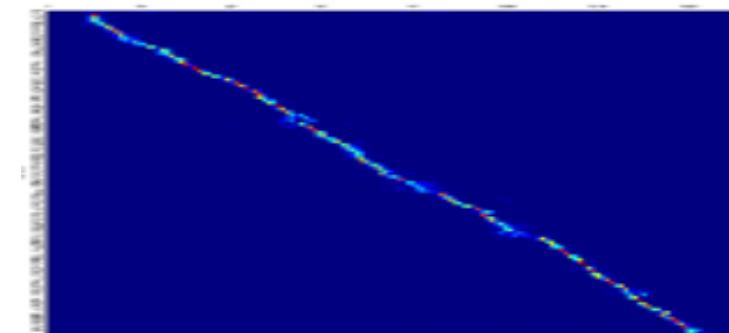
FAQ (before transformer)

- How to debug attention-based encoder/decoder?
- Please check
 - Attention pattern!**
 - Learning curves!**
- It gives you a lot of intuitive information!



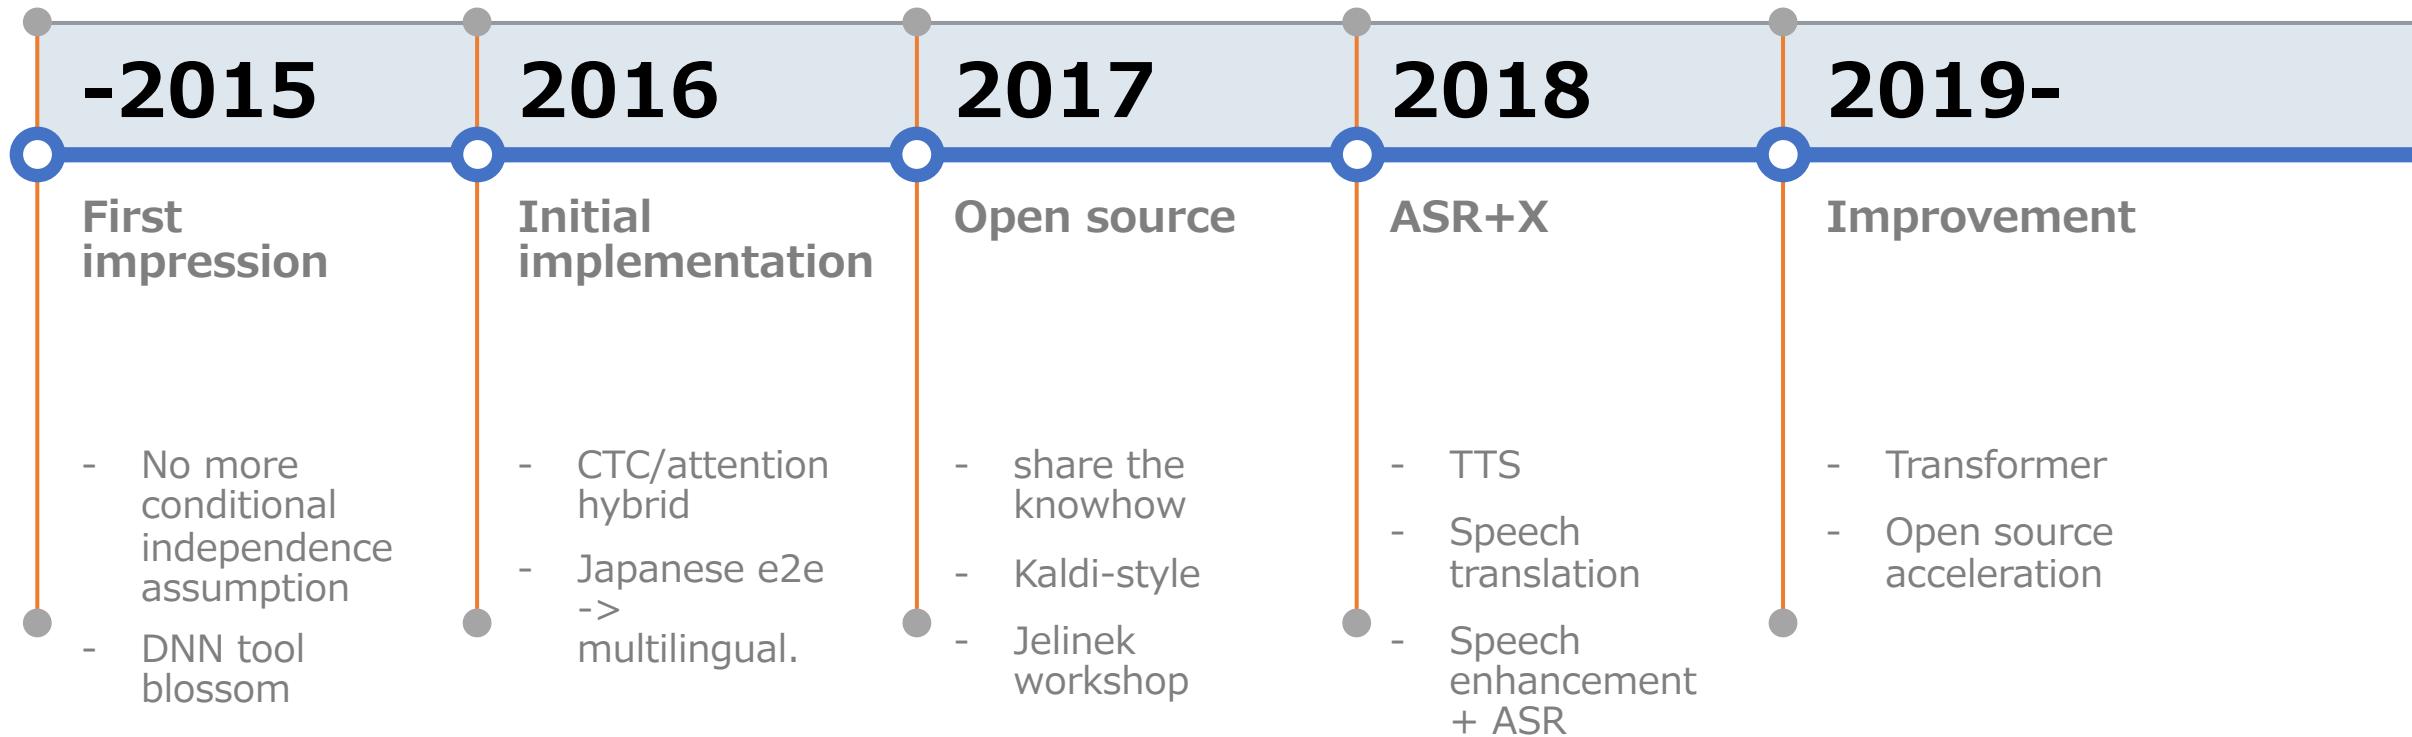
FAQ (after transformer)

- How to debug attention-based encoder/decoder?
- Please check
 - Attention pattern (including self attention)!**
 - Learning curves!**
- It gives you a lot of intuitive information!
- **Tune optimizers!**



Timeline

Shinji's personal experience for end-to-end speech processing



What's next?

- **Non autoregressive ASR**
- **Conditional chain model**
- **New architecture**
 - CNN+SA? Conformer, ???

Thanks!