

Tensor Operations: einsum & tensordot

Powerful Functions for Multi-dimensional Array Manipulation

Introduction and Mathematical Foundation

What are `einsum` and `tensordot`?

- Two powerful functions for tensor operations in NumPy and PyTorch.
- They provide concise ways to perform dot products, outer products, transpositions, matrix math, and more complex tensor contractions.

Mathematical Foundation

Einstein Summation Convention

The Einstein summation convention implies summation over repeated indices. For example, $C_{ik} = A_{ij}B_{jk}$ represents: $C_{ik} = \sum_j A_{ij}B_{jk}$. This convention simplifies tensor notation by removing explicit summation symbols.

Tensor Contraction

Tensor contraction generalizes matrix multiplication to higher-dimensional arrays. It involves element-wise multiplication followed by summation along specific dimensions.

Einsum - Core Concepts

Einstein Summation (**einsum**)

- A versatile function in NumPy and PyTorch using a special string notation.
- **Basic Idea:** `einsum('input_specs -> output_spec', tensor1, tensor2, ...)`
 - Indices repeated in inputs but not in the output are summed out (contracted).
 - Indices appearing in an input but not the output are also summed out.
 - The order of indices in `output_spec` defines the final tensor's axis order.
 - If `-> output_spec` is missing, `einsum` infers it from unsummed input indices, ordered alphabetically.

Syntax:

`numpy.einsum(subscripts, *operands, ...)`

`torch.einsum(equation, *operands)`

Einsum - Vector Operations

einsum for Vector Operations

Let $u=[1,2,3]$ and $v=[4,5,6]$.

- **Element-wise Product (Hadamard Product):** $w_i = u_i \cdot v_i$
 - `einsum('i,i->i', u, v)`
 - Result: `[4, 10, 18]`
- **Dot Product (Inner Product):** $s = \sum_i u_i v_i$
 - `einsum('i,i->', u, v)` or `einsum('i,i', u, v)`
 - Result: 32
- **Outer Product:** $M_{ij} = u_i v_j$
 - `einsum('i,j->ij', u, v)`

Result:

`[[4, 5, 6],`

`[8, 10, 12],`

`[12, 15, 18]]`

-
- **Sum of Elements:** $s = \sum_i u_i$
 - `einsum('i->', u)`
 - Result: 6

einsum - Basic Matrix Operations

einsum for Matrix Operations

Let A be (2×3) and B be (3×2).

- **Matrix Transpose:** $M_{ji} = M_{ij}$
 - `einsum('ij->ji', A_mat)`
- **Matrix-Vector Product:** $y_i = \sum_j M_{ij} u_j$
 - `einsum('ij,j->i', A_mat, u_vec_ops)`
- **Vector-Matrix Product:** $z_j = \sum_i w_i M_{ij}$
 - `einsum('i,ij->j', w_matvec, C_mat)`
- **Matrix Multiplication:** $P_{ik} = \sum_j (M1)_{ij} (M2)_{jk}$
 - `einsum('ij,jk->ik', A_mat, B_mat)` or `einsum('ij,jk', A_mat, B_mat)`

einsum - Advanced Matrix Operations

Advanced Matrix Operations with **einsum**

- **Matrix Multiplication with Reduction:**
 - Sum columns of P: `einsum('ij,jk->i', M1, M2)`
 - Sum rows of P: `einsum('ij,jk->k', M1, M2)`
 - Sum all elements of P: `einsum('ij,jk->', M1, M2)`
- **Trace of a Matrix:** $\text{tr}(M) = \sum_i M_{ii}$
 - `einsum('ii->', C_mat)`
- **Diagonal Extraction:** $\text{di} = M_{ii}$
 - `einsum('ii->i', C_mat)`
- **Frobenius Norm (Squared):** $\|A\|_F^2 = \sum_i \sum_j A_{ij}^2$
 - `einsum('ij,ij->', A_mat, A_mat)`

einsum - Axis Operations & Hadamard

Axis Operations & Hadamard Product with `einsum`

- **Sum along an Axis:**
 - Sum over columns (for each row): `einsum('ij->i', A_mat)`
 - Sum over rows (for each column): `einsum('ij->j', A_mat)`
- **Element-wise Matrix Product (Hadamard):** $P_{ij} = (M1)_{ij}(M2)_{ij}$
 - `einsum('ij,ij->ij', C_mat, D_mat)`

einsum - Chained & Tensor Operations

Chained Operations and Higher-Rank Tensors with `einsum`

- **Chained Operations:** `einsum` can combine multiple tensor operations efficiently.
 - Example: $(M1M2)w$: `einsum('ij,jk,k->i', M1, M2, w)`
 - Example: $M1M2M3$: `einsum('ij,jk,kl->il', M1, M2, M3)`
- **Tensor Operations (Rank 3+):** `einsum` excels here.
 - **Ellipsis (...) Operator:** Placeholder for leading batch dimensions.
 - Example: Batch matrix multiplication `einsum('...ij,...jk->...ik', T_A, T_B)` for `T_A` and `T_B` with shared batch dimensions.
 - **Permutation of Axes:** Reorders tensor dimensions.
 - Example: $T_{ijk} \rightarrow T_{kij}$: `einsum('ijk->kij', Tensor_permute)`

tensor dot - Introduction & Examples

tensor dot

- Performs tensor contractions by summing over specified axes.
- Available in NumPy and PyTorch.

Syntax:

`numpy.tensordot(a, b, axes)`

`torch.tensordot(a, b, dims)`

axes / dims Parameter is Key:

- **Integer N:** Sums over the last **N** axes of tensor **a** and the first **N** axes of **b**.
 - Example: `torch.tensordot(A_ex3, B_ex3, dims=2)` where **A_ex3** ends with matching dimensions to **B_ex3**'s start (e.g., `A(..., 4, 5)`, `B(4, 5, ...)`). Result shape: `(2, 3, 6, 7)`.
- **Tuple of two lists ([a_axes_list], [b_axes_list]):**
 - Sums over specified axes: `a_axes_list[k]` is contracted with `b_axes_list[k]`.
 - Resulting tensor has **a**'s remaining axes, then **b**'s remaining axes, in that order.
 - Example: `torch.tensordot(A_td_torch, B_td_torch, dims=([0,3], [0,2]))` (contract `A[0]` with `B[0]`, `A[3]` with `B[2]`). Result shape: `(3, 4, 3, 6)`.

Practical Applications

1. Batch Processing in Neural Networks (Convolution-like Operation)

`einsum` can express parts of convolutions.

Simulate batch of images and a convolutional kernel

Input patch: bchw (batch, channels_in, kernel_h, kernel_w)

Kernel: ochw (kernel_out_channels, channels_in, kernel_h, kernel_w)

`conv_like_op = np.einsum('bcij,ocij->boij', image_patch, kernel_conv)`

Output: boh (batch, kernel_out_channels, kernel_h, kernel_w)

Practical Applications - Cont.

2. Attention Mechanisms in Transformers

`einsum` is crucial for attention scores and weighted sums.

Compute attention scores: $Q @ K^T$

Q: bqd, K: bkd -> Scores: bqk

`attention_scores = np.einsum('bqd,bkd->bqk', Q_att, K_att)`

Apply attention to values: weights @ V

weights: bqk, V: bkd -> Output: bqd

`attention_output = np.einsum('bqk,bkd->bqd', attention_weights_simplified, V_att)`

Practical Applications - Cont.

3. Statistical Operations (Covariance Matrix)

`einsum` simplifies statistical computations.

Covariance matrix: $(1/(N-1)) * X_centered^T @ X_centered$

Data_centered: sn (samples, features)

`cov_matrix_einsum = np.einsum('sn,sm->nm', data_centered, data_centered) / (num_samples - 1)`

Output: nm (features, features)

Performance & Best Practices

Performance Tips and Best Practices

1. Memory Layout and Performance Considerations

For very large arrays, memory layout (Fortran vs. C order) can influence performance. `einsum` operations are generally efficient, but specific performance can vary.

2. `optimize=True` Parameter in `numpy.einsum`

NumPy's `einsum` has an `optimize` parameter to significantly improve performance for complex operations (three or more operands) by finding an optimal contraction order.

```
# Complex multi-tensor operation: 'ijk,jkl,klm->ilm'
```

```
# Use optimize=True for potential speedups:
```

```
np.einsum('ijk,jkl,klm->ilm', A_complex_perf, B_complex_perf, C_complex_perf, optimize=True)
```

Performance & Best Practices - Cont.

3. Memory Efficiency

Be mindful of memory implications for intermediate arrays, especially with large tensors. `einsum` strives for efficiency.

Common Patterns

`einsum` excels with clear, descriptive index strings.

- **Matrix transpose:** `'ij->ji'`
- **Batch matrix multiplication:** `'bij,bjk->bik'`
- **Vector dot product:** `'i,i->'`
- **Matrix trace (sum of diagonal):** `'ii->'`
- **Sum all elements (tensor):** `'...->'`
- **Quadratic form (vector-matrix-vector):** `'i,ij,j->'`

Conclusion

Key Takeaways:

1. **Concise Notation:** `einsum` offers a powerful and concise domain-specific language for tensor operations using Einstein summation convention.
2. **Versatility:** Both functions perform a wide array of operations including dot products, outer products, transpositions, permutations, and complex tensor contractions.
3. **Performance:** `einsum` can be highly performant, especially with the `optimize` flag for complex contractions. `tensordot` is also optimized for its specific tasks.
4. **Readability:** While initially cryptic, `einsum` strings become very readable for those familiar with the notation, clearly expressing tensor manipulation intent.
5. **Framework Support:** Both are available in NumPy for CPU-bound tasks and PyTorch (often with GPU acceleration) for deep learning.

When to Use Each:

- **`einsum`:** For flexible control over indices (summing, keeping, reordering), complex multi-tensor contractions, or operations hard to express concisely with standard routines (e.g., attention mechanisms).
- **`tensordot`:** When contracting specific pairs of axes between two tensors. It's a good general-purpose tensor dot product when the `axes` argument clearly defines the contraction.
- **Standard operators (`@`, `np.dot`, `np.multiply`, etc.):** For very common and simple operations, as they are often the most readable and highly optimized.

Best Practices:

1. **Clarity:** Add comments to complex `einsum` strings.
2. **Optimization:** For `numpy.einsum` with three or more tensors, use `optimize=True`.
3. **Profiling:** For critical sections, profile `einsum` against alternatives.
4. **Readability vs. Conciseness:** Balance `einsum`'s conciseness with the readability of standard operations.
5. **Start Simple:** Begin by reformulating simple operations with `einsum` and gradually move to complex ones.