



Carnegie Mellon University

## Recitation 0.24: Saving and Loading Models






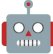
Introduction to Deep Learning  
11-785 / 685 / 485

Mengchun Zhang, Nishoak Kosaraju




---

# Today's Agenda

2

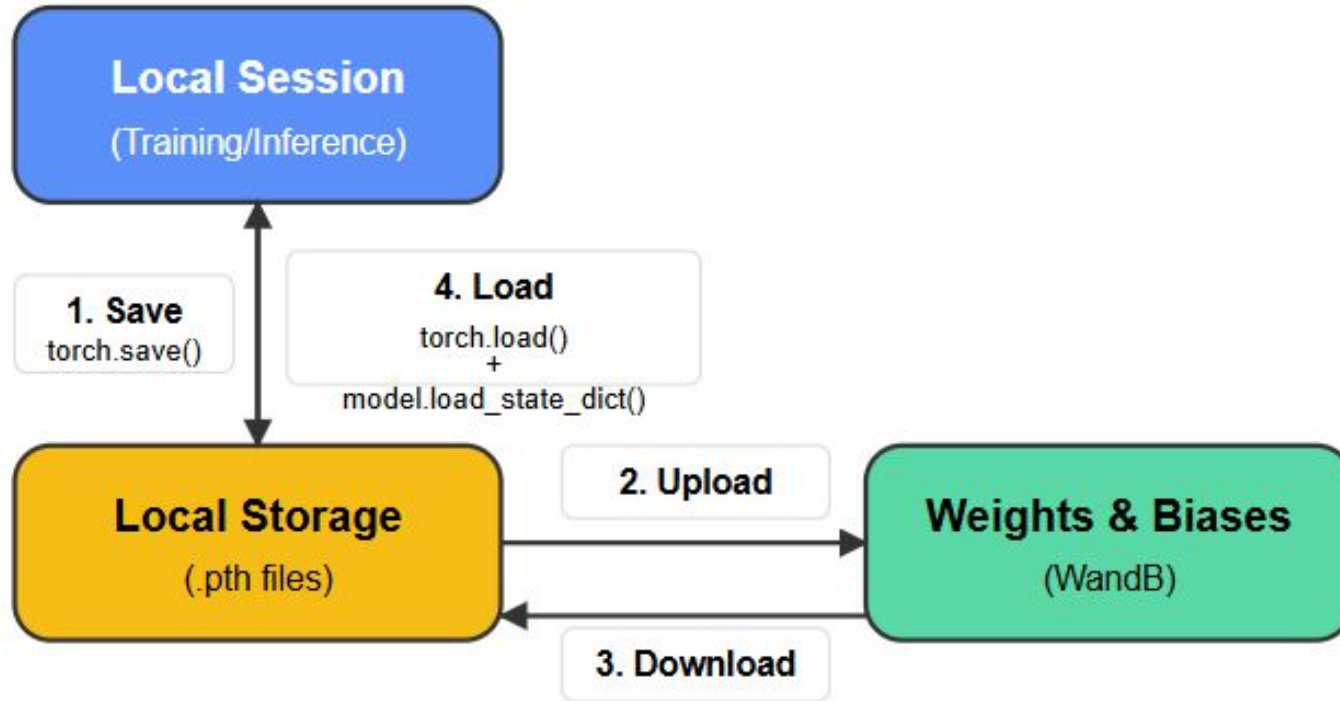
-  What is Model Checkpointing?
-  Flow of a Checkpoint (Save → Store → Load)
-  Saving Locally with `torch.save()`
-  Retrieving from Weights & Biases
-  Loading Locally & Restoring Parameters
-  Dummy Model Example

# What is Checkpointing?

- Training models **takes time...**
- Checkpointing = saving a snapshot of your full training state at a given epoch:
  - Model weights
  - Optimizer & scheduler states
  - Current epoch & metrics
- Why Checkpoint?
  -  Resume after **crashes**
  -  **Rollback** to your **best performing epoch**
  -  Branch experiments from any saved point (try new optimizers/schedulers)

# Flow of Model Checkpointing

4



# How to create a Checkpoint

# How to Save a model (Locally)


Use torch.save  
(Uses Python's  
Pickling to serialize  
the data)

```
checkpoint_path=f"<run_name>_{epoch}.pth"

# Saving your states locally with torch.save
torch.save({
    'model_state_dict': model.state_dict(), # saving the model state
    # if isinstance(model, nn.DataParallel) 'model_state_dict': model.module.state_dict()
    'optimizer_state_dict': optimizer.state_dict(), # saving the optimizer state
    'scheduler_state_dict': scheduler.state_dict(), # saving the scheduler state
    'epoch': epoch,
    'current_loss': loss
}, checkpoint_path
)
```


Saves a dictionary of model  
states and other information  
Note: If using nn.DataParallel for  
training, use  
model.module.state\_dict()

Saves the dictionary  
into a .pth file



# How to Save a Model (to Wandb)

Make sure to begin a run before your training loop



```
Before the run, you need to have started a run like so...
run = wandb.init(
    project="wandb-quickstart",
    name="<run_name>",
)

# ...
# ...
# ...
# Within a training loop (or wherever else you want)...

# Option 1:
# create artifacts (keeps track of versioning, and is much more organized to work with between collaborators)
checkpoint_artifact = wandb.Artifact("<run_name>", type="checkpoint") # You can switch type="model if you only want to save a model"

checkpoint_artifact.add_file(checkpoint_path)

run.log_artifact(checkpoint_artifact)

# Option 2:
# directly save the model to wandb
wandb.save(checkpoint_path, base_path=os.path.dirname(checkpoint_path))
```

Create the artifact and add the checkpoint file

Upload the artifact to wandb


Option 2: directly save the checkpoint file

# How to load a Checkpoint



# How to retrieve a checkpoint from Wandb

Obtain the run  
object if needed



```
# METHOD 1: Download from wandb Artifact
# If you need to re-obtain the run, you can do the following...
api = wandb.Api()
# information can be obtained from the wandb link address as follows:
# https://wandb.ai/<USERNAME>/<PROJECT_NAME>/runs/<RUN_ID>?nw=nwuser<USERNAME>
run = api.run("<USERNAME>/<PROJECT_NAME>/<RUN_ID>")
```

```
# To retrieve the artifact....
# Get the artifact (choose which version of the model you want)
artifact = run.use_artifact('<run_name>:latest')
# Downloading the artifact
artifact_dir = artifact.download()
# Loading the model dict
checkpoint_dict = torch.load(os.path.join(artifact_dir, '<run_name>'))
```

```
# METHOD 2: Download the directly saved file from wandb to Local File
checkpoint_file = wandb.restore('<run_name>', run_path="<USERNAME>/<PROJECT_NAME>/<RUN_ID>").name
checkpoint_dict = torch.load(checkpoint_file)
```

Retrieve and  
download the  
artifact


Load the  
checkpoint with  
torch.load

Option 2: restore  
the file, then load  
the file

## How to organize a model And load specific parts


# How to load a model (locally) and load the parameters

Use torch.load on the locally-saved file if you want



```
# .pth checkpoint file path can also be obtained from a locally saved .pth file.
checkpoint_path = "<checkpoint_file_path>"
checkpoint_dict = torch.load(checkpoint_path)

# Loading model weights
model.load_state_dict(checkpoint_dict['model_state_dict'])
# Loading optimizer state
optimizer.load_state_dict(checkpoint_dict['optimizer_state_dict'])
# Loading the scheduler state
scheduler.load_state_dict(checkpoint_dict['scheduler_state_dict'])
```



Load the checkpoint's model state dict into our model.

Load optimizer and scheduler state dicts

**IMPORTANT!!!!!! Your model must have the same parameter names and dimensions!!!!**

**Your optimizer/scheduler must be the same type of optimizer and scheduler as checkpoint's if you load them**

**Pytorch will simply match the keys from the checkpoint (i.e. layer/parameter names) to what it can find in your model**

# Our Dummy Model

Organize in  
submodules

```
# A simple submodule
class DummySubmodule(nn.Module):
    def __init__(self):
        super(DummySubmodule, self).__init__()
        self.layer = nn.Linear(in_features = 32, out_features = 32)

    def forward(self, x):
        return self.layer(x)

# A simple network
class DummyNetwork(nn.Module):
    def __init__(self):
        super(DummyNetwork, self).__init__()

        self.lower_layer = nn.Sequential(
            nn.Linear(in_features = 32, out_features = 64),
            nn.ReLU(),
            nn.Linear(in_features = 32, out_features = 64),
            nn.ReLU(),
        )
        self.upper_layer = nn.Sequential(
            nn.Linear(in_features = 64, out_features = 32),
            nn.ReLU(),
        )

        self.module1 = DummySubmodule()

    def forward(self, x):
        res = self.lower_layer(x)
        res = self.upper_layer(res)
        res = self.submodule(res)
        return res
```

```
# Print model's state_dict
print("Model's state_dict:")
for param_tensor in model.state_dict():
    print(param_tensor, "\t", model.state_dict()[param_tensor].size())
```

Model's state_dict:	
lower_layer.0.weight	torch.Size([64, 32])
lower_layer.0.bias	torch.Size([64])
lower_layer.2.weight	torch.Size([64, 32])
lower_layer.2.bias	torch.Size([64])
upper_layer.0.weight	torch.Size([32, 64])
upper_layer.0.bias	torch.Size([32])
module1.layer.weight	torch.Size([32, 32])
module1.layer.bias	torch.Size([32])


Keys

Values

Organize by class  
names and  
Sequential

# Loading the specific weights

Filter out the weights you want into a dictionary



```
# If you want to load specific parts of your model (in our case, we can load just the lower layers or just the upper layers)
specific_weights = { # Creates dictionary of only desired weights
    key: value
    for key, value in checkpoint_dict['model_state_dict'].items()
    if 'lower_layer' in key
}

model.load_state_dict(specific_weights, strict=False)
```



Load your weights like normal

**IMPORTANT!!!!!!** For the parameters you load, your model needs to have the same name and dimensions.