

Graph Neural Networks

John Shi and Shreyas Chaudhari
Intro to Deep Learning - Lecture
April 21, 2025

Our Research Team

Undergraduate Interns:

Yao (Lavender) Jiang
Xujin (Chris) Liu
Wendy Summer
Austin Lin
Farida Abdelmoneum

Masters/Ph.D. Students:

Mark Cheung
Oren Wright (SEI)
Mayur Gowda

Postdoc:

Jian Du

Advisor:

José Moura (ECE)

**Graph Signal Processing and Deep Learning:
Convolution, Pooling, and Topology:**
<https://arxiv.org/abs/2008.01247>

**18-898D: Special Topics in Signal Processing:
Graph Signal Processing and Learning**

Our Lecture Last Year



Carnegie Mellon

Graph Convolutional Layers

1. Aggregate node features of neighbors
2. Combine current node feature with aggregation from prior step

$$\mathbf{a}_i^{(\ell)} = AGG^{(\ell)} \left(\{\mathbf{x}_j^{(\ell-1)}, j \in \mathcal{N}(i)\} \right)$$

$$\mathbf{x}_i^{(\ell)} = COMB^{(\ell)} \left(\mathbf{x}_i^{(\ell-1)}, \mathbf{a}_i^{(\ell)} \right)$$

33

Claude Kwizera

Paul Magalotti

Abhishek



Day after ICASSP 2024...

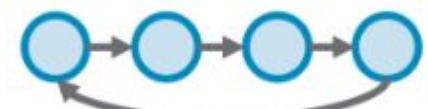
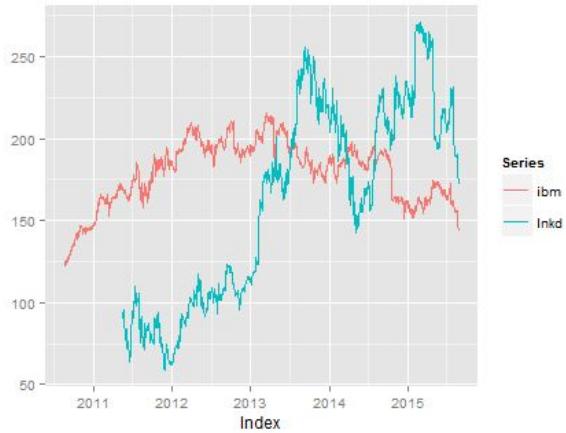
Got back from Korea Sunday night and got up Monday morning for lecture...

Luckily this year, we had a week to recover :P
(Though we were offered the day after ICASSP again...)

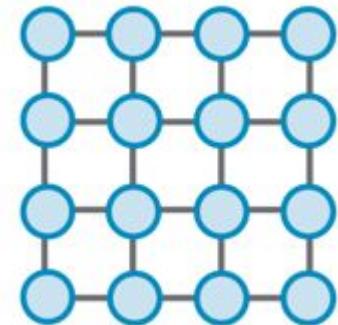
Grid-Structured Data

1-D and 2-D data

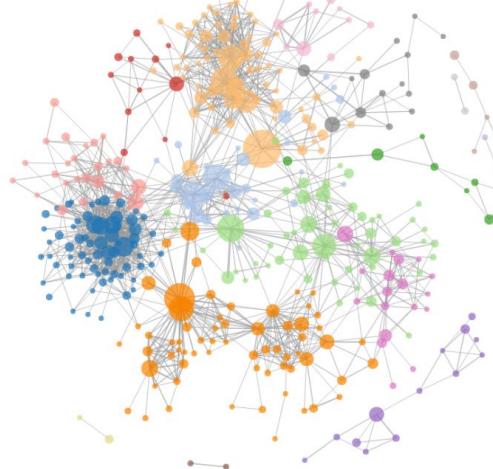
Time Series



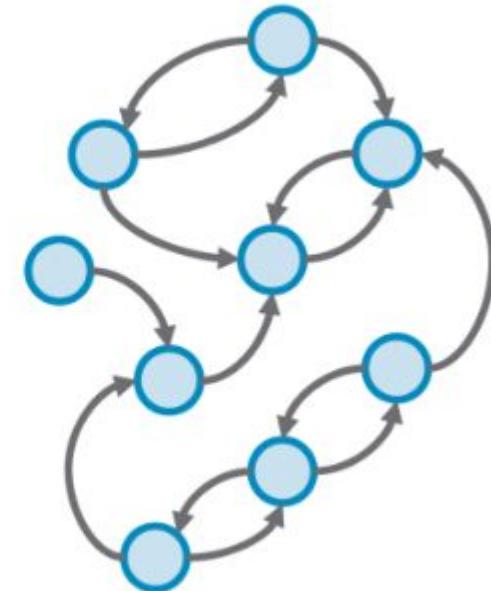
Images



Data defined on an Irregular Graph



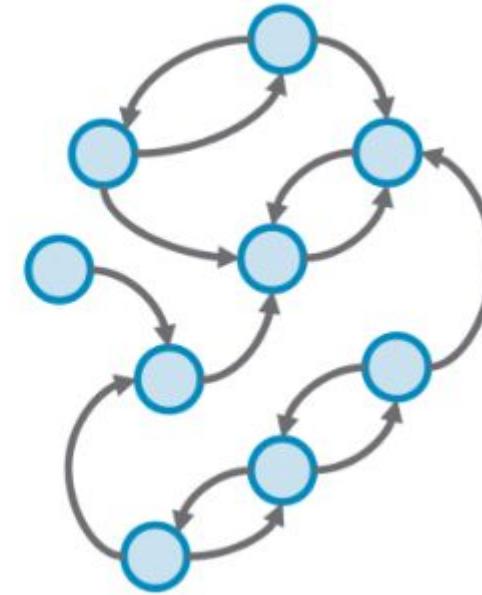
Social Networks
Sensor Feeds
Web Traffic
Supply Chains
Biological Systems
...



Challenges in Graph Data



Social Networks
Sensor Feeds
Web Traffic
Supply Chains
Biological Systems
...



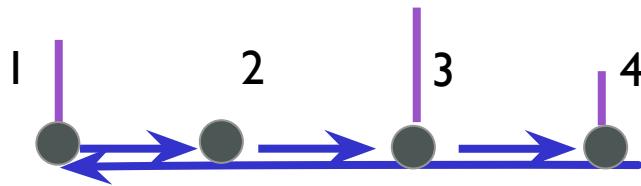
Has a Fixed, Intuitive Ordering

No Ordering (Which node is node 1 here?)

How do we incorporate the graph information?

Challenges in Graph Data

Discrete Signal Processing (DSP) works well on time and images



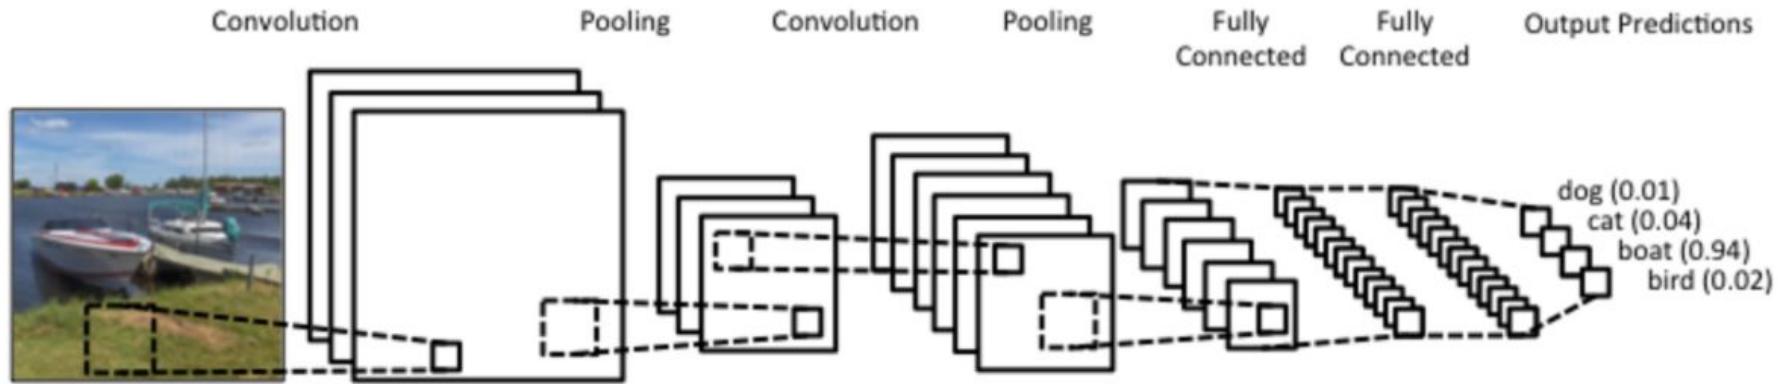
Time Series

But fail for graph data

- No fixed ordering
- No way to incorporate graph data.

Challenges in Graph Data

Classical CNNs perform well on Grid-Structured (Euclidean) Data



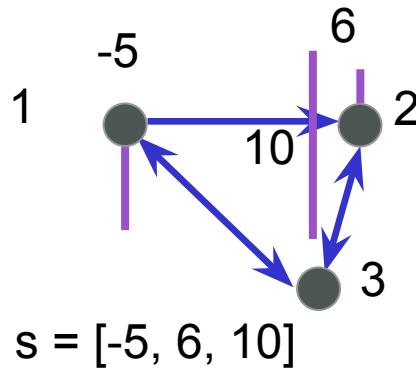
But fail for graph data

- No fixed ordering
- No way to incorporate graph data.

Challenges in Graph Data - Ordering

Ordering – Operations need to be *permutation invariant*

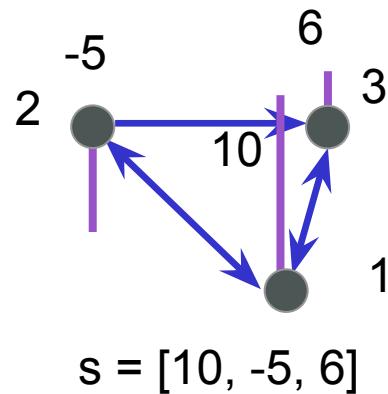
In other words, we want the operation to produce the same (or permuted) numerical results *regardless of ordering*



$$\text{Sum: } -5 + 6 + 10 = 11$$

$$\text{First Difference: } -5 - 6 - 10 = -21$$

$$\text{ReLU: } [0, 6, 10]$$



$$\text{Sum: } 10 - 5 + 6 = 11$$

$$\text{First Difference: } 10 - (-5) - 6 = 9$$

$$\text{ReLU: } [10, 0, 6]$$

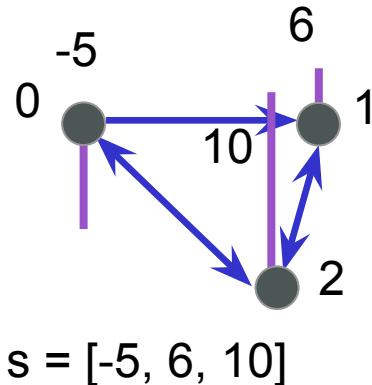
Permutation Invariant

NOT Permutation Invariant

Permutation Invariant 9

Challenges in Graph Data – Incorporating Graph Data

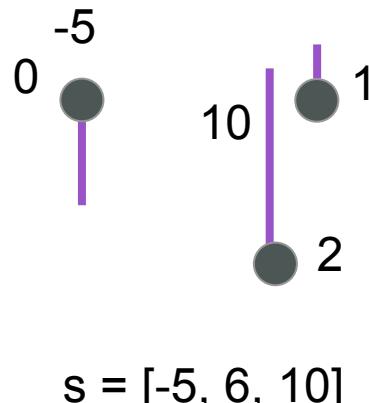
In the previous example, none of the operations actually used the graph!



Sum: $-5+6+10 = 11$

First Difference: $-5 - 6 - 10 = -21$

ReLU: [0, 6, 10]



Same Results

Main Question

How can we design deep learning techniques for graphs?

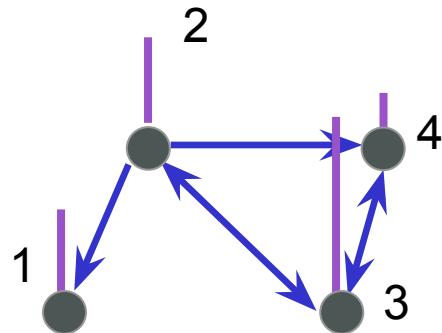
- Must incorporate graph information
- All operations must be permutation invariant

This field is called Geometric Deep Learning

Graph Signal Processing

Discrete Signal Processing works well for grid-structured data

Graph Signal Processing tries to do everything you can do in DSP, but for graph data.



$$A = \begin{bmatrix} 0 & 0 & 1 & 0 \\ 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 1 \\ 1 & 0 & 1 & 0 \end{bmatrix}$$

Adjacency Matrix
Ordering is

Node I \rightarrow Node j means

$$A_{ji} = 1$$

In GSP!

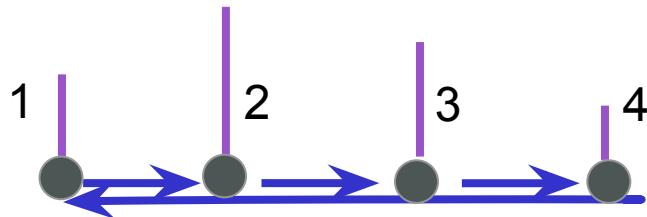
Different from Comp Sci

CMU Course: 18-898D: Special Topics in Signal Processing: Graph Signal Processing and Geometric Learning (offered most Fall semesters)

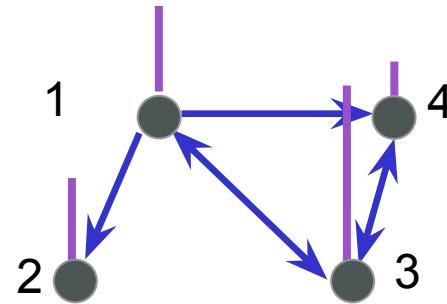
DSP to GSP – Structure, Data

DSP:

Assume the signal is periodic with period T



GSP:



time: path + boundary condition

$$A_c = \begin{bmatrix} 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix}$$

general graph G (fixed & given)

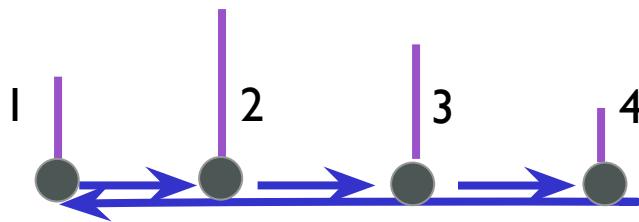
$$A = \begin{bmatrix} 0 & 0 & 1 & 0 \\ 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 1 \\ 1 & 0 & 1 & 0 \end{bmatrix}$$

Graph signal: $s : v_n \rightarrow x_n$
 $\mathcal{V} \rightarrow \mathbb{C}$

- attributes on knowledge graph
- measurements on sensor network
- voltages on power grid
- ...

DSP to GSP – Graph Shift

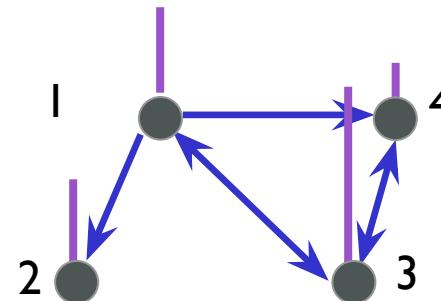
DSP:



time: line graph

$$A_c = \begin{bmatrix} 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix}$$

GSP:

general graph G (fixed & given)

$$A = \begin{bmatrix} 0 & 0 & 1 & 0 \\ 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 1 \\ 1 & 0 & 1 & 0 \end{bmatrix}$$

Graph Shift: As

$$s = [2, 5, -1, 10]^T$$

$$As = [10, 2, 5, -1]^T$$

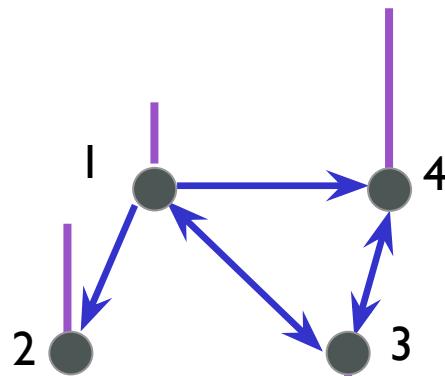
$$s = [2, 5, -1, 10]^T$$

$$As = [-1, 2, 2 + 10, 2 - 1]^T = [-1, 2, 12, 1]^T$$



Graph Shift – Permutation Invariance

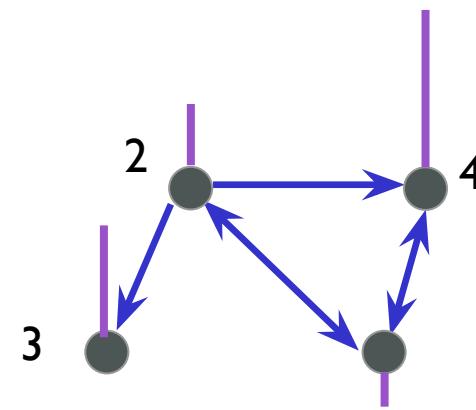
GSP:

general graph G (fixed & given)

$$A = \begin{bmatrix} 0 & 0 & 1 & 0 \\ 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 1 \\ 1 & 0 & 1 & 0 \end{bmatrix}$$

$$s = [2, 5, -1, 10]^T$$

$$As = [-1, 2, 2 + 10, 2 - 1]^T = [-1, 2, 12, 1]^T$$

Graph Shift: A_S 

$$s_2 = [-1, 2, 5, 10]^T$$

$$A_2 = \begin{bmatrix} 0 & 1 & 0 & 1 \\ 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 1 & 1 & 0 & 0 \end{bmatrix}$$

$$\begin{aligned} A_2 s_2 &= [2 + 10, -1, 2, -1 + 2]^T \\ &= [12, -1, 2, 1]^T \end{aligned}$$

Permuted As!

Graph Shift – Permutation Invariance Proof

Permutation Matrix Π : Obtained by permuting the rows of the identity matrix

Each row has 1 one and N-1 zeros. Each row and column only has one 1

Every Permutation matrix is orthogonal: $\Pi^{-1} = \Pi^T$

$$\Pi = \begin{bmatrix} 0 & 0 & 1 & 0 \\ 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \quad s = [2, 5, -1, 10]^T \quad \Pi s = [-1, 2, 5, 10]^T = s_2$$

Graph Shift: As

$$A_2 = \Pi^T A_1 \Pi \quad s_2 = \Pi s_1$$

$$A_2 s_2 = \Pi A_1 \Pi^T \Pi s_1 = \Pi A_1 s_1$$

GSP – Filtering and Convolution

- In GSP, a Linear, Shift-Invariant Graph Filter
(max degree: N-1, why?)

$$P(A) = \sum_{i=0}^{N-1} p_i A^i$$

- Convolution: Matrix-Vector product of a polynomial of the Adjacency Matrix and the signal s

$$P(A)s$$

- Graph Fourier Transform: Found using eigendecomposition of A

$$\begin{aligned}A &= \text{GFT}^{-1} \Lambda \text{GFT} \\ \hat{s} &= \text{GFT} s\end{aligned}$$

GSP – Permutation Invariant Proofs

$$A_2 = \Pi A_1 \Pi^T \quad s_2 = \Pi s_1$$

Convolution is Permutation Invariant

$$\begin{aligned} P(A_2)s_2 &= P(\Pi A_1 \Pi^T)(\Pi s_1) \\ &= \Pi P(A_1)\Pi^T\Pi(s_1) \\ &= \Pi(P(A_1)s_1) \end{aligned}$$

GFT is Permutation Invariant

$$A_1 = \text{GFT}_1^{-1} \Lambda_1 \text{GFT}_1$$

$$A_2 = \Pi A_1 \Pi^T = \Pi \text{GFT}_1^{-1} \Lambda_1 \text{GFT}_1 \Pi^T = \underbrace{\Pi \text{GFT}_1^{-1}}_{\text{GFT}_2^{-1}} \Lambda_1 \underbrace{\text{GFT}_1 \Pi^T}_{\text{GFT}_2}$$

$$\widehat{s}_2 = \text{GFT}_2 s_2 = \text{GFT}_1 \Pi^T \Pi s_1 = \text{GFT}_1 s_1 = \widehat{s}_1$$

Graph Neural Networks Tasks

Types of Experiments

1. Node Classification

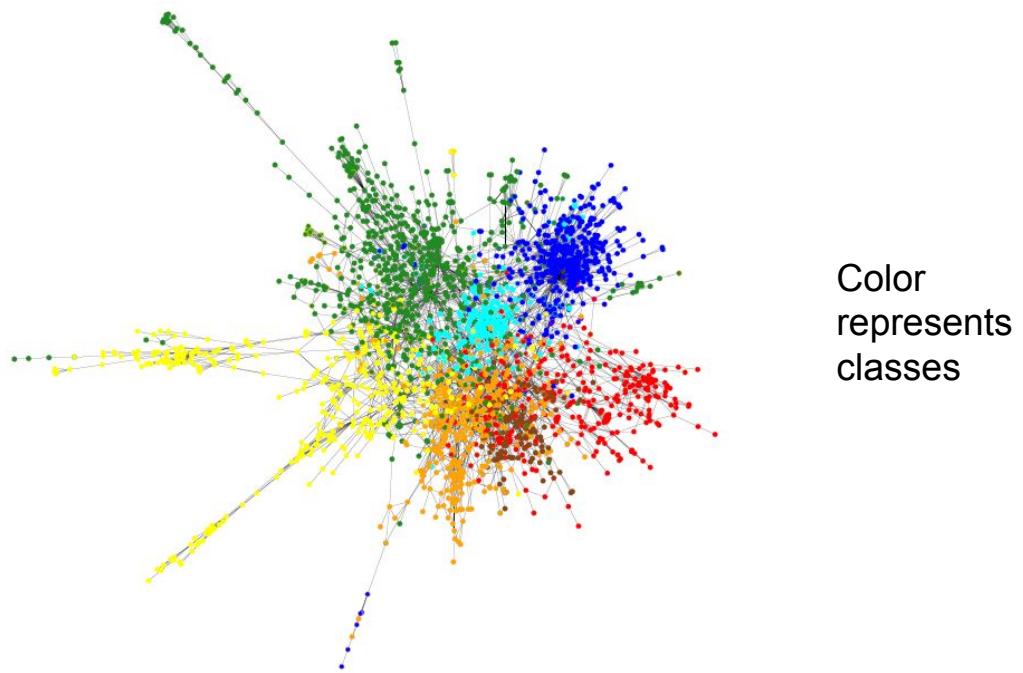
- Classify each node in the graph
- Graph shows connections between each data point
- No pooling
- GCNNs vs. MLPs

2. Graph Signal Classification

- Classify the entire graph signal
- Graph shows connections between the **features**
- Graph can be different for each data point
- Can do pooling

Node Classification: Citation Network

Classify the type of paper using the citation network (**graph**) and a bag of words (**data**) for each paper (**node**)



Dataset	Nodes	Labelled Nodes	Label Rate	Edges	Connectivity	Features	Classes
Citeseer	3327	120	0.036	4732	4e-4	3703	6
Cora	2708	141	0.052	5429	7e-4	1433	7
Pubmed	19717	59	0.003	44338	1e-4	500	3

Node Classification

Why do we need a graph? Why not just use an MLP?

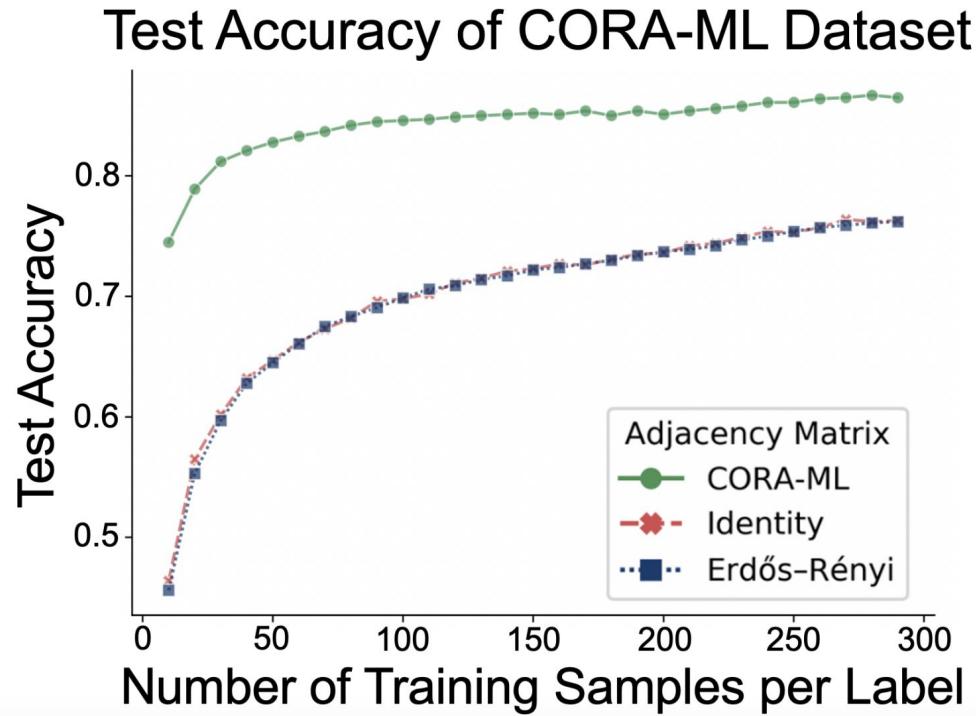
The graph allows for smaller training set.

Graph is like an “answer key” - Helps with classification.

Dataset	Nodes	Labelled Nodes	Label Rate	Edges	Connectivity	Features	Classes
Citeseer	3327	120	0.036	4732	4e-4	3703	6
Cora	2708	141	0.052	5429	7e-4	1433	7
Pubmed	19717	59	0.003	44338	1e-4	500	3

How good is my answer key?

- Graph provides “answer key” for classification problem, but how good is it?
- Graph greatly affects accuracy.



How good is my answer key?

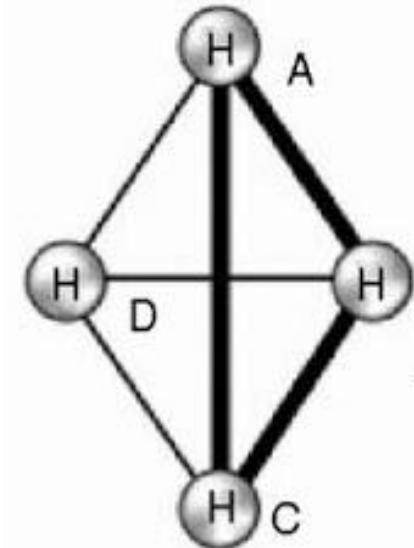
- How to form a graph:
 - Use some intuitive heuristic believed to be related to the classification
 - Use cross-correlation of the features.
- Many factors about the graph affect accuracy:
 - Homophilly/Heterophilly: Are nodes with the same classes connected? -> High homophilly
 - Features: Are nodes with similar features connected?
 - Clustering: Are nodes with the same class clustered together? Same Features?

Graph Classification: Biological Network

- Each protein, chemical and enzyme structure is a **different** graph
- Data is the elements of each node in the structure
- Predict properties based on the chemical structure

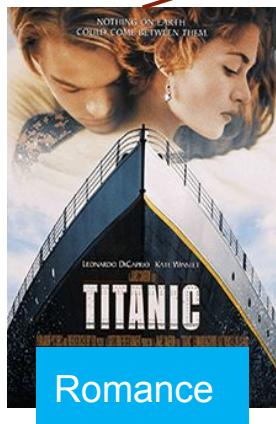
Dataset	#Graphs	Classes	Avg. #nodes	Avg. # edges
MUTAG	188	2	17.7	38.9
PTC	344	2	26.7	50.7
ENZYME	600	6	32.6	124.3
D&D	1178	2	284.4	1921.6

Graph of protein structure 1



Graph Classification: More Examples

- **IMDB-binary:**
ego-networks of actors who have appeared together in movies.
classify network into action/romance movie
- **Reddit-binary:**
does the user come from Q&A forum or discussion forum?



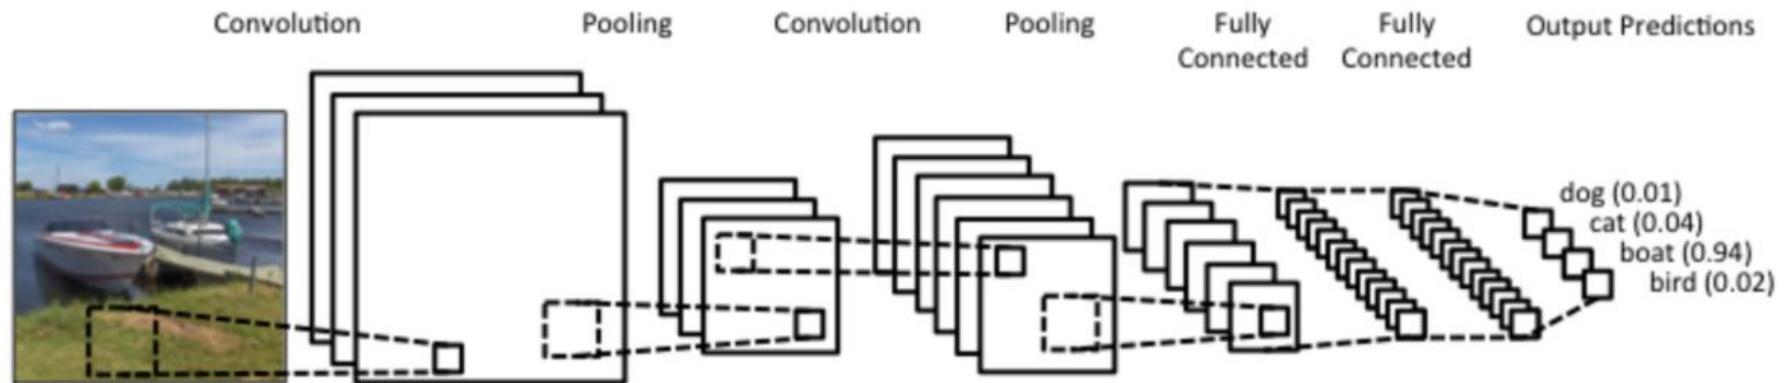
Romance



CNNs to GNNs

DSP -> GSP

How do we change the CNN to Graph CNN?



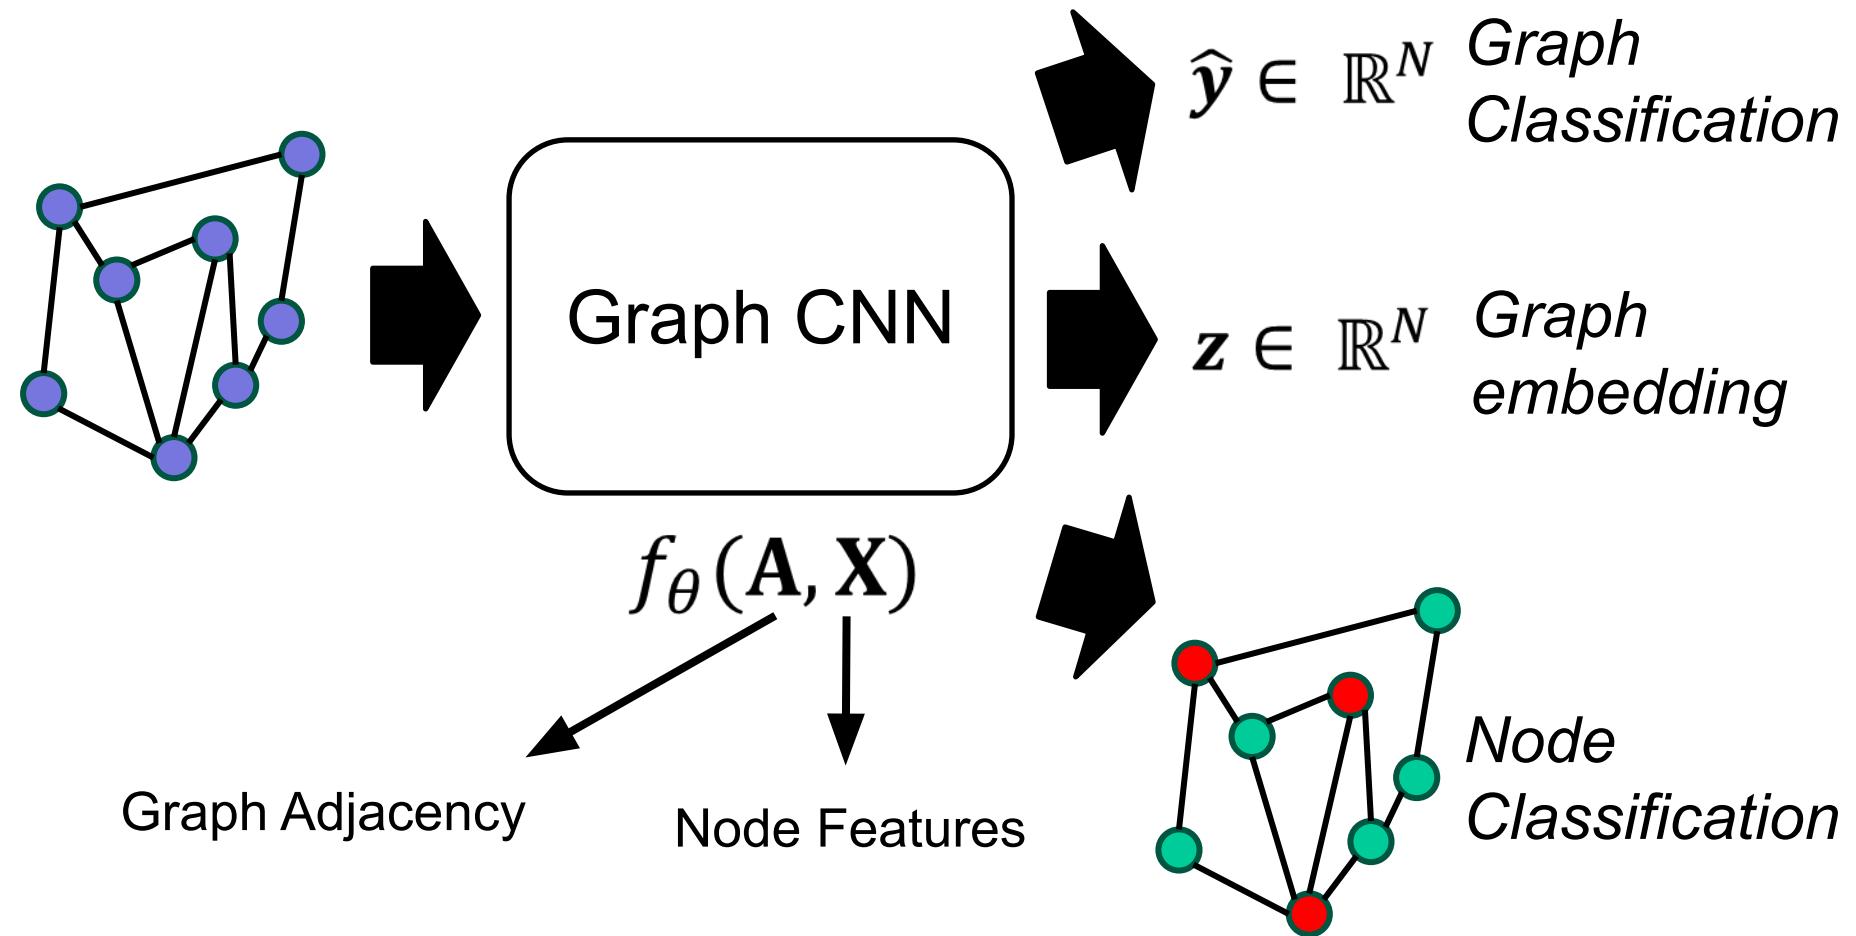
Challenges of Graph CNN

Traditional CNNs do not perform well on Graph data.

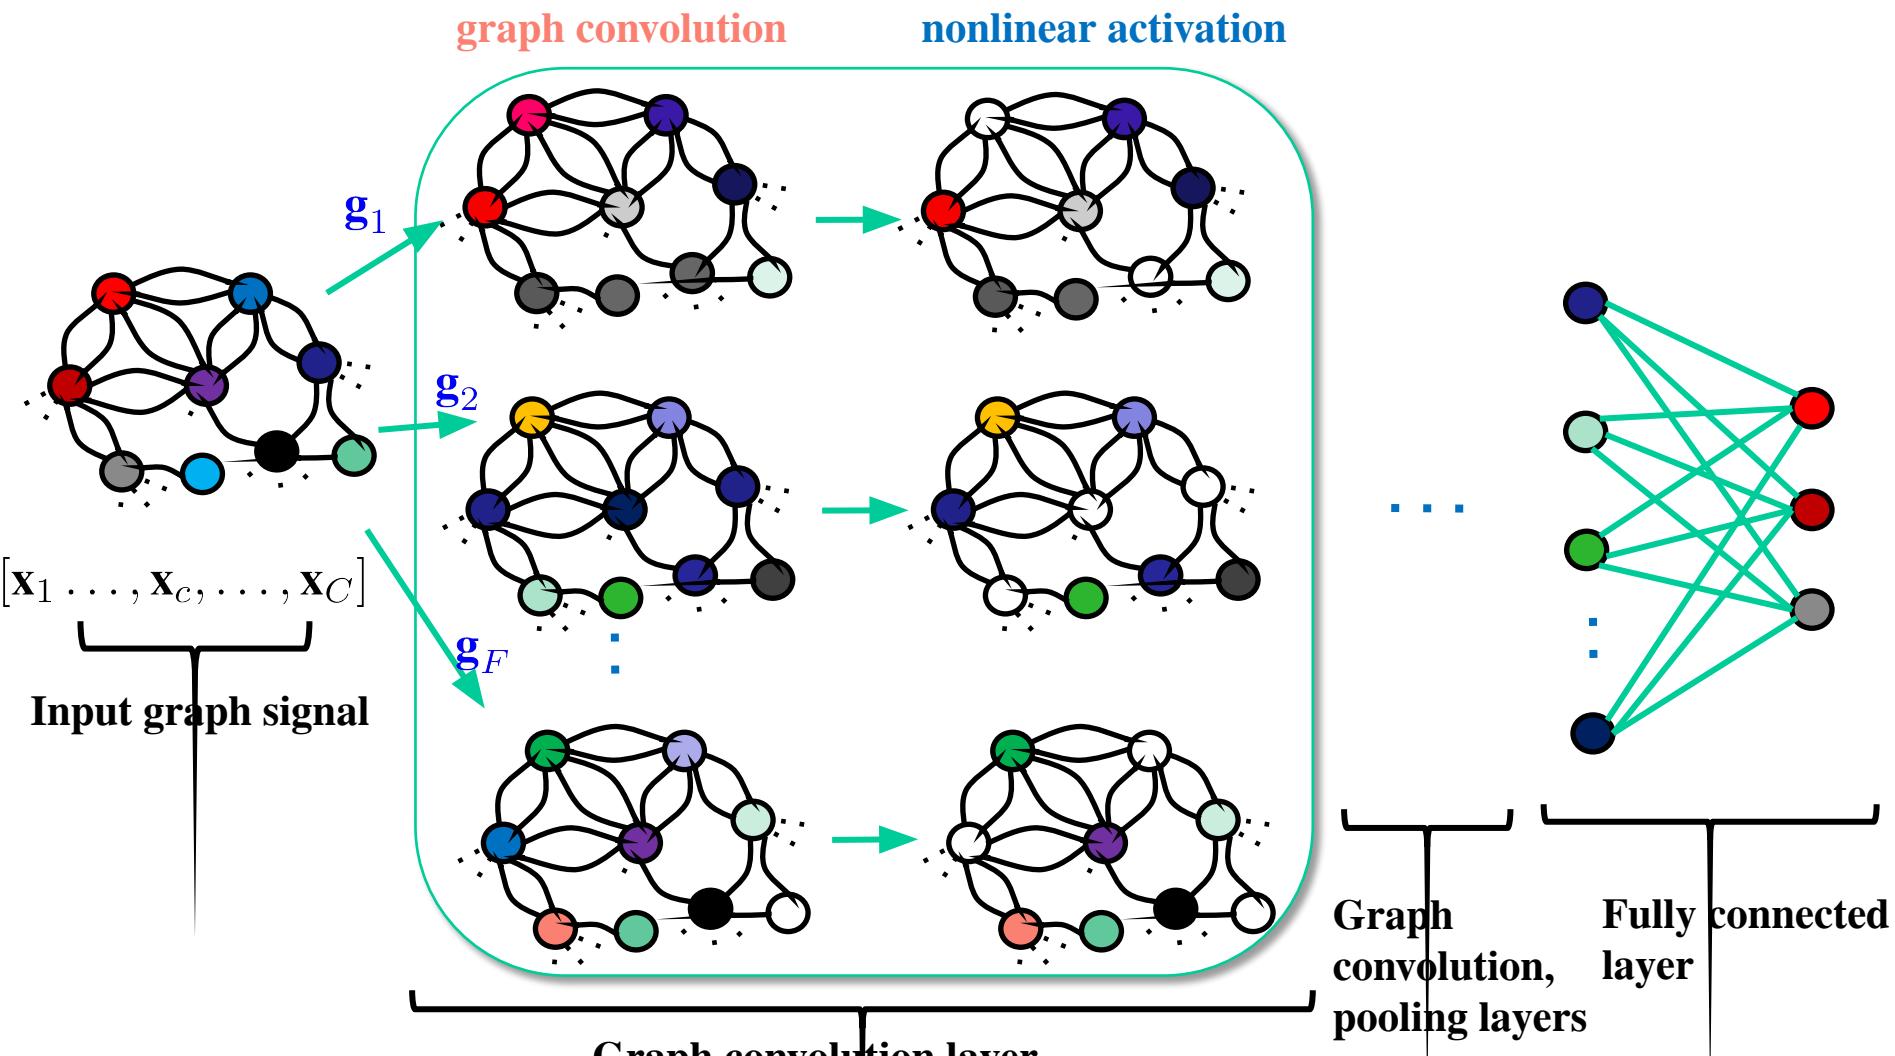
- What *graph* problems can be solved?
- What are *graph* convolutional layers?
- What is *graph* pooling?
- Which *graph* architecture should be used?

Graph Neural Network Architectures

Graph CNNs



Architecture of Graph CNN



Why do we need Graph CNNs?

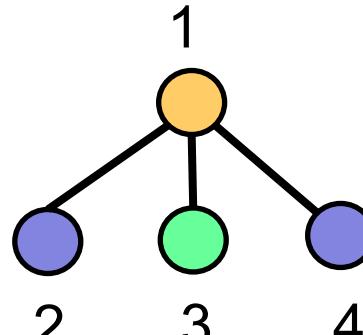
- Varying input size
 - Permutation invariance – graph representation not unique

$$f_{\theta}(\mathbf{P}\mathbf{A}\mathbf{P}^T, \mathbf{P}\mathbf{X}) = f_{\theta}(\mathbf{A}, \mathbf{X}) \quad \text{Permutation Invariance (Graph Classification)}$$

$$f_{\theta}(\mathbf{P}\mathbf{A}\mathbf{P}^T, \mathbf{P}\mathbf{X}) = \mathbf{P}f_{\theta}(\mathbf{A}, \mathbf{X}) \quad \text{Permutation Invariance (Node Classification)}$$

(Some works refer to this as Permutation Equivariance)

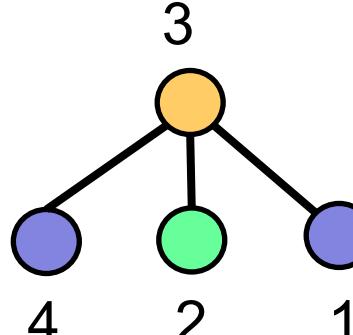
Compositions of permutation equivariant functions are also permutation equivariant



$$A = \begin{pmatrix} 0 & 1 & 1 & 1 \\ 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 \end{pmatrix}$$

$X =$



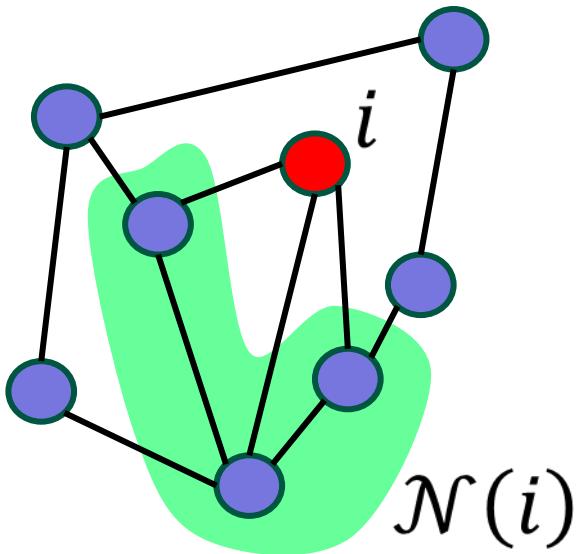


$$A = \begin{pmatrix} 0 & 0 & 1 & 0 \\ 0 & 0 & 1 & 0 \\ 1 & 1 & 0 & 1 \\ 0 & 0 & 1 & 0 \end{pmatrix}$$

X =

Graph Convolutional Layers

1. Aggregate node features of neighbors
2. Combine current node feature with aggregation from prior step



$$\mathbf{a}_i^{(\ell)} = AGG^{(\ell)} \left(\{\mathbf{x}_j^{(\ell-1)}, j \in \mathcal{N}(i)\} \right)$$
$$\mathbf{x}_i^{(\ell)} = COMB^{(\ell)} \left(\{\mathbf{x}_i^{(\ell-1)}, \mathbf{a}_i^{(\ell)}\} \right)$$

Simple Graph Convolution Layer

$$\mathbf{x}_i^{(\ell+1)} = \mathbf{W}_1^\top \mathbf{x}_i^{(\ell)} + \mathbf{W}_2^\top \sum_{j \in \mathcal{N}(i)} e_{j,i} \mathbf{x}_j^{(\ell)}$$

Learned at each layer

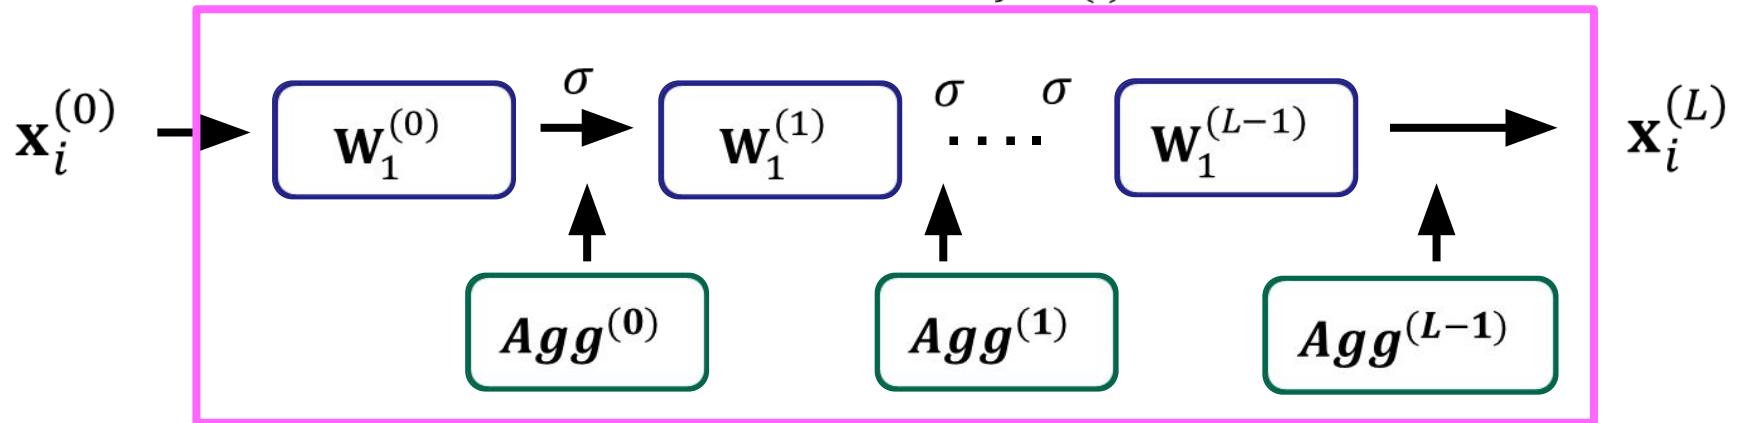
Vectorized Representation:

$$\mathbf{X}^{(\ell+1)} = \mathbf{X}^{(\ell)} \mathbf{W}_1 + \mathbf{A} \mathbf{X}^{(\ell)} \mathbf{W}_2$$

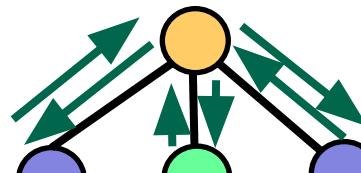
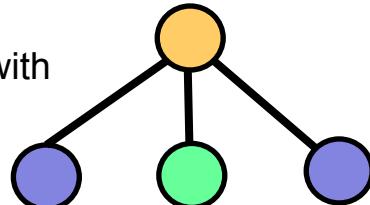
Morris et. Al., “Weisfeiler and Leman Go Neural: Higher-order Graph Neural Networks”, 2019.

Graph CNNs as Distributed MLPs

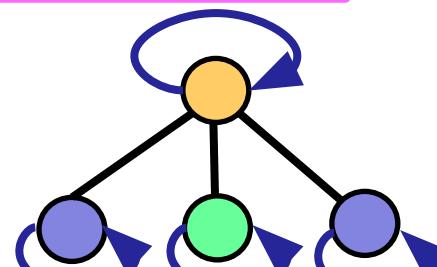
$$\mathbf{x}_i^{(\ell+1)} = \mathbf{W}_1^\top \mathbf{x}_i^{(\ell)} + \mathbf{W}_2^\top \sum_{j \in \mathcal{N}(i)} e_{j,i} \mathbf{x}_j^{(\ell)}$$



1. Aggregate
2. Single Linear layer with aggregation as bias
3. Nonlinearity
4. Repeat



Aggregate



Combine (Linear layer with aggregation as bias)

Graph Convolutional Layers

GCN: $\mathbf{X}^{(\ell+1)} = \widehat{\mathbf{D}}^{-\frac{1}{2}} \widehat{\mathbf{A}} \widehat{\mathbf{D}}^{-\frac{1}{2}} \mathbf{X} \mathbf{W}$

$$\mathbf{x}^{(\ell+1)} = \mathbf{W}^T \sum_{j \in \mathcal{N}(i)} \frac{e_{j,i}}{\sqrt{\widehat{d}_j \widehat{d}_i}} \mathbf{x}^{(\ell)}$$

Kipf, Welling, 2017.

TAGCN: $\mathbf{X}^{(\ell+1)} = \sum_{k=0}^K \left(\mathbf{D}^{-\frac{1}{2}} \mathbf{A} \mathbf{D}^{-\frac{1}{2}} \right)^k \mathbf{X} \mathbf{W}_k$

Du, Zhang, Wu, Moura, Kar, 2018.

$$\bar{\mathbf{A}} = \mathbf{D}^{-\frac{1}{2}} \mathbf{A} \mathbf{D}^{-\frac{1}{2}}$$

Normalized Adjacency matrix

Why normalize?

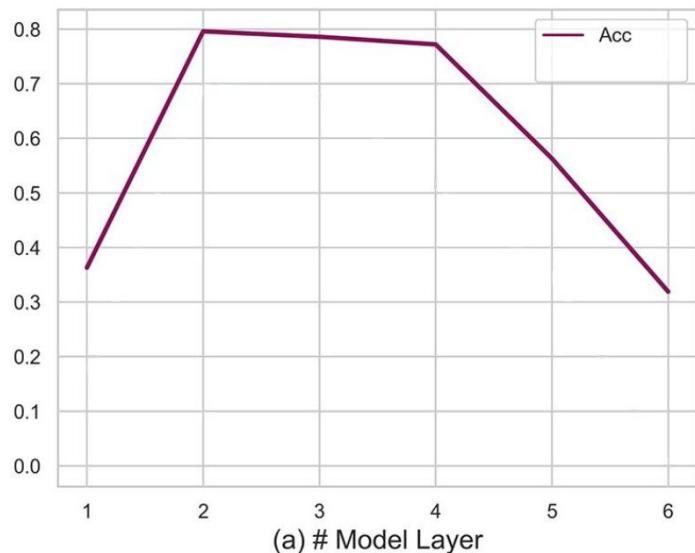
$$\mathbf{X}^{(L)} = \mathbf{A}^L \mathbf{X}^{(0)} \mathbf{W}_1 \mathbf{W}_2 \dots \mathbf{W}_L = \underbrace{(\text{GFT}^{-1}) \boldsymbol{\Lambda}^L (\text{GFT})}_{\text{Blows up if largest eigenvalue is greater than 1}} \mathbf{X}^{(0)} \mathbf{W}_1 \mathbf{W}_2 \dots \mathbf{W}_L$$

Blows up if largest eigenvalue is greater than 1

Over-Smoothing Problem

Chen et. al., "Measuring and recliving over smoothing..." AAAI 2020

Kei Ishikawa, "GNN Oversmoothing", ETH Zurich Course Slides.



The node classification accuracy (Acc) of GCNs on the CORA dataset.

Increased depth of Graph CNNs can lead to lower accuracy

$$\begin{aligned}\mathbf{X}^{(L)} &= \bar{\mathbf{A}}^L \mathbf{X}^{(0)} \mathbf{W}_1 \mathbf{W}_2 \dots \mathbf{W}_L \\ &= (\mathbf{U}^{-1}) \underbrace{\Lambda^L (\mathbf{U})}_{\text{Converges to dominant eigenvector of } \bar{\mathbf{A}}} \mathbf{X}^{(0)} \mathbf{W}_1 \mathbf{W}_2 \dots \mathbf{W}_L\end{aligned}$$

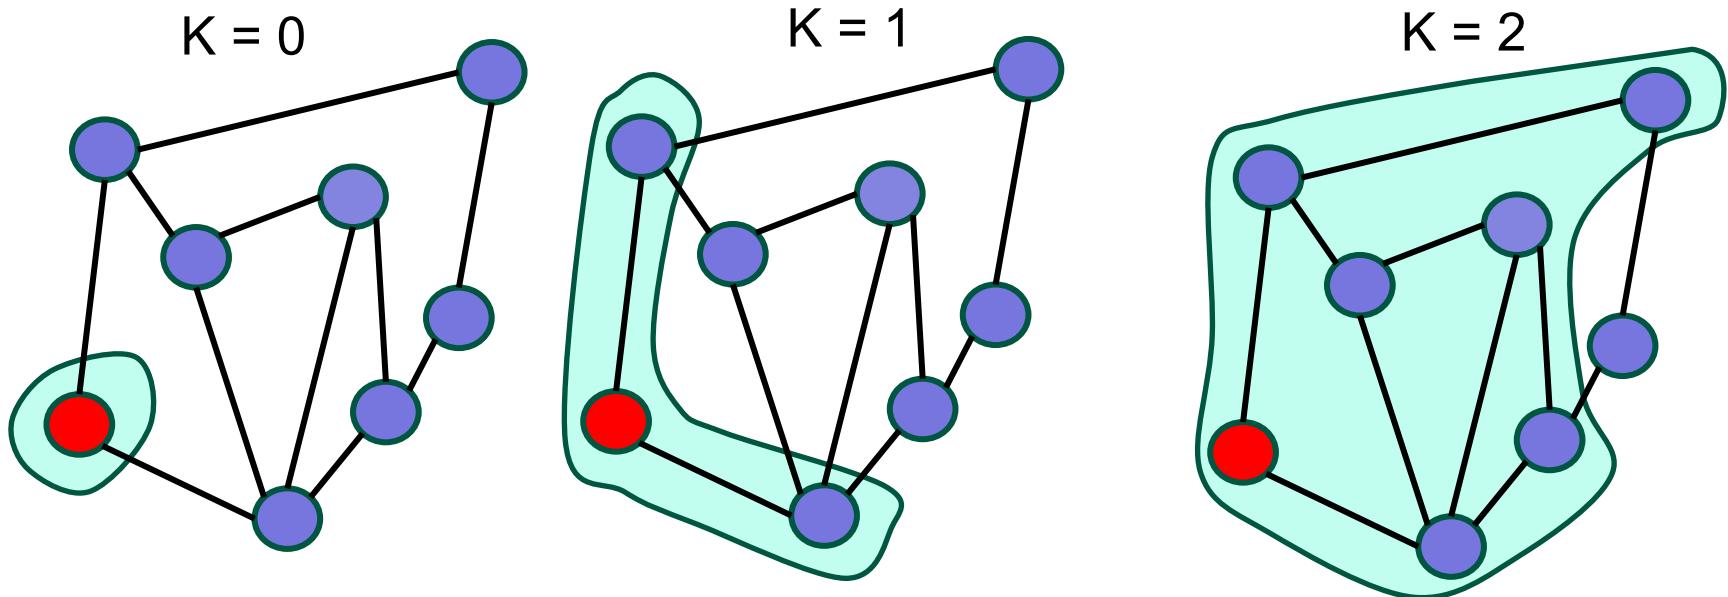
Converges to dominant eigenvector of $\bar{\mathbf{A}}$ by power iteration

TAGCN

$$\text{TAGCN: } \mathbf{X}^{(\ell+1)} = \sum_{k=0}^K \bar{\mathbf{A}}^k \mathbf{X} \mathbf{W}_k$$

Du, Zhang, Wu, Moura, Kar, 2018.

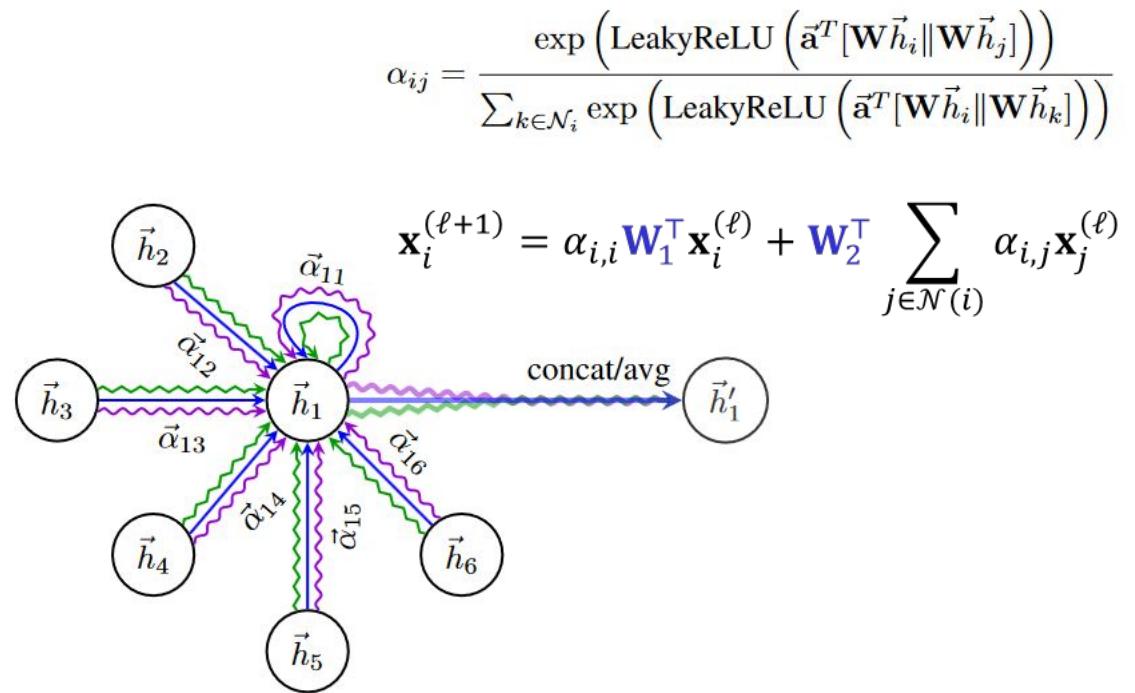
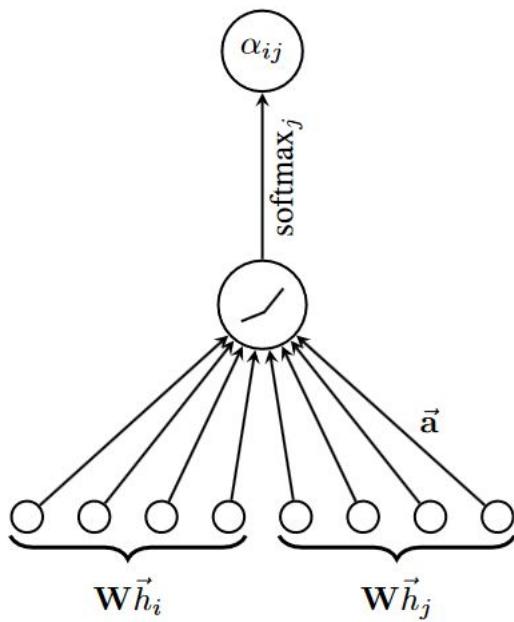
Hyperparameter K is the K-hop neighborhood over which information is aggregated



Kind of like how a CNN aggregates its neighbors!

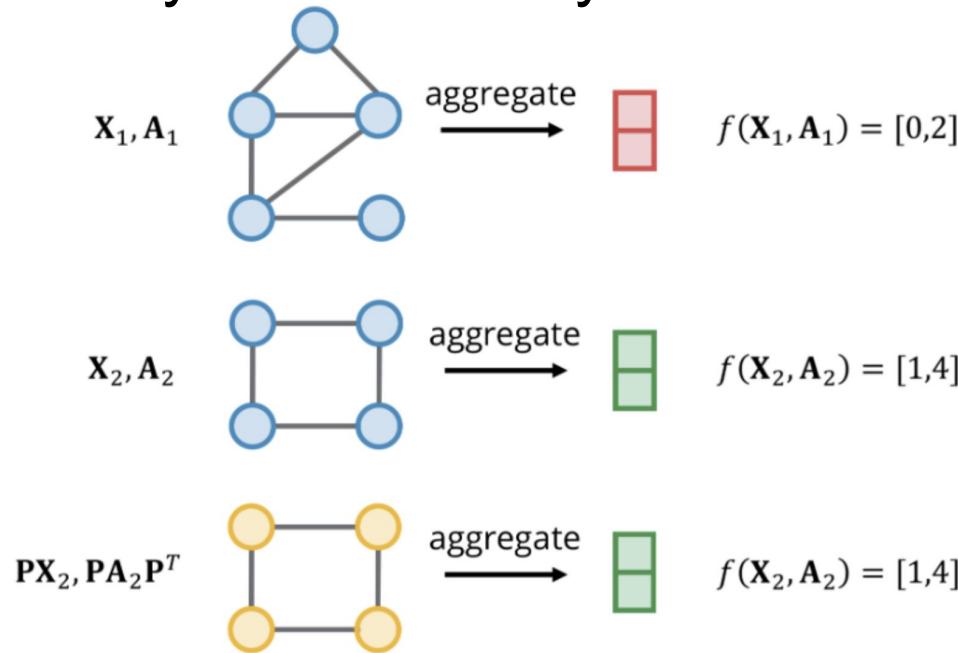
Graph Attention Layers

- Aforementioned graph CNNs use fixed edge weights of \mathbf{A}
- Graph attention networks: First learn proper edge weights, then convolve



Graph Aggregation Layers

After convolutional layers, the output is $N \times C$, which is then put into a fully connected layer.



Goal: Get the same dimension for all graphs. Make sure graphs and permuted graphs produce the same result.
This is the input to the FC layer.

Graph Aggregation Layers

- Architectures thus far yield node embeddings
- Assign each node of the graph a representation $\mathbf{z}_i \in \mathbb{R}^n$
- **How do we obtain a global graph representation?**

Aggregation Layer: $f : \mathbb{R}^{N \times C} \rightarrow \mathbb{R}^{1 \times C}$

- **Global add pooling:** $\mathbf{z}_G = \sum_{i=1}^N \mathbf{z}_i$
- **Global mean pooling:** $\mathbf{z}_G = \frac{1}{N} \sum_{i=1}^N \mathbf{z}_i$
- **Global max pooling:** $\mathbf{z}_{G,j} = \max_{i=1,\dots,N} \mathbf{z}_{i,j}$

Comparison of Aggregation Methods

Best aggregation method is data-dependent (Mark Cheung PhD Thesis, 2023)

	MUTAG	PROTEINS	IMDB-B	REDDIT-B	COLLAB
TAGCN (mean)	75.1 ± 8.2	72.4 ± 2.9	73.3 ± 5.3	91.6 ± 2.6	81.0 ± 1.1
TAGCN (var)	79.3 ± 4.2	73.5 ± 2.9	67.8 ± 2.3	91.8 ± 1.3	78.5 ± 0.9
TAGCN (max)	76.1 ± 5.5	73.0 ± 2.0	72.3 ± 2.7	90.3 ± 1.3	76.3 ± 2.0
TAGCN (random)	75.5 ± 1.1	67.1 ± 2.6	73.1 ± 2.8	85.8 ± 1.4	76.0 ± 1.7
TAGCN (mean+var)	76.1 ± 6.2	72.9 ± 2.4	74.0 ± 4.3	91.5 ± 1.7	79.5 ± 1.3
TAGCN (mean+max)	74.5 ± 8.3	74.6 ± 2.9	71.6 ± 2.7	90.6 ± 1.7	78.7 ± 2.6

Graph Pooling

Traditional CNN Pooling

0	4	0	2
2	0	1	8
1	2	3	4
5	6	7	8

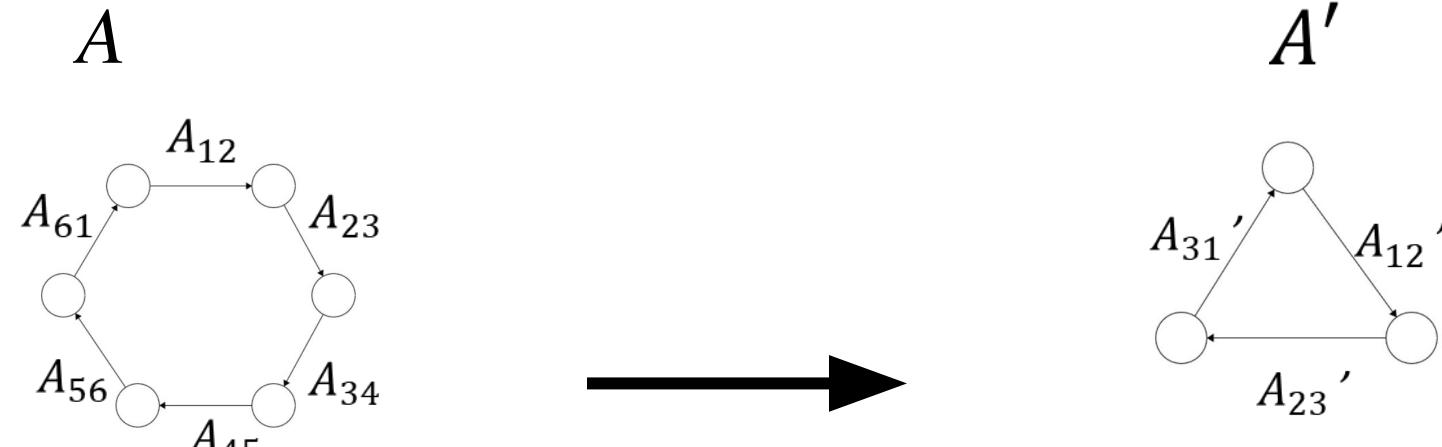


Max pool with
2x2 filters
and stride of 1

4	4	8
2	3	8
6	7	8

Graph Pooling

Graph Structure:



Data: x x_1
 x_2
 x_3
 x_4
 x_5
 x_6

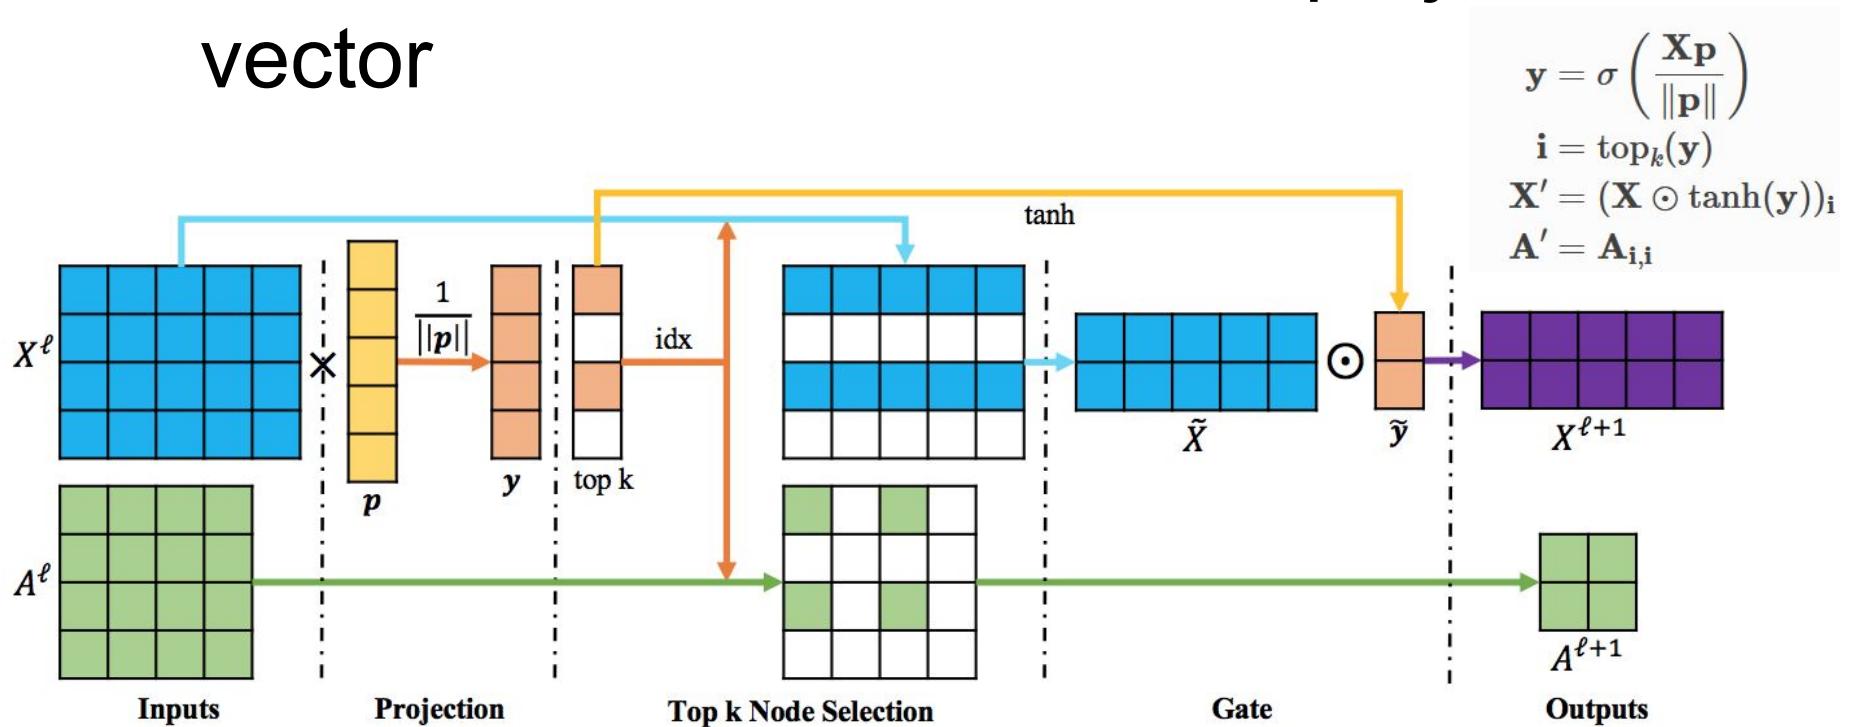
Challenge:
Need to know both the new
data and the new graph

x' x_1'
 x_2'
 x_3'

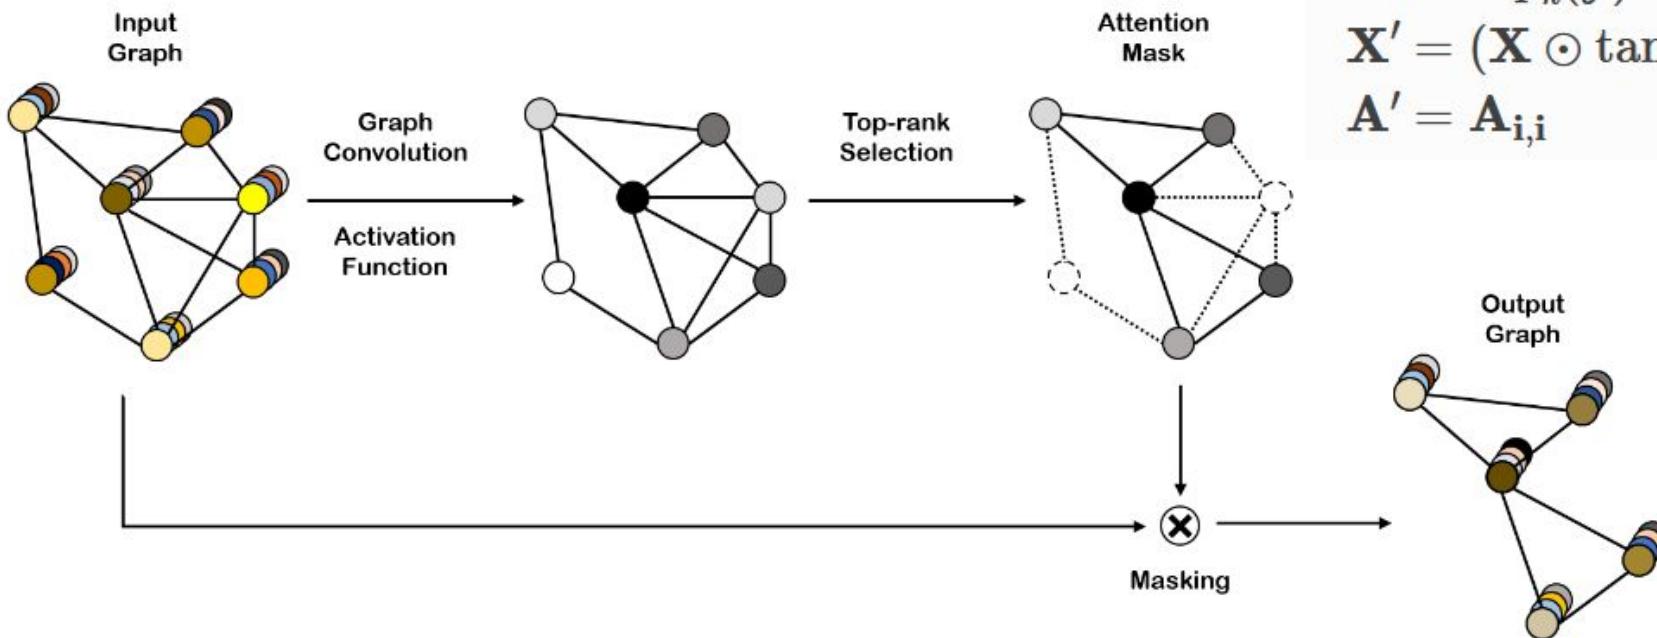
Must be permutation invariant

Top-k Pool (originally gpool)

- Autoencoder with a trainable projection vector



Self-attention Graph Pooling (sagpool)



$$\begin{aligned}\mathbf{y} &= \text{GNN}(\mathbf{X}, \mathbf{A}) \\ \mathbf{i} &= \text{top}_k(\mathbf{y}) \\ \mathbf{X}' &= (\mathbf{X} \odot \tanh(\mathbf{y}))_{\mathbf{i}} \\ \mathbf{A}' &= \mathbf{A}_{\mathbf{i}, \mathbf{i}}\end{aligned}$$

Lee, Lee, Kang: Self-Attention Graph Pooling (ICML 2019)

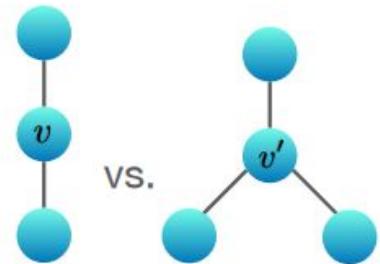
Graph CNN Implementation

- <https://pytorch-geometric.readthedocs.io/en/latest/>
- Pytorch geometric – based on standard pytorch, compatible with pytorch lightning
- Provides:
 - Efficient graph data handling
 - Provides implementations of several standard graph convolution and pooling layers
 - Provides base classes to define your own graph convolution and pooling layer
 - Access to a wide variety of graph datasets, e.g., for graph and node classification

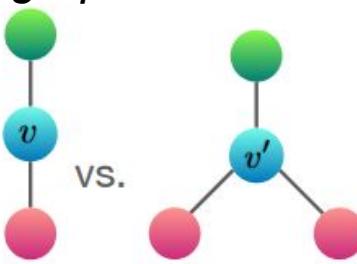


How expressive are Graph NNs

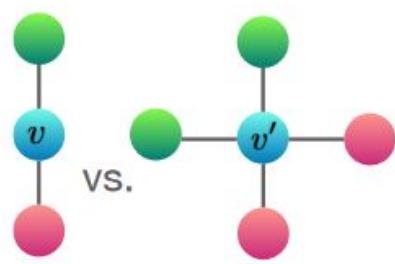
Mean/max aggregation operators assign v, v' same embedding even though the graph structures are different



(a) Mean and Max both fail

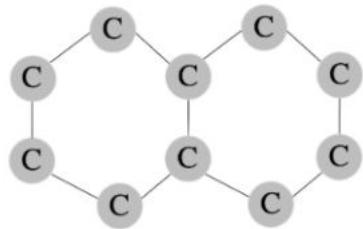


(b) Max fails

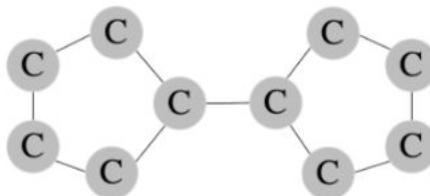


(c) Mean and Max both fail

Xu et. Al., "How Powerful Are GNNs?", 2019.



A_1 (Decalin)



A_2 (Bicyclopentyl)

Shi, et. Al. "A dual approach to graph CNNs", 2020.

How expressive are Graph NNs

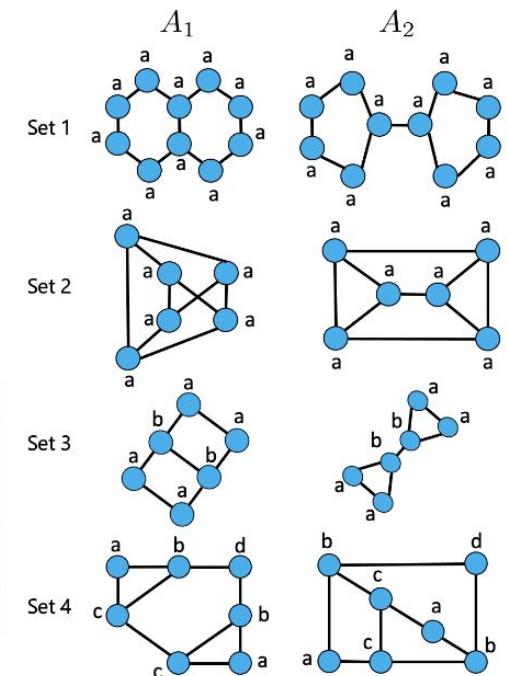
Possible Solution: Use both the nodal shift (A) and spectral shift (M)

$$\left. \begin{array}{l} \mathbf{A} = \text{GFT}^{-1} \boldsymbol{\Lambda} \text{GFT} \\ \mathbf{M} = \text{GFT} \boldsymbol{\Lambda}^* \text{GFT}^{-1} \end{array} \right\}$$

Perform graph convolutions using both shifts

SUMMARY OF RESULTS IN TERMS OF CLASSIFICATION ACCURACY FOR SYNTHETIC DATASETS

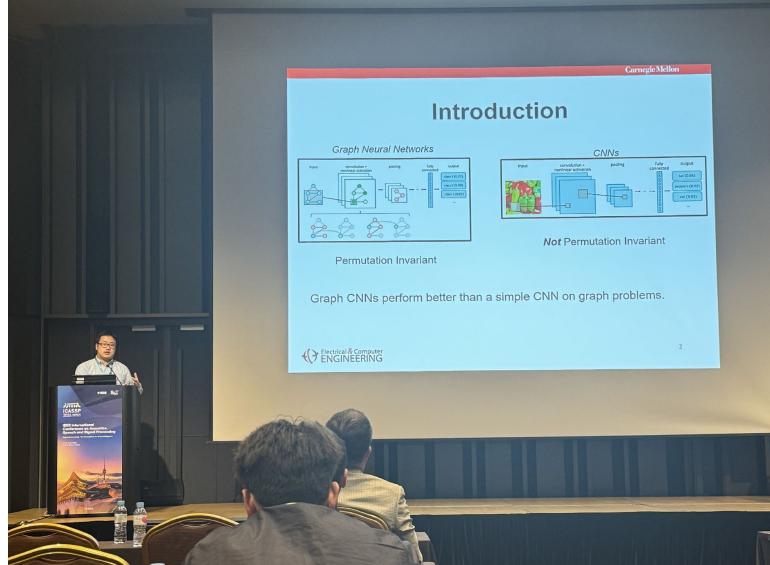
		Set 1	Set 2	Set 3	Set 4
GCN	A	0.495 ± 0.039	0.495 ± 0.014	0.491 ± 0.009	0.493 ± 0.011
	M	0.980 ± 0.008	0.580 ± 0.017	0.997 ± 0.003	0.819 ± 0.027
	$A + M$	0.983 ± 0.005	0.602 ± 0.036	0.998 ± 0.002	0.860 ± 0.031
TAGCN	$P(A)$	0.488 ± 0.028	0.484 ± 0.022	0.493 ± 0.007	0.490 ± 0.019
	$P(M)$	0.981 ± 0.007	0.712 ± 0.062	0.998 ± 0.002	0.993 ± 0.006
	$P(A + M)$	0.983 ± 0.005	0.698 ± 0.023	0.998 ± 0.003	0.993 ± 0.006



Shi, et. Al. "A dual approach to graph CNNs", 2020.

Last Year: Case Study: GNNs = CNNs + B.C.

- Presented Thursday (4/18/24) at ICASSP 2024 in South Korea
- Combines GSP and GCNNs



J. Shi, Shreyas Chaudhari, and J. M. F. Moura, “Graph Convolutional Neural Networks in the Companion Model,” International Conference on Acoustics, Speech, and Signal Processing (ICASSP) (2024).

This Year: Case Study: Inferring the Graph Structure of Images for Graph Neural Networks

- Presenting at GSP Workshop 2025 in Montreal, Canada

M. Gowda, J. Shi, A. Santos, J. M. F. Moura, “Inferring the Graph Structure of Images for Graph Neural Networks,” GSP Workshop 2025.

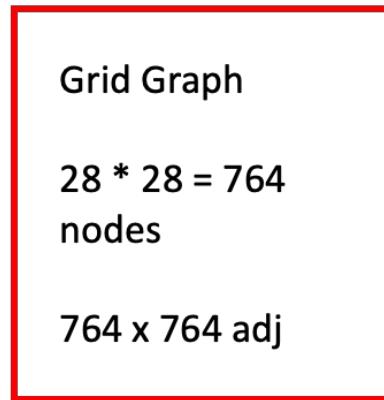
Mayur Gowda is currently an IDL student!!

Main Idea: How do we represent images using graphs?

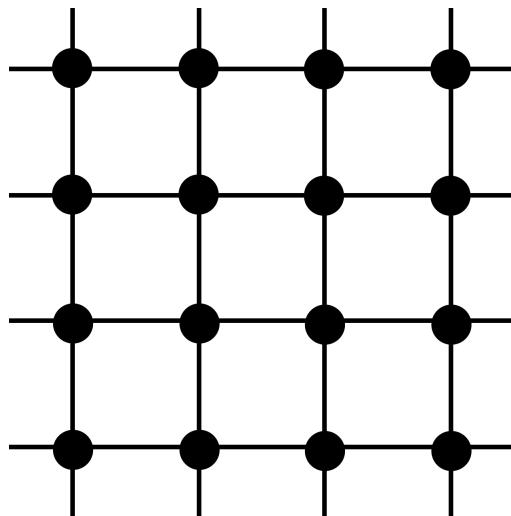


MNIST

28 x 28 pixel
images



Graph CNNs
(GCN, GAT, ...)



Most
intuitive:

Nodes = pixels,
Edges connect
Adjacent Pixels,

Features = pixel
intensities

Main Idea: How do we represent images using graphs?

Can we produce a better **graph** representation of images to increase the accuracy for downstream geometric deep learning tasks?

Capturing pixel-wise similarities. Engineering features.

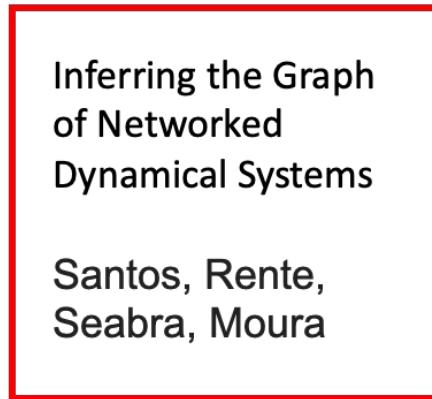
Many methods in the literature: Superpixels, Color Similarity

Main Idea: How do we represent images using graphs?



MNIST

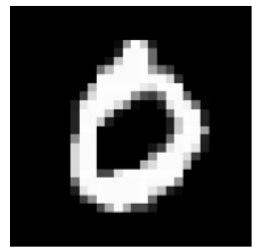
28 x 28 pixel
images



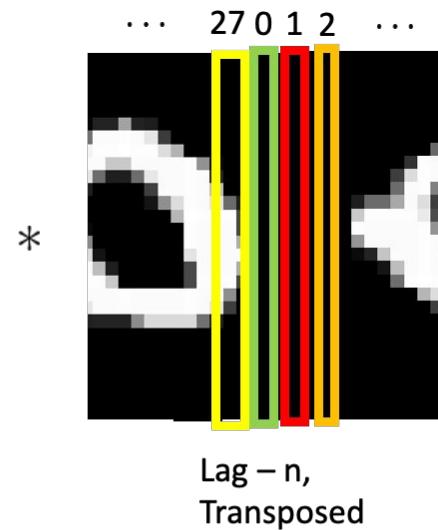
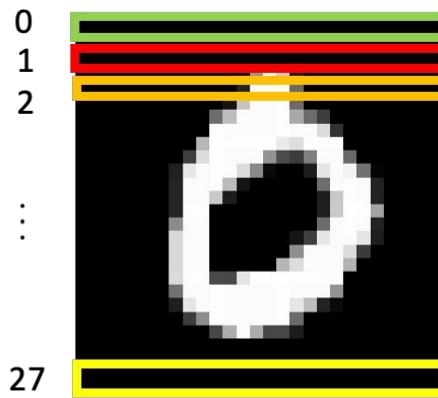
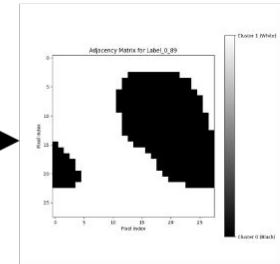
Graph CNNs
(GCN, GAT, ...)

Our method uses the method developed by our collaborators for time series correlation.

Clustering MNIST Using Rows



MNIST

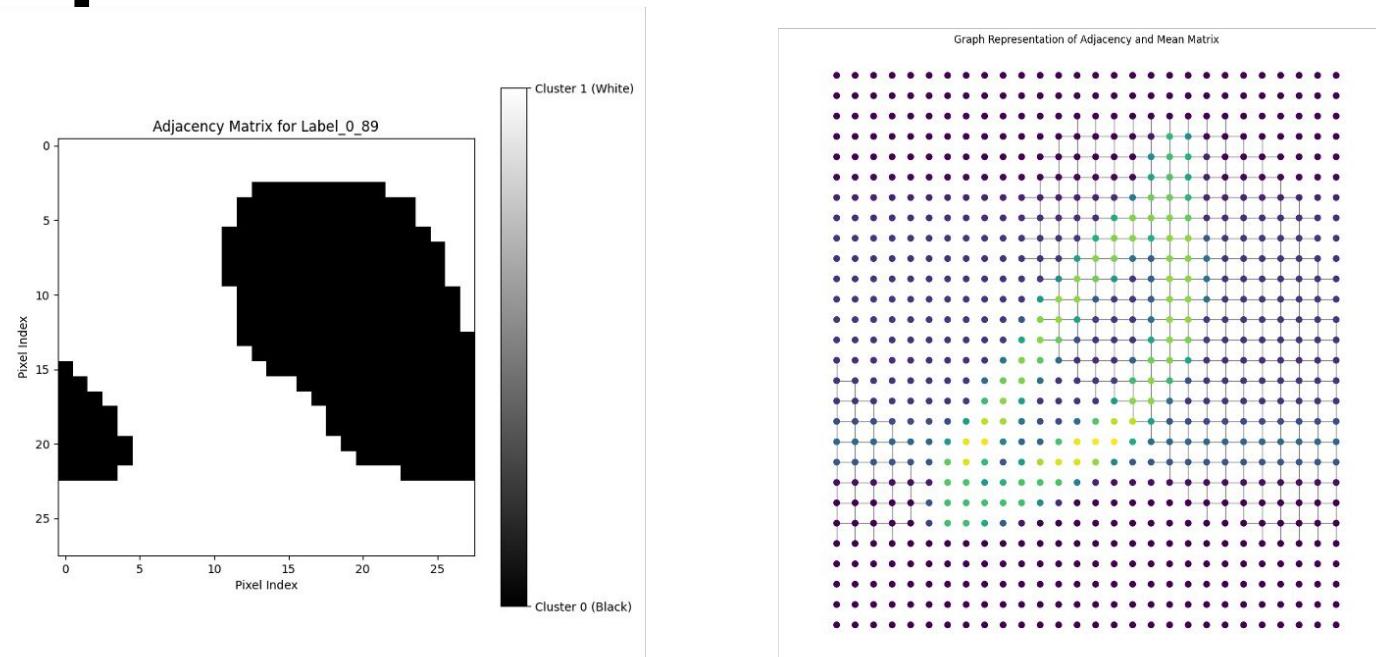
28 x 28 pixel
imagesK Means
Clustering

28 node graph

Compare each row of image with every other row

Redundant, non-local for n between 0 to 27

Representation after Clustering



28 rows = 28 nodes

28 x 28 adjacency matrix, which we superimpose the image on

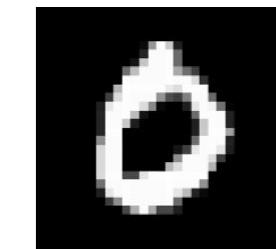
Problems

Directly using the 28×28 adjacency matrix for the row graph as the image?

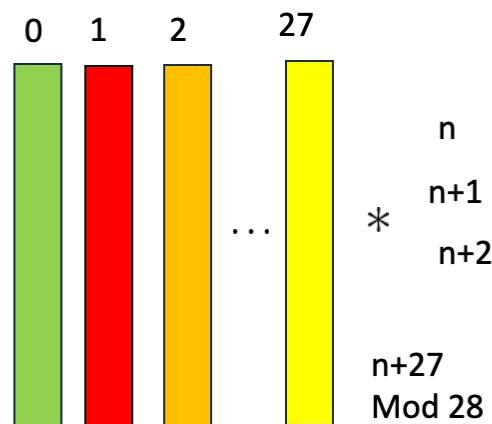
The graph only relates the rows not the relationships between the 784 pixels!

Should have 784×784 adjacency matrix like the grid graph!

Column Graph



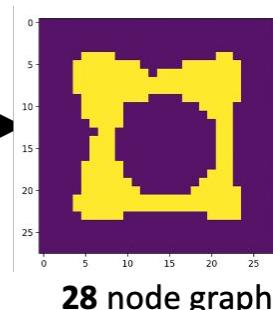
MNIST

28 x 28 pixel
images

Lag – n,
Transposed



K Means
Clustering



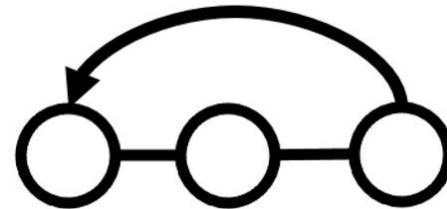
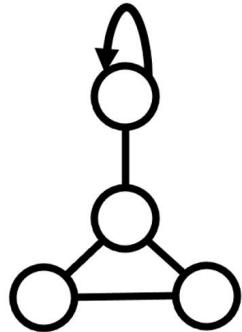
28 node graph

Compare each column of image with every other column

Same issues as row. Only relates the 28 columns!

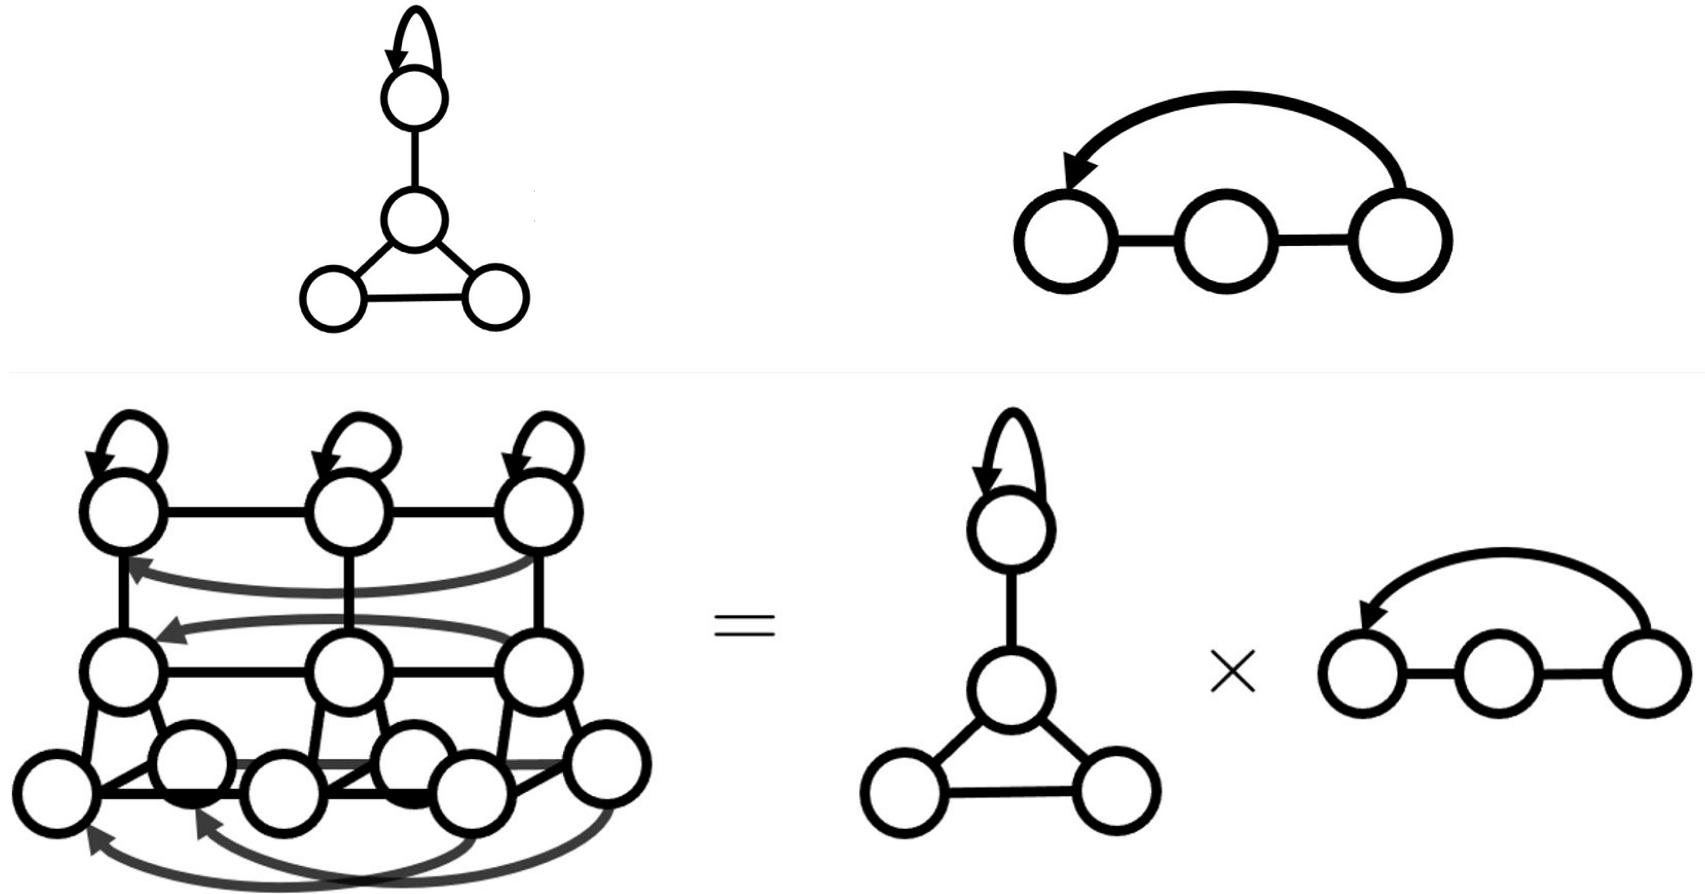
Kronecker Products in GSP

Suppose we have data defined on a graph (like an electrical grid), but it also is a time series...



Two Methods: Represent it on the space graph (left) where each node is a vector of time values. Or Represent it on a time graph (right) where each node is a vector of space values (rarely done).

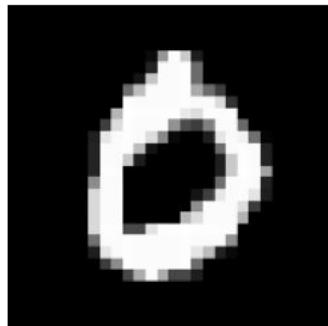
Kronecker Products in GSP



Taking the product graph gives one value per node -> very natural and intuitive

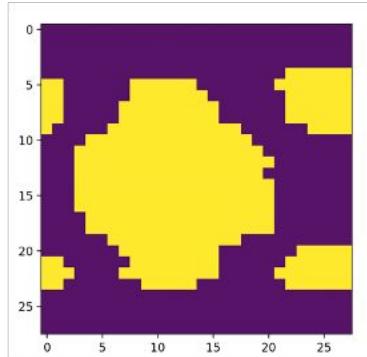
Kronecker Products in GSP

The (Cartesian) product of a line graph (row) and a line graph (column) is a grid graph



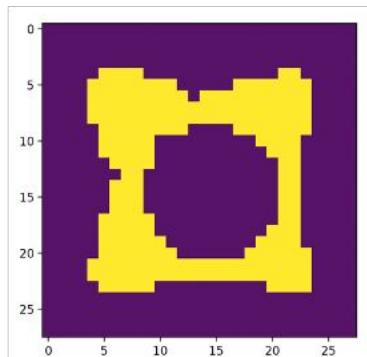
MNIST

28 x 28 pixel
images



Row Graph
28 x 28

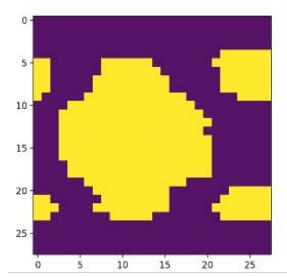
Use Kronecker
Product



Column
Graph
28 x 28

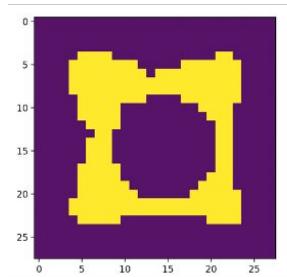
Produces 784 x
784

Clustering MNIST



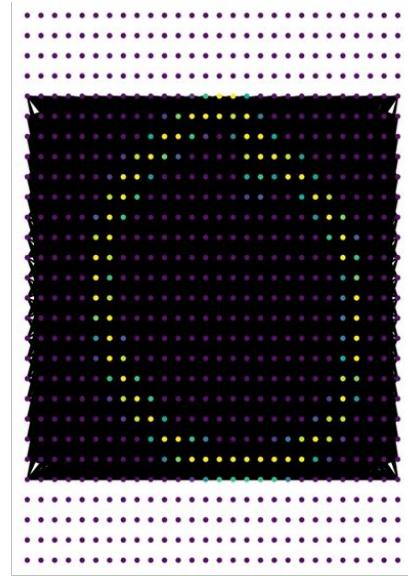
Row Graph
28 x 28

$$A_r$$



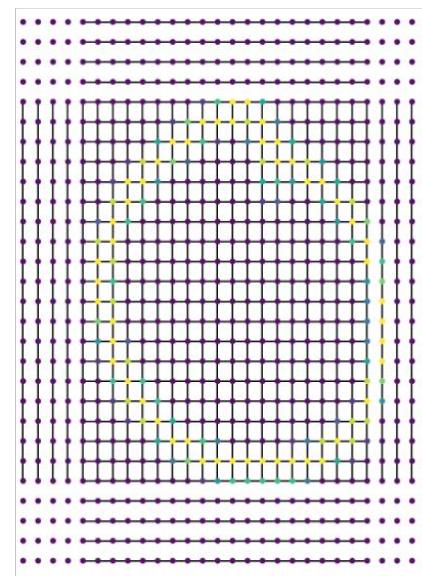
Column Graph
28 x 28

$$A_c$$



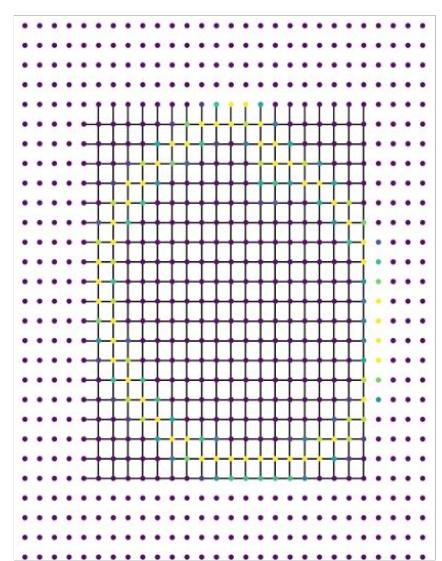
$$A_1 = A_r \otimes A_c$$

Kronecker Product
Very Dense



$$A_2 = A_r \otimes I + I \otimes A_c \quad A_2 \odot (A_r \otimes A_c + A_c \otimes A_r)$$

Cartesian Product
Grid Graph?



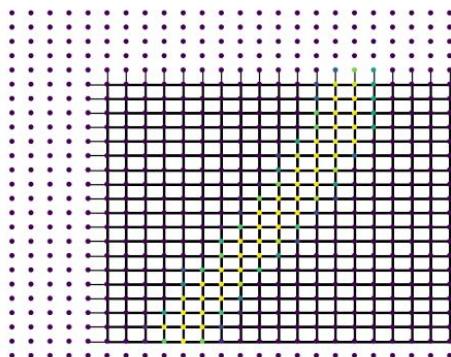
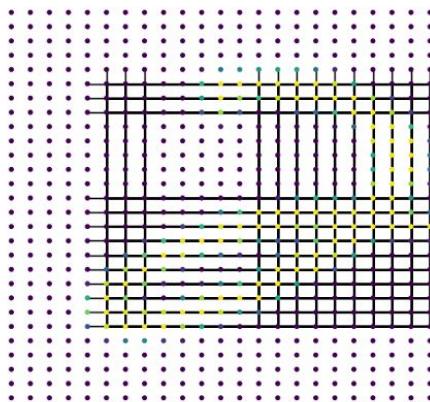
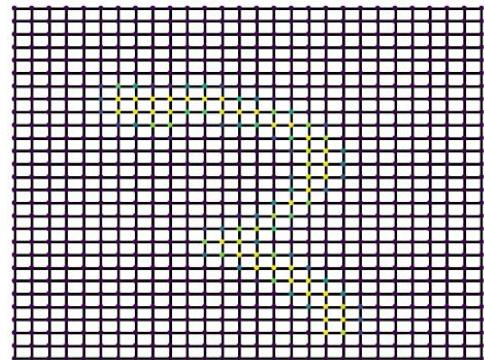
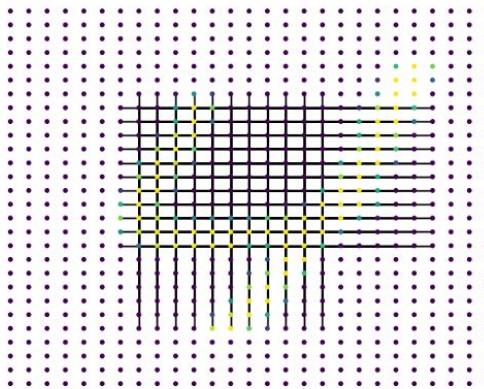
We use this
one!

Examples:

$$A_2 = A_r \otimes I + I \otimes A_c$$

$$\mathbf{G} = A_2 \odot (A_r \otimes A_c + A_c \otimes A_r)$$

Few Visual Examples of Graphs for MNIST:

Graph Representation ($G_2 = k * (k_2 + k_3)$)Graph Representation ($G_2 = k * (k_2 + k_3)$)Graph Representation ($G_3 = (k_2 + k_3) * k$)Graph Representation ($G_2 = k * (k_2 + k_3)$)

What about the features?

1 Feature: Pixel values.

Can we use more descriptive features/graph signals besides pixel values?

What about the features?

- Mean at each pixel position:

$$M(i, j) = \frac{1}{9} \sum_{m=-1}^1 \sum_{n=-1}^1 I(i + m, j + n)$$

- Variance at each pixel position:

$$\text{Var}[X] = \text{E}[X^2] - (\text{E}[X])^2.$$

What about the features?

- Gradient Magnitude:

$$G_x = S_x * I$$

$$G_y = S_y * I$$

where:

- I is the input image.
- S_x and S_y are the Sobel kernels applied along the x -axis and y -axis.
- $*$ represents the convolution operation.

The Sobel kernels:

$$S_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

$$S_y = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$

Gradient magnitude:

$$G = \sqrt{G_x^2 + G_y^2}$$

where:

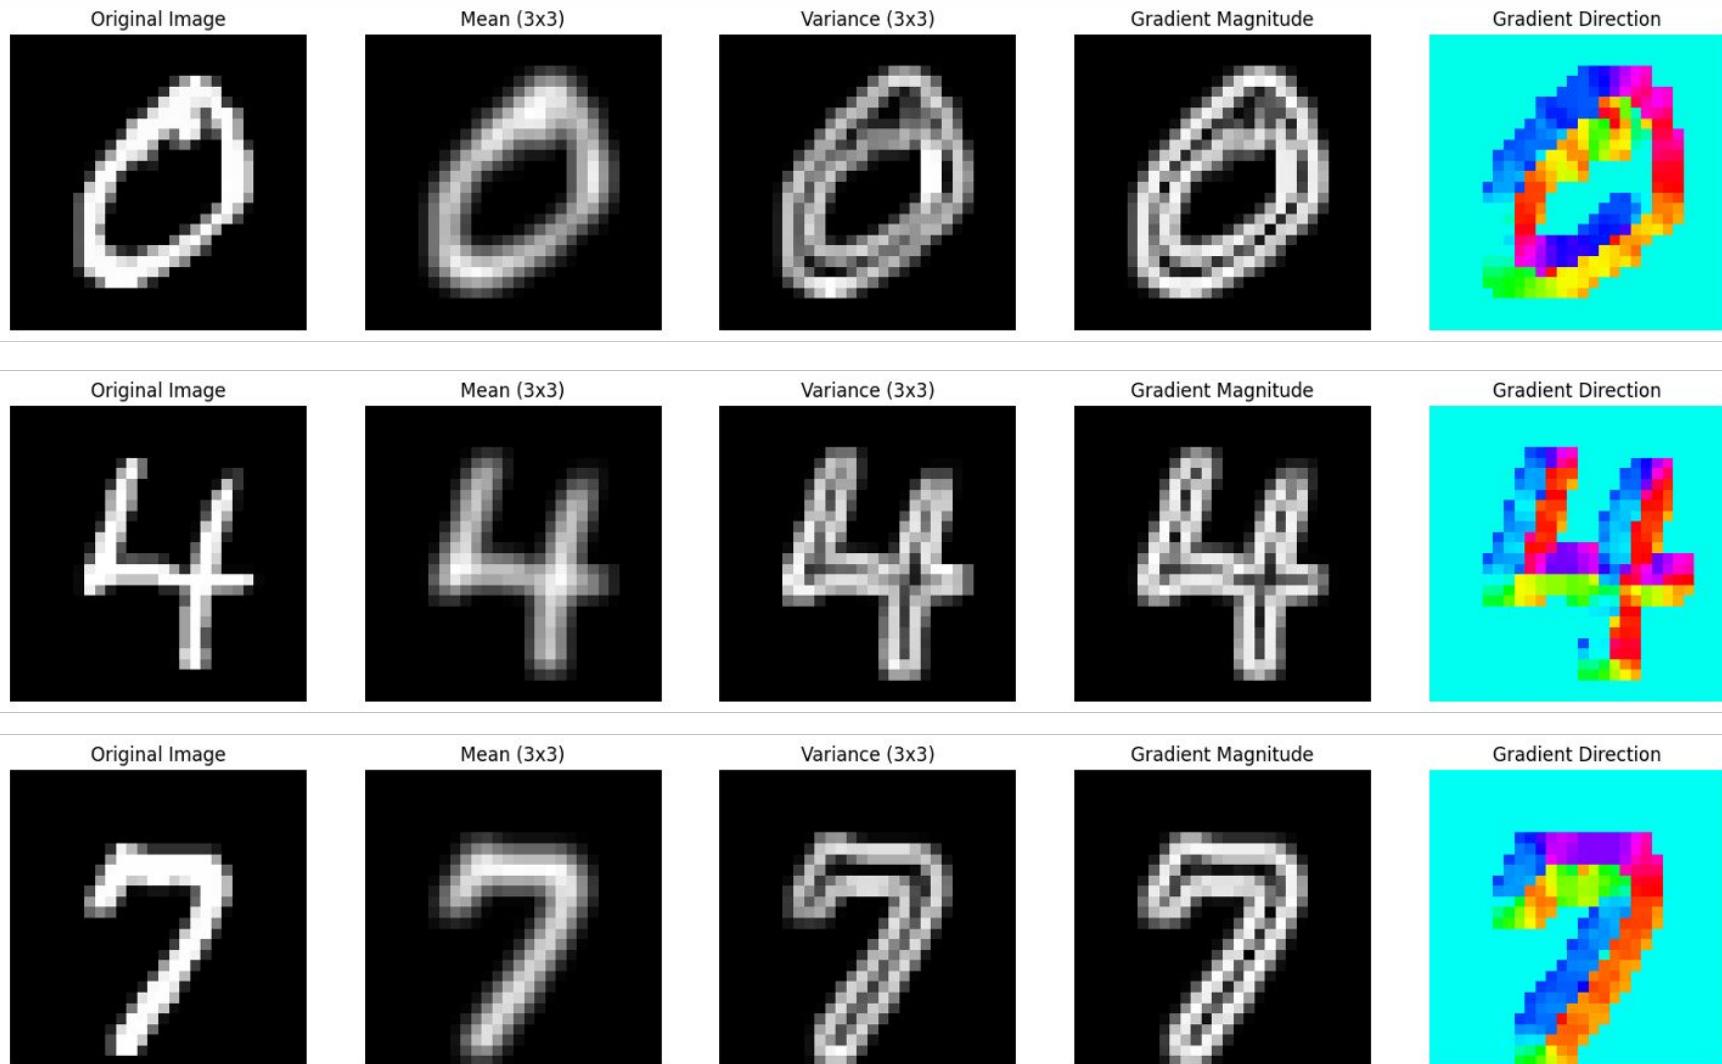
- G_x is the gradient in the horizontal direction.
- G_y is the gradient in the vertical direction.
- G represents the final gradient magnitude image.

- Gradient Direction:

The gradient direction $\Theta(i, j)$ at a pixel (i, j) is:

$$\Theta(i, j) = \tan^{-1} \left(\frac{G_y(i, j)}{G_x(i, j)} \right)$$

Example Features



Features also affect accuracy!

GNN Model	Number of Node Features	Test Accuracy (mean_acc \pm std_deviation)
GCN	1 (Pixel Intensity)	0.5505 ± 0.0289
GCN	4 (Mean, Variance, Gradient Magnitude, Gradient Direction)	0.7206 ± 0.0401
GAT	4 (Mean, Variance, Gradient Magnitude, Gradient Direction)	0.7818 ± 0.0006

Node Features for Row and Column Graphs

I is (M, M) the original Image. For MNIST and Fashion MNIST $M=28$

$$C = \begin{bmatrix} 0 & 0 & 0 & \cdots & 0 & 1 \\ 1 & 0 & 0 & \cdots & 0 & 0 \\ 0 & 1 & 0 & \cdots & 0 & 0 \\ 0 & 0 & 1 & \cdots & 0 & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & 0 & \cdots & 1 & 0 \end{bmatrix}$$

For a given lag l , the lagged transformation of the image matrix I is computed as: We have selected lag $l = 0$ to 27

$$I_l = C^l I$$

where I_l represents the lagged version of the image matrix.

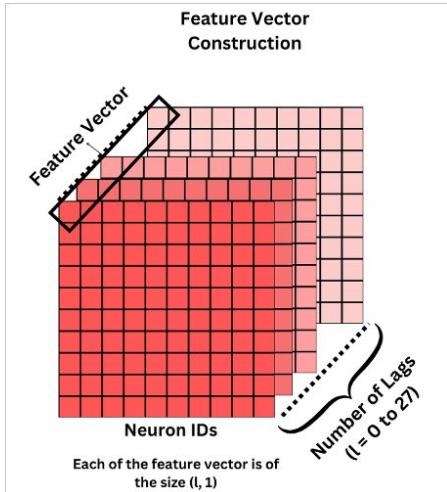
Node Features for Row and Column Graphs

The node feature matrix for lag l is given by:

$$F_l = \frac{I + I_l^T}{2}$$

where: I is the original image matrix; I_l^T is the transposed lagged version;
The result is an element-wise average.

Feature Vectors:



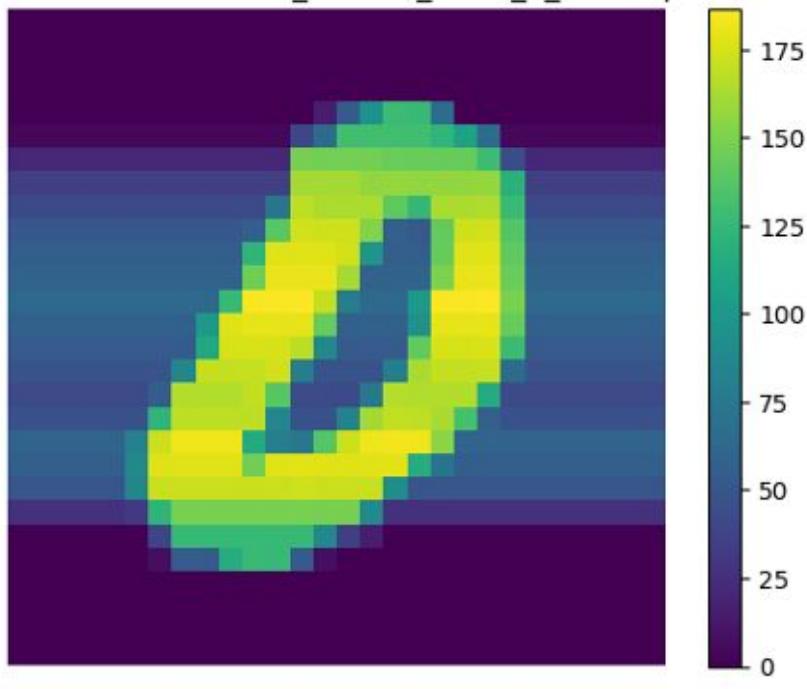
Node Feature Matrix :

$$\frac{1}{N} \sum_{l=0}^{27} F_l$$

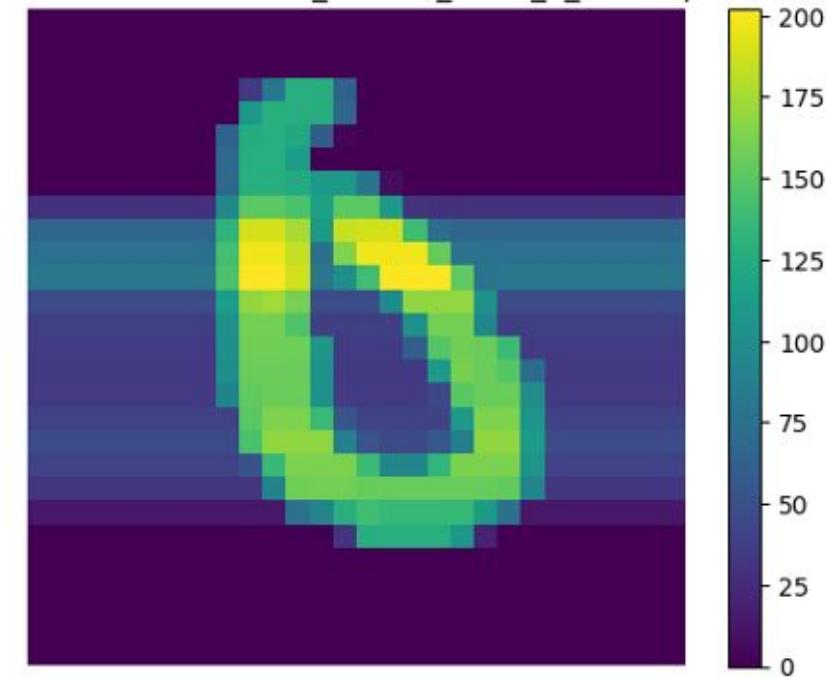
Where N is the total number of lags we are taking $N = 28$

Node Features for Row and Column Graphs

Mean Matrix - feature_vectors_Label_0_1096.npz



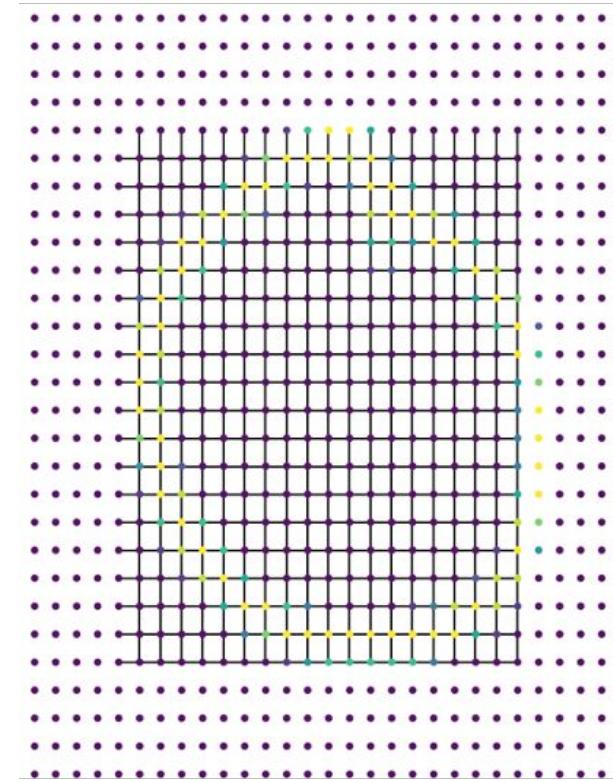
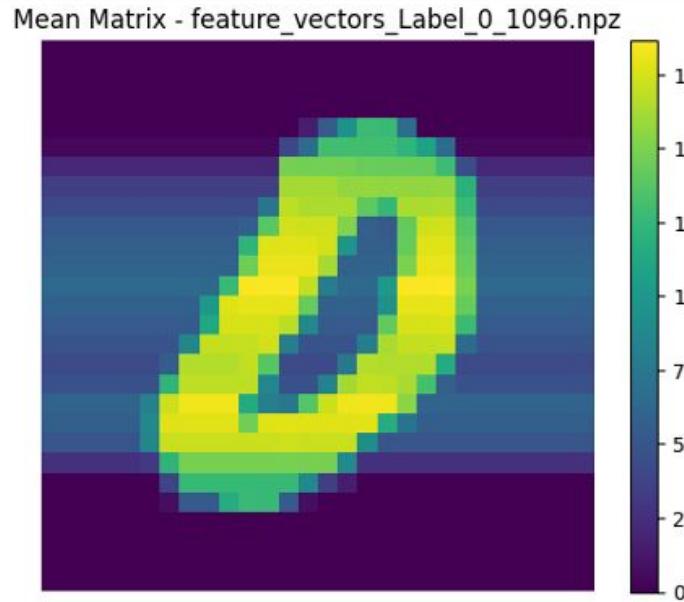
Mean Matrix - feature_vectors_Label_0_1097.npz



- Each row in the above image is fed as a feature vector to the nodes in the corresponding graph representation of the original image.
- i.e. each node as a node feature vector of size (28,1)

Similar process for column.

How do we map the (28,1) node feature for row and column to the 784 node Kronecker?



Why not just Kronecker the row and column feature matrices as well?

Results

Dataset Name	Grid Graph		Row Graph		Column Graph		Modified Cartesian Graph (Single Node Feature = Pixel Intensity)		Modified Cartesian Graph (4 Features/node, formed using modified cartesian method)		Modified Cartesian Graph (784 Features/node, formed using modified cartesian method)	
	GCN	GAT	GCN	GAT	GCN	GAT	GCN	GAT	GCN	GAT	GCN	GAT
MNIST	66.33	87.95	71.35	92.58	87.18	95.07	0.5505 ± 0.0289		0.7206 ± 0.0401	0.7818 ± 0.0006	0.9497 ± 0.00245	0.9691 ± 0.00016

- 1) GAT is better than GCN.
- 2) Kronecker does better than row/column, which does better than grid.
- 3) More useful features = better accuracy
- 4) More representative graph = better accuracy