

Introduction to Deep Learning

Recitation 0.4




Manigandan Ramadasan

Introduction

What is PyTorch?

- Open-source deep learning framework
- *Tensor* library with CPU and GPU support




Why do we need PyTorch?

- Rich Ecosystem
- Used widely in industries and research
- Makes life easier!
 - Has python APIs for building models, automatic differentiation etc.




Check out the following URL:
<https://landscape.pytorch.org/>

A Typical Model Training Process



(1) Preparing Dataset



(More about this in future Recitation 0)

torch.utils.data

Created On: Jun 13, 2025 | Last Updated On: Jun 13, 2025

At the heart of PyTorch data loading utility is the [torch.utils.data.DataLoader](#) class. It represents a Python iterable over a dataset, with support for

`class torch.utils.data.Dataset`

[\[source\]](#)


An abstract class representing a [Dataset](#).

All datasets that represent a map from keys to data samples should subclass it. All subclasses should overwrite `__getitem__()`, supporting fetching a data sample for a given key. Subclasses could also optionally overwrite `__len__()`, which is expected to return the size of the dataset by many [Sampler](#) implementations and the default options of [DataLoader](#). Subclasses could also optionally implement `__getitems__()`, for speedup batched samples loading. This method accepts list of indices of samples of batch and returns list of samples.


`class torch.utils.data.DataLoader(dataset, batch_size=1, shuffle=None, sampler=None, batch_sampler=None, num_workers=0, collate_fn=None, pin_memory=False, drop_last=False, timeout=0, worker_init_fn=None, multiprocessing_context=None, generator=None, *, prefetch_factor=None, persistent_workers=False, pin_memory_device='', in_order=True)`

[\[source\]](#)

Data loader combines a dataset and a sampler, and provides an iterable over the given dataset.



(2) Building the Neural Network



Build neural networks by stacking layers!
(and train their connection weights from data)

For more details, refer to:

<https://docs.pytorch.org/docs/stable/nn.html>


Linear Layers

nn.Identity	A placeholder identity operator that is argument-insensitive.
nn.Linear	Applies an affine linear transformation to the incoming data: $y = xA^T + b$.
nn.Bilinear	Applies a bilinear transformation to the incoming data: $y = x_1^T Ax_2 + b$.
nn.LazyLinear	A torch.nn.Linear module where <i>in_features</i> is inferred.

Convolution Layers

nn.Conv1d	Applies a 1D convolution over an input signal composed of several input planes.
nn.Conv2d	Applies a 2D convolution over an input signal composed of several input planes.
nn.Conv3d	Applies a 3D convolution over an input signal composed of several input planes.
nn.ConvTranspose1d	Applies a 1D transposed convolution operator over an input image composed of several input planes.
nn.ConvTranspose2d	Applies a 2D transposed convolution operator over an input image composed of several input planes.
nn.ConvTranspose3d	Applies a 3D transposed convolution operator over an input image composed of several input planes.

(3) Loss Functions




- Measure how far predictions are from the true labels
- Compute gradients – `loss.backward()`

More on loss functions at:
<https://docs.pytorch.org/docs/table/nn.html#loss-functions>

Loss Functions

nn.L1Loss	Creates a criterion that measures the mean absolute error (MAE) between each element in the input x and target y .
nn.MSELoss	Creates a criterion that measures the mean squared error (squared L2 norm) between each element in the input x and target y .
nn.CrossEntropyLoss	This criterion computes the cross entropy loss between input logits and target.
nn.CTCLoss	The Connectionist Temporal Classification loss.
nn.NLLLoss	The negative log likelihood loss.
nn.PoissonNLLLoss	Negative log likelihood loss with Poisson distribution of target.
nn.GaussianNLLLoss	Gaussian negative log likelihood loss.
nn.KLDivLoss	The Kullback-Leibler divergence loss.
nn.BCELoss	Creates a criterion that measures the Binary Cross Entropy between the target and the input probabilities:
nn.BCEWithLogitsLoss	This loss combines a <i>Sigmoid</i> layer and the <i>BCELoss</i> in one single class.
nn.MarginRankingLoss	Creates a criterion that measures the loss given inputs x_1, x_2 , two 1D mini-batch or 0D <i>Tensors</i> , and a label 1D mini-batch or 0D <i>Tensor</i> y (containing 1 or -1).

(4) Optimizers



Take the calculated gradients and update model parameters – `optimizer.step()`

More about `torch.optim` at:

<https://docs.pytorch.org/docs/stable/optim.html>

torch.optim

Created On: Jun 13, 2025 | Last Updated On: Aug 24, 2025

[torch.optim](#) is a package implementing various optimization algorithms.


Most commonly used methods are already supported, and the interface is general enough, so that more sophisticated ones can also be easily integrated in the future.

Algorithms

Adadelta	Implements Adadelta algorithm.
Adafactor	Implements Adafactor algorithm.
Adagrad	Implements Adagrad algorithm.
Adam	Implements Adam algorithm.
AdamW	Implements AdamW algorithm, where weight decay does not accumulate in the momentum nor variance.
SparseAdam	SparseAdam implements a masked version of the Adam algorithm suitable for sparse gradients.

Putting everything together

1. Data Preparation
 - a) Load Data
 - b) Dataset (apply transforms, etc.)
 - c) DataLoader – gives batches of (*inputs, labels*)
2. Build the neural network – Stack layers!
3. Define loss and optimizer
4. Train the neural network by going through the dataset *multiple times*
 - a) Get model predictions – *model(one batch of inputs)*
 - b) Compute the loss – *loss(prediction, labels)*
 - c) Clear/Zero out old gradients – *optimizer.zero_grad()*
 - d) Compute new gradients – *loss.backward()*
 - e) Tell optimizer to update model weights – *optimizer.step()*
5. Test the trained neural network



References

- [1] Previous course iteration's slides
- [2] PyTorch. (2024). PyTorch documentation — PyTorch 2.7 documentation. Pytorch.org. <https://docs.pytorch.org/docs/stable/index.html>
- [3] Roman, V. (2020, January 19). Deep Learning: Introduction to Tensors & TensorFlow | Towards Data Science. Towards Data Science. <https://towardsdatascience.com/deep-learning-introduction-to-tensors-tensorflow-36ce3663528f/>
- [4] Neural Network Training - an overview | ScienceDirect Topics. (n.d.).
Www.sciencedirect.com. <https://www.sciencedirect.com/topics/engineering/neural-network-training>
- [5] Päpper, M. (2022, February 28). Rethinking Batch in BatchNorm. Read. Hack. Learn. Repeat. <https://www.paepper.com/blog/posts/rethinking-batch-in-batchnorm/>
- [6] *Deep Learning: How do deep neural networks work?» Lamarr-Blog.* (2021, April 21). *Lamarr Institute for Machine Learning and Artificial Intelligence.* <https://lamarr-institute.org/blog/deep-neural-networks/>