INTRODUCTION TO WANDB

Weights & Biases

By: N.S.Akshara

WHAT IS WANDB?

WandB (Weights and Biases) is a powerful ML Operations platform that helps data scientists and ML engineers track, visualize, and manage their experiments seamlessly.

📊 Track Everything

Metrics, hyperparameters, model versions, and system resources

📈 Visualize in Real-Time

Interactive dashboards updated during training

🤝 Collaborate

Share experiments with teammates easily

🔧 Integrates Seamlessly

Works with PyTorch, TensorFlow, scikit-learn, and more

WANDB DASHBOARD

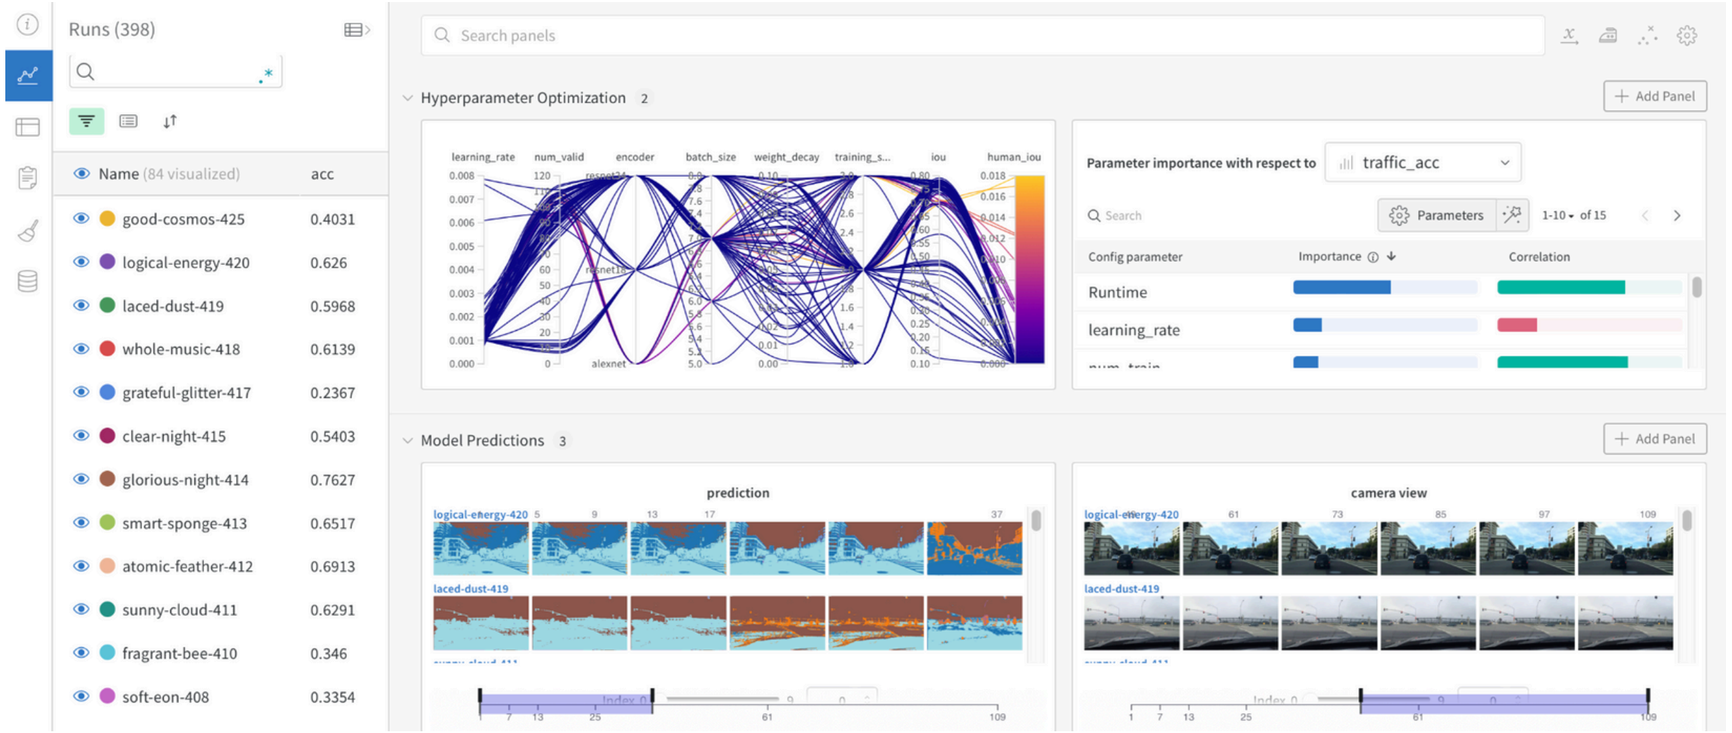

KEY FEATURES OF WANDB

Real-time Performance Logging

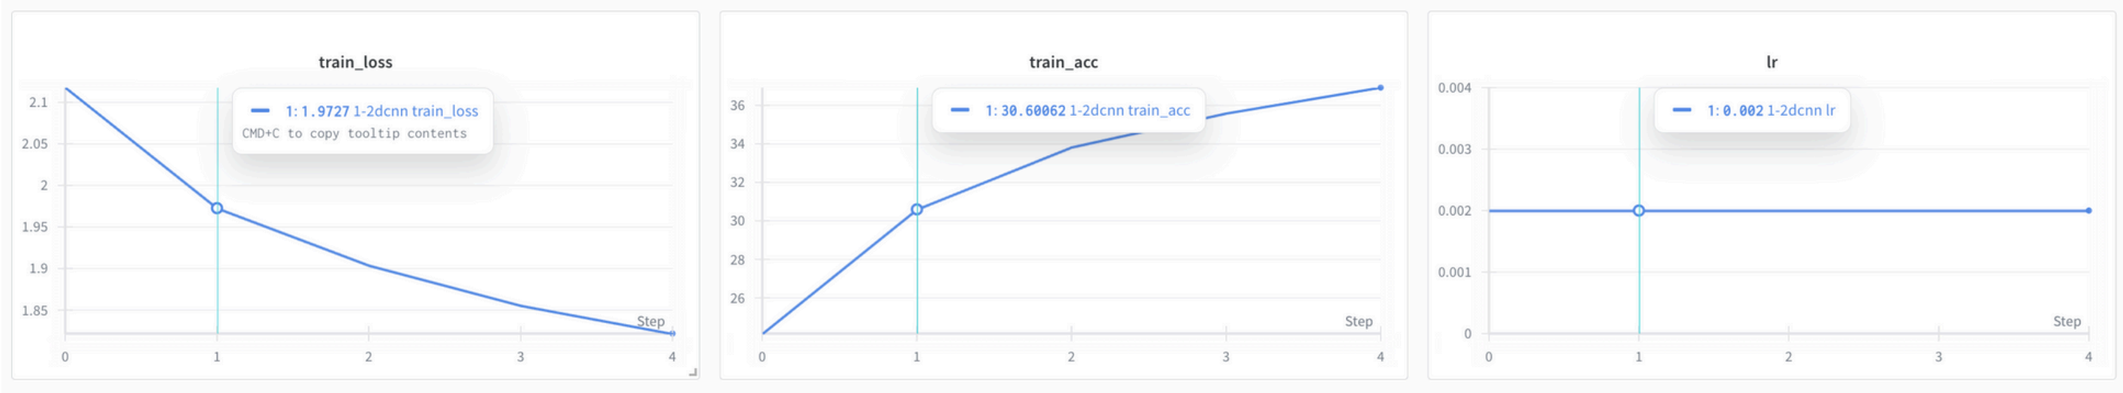

- Visualize metrics like Loss and Accuracy F1 in real-time as the model trains.
- Keep track of different models and training runs.
- Spot overfitting or exploding gradients instantly using interactive charts.

KEY FEATURES OF WANDB

Crash Recovery & Resumption

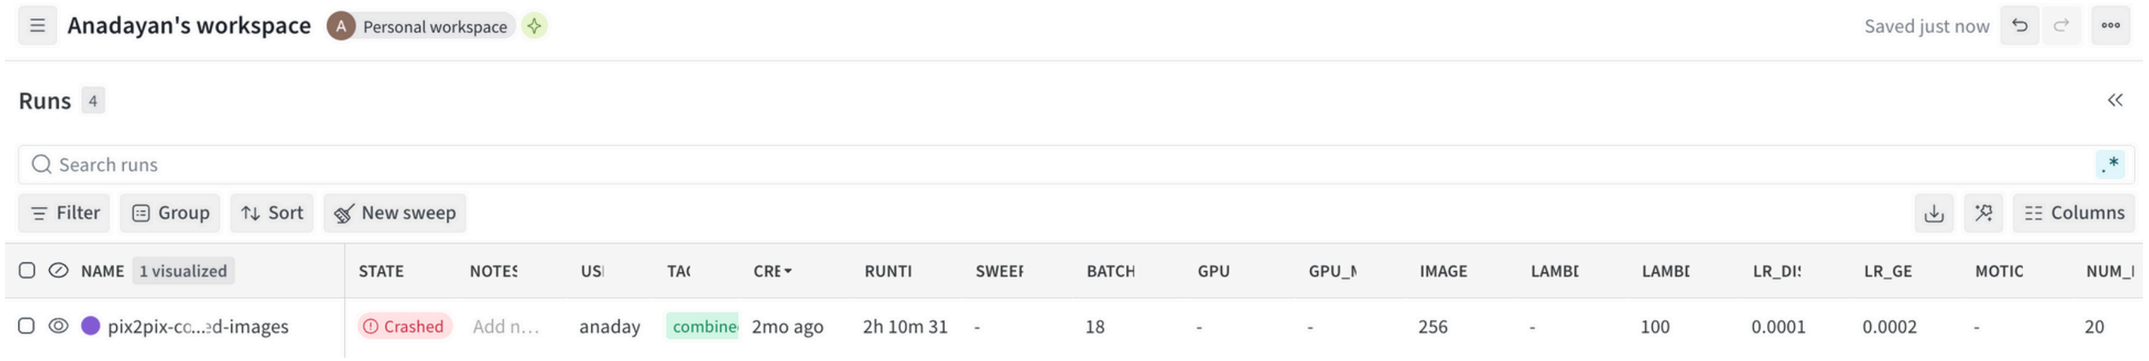

- WandB can save your model's checkpoints to the cloud.
- By passing the specific Run ID, you can download the weights and resume the epoch exactly where you left off.

KEY FEATURES OF WANDB

Hyperparameter Tuning

- Automating the search to find the best configuration (Learning Rate, Batch Size, Optimizer).
- Let WandB try different configurations while you sleep!

Collaboration

- All team members log runs to the same "Project" bucket.
- See everyone's runs side-by-side to determine whose model is performing best.

SET UP WANDB ACCOUNT

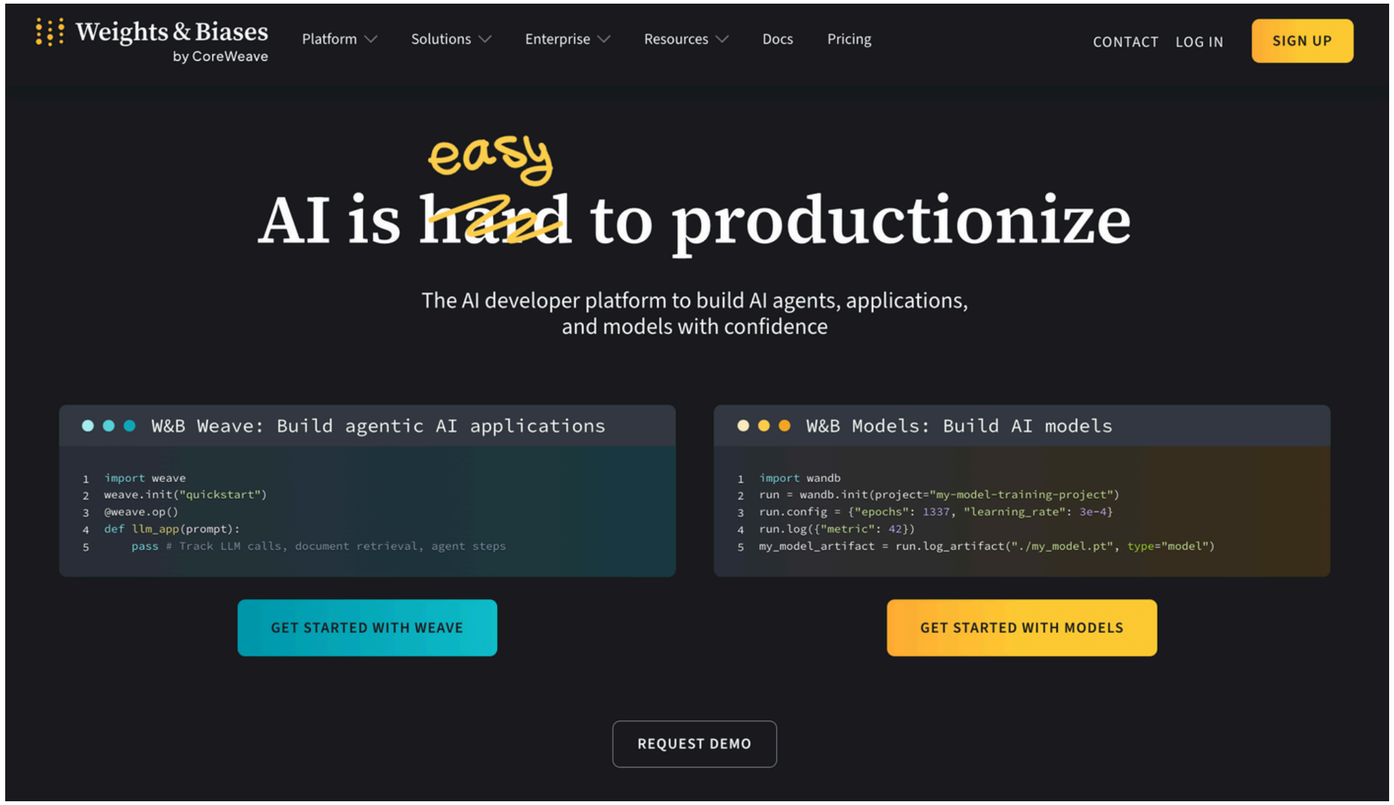

https://wandb.ai/site/

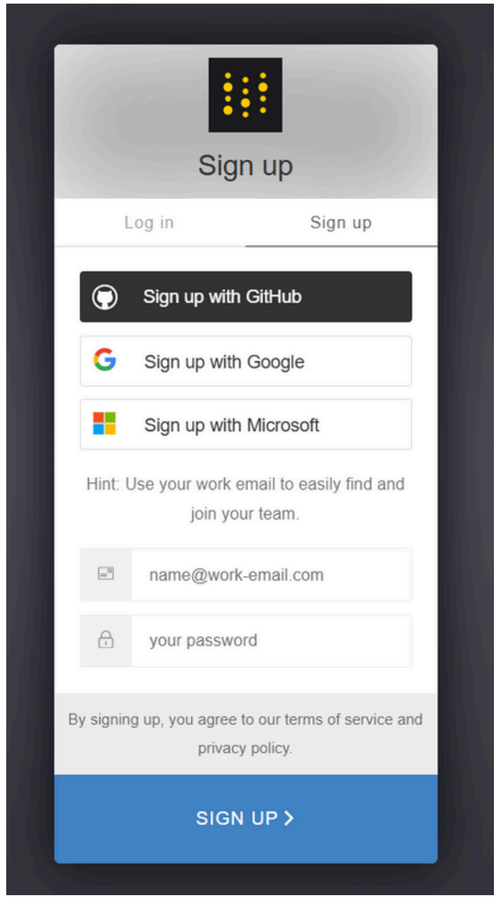

SET UP WANDB API KEY

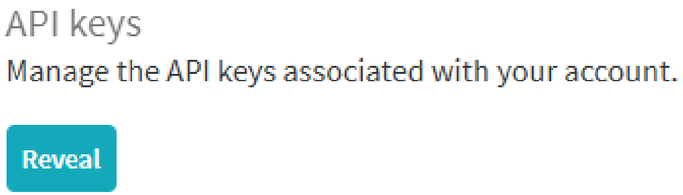

- Your API key authenticates your machine/notebook so runs can upload to your WandB dashboard.
- You can find your key in the “API Keys” section in your WandB user settings.

Never post your API key in GitHub, Piazza, screenshots, or shared notebooks!
→ If it does get leaked, please rotate the key immediately in WandB settings.

INSTALLATION & NOTEBOOK SET UP

Wandb Installation for Python

Installation and Libraries

```
[ ]  ## Installing WandB
     !pip install wandb -qqq
```

Notebook Setup

- Dataset: CIFAR10
- Neural Network: Convolutional Neural Network
- All functionalities are built in functional blocks for automated access.

Credits to Fall 2024 TAs for the Notebook!

NOTEBOOK: WANDB LOGIN

```
[ ]  import wandb, os
     os.environ['WANDB_API_KEY'] = "7328d336610dec777bdecff906dc57e2a464d0b4"#your key here
     wandb.login()
```

```
wandb: Using wandb-core as the SDK backend.  Please refer to https://wandb.me/wandb-core for more information.
wandb: Logging into wandb.ai. (Learn how to deploy a W&B server locally: https://wandb.me/wandb-server)
wandb: You can find your API key in your browser here: https://wandb.ai/authorize
wandb: Paste an API key from your profile and hit enter, or press ctrl+c to quit: ··········
wandb: Appending key for api.wandb.ai to your netrc file: /root/.netrc
True
```

NOTEBOOK: INITIALIZING PROJECT

```
run_config = {
    'model': '1-2dcnn',
    'optimizer':'sgd',
    'lr': 2e-3,
    'batch_size':64,
    'epochs': 5
}

train_loader, test_loader = build_data(run_config['batch_size'],
                                       data_train, data_test)

optimizer = get_optim(run_config['optimizer'], run_config['lr'], model)

criterion = nn.CrossEntropyLoss()

scaler = torch.cuda.amp.GradScaler()
```

```
run = wandb.init(
    #entity="wandb-starter",
    project="wandb-quickstart",
    #job_type="model-training",
    name=run_config['model'],
    config=run_config
    )
```

Tracking run with wandb version 0.23.1
Run data is saved locally in `/content/wandb/run-20251214_062711-h4lcg101`
Syncing run **1-2dcnn** to Weights & Biases (docs)
View project at https://wandb.ai/anadayan-cmu/wandb-quickstart
View run at https://wandb.ai/anadayan-cmu/wandb-quickstart/runs/h4lcg101

NOTEBOOK: LOGGING & SAVING MODELS

```
metrics = {
    "train_loss":train_loss,
    "train_acc": train_acc,
    'lr': lr
}
```

Logging metrics

```
        torch.save({
                'model_state_dict': model.state_dict(),
                'optimizer_state_dict': optimizer.state_dict()
                }, "Model.pth")

        # ALTERNATIVE 1: Saving Files as Artifacts
        # Creating Artifact
        model_artifact = wandb.Artifact(run_config['model'], type='model')

        # Adding model file to Artifact
        model_artifact.add_file("Model.pth")

        # Saving Artifact to WandB
        run.log_artifact(model_artifact)

        # ALTERNATIVE 2: Saving Files as Files
        wandb.save("Model.pth")

if finish:
  wandb.finish()
```

Saving models

NOTEBOOK: LOGGING & SAVING MODELS

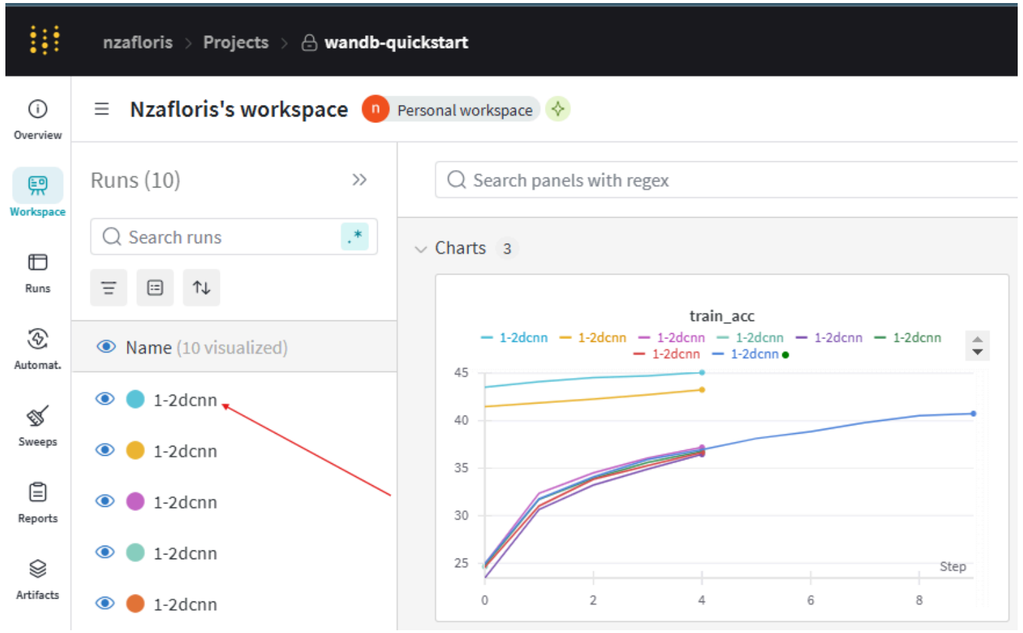

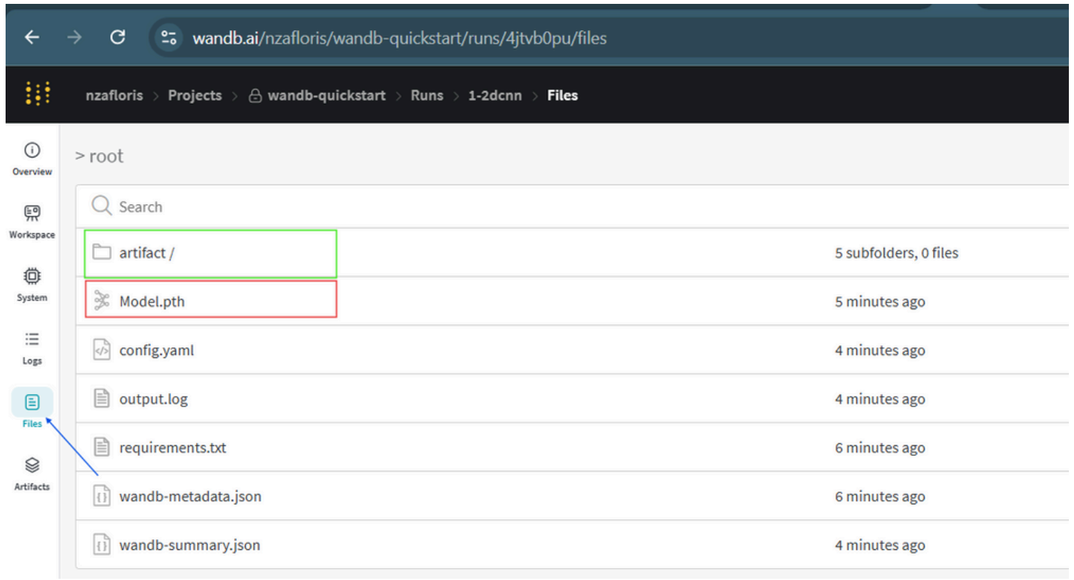

NOTEBOOK: GET RUN ID

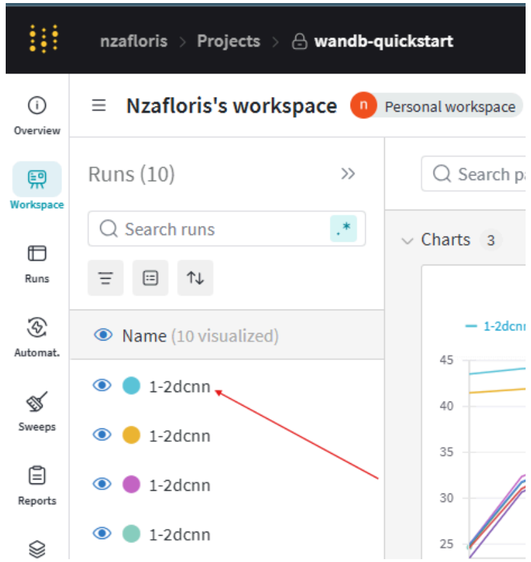

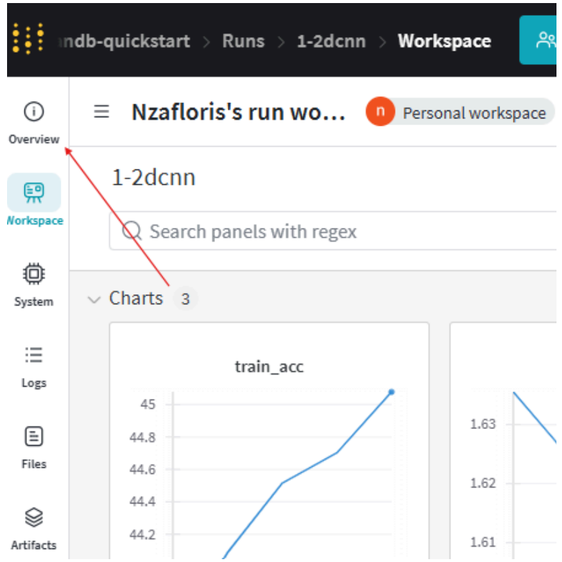

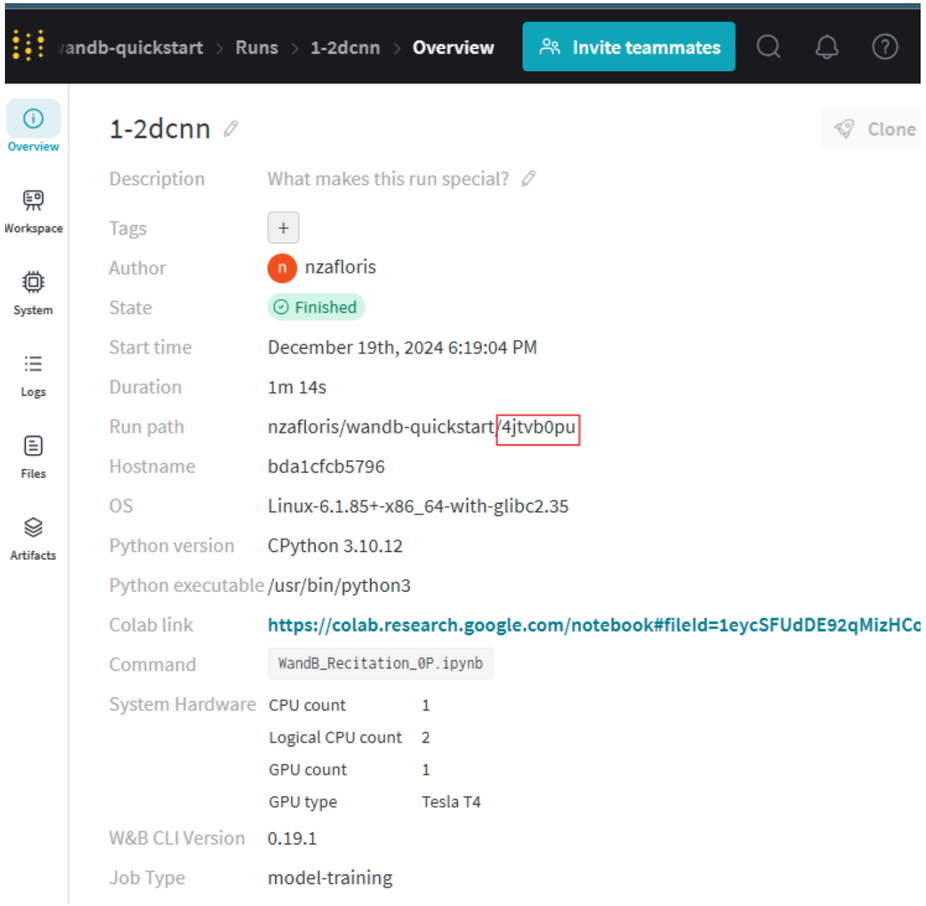

NOTEBOOK: RESUMING RUN

```
run_id = "4jtvb0pu" ### Replace with run id string
run = wandb.init(
    id     = run_id, ### Insert specific run id here if you want to resume a previous run
    resume = "must", ### You need this to resume previous runs, but comment out reinit = True when using this
    project = "wandb-quickstart", ### Project should be created in your wandb account
)
```

NOTEBOOK: HYPERPARAMETER SWEEPING

```
[ ]  # Hyperparameters to work with

     parameters_dict = {
         'optimizer':{
             'values': ['sgd', 'adam']
         },
         'learning_rate':{
             'distribution':'uniform',
             'min':2e-4,
             'max':1e-1
         },
         'batch_size': {
             'distribution': 'q_log_uniform_values',
             'q':4,
             'min': 16,
             'max': 128
         },
         'epochs':{
             'value': 5
         }
     }
     sweep_config['parameters'] = parameters_dict
```

```
[ ]  # Initalizing the sweep

     sweep_id = wandb.sweep(sweep_config, project="CIFAR-Sweep2")
```

```
def train_sweep(config = None):
    with wandb.init(config=config) as run:
        run.name=f"Jeel_{wandb.config.learning_rate}_{wandb.config.batch_size}_{wandb.config.optimizer}"
        config = wandb.config

        train_loader, test_loader = build_data(config.batch_size, data_train, data_test)

        model = Network().to(device)

        optimizer = get_optim(config.optimizer, config.learning_rate, model)

        criterion = nn.CrossEntropyLoss()

        scaler = torch.cuda.amp.GradScaler()

        for epoch in range(config.epochs):

            model, loss = train_epoch(model, train_loader, optimizer, criterion, scaler)

            wandb.log({'loss': loss})
```

```
[ ]  # Running the sweep

     wandb.agent(sweep_id, train_sweep, count=2)
```

MAKING YOUR WANDB PROJECT PUBLIC

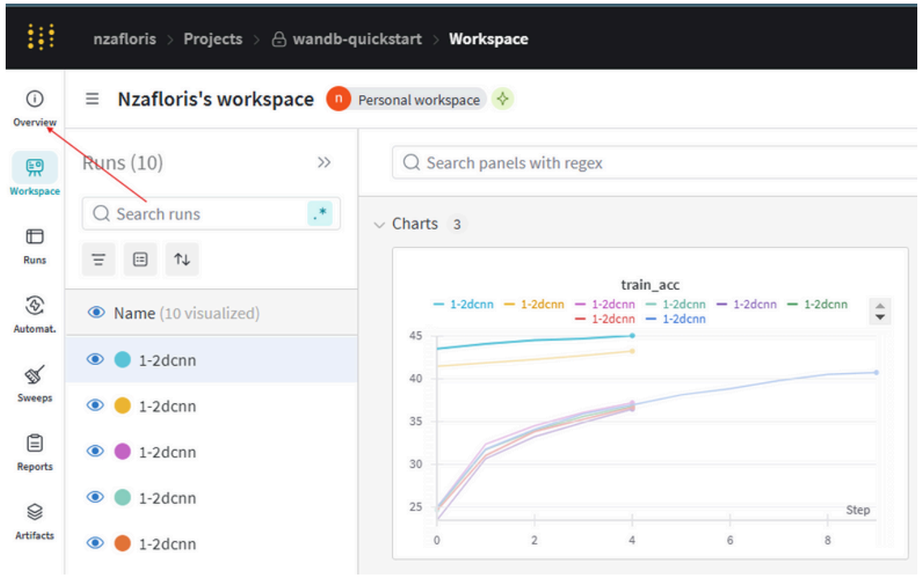

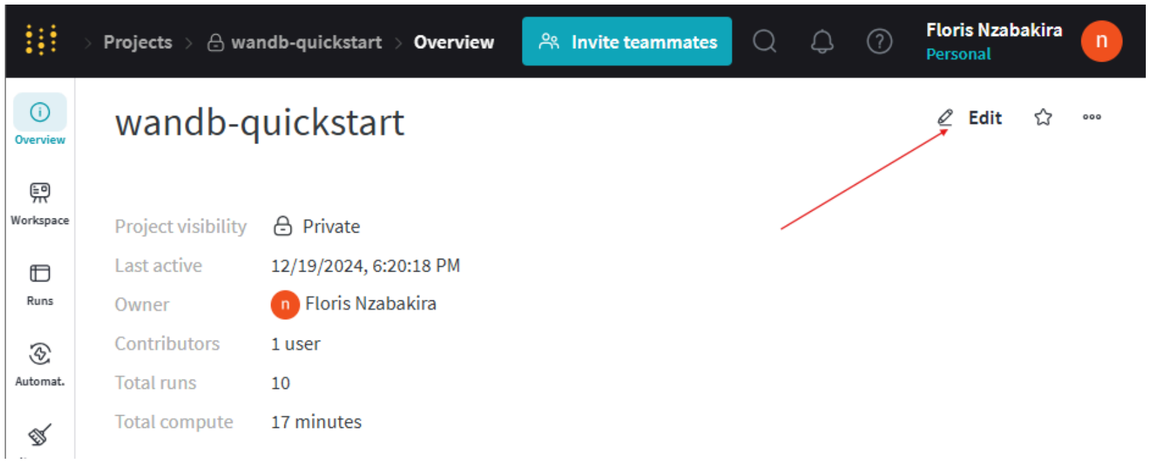

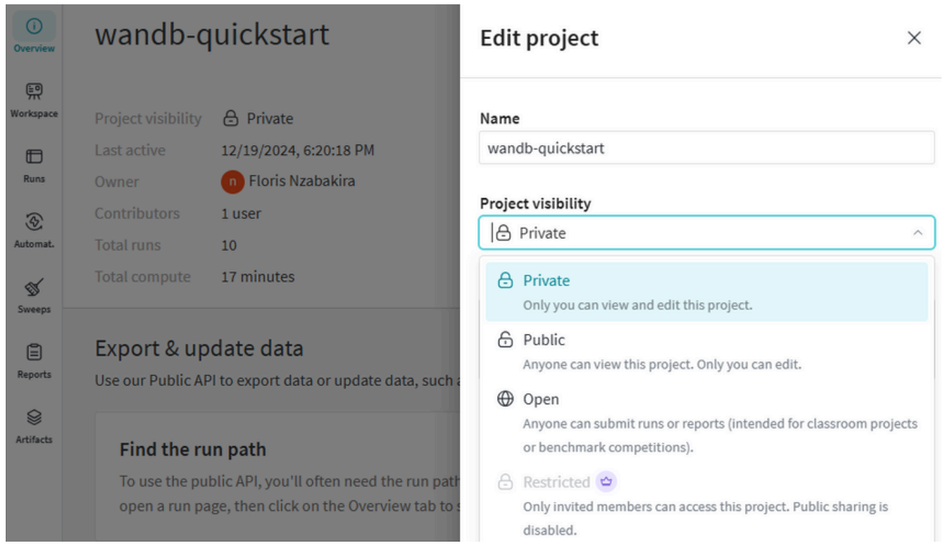

WANDB STORAGE FOR ACCOUNTS

Once your WandB free storage is full, you won't be able to access your runs!

- Individual Account:
  - Provides 100GB free
  - Can be accessed through https://wandb.ai/subscriptions

- Group Account:
  - Provides 5GB free per team
  - Can be accessed through: https://wandb.ai/account-settings/<Team-Name>/billing

THANK YOU!