

HW1P1 Bootcamp

Akshara, Anurag, Ron
01-17-26

Checklist

1. Getting Started

- Watch the recitation 0 lectures (like Python, PyTorch, Debugging, HW Workflows), if required
- Download code handout and extract the file
- Setup new environment and Install required python libraries
- Read the whole assignment write-up for an overview
- Once read the write-up, you can find all the formulas in the [Appendix](#) section for easy access
- Watch the 10-min video [Backpropagation](#) by 3B1B to understand the `backward` methods ¹.

2. Complete the Components of a Multilayer Perceptron Model

- Revisit lecture 2 about linear classifier, activation function, and perceptron
- Complete the linear layer class
- Complete the 4 activation functions

3. Complete 3 Multilayer Perceptron Models using Components Built

- Revisit lecture 2 about MLP
- Write a MLP model with 0 hidden layers
- Write a MLP model with 1 hidden layers
- Write a MLP model with 4 layers

4. Implement the Criterion Functions to evaluate a machine-learning model

- Revisit lecture 3 about Loss
- Implement Mean Squared Error (MSE) Loss for regression models
- Implement Cross-Entropy Loss for classification models

5. Implement an Optimizer to train a machine learning model

- Revisit lecture 6 about momentum, and lecture 7 about SGD
- Implement SGD optimizer

6. Implement a Regularization method: Batch Normalization

- Revisit lecture 8 about Batch normalization
- Translate the element-wise equations to matrix equations
- Write the code based on the matrix equations you wrote

7. Hand-in

- Set all flags to `True` in `hw1p1_autograder_flags.py`
- Make sure you pass all test cases in the local autograder
- Make the `handin.tar` file and submit to autolab

Please track your progress using the Checklist provided in the Writeup!

System Setup

- Step 1: Download the Anaconda installer specific to your OS.
- Step 2: Open Anaconda Prompt (Windows) or Terminal (macOS/Linux).
- Step 3: File Extraction: `tar -xvf HW1P1_S26_handout.tar`
- Step 4: Navigate to your HW1P1 directory and create Environment:

```
conda create -n idls26 python=3.13 -y
```

- Step 4: Activate the Environment: `conda activate idls26`
- Step 5: Install the Requirements: `pip install -r requirements.txt`

Local Development & Submission Workflow

- Local Testing:
 - Set flags to `True` in `hw1p1_autograder_flags.py` to test individual components.
 - Execute autograder:
`(idls26) python3 autograder/hw1p1_autograder.py`
- Final Submission: Create submission file from the top-level directory:
`tar -cvf handin.tar models mytorch`

The Big Picture - Main Components

- Linear layer: Computes an affine transformation.
- Activation: Introduces non-linearity to allow the model to approximate complex functions.
- Loss: Quantifies the mismatch between prediction and ground truth.
- Optimizer: Adjusts model parameters to minimize the loss function based on gradients.
- Regularizer: Used to prevent Overfitting.




Figure A: An abstraction of an end-to-end topology.

Neural Network Layer

Linear Layer

- The Linear layer connects every input neuron to every output neuron via learnable weights and biases.
- Forward pass equation: $Z = A \cdot W^T + \iota_N \cdot b^T$
- Backward pass equations [Backpropagation]:

$$\frac{\partial L}{\partial A} = \left(\frac{\partial L}{\partial Z} \right) \cdot W$$


$$\frac{\partial L}{\partial W} = \left(\frac{\partial L}{\partial Z} \right)^T \cdot A$$

$$\frac{\partial L}{\partial b} = \left(\frac{\partial L}{\partial Z} \right)^T \cdot \iota_N$$

$$A \cdot W^T + \iota \cdot b^T = Z$$

-4	-3	-2	-1	1	-1	10	-3	-16
-2	-1	0	1	1	0	4	-1	-6
0	1	2	3	1	1	-2	1	4
2	3			1	1	-8	3	14

Figure B: Linear Layer Forward Example








1. ACTIVATIONS INTRO

The “Why”: Neural networks need nonlinear components to approximate nonlinear functions. Without activations, the network remains linear regardless of depth.

Types:

- Scalar Activation:** Applied element-wise like Sigmoid, ReLU and Tanh (each output depends on single input only)
- Vector Activation:** Each output element depends on all input elements (e.g. Softmax)

2. ALL ACTIVATIONS YOU IMPLEMENT



3. RECOMMENDATIONS

- Use the hints:
 - For ReLU :
 - For GeLU:

Hint: For coding, search and read the docs on `np.where`.

Hint: Search the docs of the `math` and `scipy` libraries

BUILDING THE MLPs:




Figure H: MLP 1 Example Topology (Hidden Layers = 1)

COMMON PITFALLS & FIXES:

Mistake 1: Treating Linear and Activation Differently

- Don't write separate logic for Linear vs Activation layers
- Solution: Both are just `self.layers[i]` - call their forward/backward methods the same way

Mistake 2: Backward in Forward Order

- Don't loop through layers 0 to end during backward pass
- Solution: MUST use `reversed(range(len(self.layers)))` to go from end to 0

Mistake 3: Wrong Dimensions

- Don't guess the input/output sizes for each Linear layer
- Solution: Draw the network on paper! Follow the arrows in Figures G, H, and I from the writeup

What is a Loss Function?

A loss function measures how wrong your model's predictions are. It outputs a single scalar value that quantifies the mismatch between:

- Model Output (A): What your network predicted
- Ground Truth (Y): What the correct answer should be

Lower loss equals better model! We use this value to update network parameters via optimization.

MSE Loss (Mean Squared Error) - For Regression

- Use Case: Predicting continuous values (house prices, temperature, stock prices)

$$\text{MSE} = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

$$SE(A, Y) = (\textcolor{blue}{A} - Y) \odot (\textcolor{blue}{A} - Y)$$

$$SSE(A, Y) = \textcolor{blue}{\iota_N^T} \cdot SE(A, Y) \cdot \iota_C$$

$$MSELoss(A, Y) = \frac{SSE(A, Y)}{N \cdot C}$$



Cross-Entropy Loss - For Classification

- Use Case: Predicting categories (cat vs dog, digit 0-9, spam vs not spam)
- Negative log probability of correct class

$$L = - \sum_{k=1}^K y_k \log(p_k)$$

Optimizers

Stochastic Gradient Descent (SGD)



$$W := W - \lambda \frac{\partial L}{\partial W}$$
$$b := b - \lambda \frac{\partial L}{\partial b}$$

$$v_W := \mu v_W + \frac{\partial L}{\partial W} \quad W := W - \lambda v_W$$
$$v_b := \mu v_b + \frac{\partial L}{\partial b} \quad b := b - \lambda v_b$$

Without Momentum

With Momentum

Regularization

$$\mu_j = \frac{1}{N} \sum_{i=1}^N Z_{ij}$$

$$\sigma_j^2 = \frac{1}{N} \sum_{i=1}^N (Z_{ij} - \mu_j)^2$$

$$\hat{Z}_i = \frac{Z_i - \mu}{\sqrt{\sigma^2 + \epsilon}}$$




Figure L: Batchnorm Topology


Batch Normalization

$$E[Z] = \alpha * E[Z] + (1 - \alpha) * \mu$$

$$Var[Z] = \alpha * Var[Z] + (1 - \alpha) * \sigma^2$$

$$\hat{Z}_i = \frac{Z_i - E[Z]}{\sqrt{Var[Z] + \epsilon}}$$

$$\tilde{Z}_i = \gamma \odot \hat{Z}_i + \beta$$



Troubleshooting